

AI-Attribution Statement

Prepared 14 September 2026, covering the project from its 26 August 2026 ratification of the Engineering Constitution to date.

What this document is.

Item 5 of the six criteria docs/PHASE7_DECISIONS.md sets for calling this project “portfolio-ready.” Its own words for the job: “What Viktor designed, decided, tested and ruled on, versus what Claude produced... never assembled into one place.” This is that assembly. It is a separate document from Phase7_Portfolio_Document.pdf (item 4), which narrates the engine and its history; this document instead accounts for authorship -- who decided, who wrote, who was overruled, and by what evidence each of those claims can be checked.

Method

Every claim in this document is checkable against something outside it: an inline code comment (grep-able by the tag "VIKTOR'S RULING"), a dated entry in the engineering log (docs/Phase7_Engineering_Notes.pdf, itself append-only and never retroactively edited), a commit message, or the rulings recorded in docs/PHASE7_DECISIONS.md and docs/PHASE7_HISTORY.md. Nothing here is asserted from memory alone. Where the underlying evidence is a memory of a conversation rather than a committed artifact, this document says so rather than presenting it with the same weight as a commit hash.

The engineering log alone -- docs/build/build_engineering_notes.py, counted directly from its own source rather than estimated -- carries 118 numbered entries as of this writing, tagged by kind:

Tag	Count	What it records
DECISION — ADOPTED / ON ICE	30	A ruling Viktor made and acted on (or deliberately deferred).
PROCESS — RECORDED	20	A change to how the work itself is done.
OBSERVATION — FILED FOR THE RECORD	17	Noted, not (yet) acted on.
FINDINGS — FIXED / FROM EXECUTION / FROM AN UNFINISHED AUDIT	22	A defect: found, and (in 11 of the 22) closed.
MILESTONE — RECORDED / PHASE A COMPLETE	7	A phase or step completed.
AUDIT RUN/STEP — RECORDED or VERIFIED	7	An independent model's audit pass, logged.
EXTERNAL ASSESSMENT	6	An outside party's judgment on the project itself (Grok, Gemini, GPT-6 Astra, Meta Muse Spark).
REFERENCE — FILED FOR THE RECORD	5	Background material kept for context.
CANDIDATE — NOT ADOPTED	2	Proposed, considered, and deliberately not taken.
IDEA / KEY INSIGHT	2	A single unevaluated idea, and one named structural insight.

118 entries, verified by parsing docs/build/build_engineering_notes.py's own source (every entry_box()/highlighted_entry_box() call, numbered 1 through 118 with no gap and no duplicate) rather than counted by eye.

Roles, as the Constitution itself defines them

The Phase-7 Engineering Constitution states the working relationship directly, in its own "Roles & Authority" section: this document assigns Claude the technical and architectural judgment calls -- not blind authority. "The working relationship is: Claude proposes, reasons, challenges, implements, and tests; Viktor reviews, questions, performs QA, and ultimately accepts or rejects the change." Audit independence is structured the same way: "Claude prepares the audit package... the independent auditor writes none of the findings" -- a third party, neither Viktor nor Claude, evaluates the engine against the register both of them helped write.

In practice this produced three distinct roles rather than two: Viktor, who ratifies, rules, and accepts or rejects; Claude, who drafts, implements, tests, and researches, under review; and a rotating set of independent AI models, commissioned specifically because neither Viktor nor Claude can certify Claude's own work as compliant with rules Claude helped write.

Decisions Viktor made himself

A representative set, not an exhaustive one -- the engineering log's 29 "DECISION — ADOPTED" entries carry the complete record. Each of these below is independently checkable: the inline comments cited exist verbatim in the current source.

Date	Ruling	Where it is checkable
27 Aug 2026	Item 6 (Traceability) raised from Major to Critical, on the principle that severity reflects consequence rather than repair effort.	docs/PHASE7_HISTORY.md
29 Aug 2026	"Degrade, don't halt" — a failed input is recorded and confidence capped, rather than the run stopping.	models/decision_model.py's _apply_degradation()
29 Aug 2026	The engine must not compute monetary position sizing; a portfolio layer's job, not the engine's.	core/config.py, SEQUENCE ITEM 13 comment
31 Aug 2026	Items 3, 11, and 14's re-audit rulings, decided personally rather than delegated ("decide items 3, 11, 14 myself").	docs/PHASE7_DECISIONS.md, commit c4dfcc7
2 Sep 2026	Bias is the sole source of trade direction; no other signal may open or override it.	models/decision_model.py, models/signal_router.py
5 Sep 2026	"A lean is not a case" — a directional bias below a fixed strength floor (MIN_ACTION_BIAS) waits rather than acting.	models/decision_model.py
5 Sep 2026	trend_health is an unsigned magnitude and may not be read as a directional signal on its own.	models/decision_model.py
14 Sep 2026	The setup's direction (long/short) must be stated in words on the panel, not only inferable from the targets — specified placement (under the DECISION line) himself.	core/panel_render.py, commit 66f1479
14 Sep 2026	A NEUTRAL bias prints no direction line at all, overruling Claude's own unrequested "no directional lean" addition.	core/panel_render.py, commit 39e0e79
14 Sep 2026	Closing round 6's findings needs the same auditor confirming its own fixes, not a fresh independent model — a genuinely independent pass is saved for immediately before backtesting.	docs/PHASE7_HISTORY.md

What Claude produced

Most of the code and most of the prose in this repository, under the review relationship described above -- proposed, reasoned about, implemented, and tested, then accepted, corrected, or rejected by Viktor. Concretely, across this project: the majority of the engine's Python implementation across `core/`, `data/`, `indicators/`, `models/`, `structure/`, and `utils/`; the drafting of the Engineering Constitution's prose (from a joint principles discussion, then revised across eight revisions in response to outside review); the engineering log and its 118 entries; the test suite's design and the golden-path/pinned-data reproducibility mechanism; the git archaeology and cross-referencing used to keep documents honest against the actual commit history (including, twice, catching this project's own documentation going stale); and the mechanics of every patch delivered into this repository -- verification against a fresh clone, `code_fingerprint` and `golden-snapshot` checks, and the commit messages themselves.

None of that authorship implies authority. Every one of those outputs reached the repository because Viktor reviewed and accepted it, and a meaningful fraction did not reach it unchanged -- Section 5 accounts for those.

Where Claude was wrong, and Viktor caught it

An attribution statement that only credits Viktor's decisions and Claude's execution would misstate the actual working relationship, which included real errors, caught by review, not by construction. Recorded here rather than omitted, per this project's own standing practice of keeping mistakes in the log rather than quietly correcting them.

What happened	How it was corrected
1 Sep 2026: Claude told Viktor "nine Major and Moderate findings still open," a count taken from an audit report's verdicts rather than checked against the actual code.	Checking each location the report quoted found three already closed by unrelated work; the true count was different. Recorded in <code>docs/PHASE7_NEXT.md</code> (now in <code>docs/PHASE7_HISTORY.md</code>) rather than silently revised.
3 Sep 2026: five separate errors in one session, all the same shape — a conclusion stated with the same confidence whether or not it had actually been checked (two incomplete change predictions, a shipped defect, three unread-first documentation alarms, one vacuous test assertion).	Viktor named the pattern explicitly and asked for a standing practice: state a claim's verification status inside the sentence making the claim, not as a caveat after it. Adopted as a working rule, still in force.
Some point before 29 Aug 2026: the engine computed monetary position sizing (account balance and risk-percent constants, read only by a sizing block in <code>engine_core.py</code>).	Viktor ruled the engine must not size positions — a portfolio layer's job, not an analytical tool's. The computation and its constants were removed entirely, not gated behind a flag.
Constitution drafting: a rationale text drifted toward implying the engine should size positions, introduced during an earlier revision.	Viktor caught the drift on his own reading and it was corrected in Revision 5, before ratification — recorded in the Constitution's own revision history rather than silently rewritten.
14 Sep 2026: Claude added an unrequested "NEUTRAL — no directional lean this run" line to the panel, reasoning it matched the style of other "not located" messages elsewhere.	Viktor ruled a NEUTRAL bias is a correct reading, not a failed measurement, and should print nothing — overruling the addition. Landed at commit 39e0e79.

Independent third parties

Neither Viktor nor Claude certifies the engine compliant with the Constitution -- that is the audit's entire reason to exist, and it is why this section is neither "Viktor" nor "Claude." Multiple AI models across multiple labs (at various points: DeepSeek, Kimi K3, GLM 5.3, GPT-5.6 Luna Pro, GPT-6 Astra, and Meta's Muse Spark, among others, per README.md and the engineering log) have reviewed the Constitution, audited the engine's source, or verified a specific set of fixes. Which document each model has seen is tracked deliberately, because a model that has read the standard it is later asked to audit against is no longer independent of it -- a distinction this project treats as load-bearing rather than a formality, and checks explicitly before naming an auditor for a given round.

One exposure is stated rather than glossed over: three models (Claude Sonnet 4, DeepSeek V3, and DeepSeek R1) touched this codebase during the build itself via an AI coding tool, which means a later, nominally blind audit run by a DeepSeek-lineage model carried weaker independence than it first appeared. Independence is tracked at the lab level for exactly this reason, not at the model-version level, since treating a version bump as a reset would make the safeguard ceremonial.