

The Phase-7 Documentation & Change-Log Standard

A Practical Companion to the Engineering Constitution — Recording What Changed, What It Did to the Engine, and Whether It Was Progress

Status: **Working Standard, v1.0** — in effect immediately

Date: **August 25, 2026**

Relationship to the Constitution: **a separate, practical companion document — not an edit to the Constitution's frozen 17 / 7 / 10 / 6 rules.** See the closing section for exactly why that distinction is being kept.

"I do not want to sell snake oil. I want to sell a documented and certified good product, as good as it can be for its specific purpose." — Viktor, on why this document exists.

Contents

- Why This Document Exists
- What the Constitution Already Requires — and Where It Stops
- The Change Impact Record
- The Record-Keeping Standard
- Worked Example: The Confidence-Score Rename
- Template — Start Using This Today
- How This Feeds a Future Customer-Facing Disclosure
- Relationship to the Constitution
- Version History

Why This Document Exists

This document exists because of one message. Viktor pointed out that if Phase-7 is ever sold to real customers — even at a low, accessible price — it can't be sold on confidence alone. It has to be sold on evidence: a documented, certified product, as good as it can be for its specific purpose, not something dressed up to look more capable than it's been shown to be.

That's a business concern and an engineering concern at the same time, and they turn out to require the same thing: a habit of recording, for every meaningful change to the engine, what changed, why, what effect it actually had once measured, and whether that effect was progress, regression, or still unclear. Not a memory of that — a written record. This document defines exactly what that record looks like and starts it today.

What the Constitution Already Requires — and Where It Stops

The Engineering Constitution already asks for most of the raw ingredients here. It's worth being precise about what's already covered before describing what's missing — the gap is real, but it's narrower than it might first sound.

Already required by the Constitution	What it covers
Traceability (Tier 1, Item 6)	Any output must be explainable, step by step, back to the raw data that produced it.
Reproducibility (Tier 1, Item 5)	Any past result should be reconstructable later — same data, same code version, same configuration.
Documentation of significant decisions (Tier 3)	Why a change was made, not just what changed — usually as a short Architectural Decision Record.
Fixed evaluation datasets (Tier 3)	A stable dataset changes get measured against, so “better” and “worse” mean the same thing from one change to the next.
Hypothesis-driven development (Tier 3)	Every change already has a measurement and an evaluation step built into its process, before it's accepted or rejected.
Regression tests & known-good checkpoints (Tier 3)	A change isn't allowed to silently destroy something that used to work, and every major version stays recoverable.

Taken together, these already mean the engine is *supposed to* be measured and documented at every change. What's missing is narrower: none of these require that record to live anywhere as one running, referenceable log. Right now, the evidence that a given change helped, hurt, or did nothing measurable could exist only in a chat history, a memory, or nowhere at all. That's the specific gap this document closes — not by adding a new rule to the Constitution today, but by starting the practice now.

The Change Impact Record

A Change Impact Record is a short, dated entry created every time something meaningful changes about the engine — closing the loop between “we made a change” and “here's what actually happened, measured.”

It's a close cousin of an Architectural Decision Record, but answers a different question. An ADR records the decision — what was decided, and why, at the moment the change was made. A Change Impact Record records the outcome — what was actually observed once the change existed, checked against a fixed dataset, and given an honest verdict. A significant change deserves both: the reasoning going in, and the evidence coming out.

The four verdict categories below are deliberately built the same way the Constitution already treats audit findings: real categories, including one that means “we don't know yet,” because guessing at a verdict would violate the same epistemic honesty the Constitution already requires of the engine itself.

1 Improvement

IMPROVEMENT

The change was measured against the fixed evaluation dataset (or otherwise verified) and produced a better, more reliable, or more correct result than before — with the evidence attached, not just the impression that it went well.

2 Regression

REGRESSION

The change made something measurably worse — even if it fixed the thing it was meant to fix. A regression isn't a failure to hide; it's exactly the kind of finding this log exists to catch before it reaches a customer.

3 Neutral / No Measurable Effect

NEUTRAL

The change did what it was supposed to do without making anything measurably better or worse — a legitimate, common outcome, especially for refactors, cleanups, or safety changes that aren't meant to move a metric.

4 Not Yet Measurable

NOT YET

The honest answer when there isn't yet enough data, time, or a suitable test to know the effect. This is not a placeholder for laziness — it's the same “Unknown” principle the Constitution already uses for audit findings, applied here to change impact.

The Record-Keeping Standard

Every Change Impact Record should carry the same fields, so entries stay comparable to each other over the life of the project — the same reasoning behind the audit-finding structure already used in the Constitution's Next Steps section.

Field	What goes here
Date & Change ID	When the change happened, and a short reference id so it can be cited elsewhere (a bug report, a future audit finding, a customer-facing claim).
What changed	A plain description — one or two sentences, no jargon required.
Why	The reasoning behind the change. Link to the ADR if one exists for this change.

Affected module(s)	Which part of the engine this touched — helps a future reader (or audit) understand blast radius.
How it was measured	Which fixed dataset, test, or comparison was used to check the effect. “I ran it and it looked fine” is not a measurement.
Before / after comparison	The actual numbers, behavior, or output difference observed — the evidence itself, not a summary of it.
Verdict	Improvement / Regression / Neutral / Not Yet Measurable — one of the four, chosen honestly, not the one that reads best.
Notes	Anything else worth knowing later — surprises, caveats, follow-up needed.

When an entry is required: anything that could plausibly affect what the engine outputs — a new indicator, a model tweak, a bug fix, a changed default, a reworked calculation. Formatting, comments, and other changes with no behavioral effect don't need an entry — the standard is “could this change what a customer would see,” not “did a file change.”

Worked Example: The Confidence-Score Rename

The Constitution already references a real incident from this project's history: a code path silently renamed `confidence_score` to `confidence`, and the mismatch broke fourteen downstream modules before it was caught. Here's what a Change Impact Record for that change would have looked like, had this standard existed at the time — reconstructed here for illustration, not as a claim about exactly how or when it was actually caught.

Field	Entry
Date / Change ID	(illustrative) — CHG-EXAMPLE-01
What changed	A code path's output field was renamed from <code>confidence_score</code> to <code>confidence</code> .
Why	Naming cleanup, intended to be purely cosmetic.
Affected module(s)	The renaming module, plus every downstream module reading the old field name.
How it was measured	Not measured before merging — this is precisely the gap that let the break happen silently.
Before / after comparison	Fourteen downstream modules began silently reading a missing or mismatched field.
Verdict	Regression — severe. Would have been caught immediately had a Change Impact Record been required before the change was considered complete.

This is exactly the kind of finding the standard exists to surface early rather than after the fact — and it's also why Tier 2 of the Constitution already contains a rule (“explicit, evaluated changes to interfaces or behavior”) written in direct response to this same incident. The two documents are reinforcing the same lesson from two different angles.

Template — Start Using This Today

A blank Change Impact Record, ready to copy for the next meaningful change made to the engine. No tooling required to start — a plain text file or spreadsheet row is enough; what matters is that it exists and gets filled in honestly.

Field	Entry
Date / Change ID	—
What changed	—
Why	—
Affected module(s)	—
How it was measured	—
Before / after comparison	—
Verdict	Improvement / Regression / Neutral / Not Yet Measurable
Notes	—

How This Feeds a Future Customer-Facing Disclosure

This is the point that connects the practice directly back to the reason Viktor raised it in the first place.

When it's time to write the document a future customer would actually read — something explaining what the engine does, how it does it, and how it was validated — every claim in that document should be able to point to something. Not “it works well,” but “here's the dated record showing what changed, how it was tested, and what the measured effect was.” A product described this way isn't snake oil, because the description isn't the only evidence for itself — there's a trail behind it. That trail is exactly what a running set of Change Impact Records becomes, almost as a side effect of just keeping honest records along the way.

This also means the customer-facing disclosure document, whenever it gets built, doesn't have to be written from memory or reconstructed after the fact. It gets pulled from this log — which is a much stronger position to sell from than trying to recall, months later, what was actually tested and what wasn't.

Relationship to the Constitution

Worth stating plainly, since the Constitution and this document sit right next to each other in purpose.

This standard is deliberately **not** an edit to the Engineering Constitution. The Constitution's scope — 17 Tier 1 invariants, 7 Tier 2 principles, 10 Tier 3 process items, 6 Tier 4 preferences — is frozen, on purpose, until the audit described in its Next Steps section actually runs. Adding a new rule to that document right now, even a good one, would break the exact discipline the Constitution's own “Scope Freeze” section exists to protect. Viktor's own instruction when this came up was explicit: this doesn't need to go into the Constitution right now.

What this document does instead is operationalize principles the Constitution already contains — Traceability, Reproducibility, Documentation of Significant Decisions, Fixed Evaluation Datasets — into one concrete, usable habit, starting immediately, without touching the frozen document at all. If it proves useful in practice, formalizing “Change Impact Logging” as an actual Tier 3 item is a natural, well-evidenced candidate for consideration once the audit runs and the freeze is revisited — but that’s a decision for later, made with real experience behind it rather than as a guess made today.

Version History

Version	Date	Status	Notes
v1.0	August 25, 2026	Working standard — in effect immediately	Initial version, written directly from Viktor's stated concern about selling a documented, certified product rather than an undocumented one. Not part of the Constitution's frozen scope; usable starting today.

This document is a working tool, not a governing one — it can be revised freely as the practice of using it reveals what works and what doesn't. The one thing worth keeping constant is the reason it exists: an engine sold honestly needs a paper trail, not just a pitch.