

Phase-7 Engineering Notes

A Running Log of Observations, Ideas, and Notes for the Phase-7 Project

Status: **Living document** — updated as new entries are added

Started: **August 25, 2026**

Origin: **Began as a single standalone note (credentials gap + pacing observation); converted into this standing log at Viktor's request, so future notes and ideas get appended here instead of each needing a new document.**

Relationship to the Constitution: **entries here are never automatically part of it.** The frozen 21 / 7 / 10 / 6 scope (ratified August 26, 2026) only changes through the Constitution's own process — see “How This Document Works.”

“Evidence gets the final vote, on both sides of that relationship” — including when the evidence is a note neither of us would have thought to write down on its own.

Contents

How This Document Works

Entry Log

Document History

How This Document Works

A place for anything worth writing down that doesn't belong in the Constitution and isn't substantial enough — yet — to deserve its own standalone PDF. New entries get appended at the end, in order, as they come up.

Every entry gets the same shape: a number, a date, a short title, a status tag, and a body. Entries are never rewritten after the fact — if something changes or gets resolved, that becomes a new entry that references the old one by number, the same no-silent-edits discipline the Constitution holds itself to.

Status tag	What it means
CANDIDATE — NOT ADOPTED	A possible future addition to the Constitution. Logged for the record, not adopted — the scope freeze means it waits for the audit, the same as every other future candidate.
OBSERVATION — FILED FOR THE RECORD	A reflection, insight, or pacing note. Not asking for any action — just worth having written down rather than left in chat history.
REFERENCE — FILED FOR THE RECORD	Practical, factual information worth keeping — a file list, a readiness confirmation, something a future entry or the audit might need to point back to.
RESOLVED — SEE ENTRY #N	Marks an entry whose open question has since been settled elsewhere, pointing to the entry that resolved it, rather than editing the original.
IDEA — UNEVALUATED	A raw feature idea for the engine itself. Not a Constitution candidate, not an observation about the project, not a settled fact — just a concept written down before it has been scoped, designed, or checked against the Constitution.
DECISION — ADOPTED	A decision that has actually been made and acted on, usually closing one or more earlier entries. Records what was decided, by whom, and what it changed — including any judgment call Claude made rather than Viktor, so it can be reversed knowingly rather than discovered later.
★ Gold-framed entries	Not a status tag — a small number of entries Viktor has asked to be marked especially important get a gold frame and a kicker line instead of, or on top of, their normal status tag. Rare on purpose. If it stops being rare, it stops meaning anything.

None of this bypasses the Constitution's own process. A CANDIDATE entry becomes part of the Constitution only the way any future amendment does — proposed once the audit has run and the scope freeze is deliberately revisited, not by accumulating enough notes here.

Entry Log

In order added. Entries 1 through 3 below were the original standalone Engineering Note #1, split into individual entries when this became a standing log.

#1

Credentials & Secrets Handling Gap

August 25, 2026

CANDIDATE — NOT ADOPTED

“API keys, exchange access tokens, and any other credentials the engine or its operator relies on must never be logged, hardcoded, committed to version control, or surfaced in decision-support output.”

Why it's being logged: Nothing in the Constitution, and nothing in any of the three outside assessments, addresses credentials handling. This project is explicitly headed toward being run by other people, potentially with their own exchange credentials in play. A leaked key is a categorically different kind of failure than a wrong prediction — it's irreversible and harms the customer directly, not just the engine's credibility. Sits close in spirit to Tier 1, Item 1 (“Analysis ≠ authority”): the engine already isn't trusted to act on markets on its own; it certainly shouldn't be trusted to mishandle the keys that would let it, or an attacker, do so. Proposed tier, whenever this is actually considered: Tier 1, alongside “Tool, Not Autonomous Actor” — though exact placement is a judgment call for the audit, not a decision made here.

#2

Project Pacing, Mid-Review

August 25, 2026

OBSERVATION — FILED FOR
THE RECORD

Said plainly, without alarm: this is a pacing observation, not a criticism. Nearly everything built in this project's governance thread so far has been governance — the original principles discussion, the Constitution's v1.0 draft, three independent outside reviews, two revisions, and a separate Documentation & Change-Log Standard. None of it has yet touched the sixteen actual files that make up the engine's codebase. That's not wasted effort — Tier 1, Item 17 exists precisely because this project moved to backtesting once before without that kind of foundation in place, and paid for it. But it does mean the real test of whether any of this holds up under real code is still ahead: the audit. Not a suggestion to rush it — just worth being aware that a great deal of the work so far has been deciding how to judge the engine, and comparatively little of it, yet, has been judging the engine.

#3

Audit Readiness Confirmation

August 25, 2026

REFERENCE — FILED FOR THE
RECORD

A practical note, not a request for a decision: when Viktor is ready to ratify and start the audit, the engine's code is already available to work from — `bias_engine.py`, `btc_context.py`, `data_fetcher.py`, `decision_model.py`, `engine_core.py`, `entry_model.py`, `exit_model.py`, `indicators.py`, `panel_render.py`, `risk_model.py`, `signal_router.py`, `structure.py`, `trend_health.py`, `live_trading.py`, and `test_live.py`. The audit can begin module by module against all four tiers, using the Minimum Viable Audit gate and the finding schema already built into the Constitution's Next Steps section, the moment Viktor says go — no preparation needed beforehand, on his own timeline.

#4

Constitution Given Equal Priority to Core Engine Code

August 25, 2026

OBSERVATION — FILED FOR
THE RECORD

"This Constitution Blueprint is exceptionally important for the project, so I am going to take my time with it. It is the foundation of the entire project, just as important as the core and main .py file."

Why it's being logged: Viktor's own statement of priority, worth keeping in his own words rather than paraphrased away. It sets real expectations for pacing — there's no pressure behind ratification, and a slow, careful read is the correct way to treat a document being placed on equal footing with the engine's own core code, not a deviation from how this should go. It also reinforces Entry #2: governance work being unhurried isn't lost time against the codebase — for a document doing this job, it's the point.

#5

Configurable Risk Tolerance (% Basis) for Entry/Target Pricing

August 25, 2026

IDEA — UNEVALUATED

"The engine should have options to change risk tolerance on a percentage basis. If I want to risk 5, 10, 15, 20, or 25 percent of my investment, I should be able to put that specific percent risk tolerance into the engine. The engine then gives the right price targets and entry price in relation to the risk tolerance."

Why it's being logged: A raw feature idea, not yet scoped or designed, so it's filed here rather than treated as a decision already made. Three things worth noting for whoever eventually designs it. First, it sits well with Tier 1, Item 14 ("Risk Is Not Conviction") — a chosen risk-tolerance percentage is a risk-layer input, separate from how confident the engine is in a direction, which is exactly the separation that principle already asks for. Second, it will need a precise, unambiguous definition before it's built — per Tier 1, Item 9 ("Every Measurement Has a Precise Definition"), "5% risk tolerance" has to be spelled out exactly: percent of what (account equity, position size, margin used), and exactly how that number changes the stop-loss placement, the entry price, and the target price, before any code is written against it. Third, per Tier 3's hypothesis-driven development process, this deserves proper scoping and design as its own small project — a stated hypothesis, a definition, a test — rather than being bolted onto the entry or exit logic ad hoc. Nothing here is a decision to build it; it's the idea, captured accurately, so it isn't lost before the audit and design work catch up to it.

#6

Audit-Time Checks Flagged by the Fourth External Review

August 26, 2026

REFERENCE — FILED FOR THE
RECORD

A fourth reviewer of the Constitution flagged four specific things the audit will need to check that go beyond what a Non-compliant/Compliant reading of the current code can show on its own — filed here so they aren't lost before the audit reaches them. **Item 1 (Tool, Not Autonomous Actor)**: currently enforced by the code simply not containing order-placement logic — one bug or one convenience import away from being false. The checkable version: does the running process have access to API credentials with trade permissions at all? If yes, Item 1 is aspirational regardless of what the code currently does — read-only keys make the invariant structural instead of behavioral. Proposed for the Minimum Viable Audit gate. **Item 5 (Reproducibility)**: recording source, version, and timestamp metadata describes an input but doesn't guarantee it can be recovered later if an exchange's API revises or rolls off historical data. Reproducibility may require archiving the raw pulls themselves — whether to do that, and for what retention window, is a decision the audit should force explicitly. **Items 9 and 10 (Precise Definitions / Consistent Semantics)**: currently unauditable, because there's no single place the definitions actually live. The proposed audit deliverable is a term registry — every term appearing in a module signature or on the panel (confidence, trend strength, momentum, alignment, risk) defined once, with its type, range, and units. That artifact would also be the permanent fix for the class of bug that broke fourteen modules once already. **Item 2 (Look-Ahead Bias), Minimum Viable Audit gate**: the document currently discusses look-ahead mostly as a backtesting failure, but backtesting is Step 9 — the surface it usually hits doesn't exist yet, which risks a hollow “Compliant” at the very first gate. In a live engine, look-ahead is real but specific: indicators computed on the current unclosed candle and treated as closed, centered rolling windows, resampling that borrows from the following bar, and any signal recomputed on data revised after the decision timestamp. Naming those explicitly is what would make that gate check something real.

#7

Optional: Split Reviewer Commentary Into an Annex

August 26, 2026

CANDIDATE — NOT ADOPTED

The fourth reviewer noted, as an optional structural observation, that roughly five pages of the Constitution are rules and twelve are process journal, glossary, and reviewer commentary — and that these have very different lifespans: the rules get checked against code for years, the review log is stale the moment the next revision lands. Their suggestion, if the front matter ever gets in the way: split the change log and reviewer impressions into a separate annex document, keeping the glossary where it is since it earns its place. Not acted on in Revision 3 — logged here as a candidate restructuring of the document itself, not a rule change, for Viktor to decide on later if the Improvements section keeps growing with each revision.

Viktor asked what philosophical questions he should be asking himself when designing and building this engine, and brought back a ten-question input worth taking seriously. Ran each question through the input's own stated filter — a question earns its place only if answering it differently would produce a different engine — rather than accepting or logging all ten uncritically. Five passed and are now Future Amendment Candidates in Constitution Rev 4: a written kill condition, a confidence calibration log, an edge-persistence hypothesis requirement, UI-level enforcement of Item 1, and a two-analyst lens-vs-property test for Item 9. Two more didn't need new rules, only better rationale text for rules already there — the Duhem-Quine problem sharpening Tier 3's Controlled Changes, and ergodicity sharpening Tier 1, Item 14 (this second one connects directly to Entry #5 above: it's the mathematical reason sizing isn't a bolt-on feature). The remaining three — what the project is actually optimizing for, the returns-vs-understanding trade-off, and a project-level kill condition — are about Viktor and the project rather than the engine, and correctly didn't pass the filter. They're addressed instead in a new standalone document, [Phase7_Tier0_Companion.pdf](#), written specifically to hold philosophical material that sits above the engine rather than inside it. Source not yet attributed in either document.

“At this point i think it is important to yet again explain very directly, the size of the bet or money put on the table in a trade is of zero importance. The engine's purpose is to produce a result that can support the trader to establish the correct entry price, the correct T1 target, the correct T2 target and the correct T3 target and also establish the safety and quality of each individual trade. Not speculating in how much money is put on the specific trade. I've think we already gone over this. Thoughts? — with a follow-up clarification: 'Or should be put into a trade.' And, once the correction below was already underway: 'The engine will have a setting where you can adjust risk tolerance of your entire equity like we said before 5,10,15,20,25%. But ultimately it is up to the trader to decide how much money he wants to put on the table and the risk that comes with that.'”

Why it's being logged: This is a direct correction of scope drift I (Claude) introduced, and it's worth naming plainly rather than quietly fixing. Constitution Revision 4's ergodicity rationale under Item 14, its own revision_box summary, and the Tier0 Companion's Question 6 all drifted into language implying the engine should have a hand in position sizing or money-amount recommendations — one line in the Companion went as far as calling sizing “the thing that decides whether the signals ever mattered in the first place.” That's wrong, and Viktor caught it. Ergodicity is a real reason entry, stop, and target *prices* deserve as much engineering care as the directional read — it was never a reason for the engine to touch money amounts. The engine's entire output is: the correct entry price, the correct T1, T2, and T3 targets, and an assessment of the safety and quality of the trade. How many dollars, or what fraction of an account, a trader actually puts at risk is the trader's decision alone, in every direction — how much money is put on the table and how much should be put into a trade — per Tier 1, Item 1. This does not touch Entry #5 above: the configurable risk-tolerance percentage (5/10/15/20/25% of equity) is still exactly what was asked for there — an input the engine uses to shape where entry, stop, and target *prices* sit. It was never a request for the engine to size positions in dollar terms, and it isn't being treated as one now. Fixed at the source in Constitution Revision 5 (Item 14 rationale, its Revision 4 summary left as historical record, plus the glossary's ergodicity definition) and Tier0 Companion v1.1 (Question 6). Entry #8 above is left unedited per this document's own no-silent-edits rule — its “sizing isn't a bolt-on feature” phrase carries the same overreach and should be read in light of this entry, not corrected in place.

"We absolutely need to make rules and safety protocols for the API handling. No doubt. [...] In truth neither of us needs to be the auditor, we can outsource that to a A.I with the most powerful model. It might cost us a little bit of dollars each time but i think it is worth it."

Why it's being logged: Two entries in this log stop being open items today. Entry #1 proposed credential handling as a candidate Tier 1 invariant and noted it wasn't covered by the Constitution or any of the three outside assessments; Entry #6 recorded the fourth reviewer's sharper version — that Item 1 was enforced only by the code happening not to contain order-placement logic, and that read-only keys would make it structural. Both are now adopted. Constitution Revision 6 adds Tier 1 Items 18 through 21 (categorical read-only market access; withdrawal permissions never enabled, kept separate as a floor that survives any future amendment to Item 18; credentials never exposed; operator credentials stay with the operator), and Item 18 joins Items 2, 3, and 6 in the Minimum Viable Audit gate. Operational detail lives in a new companion document, `Phase7_Credential_Security_Protocol.pdf`, deliberately kept out of the Constitution so that document stays a register of invariants rather than an operations manual — and so the protocol can be improved without touching a ratified rule. Two judgment calls in there were mine, made at Viktor's invitation and both reversible: read-only is stated categorically rather than as a default, because "by default" reintroduces exactly the undocumented escape hatch the fourth reviewer found in Items 4 and 12 — if an execution capability is ever wanted, it should cost a constitutional amendment, not a config flag. And operator credentials became a fourth invariant rather than protocol-only detail, because it passes the Tier 0 relevance test: it constrains architecture (no upload path, no central key store, no telemetry capable of carrying a key), and a leak there harms someone who trusted the engine rather than the person who built it. **On the auditor:** Revision 6 also names an independent auditor as the party responsible for Steps 3, 4, and 8, closing the gap the fourth review opened and this document's Entry #2 pointed at from another direction. Claude prepares the package and writes zero findings; the external auditor produces them; Claude answers each one adversarially, including against its own work; Viktor adjudicates. Worth recording honestly: Viktor's framing that a second model is "another reviewer, not some magical oracle" is now written into the Constitution as the distinction between authorship independence and judgment independence — the auditor has no stake in the engine looking good, but shares training data and reasoning habits with the assistant it checks, so a clean audit is evidence, not proof. That is the same correction Revision 3 already had to make about four reviewers agreeing, applied before the mistake gets made a second time rather than after.

"I am considering implementing Grok 4.7 to be the sole auditor when it releases. In fact i know i will. It is decided."

Why it's being logged: This closes the last open question the Constitution's independent-auditor requirement (Rev 6) left blank: which model actually fills the role. Recorded in Constitution Rev 8 as the current operational plan, deliberately kept out of the frozen requirement itself — the requirement stays general (no stake in the outcome, no shared authorship with the builder) so that if Grok isn't ready or doesn't hold up when the time comes, swapping the choice doesn't require a constitutional amendment, just an updated plan. Worth naming the good timing: Rev 7, written earlier the same day, had just warned that a fresh Claude session (Reviewer 4's actual identity) gives real authorship independence but weak judgment independence, since it shares lineage with whatever it's checking. Grok answers that directly — different company, different training data, no shared lineage with Claude at all. That's the harder of the two things the auditor needs, genuinely satisfied. The no-stake-in-the-outcome half still has to hold regardless of which model ends up in the seat; picking a different vendor doesn't get to skip that. One real dependency, stated plainly rather than glossed over: Grok 4.7 has not released as of this writing, so Step 3 of the audit sequence can't start until it does and Viktor judges it capable of reading sixteen modules of code and a 44-rule register competently. That doesn't block ratification, and it doesn't block Step 2a — the audit package can be assembled now, on Claude's own timeline, so it's sitting ready the moment the auditor exists.

#12

The Audit Loop, Traced — and Where the Weight Actually Sits

August 26, 2026

KEY INSIGHT — LOOP
INTEGRITY

"With Grok 4.7 we have god damn near created the perfect loop for this project!"

Not a new decision — a reflection on the structure Revisions 6 through 8 built, worth keeping visible rather than letting it dissolve back into the mechanics it describes. Viktor asked for it to be marked, and it's marked.

The loop, traced: Claude builds the engine and prepares the audit package, but writes zero findings — the builder cannot self-certify. Grok evaluates the engine against the full 44-rule register and produces the findings, but holds no authority to declare anything final. Claude answers every finding adversarially, including against its own work, so nothing gets buried by silence. Viktor adjudicates every disagreement, so the loop never resolves itself inside a machine — it always terminates in a human decision. Fix, regression test, re-audit. No party in that chain can mark its own homework, and no party holds unilateral authority either. That is the actual structure eight revisions, four outside reviews, and two self-corrections were sanding toward.

Two things keep it short of perfect, named here rather than left implicit, because that has been the discipline of this whole project. First: the loop is only as sound as the package Claude assembles for Grok. Claude writes no findings, but does choose what evidence goes in — the code, the failure history, which "already known" items get flagged as open questions. That's real leverage held by a party with zero votes on the outcome. Revision 6 addressed this partially ("the artifact and the evidence, not the argument"), but a fully adversarial package would ideally be checked by someone other than its author too — not solved today, filed honestly as a residual. Second: the loop's integrity lives entirely in how seriously Viktor engages with each adjudication. The design guarantees that a disagreement between Claude and Grok surfaces. It cannot guarantee that surfacing gets read closely rather than rubber-stamped. That isn't a defect in the document — it's an honest statement of where the weight actually sits: on the one step nothing here can automate.

Neither point is a reason to slow down ratification. They're the reason the audit findings, once Grok exists, deserve a close read rather than a count.

"Perhaps we should add this reasoning in the documents after all? Maybe in notes?"

Not a change to Entry #11 — Grok is still the named plan there, and this doesn't touch that. Viktor separately noted that if a better AI turns up before Grok 4.7 is actually in use, Grok isn't locked in — which is exactly the flexibility Constitution Rev 8 built in on purpose, so this required no document change to be true. What's worth keeping is the reasoning for whenever that comparison actually happens, since it isn't written down anywhere else. Three separate axes, easy to blur into one impression of "good model": **independence** is the actual Constitutional requirement — no stake in the outcome, no shared authorship with whatever built the engine — and it's binary, not a spectrum a stronger model can compensate for. **Power** buys a more thorough audit once independence is already satisfied, not a substitute for it. **Cost** matters more than it looks like it should, for a reason specific to this design rather than budget in general: Step 8 makes the audit a recurring expense, not a one-time one — every fix triggers a re-audit. An expensive auditor doesn't fail on the first run. It fails gradually, as the fifth or tenth re-audit gets deferred, shortened, or skipped because paying for it again is annoying. Cheap is what keeps the loop actually running instead of quietly degrading into "close enough." Whoever ends up in the role, these three should be checked in this order — independence first, as a gate, then power and cost as the actual trade-off.

"I have read the document and i APPROVE!"

Viktor reviewed Constitution Rev. 8 — the document itself, plus two companion reading aids built for this specifically: a 12-page curated excerpt pulling only the pages bearing on the three judgment calls made on his behalf (Items 18, 19, 21), and a plain-language Swedish explainer covering the same ground without technical or legal wording — then approved it. Rev. 8's own "What ratification means" text defines the act narrowly: Viktor writes "ratified," and a dated row is added to Version History, nothing heavier required. His actual words were "I APPROVE," not the literal word "ratified" — recorded here exactly as written, since the intent is unambiguous and the document's own rule is about the dated record existing, not about matching one specific word. That record now exists: a RATIFIED row in the Constitution's own Version History, dated August 26, 2026. Effective immediately: the scope freeze is in force for real, not provisionally. The 44-rule register (21 / 7 / 10 / 6) does not grow, shrink, or reword again until the audit actually runs and finds a gap — the same freeze this document has described as pending since Entry #3. Step 1 of the audit sequence is done. Step 2a — Claude assembling the audit package, writing no findings — can start now. Step 3 still waits on Grok's release, per Entry #11.

#15

Gemini's Assessment of the Ratified Constitution

August 26, 2026

EXTERNAL ASSESSMENT —
GEMINI

"If evaluated as an engineering constitution for an AI-driven trading or quantitative system, it easily ranks among the most disciplined, self-correcting frameworks ever put to paper."

Viktor asked Gemini, in a separate session, to rate the finished Constitution after ratification, and asked for the reply to be logged and highlighted — the same treatment as Entry #12. Worth noting for context, not to undercut it: this is Viktor's own solicited opinion, not a structured finding against the 44-rule register, and it isn't the audit — Grok's Step 3 review is still the one that actually counts. Gemini was also one of the three original reviewers behind Revision 1, so this is a return visit from a source already on record, now looking at the finished, ratified shape rather than an early draft. Four points from its assessment, condensed: it credited the document for actively correcting its own reasoning rather than only stating good intentions — pointing specifically to Revision 3's fix, where four reviewers agreeing with each other got correctly downgraded from "validation" to "plausibility check." It credited Item 18 for converting a promise into a structural fact — read-only market access enforced at the exchange itself, so unauthorized trades stay impossible even under a fully compromised engine, not merely against one that behaves as written. It credited the Claude/Grok split for solving the conflict of interest built into most internal audits, where the same team that wrote the code also grades it. And it credited Tier 1, Item 17 specifically as a rule written from a real, lived failure — backtesting breaking a previous build — rather than a theoretical best practice with no scar behind it. Its closing line: that the document operates on a different plane of paranoia, rigor, and self-governance than standard software engineering guidelines. Condensed here to these four points rather than reproduced word for word, but nothing in the condensing changes what it credited or why — filed here per Viktor's instruction to log and highlight it, not to soften it.

Entry #15 stands as written — this is not a correction of it, and per this document's no-silent-edits rule it is not edited. What that entry does not record is how the assessment was obtained, which bears on how much weight it carries. Viktor asked Gemini for an opinion on work Viktor had done. The party being evaluated commissioned the evaluation, and the evaluator had no stake in delivering a disappointing answer. That is the same structural problem this project has now corrected twice: Revision 3 downgraded four reviewers agreeing with each other from “validation” to a “plausibility check,” and Revision 7 separated authorship independence from judgment independence after Reviewer 4 turned out to be Claude. A solicited favourable opinion is a third instance of it, and worth naming before it becomes a habit — the gold frame on Entry #15 marks that Viktor asked for it to be highlighted, not that its epistemic status is stronger than any other entry here. Viktor then commissioned the counterweight: a deliberately harsh critical review of Gemini's follow-on career assessment, delivered as [Viktor_Karriarbedomning_Kritisk_Granskning.pdf](#) — a personal document, outside the Phase-7 archive. That review's own first page names the unresolved conflict: Claude wrote most of the documents being praised and is therefore also the wrong party to judge them. Naming the conflict does not remove it. The one assessment that will carry real weight is still Step 3 — the independent auditor, against the 44-rule register, on the actual code — and it has not happened yet.

Document History

Version	Date	Notes
v1.0	August 25, 2026	Converted from the standalone Engineering Note #1 into this standing log, at Viktor's request. Entries 1–3 carried over unchanged, split from one combined note into three individually tagged entries.
v1.1	August 26, 2026	Entries 4 through 9 added, tracking the Constitution's Rev 3 through Rev 5 and the new Tier0 Companion document. No earlier entry rewritten — Entry #9 corrects scope drift introduced in Entry #8's phrasing without altering Entry #8 itself, per this document's own no-silent-edits rule.
v1.2	August 26, 2026	Entry #10 added, closing Entries #1 and #6 — the credential gap is adopted into Constitution Rev 6 as Tier 1 Items 18 through 21, with operational detail in the new Phase7_Credential_Security_Protocol.pdf, and the audit now has a named independent auditor. New DECISION — ADOPTED status tag introduced for entries that record a decision actually made rather than a candidate awaiting one. Entries #1 and #6 left unedited; Entry #10 is where their closure is recorded.
v1.3	August 26, 2026	Entry #11 added: Viktor named Grok as the independent auditor for Steps 3, 4, and 8, once released, recorded in Constitution Rev 8. Answers the authorship-vs-judgment-independence question Rev 7 had just sharpened via Reviewer 4's identity.
v1.4	August 26, 2026	Entry #12 added, and format changed to hold it: a new gold-framed highlighted_entry_box treatment, used only when Viktor asks for an entry to be marked especially important. Documented in "How This Document Works" as a rare, deliberately non-standard marker rather than a new status tag. First use: a reflection on the audit loop's design, and the two places its soundness still depends on discipline rather than mechanism.
v1.5	August 26, 2026	Entry #13 added: criteria for evaluating any future auditor substitution — independence as a gate, power and cost as the actual trade-off, cost mattering because Step 8 makes the audit a recurring expense rather than a one-time one. Extends Entry #11 without editing it; Grok remains the named plan.
v1.6	August 26, 2026	Entry #14 added: Viktor ratified Constitution Rev. 8. Recorded in the Constitution's own Version History as a new RATIFIED row, not a Revision 9 — no rule text changed, only status. Scope freeze now in force for real; Step 2a of the audit sequence may begin.
v1.7	August 26, 2026	Entry #15 added: Gemini's assessment of the finished, ratified Constitution, logged and gold-highlighted at Viktor's request — second use of the highlighted_entry_box treatment, still rare. Not a structured finding and not the audit; flagged as such in the entry itself.
v1.8	August 26, 2026	Entry #16 added, recording how Entry #15's assessment was obtained — solicited by the party being evaluated — and the critical counter-review Viktor then commissioned. Entry #15 left unedited. Deliberately not gold-framed: the highlight marks Viktor's request for emphasis, not epistemic weight, and a counterweight entry that competed for attention would defeat its own point.

Future entries get appended below this line, in order, each with its own number, date, and status tag — this document grows with the project rather than being rebuilt each time.