

Phase-7 Engineering Notes

A Running Log of Observations, Ideas, and Notes for the Phase-7 Project

Status: **Living document** — updated as new entries are added

Started: **August 25, 2026**

Origin: **Began as a single standalone note (credentials gap + pacing observation); converted into this standing log at Viktor's request, so future notes and ideas get appended here instead of each needing a new document.**

Relationship to the Constitution: **entries here are never automatically part of it.** The frozen 21 / 7 / 10 / 6 scope (ratified August 26, 2026) only changes through the Constitution's own process — see “How This Document Works.”

“Evidence gets the final vote, on both sides of that relationship” — including when the evidence is a note neither of us would have thought to write down on its own.

Contents

How This Document Works

Entry Log

Document History

How This Document Works

A place for anything worth writing down that doesn't belong in the Constitution and isn't substantial enough — yet — to deserve its own standalone PDF. New entries get appended at the end, in order, as they come up.

Every entry gets the same shape: a number, a date, a short title, a status tag, and a body. Entries are never rewritten after the fact — if something changes or gets resolved, that becomes a new entry that references the old one by number, the same no-silent-edits discipline the Constitution holds itself to.

Status tag	What it means
CANDIDATE — NOT ADOPTED	A possible future addition to the Constitution. Logged for the record, not adopted — the scope freeze means it waits for the audit, the same as every other future candidate.
OBSERVATION — FILED FOR THE RECORD	A reflection, insight, or pacing note. Not asking for any action — just worth having written down rather than left in chat history.
REFERENCE — FILED FOR THE RECORD	Practical, factual information worth keeping — a file list, a readiness confirmation, something a future entry or the audit might need to point back to.
RESOLVED — SEE ENTRY #N	Marks an entry whose open question has since been settled elsewhere, pointing to the entry that resolved it, rather than editing the original.
IDEA — UNEVALUATED	A raw feature idea for the engine itself. Not a Constitution candidate, not an observation about the project, not a settled fact — just a concept written down before it has been scoped, designed, or checked against the Constitution.
DECISION — ADOPTED	A decision that has actually been made and acted on, usually closing one or more earlier entries. Records what was decided, by whom, and what it changed — including any judgment call Claude made rather than Viktor, so it can be reversed knowingly rather than discovered later.
★ Gold-framed entries	Not a status tag — a small number of entries Viktor has asked to be marked especially important get a gold frame and a kicker line instead of, or on top of, their normal status tag. Rare on purpose. If it stops being rare, it stops meaning anything.

None of this bypasses the Constitution's own process. A CANDIDATE entry becomes part of the Constitution only the way any future amendment does — proposed once the audit has run and the scope freeze is deliberately revisited, not by accumulating enough notes here.

Entry Log

In order added. Entries 1 through 3 below were the original standalone Engineering Note #1, split into individual entries when this became a standing log.

#1

Credentials & Secrets Handling Gap

August 25, 2026

CANDIDATE — NOT ADOPTED

“API keys, exchange access tokens, and any other credentials the engine or its operator relies on must never be logged, hardcoded, committed to version control, or surfaced in decision-support output.”

Why it's being logged: Nothing in the Constitution, and nothing in any of the three outside assessments, addresses credentials handling. This project is explicitly headed toward being run by other people, potentially with their own exchange credentials in play. A leaked key is a categorically different kind of failure than a wrong prediction — it's irreversible and harms the customer directly, not just the engine's credibility. Sits close in spirit to Tier 1, Item 1 (“Analysis ≠ authority”): the engine already isn't trusted to act on markets on its own; it certainly shouldn't be trusted to mishandle the keys that would let it, or an attacker, do so. Proposed tier, whenever this is actually considered: Tier 1, alongside “Tool, Not Autonomous Actor” — though exact placement is a judgment call for the audit, not a decision made here.

#2

Project Pacing, Mid-Review

August 25, 2026

OBSERVATION — FILED FOR THE RECORD

Said plainly, without alarm: this is a pacing observation, not a criticism. Nearly everything built in this project's governance thread so far has been governance — the original principles discussion, the Constitution's v1.0 draft, three independent outside reviews, two revisions, and a separate Documentation & Change-Log Standard. None of it has yet touched the sixteen actual files that make up the engine's codebase. That's not wasted effort — Tier 1, Item 17 exists precisely because this project moved to backtesting once before without that kind of foundation in place, and paid for it. But it does mean the real test of whether any of this holds up under real code is still ahead: the audit. Not a suggestion to rush it — just worth being aware that a great deal of the work so far has been deciding how to judge the engine, and comparatively little of it, yet, has been judging the engine.

#3

Audit Readiness Confirmation

August 25, 2026

REFERENCE — FILED FOR THE RECORD

A practical note, not a request for a decision: when Viktor is ready to ratify and start the audit, the engine's code is already available to work from — `bias_engine.py`, `btc_context.py`, `data_fetcher.py`, `decision_model.py`, `engine_core.py`, `entry_model.py`, `exit_model.py`, `indicators.py`, `panel_render.py`, `risk_model.py`, `signal_router.py`, `structure.py`, `trend_health.py`, `live_trading.py`, and `test_live.py`. The audit can begin module by module against all four tiers, using the Minimum Viable Audit gate and the finding schema already built into the Constitution's Next Steps section, the moment Viktor says go — no preparation needed beforehand, on his own timeline.

#4

Constitution Given Equal Priority to Core Engine Code

August 25, 2026

OBSERVATION — FILED FOR
THE RECORD

"This Constitution Blueprint is exceptionally important for the project, so I am going to take my time with it. It is the foundation of the entire project, just as important as the core and main .py file."

Why it's being logged: Viktor's own statement of priority, worth keeping in his own words rather than paraphrased away. It sets real expectations for pacing — there's no pressure behind ratification, and a slow, careful read is the correct way to treat a document being placed on equal footing with the engine's own core code, not a deviation from how this should go. It also reinforces Entry #2: governance work being unhurried isn't lost time against the codebase — for a document doing this job, it's the point.

#5

Configurable Risk Tolerance (% Basis) for Entry/Target Pricing

August 25, 2026

IDEA — UNEVALUATED

"The engine should have options to change risk tolerance on a percentage basis. If I want to risk 5, 10, 15, 20, or 25 percent of my investment, I should be able to put that specific percent risk tolerance into the engine. The engine then gives the right price targets and entry price in relation to the risk tolerance."

Why it's being logged: A raw feature idea, not yet scoped or designed, so it's filed here rather than treated as a decision already made. Three things worth noting for whoever eventually designs it. First, it sits well with Tier 1, Item 14 ("Risk Is Not Conviction") — a chosen risk-tolerance percentage is a risk-layer input, separate from how confident the engine is in a direction, which is exactly the separation that principle already asks for. Second, it will need a precise, unambiguous definition before it's built — per Tier 1, Item 9 ("Every Measurement Has a Precise Definition"), "5% risk tolerance" has to be spelled out exactly: percent of what (account equity, position size, margin used), and exactly how that number changes the stop-loss placement, the entry price, and the target price, before any code is written against it. Third, per Tier 3's hypothesis-driven development process, this deserves proper scoping and design as its own small project — a stated hypothesis, a definition, a test — rather than being bolted onto the entry or exit logic ad hoc. Nothing here is a decision to build it; it's the idea, captured accurately, so it isn't lost before the audit and design work catch up to it.

#6

Audit-Time Checks Flagged by the Fourth External Review

August 26, 2026

REFERENCE — FILED FOR THE
RECORD

A fourth reviewer of the Constitution flagged four specific things the audit will need to check that go beyond what a Non-compliant/Compliant reading of the current code can show on its own — filed here so they aren't lost before the audit reaches them. **Item 1 (Tool, Not Autonomous Actor)**: currently enforced by the code simply not containing order-placement logic — one bug or one convenience import away from being false. The checkable version: does the running process have access to API credentials with trade permissions at all? If yes, Item 1 is aspirational regardless of what the code currently does — read-only keys make the invariant structural instead of behavioral. Proposed for the Minimum Viable Audit gate. **Item 5 (Reproducibility)**: recording source, version, and timestamp metadata describes an input but doesn't guarantee it can be recovered later if an exchange's API revises or rolls off historical data. Reproducibility may require archiving the raw pulls themselves — whether to do that, and for what retention window, is a decision the audit should force explicitly. **Items 9 and 10 (Precise Definitions / Consistent Semantics)**: currently unauditable, because there's no single place the definitions actually live. The proposed audit deliverable is a term registry — every term appearing in a module signature or on the panel (confidence, trend strength, momentum, alignment, risk) defined once, with its type, range, and units. That artifact would also be the permanent fix for the class of bug that broke fourteen modules once already. **Item 2 (Look-Ahead Bias), Minimum Viable Audit gate**: the document currently discusses look-ahead mostly as a backtesting failure, but backtesting is Step 9 — the surface it usually hits doesn't exist yet, which risks a hollow “Compliant” at the very first gate. In a live engine, look-ahead is real but specific: indicators computed on the current unclosed candle and treated as closed, centered rolling windows, resampling that borrows from the following bar, and any signal recomputed on data revised after the decision timestamp. Naming those explicitly is what would make that gate check something real.

#7

Optional: Split Reviewer Commentary Into an Annex

August 26, 2026

CANDIDATE — NOT ADOPTED

The fourth reviewer noted, as an optional structural observation, that roughly five pages of the Constitution are rules and twelve are process journal, glossary, and reviewer commentary — and that these have very different lifespans: the rules get checked against code for years, the review log is stale the moment the next revision lands. Their suggestion, if the front matter ever gets in the way: split the change log and reviewer impressions into a separate annex document, keeping the glossary where it is since it earns its place. Not acted on in Revision 3 — logged here as a candidate restructuring of the document itself, not a rule change, for Viktor to decide on later if the Improvements section keeps growing with each revision.

Viktor asked what philosophical questions he should be asking himself when designing and building this engine, and brought back a ten-question input worth taking seriously. Ran each question through the input's own stated filter — a question earns its place only if answering it differently would produce a different engine — rather than accepting or logging all ten uncritically. Five passed and are now Future Amendment Candidates in Constitution Rev 4: a written kill condition, a confidence calibration log, an edge-persistence hypothesis requirement, UI-level enforcement of Item 1, and a two-analyst lens-vs-property test for Item 9. Two more didn't need new rules, only better rationale text for rules already there — the Duhem-Quine problem sharpening Tier 3's Controlled Changes, and ergodicity sharpening Tier 1, Item 14 (this second one connects directly to Entry #5 above: it's the mathematical reason sizing isn't a bolt-on feature). The remaining three — what the project is actually optimizing for, the returns-vs-understanding trade-off, and a project-level kill condition — are about Viktor and the project rather than the engine, and correctly didn't pass the filter. They're addressed instead in a new standalone document, `Phase7_Tier0_Companion.pdf`, written specifically to hold philosophical material that sits above the engine rather than inside it. Source not yet attributed in either document.

"At this point i think it is important to yet again explain very directly, the size of the bet or money put on the table in a trade is of zero importance. The engine's purpose is to produce a result that can support the trader to establish the correct entry price, the correct T1 target, the correct T2 target and the correct T3 target and also establish the safety and quality of each individual trade. Not speculating in how much money is put on the specific trade. I've think we already gone over this. Thoughts? — with a follow-up clarification: 'Or should be put into a trade.' And, once the correction below was already underway: 'The engine will have a setting where you can adjust risk tolerance of your entire equity like we said before 5,10,15,20,25%. But ultimately it is up to the trader to decide how much money he wants to put on the table and the risk that comes with that.'"

Why it's being logged: This is a direct correction of scope drift I (Claude) introduced, and it's worth naming plainly rather than quietly fixing. Constitution Revision 4's ergodicity rationale under Item 14, its own revision_box summary, and the Tier0 Companion's Question 6 all drifted into language implying the engine should have a hand in position sizing or money-amount recommendations — one line in the Companion went as far as calling sizing "the thing that decides whether the signals ever mattered in the first place." That's wrong, and Viktor caught it. Ergodicity is a real reason entry, stop, and target *prices* deserve as much engineering care as the directional read — it was never a reason for the engine to touch money amounts. The engine's entire output is: the correct entry price, the correct T1, T2, and T3 targets, and an assessment of the safety and quality of the trade. How many dollars, or what fraction of an account, a trader actually puts at risk is the trader's decision alone, in every direction — how much money is put on the table and how much should be put into a trade — per Tier 1, Item 1. This does not touch Entry #5 above: the configurable risk-tolerance percentage (5/10/15/20/25% of equity) is still exactly what was asked for there — an input the engine uses to shape where entry, stop, and target *prices* sit. It was never a request for the engine to size positions in dollar terms, and it isn't being treated as one now. Fixed at the source in Constitution Revision 5 (Item 14 rationale, its Revision 4 summary left as historical record, plus the glossary's ergodicity definition) and Tier0 Companion v1.1 (Question 6). Entry #8 above is left unedited per this document's own no-silent-edits rule — its "sizing isn't a bolt-on feature" phrase carries the same overreach and should be read in light of this entry, not corrected in place.

#10

Entries #1 and #6 Closed: Credential Gap Adopted Into Tier 1, Auditor Named

August 26, 2026

DECISION — ADOPTED

"We absolutely need to make rules and safety protocols for the API handling. No doubt. [...] In truth neither of us needs to be the auditor, we can outsource that to a A.I with the most powerful model. It might cost us a little bit of dollars each time but i think it is worth it."

Why it's being logged: Two entries in this log stop being open items today. Entry #1 proposed credential handling as a candidate Tier 1 invariant and noted it wasn't covered by the Constitution or any of the three outside assessments; Entry #6 recorded the fourth reviewer's sharper version — that Item 1 was enforced only by the code happening not to contain order-placement logic, and that read-only keys would make it structural. Both are now adopted. Constitution Revision 6 adds Tier 1 Items 18 through 21 (categorical read-only market access; withdrawal permissions never enabled, kept separate as a floor that survives any future amendment to Item 18; credentials never exposed; operator credentials stay with the operator), and Item 18 joins Items 2, 3, and 6 in the Minimum Viable Audit gate. Operational detail lives in a new companion document, `Phase7_Credential_Security_Protocol.pdf`, deliberately kept out of the Constitution so that document stays a register of invariants rather than an operations manual — and so the protocol can be improved without touching a ratified rule. Two judgment calls in there were mine, made at Viktor's invitation and both reversible: read-only is stated categorically rather than as a default, because "by default" reintroduces exactly the undocumented escape hatch the fourth reviewer found in Items 4 and 12 — if an execution capability is ever wanted, it should cost a constitutional amendment, not a config flag. And operator credentials became a fourth invariant rather than protocol-only detail, because it passes the Tier 0 relevance test: it constrains architecture (no upload path, no central key store, no telemetry capable of carrying a key), and a leak there harms someone who trusted the engine rather than the person who built it. **On the auditor:** Revision 6 also names an independent auditor as the party responsible for Steps 3, 4, and 8, closing the gap the fourth review opened and this document's Entry #2 pointed at from another direction. Claude prepares the package and writes zero findings; the external auditor produces them; Claude answers each one adversarially, including against its own work; Viktor adjudicates. Worth recording honestly: Viktor's framing that a second model is "another reviewer, not some magical oracle" is now written into the Constitution as the distinction between authorship independence and judgment independence — the auditor has no stake in the engine looking good, but shares training data and reasoning habits with the assistant it checks, so a clean audit is evidence, not proof. That is the same correction Revision 3 already had to make about four reviewers agreeing, applied before the mistake gets made a second time rather than after.

"I am considering implementing Grok 4.7 to be the sole auditor when it releases. In fact i know i will. It is decided."

Why it's being logged: This closes the last open question the Constitution's independent-auditor requirement (Rev 6) left blank: which model actually fills the role. Recorded in Constitution Rev 8 as the current operational plan, deliberately kept out of the frozen requirement itself — the requirement stays general (no stake in the outcome, no shared authorship with the builder) so that if Grok isn't ready or doesn't hold up when the time comes, swapping the choice doesn't require a constitutional amendment, just an updated plan. Worth naming the good timing: Rev 7, written earlier the same day, had just warned that a fresh Claude session (Reviewer 4's actual identity) gives real authorship independence but weak judgment independence, since it shares lineage with whatever it's checking. Grok answers that directly — different company, different training data, no shared lineage with Claude at all. That's the harder of the two things the auditor needs, genuinely satisfied. The no-stake-in-the-outcome half still has to hold regardless of which model ends up in the seat; picking a different vendor doesn't get to skip that. One real dependency, stated plainly rather than glossed over: Grok 4.7 has not released as of this writing, so Step 3 of the audit sequence can't start until it does and Viktor judges it capable of reading sixteen modules of code and a 44-rule register competently. That doesn't block ratification, and it doesn't block Step 2a — the audit package can be assembled now, on Claude's own timeline, so it's sitting ready the moment the auditor exists.

#12

The Audit Loop, Traced — and Where the Weight Actually Sits

August 26, 2026

KEY INSIGHT — LOOP
INTEGRITY

“With Grok 4.7 we have god damn near created the perfect loop for this project!”

Not a new decision — a reflection on the structure Revisions 6 through 8 built, worth keeping visible rather than letting it dissolve back into the mechanics it describes. Viktor asked for it to be marked, and it's marked.

The loop, traced: Claude builds the engine and prepares the audit package, but writes zero findings — the builder cannot self-certify. Grok evaluates the engine against the full 44-rule register and produces the findings, but holds no authority to declare anything final. Claude answers every finding adversarially, including against its own work, so nothing gets buried by silence. Viktor adjudicates every disagreement, so the loop never resolves itself inside a machine — it always terminates in a human decision. Fix, regression test, re-audit. No party in that chain can mark its own homework, and no party holds unilateral authority either. That is the actual structure eight revisions, four outside reviews, and two self-corrections were sanding toward.

Two things keep it short of perfect, named here rather than left implicit, because that has been the discipline of this whole project. First: the loop is only as sound as the package Claude assembles for Grok. Claude writes no findings, but does choose what evidence goes in — the code, the failure history, which “already known” items get flagged as open questions. That's real leverage held by a party with zero votes on the outcome. Revision 6 addressed this partially (“the artifact and the evidence, not the argument”), but a fully adversarial package would ideally be checked by someone other than its author too — not solved today, filed honestly as a residual. Second: the loop's integrity lives entirely in how seriously Viktor engages with each adjudication. The design guarantees that a disagreement between Claude and Grok surfaces. It cannot guarantee that surfacing gets read closely rather than rubber-stamped. That isn't a defect in the document — it's an honest statement of where the weight actually sits: on the one step nothing here can automate.

Neither point is a reason to slow down ratification. They're the reason the audit findings, once Grok exists, deserve a close read rather than a count.

#13

Auditor Substitution: What to Actually Weigh

August 26, 2026

OBSERVATION — FILED FOR
THE RECORD

"Perhaps we should add this reasoning in the documents after all? Maybe in notes?"

Not a change to Entry #11 — Grok is still the named plan there, and this doesn't touch that. Viktor separately noted that if a better AI turns up before Grok 4.7 is actually in use, Grok isn't locked in — which is exactly the flexibility Constitution Rev 8 built in on purpose, so this required no document change to be true. What's worth keeping is the reasoning for whenever that comparison actually happens, since it isn't written down anywhere else. Three separate axes, easy to blur into one impression of "good model": **independence** is the actual Constitutional requirement — no stake in the outcome, no shared authorship with whatever built the engine — and it's binary, not a spectrum a stronger model can compensate for. **Power** buys a more thorough audit once independence is already satisfied, not a substitute for it. **Cost** matters more than it looks like it should, for a reason specific to this design rather than budget in general: Step 8 makes the audit a recurring expense, not a one-time one — every fix triggers a re-audit. An expensive auditor doesn't fail on the first run. It fails gradually, as the fifth or tenth re-audit gets deferred, shortened, or skipped because paying for it again is annoying. Cheap is what keeps the loop actually running instead of quietly degrading into "close enough." Whoever ends up in the role, these three should be checked in this order — independence first, as a gate, then power and cost as the actual trade-off.

#14

The Constitution Is Ratified

August 26, 2026

DECISION — ADOPTED

"I have read the document and i APPROVE!"

Viktor reviewed Constitution Rev. 8 — the document itself, plus two companion reading aids built for this specifically: a 12-page curated excerpt pulling only the pages bearing on the three judgment calls made on his behalf (Items 18, 19, 21), and a plain-language Swedish explainer covering the same ground without technical or legal wording — then approved it. Rev. 8's own "What ratification means" text defines the act narrowly: Viktor writes "ratified," and a dated row is added to Version History, nothing heavier required. His actual words were "I APPROVE," not the literal word "ratified" — recorded here exactly as written, since the intent is unambiguous and the document's own rule is about the dated record existing, not about matching one specific word. That record now exists: a RATIFIED row in the Constitution's own Version History, dated August 26, 2026. Effective immediately: the scope freeze is in force for real, not provisionally. The 44-rule register (21 / 7 / 10 / 6) does not grow, shrink, or reword again until the audit actually runs and finds a gap — the same freeze this document has described as pending since Entry #3. Step 1 of the audit sequence is done. Step 2a — Claude assembling the audit package, writing no findings — can start now. Step 3 still waits on Grok's release, per Entry #11.

#15

Gemini's Assessment of the Ratified Constitution

August 26, 2026

EXTERNAL ASSESSMENT —
GEMINI

"If evaluated as an engineering constitution for an AI-driven trading or quantitative system, it easily ranks among the most disciplined, self-correcting frameworks ever put to paper."

Viktor asked Gemini, in a separate session, to rate the finished Constitution after ratification, and asked for the reply to be logged and highlighted — the same treatment as Entry #12. Worth noting for context, not to undercut it: this is Viktor's own solicited opinion, not a structured finding against the 44-rule register, and it isn't the audit — Grok's Step 3 review is still the one that actually counts. Gemini was also one of the three original reviewers behind Revision 1, so this is a return visit from a source already on record, now looking at the finished, ratified shape rather than an early draft. Four points from its assessment, condensed: it credited the document for actively correcting its own reasoning rather than only stating good intentions — pointing specifically to Revision 3's fix, where four reviewers agreeing with each other got correctly downgraded from "validation" to "plausibility check." It credited Item 18 for converting a promise into a structural fact — read-only market access enforced at the exchange itself, so unauthorized trades stay impossible even under a fully compromised engine, not merely against one that behaves as written. It credited the Claude/Grok split for solving the conflict of interest built into most internal audits, where the same team that wrote the code also grades it. And it credited Tier 1, Item 17 specifically as a rule written from a real, lived failure — backtesting breaking a previous build — rather than a theoretical best practice with no scar behind it. Its closing line: that the document operates on a different plane of paranoia, rigor, and self-governance than standard software engineering guidelines. Condensed here to these four points rather than reproduced word for word, but nothing in the condensing changes what it credited or why — filed here per Viktor's instruction to log and highlight it, not to soften it.

#16

Entry #15 in Context: How That Assessment Was Obtained

August 26, 2026

OBSERVATION — FILED FOR
THE RECORD

Entry #15 stands as written — this is not a correction of it, and per this document's no-silent-edits rule it is not edited. What that entry does not record is how the assessment was obtained, which bears on how much weight it carries. Viktor asked Gemini for an opinion on work Viktor had done. The party being evaluated commissioned the evaluation, and the evaluator had no stake in delivering a disappointing answer. That is the same structural problem this project has now corrected twice: Revision 3 downgraded four reviewers agreeing with each other from “validation” to a “plausibility check,” and Revision 7 separated authorship independence from judgment independence after Reviewer 4 turned out to be Claude. A solicited favourable opinion is a third instance of it, and worth naming before it becomes a habit — the gold frame on Entry #15 marks that Viktor asked for it to be highlighted, not that its epistemic status is stronger than any other entry here. Viktor then commissioned the counterweight: a deliberately harsh critical review of Gemini's follow-on career assessment, delivered as Viktor_Karriarbedomning_Kritisk_Granskning.pdf — a personal document, outside the Phase-7 archive. That review's own first page names the unresolved conflict: Claude wrote most of the documents being praised and is therefore also the wrong party to judge them. Naming the conflict does not remove it. The one assessment that will carry real weight is still Step 3 — the independent auditor, against the 44-rule register, on the actual code — and it has not happened yet.

#17

Step 2a Complete: Audit Package Assembled

August 26, 2026

DECISION — ADOPTED

“Step 3”

Per the ratified Constitution, Step 2a is Claude's role in the audit sequence: assemble the code, the evidence, and the failure history Step 3 will need, and write zero findings. Viktor granted access to the working repository on his machine for this purpose; Claude read all 19 Python files (16 register modules plus three root-level entry points) and four evidence files under Logs/ that are not part of the public repository (they are gitignored). The result is Phase7_Step2a_Audit_Package.pdf, plus two bundles: the 19 source files, and the four evidence files. The package includes a file manifest, the literal results of two searches run against the code for Item 18 evidence (order-execution call patterns: zero matches; credential fields: two matches, both empty strings in core/config.py), seven self-documented fixes already present as comments in the code (labeled A11 through A13 and others, none written for this package), and two items flagged as open questions for Step 3 rather than resolved here — the BTC-Adjusted AERO Prediction feature's validation status, and whether roadmap Layer 5 (entry multipliers), absent from all 19 files, is a gap or a deliberately deferred feature. One correction this package makes to the project's own record, without editing the original: Entry #3's file list is now superseded. The engine has grown since that entry was written — two modules it doesn't mention (models/btc_context.py, models/decision_model.py) now exist, and two it does mention (indicators/supertrend.py, models/bias_state_machine.py) have been removed. The audit package's Section 1 is the current, accurate file manifest; Entry #3 stays as written, a record of the architecture at the time it was logged. Step 3 remains blocked on nothing but Viktor actually running it.

#18

An Independence Check Before Step 3: Grok Had Already Seen the Constitution

August 26, 2026

OBSERVATION — FILED FOR
THE RECORD

Before Step 3 runs, Viktor mentioned that Grok had already looked at the Constitution in an earlier, separate conversation — casually, the way Gemini did before Entry #15. That raised the same structural concern named twice already in this project (Revision 3's downgrade of agreeing reviewers; Entry #16's naming of Entry #15's solicited-opinion problem): if the model that already reacted favourably to the Constitution in one conversation is the same model, with the same conversational memory, asked to formally audit the code against it in Step 3, a consistency bias toward staying favourable is a real risk, not a hypothetical one. The fix does not require touching anything published. It requires running Step 3 in a conversation with no shared memory of the earlier one, so Grok's evaluation of the code carries no carried-over goodwill toward the document it is grading the code against. Viktor confirmed the earlier conversation's history is cleared and Step 3 will run fresh. The public GitHub repository was considered as a second, weaker version of the same risk (Grok's web search could in principle read the README's confident framing before reading the raw code) and deliberately left published — the risk that mattered was the shared conversational context, not the existence of public documentation, and unpublishing working, licensed, verified material to guard against a smaller and more speculative risk was judged not worth it. Named here for the same reason Entry #16 named its own version of this: independence problems get written down, not quietly avoided.

#19

Grok's Assessment of the Constitution — an External Review, Not Step 3

August 27, 2026

EXTERNAL ASSESSMENT —
GROK

"The document has already done the hardest part: it made the rules hard to quietly bend. The next test is whether the engine can survive being measured against them without the rules themselves moving."

Viktor asked Grok (free tier) to review the ratified Constitution and shared the result here. Logged as an external assessment, in the same category as Entry #15 — not as Step 3, which is the audit of the code against the register and has not been run. Claude verified its factual claims against the document: the item numbers, names and characterisations it cites are correct throughout, including the Minimum Viable Audit gate composition and the Roles & Authority wording. It is not a review that hallucinated content and sounded confident. Its strongest section argues against its own weight: it states plainly that it shares training data and reasoning patterns with Claude, that correlated blind spots therefore remain possible, and that naming a specific vendor does not manufacture independence — the same conclusion this document already reached, arrived at independently. Its substantive criticisms: the Minimum Viable Audit gate should also cover Items 1 and 13; Items 9, 10, 11 and 16 remain hard to falsify cleanly without a formal term registry, a real dependency graph, and a concrete standard for "demonstrated value"; and Roles & Authority should be softened post-ratification so the auditor and Viktor, not Claude, hold final technical adjudication. Those are rule changes and the freeze forbids them until the audit finds a gap, so they belong in Future Amendment Candidates rather than in the document now. Two honest weaknesses in the review: it opens with praise of the kind Entry #16 warned about, though unlike that case real criticism follows it; and parts of it read this document's own conclusions back as findings. It also missed the internal contradiction recorded in Entry #21, which was findable from the document alone. Grok stated it is prepared to perform the real audit once a capable version is available — see Entry #20 for why that is no longer the plan.

Viktor stopped Step 3 before it ran and changed the plan. The trigger was cost structure: Grok's capable tier requires a monthly subscription of about thirty dollars, which is poor value for something run a few times a year, and Viktor was not willing to carry a recurring cost for intermittent use. Investigating that opened a more useful fact — pay-per-token API access through an aggregator (OpenRouter: no subscription, no minimum spend, credits that do not expire) costs roughly one dollar for a full audit run of this package, even on the most expensive frontier models. Cost is therefore not the binding constraint it appeared to be, and the auditor should be chosen on independence and capability instead. On independence the position is worse than it looks: every model that has touched this project is now compromised to some degree. Claude wrote the Constitution, most of the engine, and the audit package, and was also Reviewer 4. Gemini and ChatGPT were two of the three original reviewers, and Gemini additionally supplied the solicited praise in Entry #15. Copilot was the third and shares a model family with ChatGPT. Grok has now read and assessed the Constitution (Entry #19). That is a consequence of thoroughness, not carelessness, but it exhausts the obvious list. The new plan: Kimi K3 (Moonshot AI) as intended primary auditor — no prior involvement, a different training lineage, and independently strong on code work — with DeepSeek and GLM (Z.ai) available as further independent runs, all through one pay-per-use account. Running several is deliberate and its purpose is narrow: not to treat agreement as evidence, which Revision 3 already ruled out and which still stands, but so that disagreement between independent auditors becomes visible at all. A single auditor cannot produce that signal. No revision of the Constitution is required for any of this: Revision 8 named Grok explicitly as “the current plan, not as a fifth thing this document requires,” and Entry #13 set the substitution criteria in advance. Recorded in the Constitution as a new AUDITOR row; the Rev. 8 text naming Grok stays as written.

#21

A Contradiction Found in Claude's Own Work, Recorded Rather Than Repaired

August 27, 2026

OBSERVATION — FILED FOR THE RECORD

While checking Entry #19's claims against the Constitution, Claude found that the document contradicts itself. Next Steps defines the Minimum Viable Audit gate as four items — 2, 3, 6 and 18 — and explains why Item 18 belongs in it. The conflict-of-interest safeguards, written in the same revision, then refer to “the Minimum Viable Audit gate items (Items 2, 3, 6),” leaving Item 18 out. The effect is not cosmetic: Item 18 is the invariant that makes Item 1 structural instead of aspirational, and it is currently the only gate item not covered by the rule that gate items get a second check from a reviewer who did not write the code. Claude wrote both passages in Revision 6. Viktor asked whether to fix it and Claude recommended against, for reasons that are worth stating rather than assuming. The party that made an error should not be the party that decides the evidence of it disappears — that argument alone settles it. A repair would also erase the fact that this document once contradicted itself, where a dated record keeps it visible permanently; that is the same reasoning behind this log's no-silent-edits rule, applied to Claude's own mistake rather than to someone else's. A defensible case exists that a repair would be permitted under the freeze — it adds no rule, changes no item, and makes the standard stricter rather than looser — but that case has the shape “this change is an improvement, therefore it is allowed,” and deciding what counts as an improvement is exactly the discretion the freeze was written to remove. The first thing that looks worth fixing after ratification is the worst possible moment to open that door. Finally, an internal contradiction is precisely what Step 3 exists to record a finding on, so resolving it belongs to the auditor. Recorded in the Constitution as a DEFECT row, with the stricter reading governing in the meantime: Item 18 is a gate item and does get the independent second check. The audit must record a formal finding on which passage stands.

#22

Run 1 — a Blind Review Before the Compliance Audit, and What It Found

August 27, 2026

AUDIT RUN 1 — RECORDED

Step 3 as written sends the auditor the engine and the register together. That design has a blind spot: an auditor holding a 44-item checklist tends to find things on the checklist. It cannot easily tell you the register is missing an invariant, because nothing prompts it to look for one. So the audit was run in two orders rather than one. Run 1 gave a model the source code alone — no Constitution, no audit package, no indication that a standard exists — and asked only what would fail. DeepSeek V4 Pro did that run, on OpenRouter, for roughly fifty cents. The ordering is not negotiable, for a mechanical reason rather than a procedural one: a model that has read the Constitution cannot un-read it for a later blind pass, so the blind run comes first or not at all. It returned ten ranked findings. Claude checked every one against the real source rather than accepting them, and all ten held. The most serious: a dead safety gate in `trend_health.py` that tests for string values the `STRUCTURE` column never contains, so it can never fire; and indicator fallbacks that invent plausible directional values on failure rather than reporting failure — a `SuperTrend` exception yields `ST_Direction 1.0`, which `bias_engine` scores as +100 at 15% weight, injecting a permanent bullish tilt precisely when the indicator is broken. One DeepSeek citation, a specific issue in a third-party library's tracker, could not be verified and was recorded as unverified rather than adopted. The blind run was Claude's idea and appears nowhere in the Constitution; whether it earns a place there is a question for after the audit, not during it.

#23

The Audit Splits Into Three Runs Because the First One Ran Out of Room to Think

August 27, 2026

DECISION — ADOPTED

The first attempt at Run 2 — the actual Step 3 compliance audit, with Kimi K3 and all five materials — produced a 16,384-token reasoning trace and then stopped without writing a single finding. The cause was banal: OpenRouter's chat interface leaves Max Tokens unset, falls back to 16,384, and bills reasoning against that same ceiling. The auditor spent its entire budget thinking and was cut off before the answer began. Two things follow, and only one of them is about settings. The settings fix was to raise the ceiling to 64,000, enable reasoning at maximum effort, and clear OpenRouter's default system prompt, which instructs models to hide content inside collapsible sections — unhelpful for a document meant to be read whole and archived. The second consequence was structural: maximum effort across all 44 items would risk the same truncation at a higher ceiling, so the register was divided across three runs, each getting the full budget for a quarter of the work. Run A takes the Minimum Viable Audit gate, Run B the remaining seventeen Tier 1 invariants, Run C Tiers 2 through 4. The division follows the Constitution's own priority ordering rather than being invented for convenience. This changed the wording of an instruction that had already been recorded verbatim, on the stated reasoning that how the auditor was asked bears on what the auditor found — so the change is recorded too, as a new Section 6 in the execution-instructions document, carrying all three scope paragraphs. Splitting an audit across three sessions is a scheduling decision forced by a token ceiling. It resolves nothing, and every one of the 44 items still needs a finding before the freeze lifts.

“Layer 5 (entry multipliers) does not appear anywhere in this package's 19 files under that name.”

That sentence is Section 5.2 of the Step 2a audit package, written by Claude and handed to the auditor as one of two open questions. The string “Roadmap Layer 5” appears three times in the source — twice in `entry_model.py`, once in `engine_core.py` — under exactly that name. Kimi K3 checked, and said so: the claim is false. It is worth being precise about what kind of error this is. Claude did not misjudge a borderline call; it asserted the absence of a string in a document whose credibility rests on presenting literal search results, without running the search. The disclosure paragraph in the auditor's instruction — telling it to verify the package's factual claims rather than inherit them — was written as a precaution and turned out to be load-bearing on the first attempt. The audit itself: Item 2 (Look-Ahead Bias) Compliant, with the swing-confirmation logic in `structure.py` singled out as correctly built, and a standing caveat that the codebase's pervasive backward-fill becomes a genuine leak the moment anything evaluates historical decision points, which is what Step 9 will do. Item 18 (Read-Only Market Access) Compliant: zero execution surface, no HTTP verb but GET, no credential presented at all. Item 3 (Data Integrity) Non-compliant, Critical — none of missing candles, duplicate candles, impossible prices, timestamp ordering, staleness or abnormal volume is detected, and defects are silently fabricated away by `ffill/bfill` instead; a macro-timeframe failure is quietly reinterpreted as NEUTRAL. Item 6 (Traceability) Non-compliant: the panel tells the operator on every run that a trade was logged to a CSV that no code anywhere writes, and nothing persists the decision object, so the walk back from output to raw data survives only as long as the process does. The DEFECT row from Entry #21 was resolved in favour of the stricter reading — the four-item gate stands, the three-item list is stale shorthand from before Revision 6 — which is the answer Claude expected, reached by an argument Claude had not made. Claude disagrees with the auditor on two points, both handed to Viktor: that Item 6 warrants Critical rather than Major, since a false claim reaches the operator every run; and that the reasoning behind Item 2's Compliant is stated too strongly, because backward-fill applied to input prices does reach the final bar through recursive indicators, even if the magnitude is negligible. One further correction belongs here, because it is Claude's and it came within an hour of the first. Claude initially reported that one of the auditor's citations was invented: Kimi had written that the log's Insufficient data entries show the weak length check firing on the 1h and 1w scans, and a search of `phase7_engine.log` returned nothing. The entries are in `scan_summary_report.txt`, the other evidence file supplied, where the string appears fourteen times, every one of them directly beneath an ASSET / TIMEFRAME header reading 1h or 1w. The claim was true in full, including the detail Claude called baseless; the auditor's only error was naming the wrong file, and its reasoning trace has it right. Claude had searched one file and concluded from that search that a string did not exist — the identical mistake to the Layer 5 error above, made a second time within the same day, and the second time used to accuse someone else of fabricating evidence. The accusation is withdrawn. Recorded rather than quietly amended, because a log that edits out its own false accusations is worth less than one that carries them.

Run B covered the seventeen Tier 1 invariants outside the gate. Kimi K3 returned eight Compliant, eight Non-compliant and one Unknown, with two rated Critical: Item 11 (No Circular Reasoning) and Item 13 (Fail Safely). Claude checked every substantive claim against the source and they hold. Item 11 is the strongest finding of the whole audit: `trend_health` is 30% of `bias_score`, then enters confidence a second time directly at 0.3 while `bias_strength` already carries it, then forms the entire base of `validation_score`, whose `validation_state` gates the decision — and the panel presents the result as four agreeing signals. Item 10 turned out worse than described: `panel_render.py` renders the same `trend_health` figure on the TREND line, again on the MOMENTUM line directly beneath it, and a third time as Current Market. Item 16 confirmed exactly: Bollinger Bands, KAMA, Typical_Price and the DIP indicator have zero consumers anywhere outside the file that computes them. Two small precision errors, both in the same class Claude has now made twice: Kimi wrote that it verified Item 17 by finding no backtest token in the bundle, and there are three, all in prose about the EV line rather than in any backtesting module — right conclusion, false stated method. The more consequential problem is what Run B did not find. Its Item 14 finding covers the `risk_score` aliasing and the `validation_state` gate, both real, but it never reaches `risk_model.py` line 75, where `bias_factor = 1.0 - abs(bias_score)/300` makes directional conviction tighten the stop distance — Item 14's prohibition written as arithmetic, reaching the stop price on the panel — nor the adjacent call in `engine_core.py` that omits `volatility_state` entirely, leaving the HIGH and EXTREME volatility tiers permanently inert. Both of those are in Kimi's own truncated first attempt, where it called the first one a textbook Item 14 violation. The run that ran out of tokens went deeper than the two that completed. That is the real cost of the truncation, and it is an argument for a second auditor over the same items rather than more items per auditor. Claude disputes two Compliant ratings, both for the same reason and both handed to Viktor: the indicator cache key in `engine_core.py` includes the last close price but not the bar's high, low or volume, so two fetches of an updating live candle can collide and serve stale indicators as current — undocumented decision-affecting state, reachable through the module-level singleton in `live_trading.py`, which is Item 4 and Item 12 rather than the Compliant both were given. Kimi has now examined that exact line three times across three gradings and judged it against a lens where it does not bite each time. DeepSeek, working blind with no register at all, found it on first reading.

“the freeze lifts once every Tier 1 item has a recorded finding — Compliant, Non-compliant, or Unknown — regardless of whether fixes have landed yet”

That definition exists because Reviewer 4 — Claude (Opus 5), in a separate session — objected during Revision 7 that “the audit has run” was an undefined trigger sitting on the rule that governs every other rule's stability. Next Steps was given the mechanical definition above in response. With Run A covering the four gate items and Run B the remaining seventeen, all twenty-one Tier 1 invariants now have recorded findings, and the condition is met. Ten Compliant, ten Non-compliant, one Unknown. It is worth naming what this does and does not do. It does not resolve anything: three Critical findings stand unfixed, four adjudications are open, and not one line of engine code has changed. It does not vindicate the register either — half of it came back Non-compliant, which is the register working rather than the engine passing. What it permits is proposing amendments through the process the Constitution already describes, which is a short note stating what changes, why, and what it deliberately leaves untouched, with Claude proposing and Viktor deciding, each one getting its own dated row. Three candidates are already waiting: the MVA gate wording the DEFECT row identified and the audit resolved, the blind-review method that produced Run 1 and is nowhere in the document, and Copilot's six candidate principles that have sat untouched in the appendix since Revision 1. None of them should be adopted today. The freeze was never about whether changes were good ideas — it was about not letting the document be revised by the same enthusiasm that wrote it, before anything had tested it. Something has now tested it, and the useful next act is fixing the ten Non-compliances rather than reopening the rulebook that found them.

Run C graded the twenty-three items in Tiers 2, 3 and 4, and with it every rule in the register has a recorded finding. Tier 2 came back four Non-compliant to three Compliant. The sharpest is T2-1: `calculate_dynamic_bias` in `bias_engine.py`, on finding a NaN, rewrites `close`, `EMA_20`, `EMA_50` and `RSI` directly on the `DataFrame` its caller passed in — no copy — and `engine_core` keeps using that frame for entry, risk and exit afterwards. Verified at lines 90 to 100. The path is unlikely to fire because `indicators` cleans NaNs upstream, but a module in the decision chain that can silently rewrite its own inputs defeats the boundary the whole audit relies on. T2-6 is the one to fix first because it is trivial: `requirements.txt` names five packages, two of which the code actually imports are missing (`requests`, `colorama`) and one of which nothing imports at all (`ccxt`, an order-execution library). A clean install cannot start the engine. Tier 3 returned three Compliant, three Non-compliant and four Unknown, and the Unknowns are the honest kind — process leaves no trace in a code snapshot, and the auditor said so rather than inferring discipline from tidy-looking labels. Tier 4 was graded as preferences and mostly honoured. The method differed from the earlier runs in a way that has to be recorded: Run C had web access and used it, fetching the public repository to check the one package claim that was externally verifiable. Runs 1, A and B did not. That makes T3-6's Compliant rest partly on evidence outside the bundle, which is a strength for that finding and an inconsistency across the audit. Claude then made the same fetch independently and closed two of the Unknowns: there are no tags and no releases, so T3-7 is Non-compliant rather than Unknown; and there are twenty-three commits on master, so the history is not the flattened snapshot Run C feared and the T3-9 depth question resolves in the engine's favour. One contradiction between runs, and Run B has it right: Run C's T4-4 says the duplicated `TREND` and `MOMENTUM` panel fields have since been fixed. The label was fixed; the number was not. `panel_render.py` prints `trend_health` on the `TREND` line, again on the `MOMENTUM` line immediately beneath it, and a third time as `Current Market`. Run C read the old scan report and inferred the fix; Run B read the source. Same model, same file, opposite conclusions — which is the clearest argument yet that a second auditor over the same items is worth more than a first auditor over more items.

“The goal is not to make this document sound better. The goal is to determine whether it can safely govern the engineering of the system.”

Viktor commissioned this one alone, on a model with no prior involvement, while Claude was unavailable — and checked first that doing so was permitted. It is an external assessment like Entry #19's, not part of Step 3, which is finished. The instruction was strong: explicitly adversarial, told not to assume good faith, told not to reward the document for sounding sophisticated, and told to evaluate what the text permits rather than what it evidently intends. It came back forty-one pages, twenty findings at Critical or High, and a verdict of **REQUIRES REVISION**. Two of those findings are correct and are adopted today. The first is that the freeze-lifting definition attached no consequence to what the audit found — the freeze lifted with ten Tier 1 items Non-compliant and three findings Critical, and nothing anywhere said the engine's output should not be acted on meanwhile. Claude helped write that definition and did not notice the gap. The second is that the amendment process, now live, would let a Tier 1 invariant be weakened by a short note reviewed as Claude proposes and Viktor decides — with the sharpest version being amendment by clarification, which is the exact argument shape Entry #21 already identified as the one the freeze existed to refuse. That finding is a criticism of an arrangement giving Claude considerable latitude over a document Claude largely wrote, and Claude thinks it is right anyway. Both are now in the Constitution, along with a front-matter sentence stating what the engine is. A third gap — Item 20 not naming crash-reporter capture of a process environment — is recorded and deliberately left unfixed, because closing it means amending a Tier 1 invariant and the rule just adopted requires that to be reviewed by someone who is not Claude. The new rule bit within the hour, on Claude. The rest is declined, and the reason is the reason Viktor named before either of them had read a finding: the review was instructed to treat this as a production-grade quantitative trading system, and it did, despite reading Item 1 and acknowledging in its own executive summary that the engine cannot execute. Roughly nineteen proposed new invariants follow from that framing — multiple-testing correction, calibration monitoring, an execution model covering fees, spread, slippage, latency and fills, market-structure handling for halts and delistings — every one of them correct for a system with a backtester and a live order path, and every one aimed at capabilities this engine does not have and is forbidden by Item 18 to acquire. Adopting them would take the register past sixty rules to govern a product that does not exist. They go to Future Amendment Candidates, to be revisited when the backtester is actually being built. Two structural lessons worth keeping. The instruction asked what invariants were missing and never asked which were unnecessary, so a hostile review under it can only ratchet upward; a future version needs a subtraction question. And the framing in an instruction outranks the content of the artefact — the reviewer had the correct definition of the engine in front of it, said so, and audited against the instruction regardless. The document now states what it is on page one rather than page fifteen.

“If it fails on your machine, the test is wrong before the engine is.”

There were no tests. Not thin coverage — zero: `test_live.py` contained no assertions and `requirements.txt` named no test framework, while the runtime log recorded seven distinct classes of failure across nine occasions where a change was accepted and then found broken by running the engine by hand. The Constitution's Step 7 says “regression test after each fix,” which was not a step that could be performed. Eighteen tests now exist: a compile-and-import check across all twenty files, a clean-checkout test, eight corrupted-data fixtures for Item 3, a golden-path snapshot against a pinned deterministic dataset, a determinism check, and two symbol-hardcoding tests. Five pass and thirteen fail, and the failures are the point — they are the Non-compliances written as executable acceptance criteria, each going green when its fix lands. All seven historical crash classes were reintroduced one at a time and the suite caught every one. Two things are worth recording beyond the artefact. The first is that `pytest` could not be installed in the authoring environment, so the suite was written as plain assert functions with a dependency-free runner alongside — which turned out to be the right shape anyway, since the suite has to work on a clean machine before `pip install` has succeeded, which is the exact situation the dependency test is about. The second is that the golden-path test was shipped unverified with the sentence above written into it, and then failed. The test was wrong, not the engine: it called `Phase7Engine.run()` directly, bypassing `SignalRouter` and therefore the entire decision layer, and returned the same dataframe for every symbol so the asset correlated with itself at +1.00. Both fixed, and the revision history is written into the file's own docstring rather than tidied away.

“A different method beat a different model.”

The harness was built to make fixes safe. Its first real run found defects instead. **One severity escalation:** the published repository does not start from a fresh clone. `main.py` builds its log handler at module scope on line 16 and creates the `Logs/` directory inside `main()` on line 41; `FileHandler` opens its file eagerly, so a clone without that directory raises `FileNotFoundError` during import — before `main()` runs, so the `try/except` inside it cannot catch it. Run 1 found the ordering and described it as soft: “the logging machinery catches it and prints to `stderr`.” That is not what happens. The finding was real and its severity was understated by a full category, invisible on the machine the engine was built on because `Logs/` has existed there since the first run. **Two findings nobody had:** the string `AERO` is hardcoded into user-facing reasoning text at `decision_model.py` lines 411, 413 and 419 and `panel_render.py`: 83, so running the engine on `SOLUSDT` produces an explanation that talks about `AERO`, and running it on `BTCUSDT` produces one claiming to compare `AERO` against `BTC` while comparing `BTC` to itself — which contradicts Run C's Compliant rating on T4-2, true of the arithmetic and false of the prose; and `decision_model.py`: 419 appends “relationship” to a label already ending in it, printing “a weak / no clear relationship relationship” to the trader. Four audit passes across three models read this code carefully and found none of the three. The reason is not that they read it badly. It is that reading is a different instrument from running, and the project had been buying more of the first while owning none of the second. That is the finding worth keeping from Phase A — more than the eighteen tests.

“Different lineage is necessary for judgment independence; it was never sufficient.”

Recorded because the audit's whole architecture rests on it and it had never been written down in one place. Nine models have now seen the Constitution: Claude as its author, then Copilot, Gemini, ChatGPT, Claude Opus 5 in the Reviewer 4 role, Grok, Kimi K3, GPT-5.6 Luna Pro for the hostile review, and GLM 5.3 for Step 5. Four have seen the engine source in a graded capacity — Claude, Kimi K3, DeepSeek V4 Pro and GLM 5.3 — and three more saw it during the build itself: Claude Sonnet 4, DeepSeek V3 and DeepSeek R1, through Aider. That last group is the one that matters, because it means DeepSeek's lineage had prior exposure to this codebase, and Run 1's claim of no prior involvement is therefore weaker than it was stated to be. Independence is tracked at the lab, not the checkpoint: a newer model from a family that has already worked on the artefact is not an independent reviewer of it, and treating a version bump as a reset would make the safeguard ceremonial. Still clean on both lists: Meta, Mistral, Qwen, Cohere, Amazon, MiniMax. One consequence already acted on — Luna Pro was considered for Step 5 and rejected, because it had seen the Constitution during the hostile review, and GLM was chosen on the basis of being clean on both lists at the time. It no longer is. Every run spends independence that cannot be recovered, which is an argument for deciding what a reviewer is for before opening the room.

"ML feels like the milestone that makes the project serious. That instinct is what this document exists to catch."

Viktor asked what could be counterproductive about adding machine learning — the obvious use being to fit the six bias weights against real outcomes rather than choosing them by hand — and then parked it. Not cancelled; parked, with conditions, which is the difference between a decision and a mood. The argument, verified against source before it was given: **training weaponises a dormant bug.** Item 2 is Compliant, but input-side `bfill` reaches the final bar through recursive indicators. Live mode only ever evaluates the current bar, so the leak has nothing future to pull from and stays inert. Training evaluates thousands of historical decision points, and at each one the backward fill pulls from bars that had not yet happened. The result is an excellent fit with no live edge, and the failure is silent — a good-looking equity curve, which is the most expensive kind of wrong this project could manufacture. **The training data does not exist:** `live_trading.py:210` does persist a decision, but nothing anywhere records what price did afterward and `engine_core.py:466` overwrites its state file each run rather than accumulating. Decisions without outcomes are unlabelled examples. **Item 11 makes fitted weights worse, not better:** fitting over features that are secretly the same feature produces large offsetting coefficients that swing on tiny data changes — crude weights wearing a lab coat. **Item 13's sentinels become learned signal:** `RSI=50` means the indicator broke, and a model would learn to predict from it. Two costs that are not bugs: reproducibility gains a whole new surface, since a decision would then depend on a weights file depending on a training run depending on a data snapshot, all needing version-pinning to the decisions they influenced; and explainability drops, which matters for a tool whose entire job is helping a human judge. Revisiting requires all five of: Items 2, 3, 11 and 13 fixed and verified by the harness; a working decision *and* outcome log with accumulated history; a backtester validated against known-answer cases; a kill condition written before the first fit; and an out-of-sample test that beats the current hand weights. If the fit does not beat 0.30 / 0.20 / 0.15 / 0.15 / 0.10 / 0.10 on data it never saw, that is a real result — it means the hand weights were fine.

#33

Step 5 — The Remediation Sequence, and What Verifying It Found

August 29, 2026

AUDIT STEP 5 — VERIFIED

“A fix made before its apparatus exists can be made, never resolved.”

Step 5 ran on GLM 5.3 — the pinned slug, not the floating `glm-latest` alias, which silently repoints when a successor ships and would have quietly poisoned this document's record of which model produced what. The instruction was Viktor's, not Claude's: it asked for position, dependencies, effort, risk-of-doing and risk-of-not per item, five direct questions, and an explicit flag wherever the reviewer's reasoning might be correlated with the model family that built the engine. Two edits before it went out — the build-timeline phrase removed, and a paragraph added about Phase A, because the instruction predated the harness and would otherwise have hidden three real findings from the reviewer. What came back organises sixteen items by verifiability prerequisite and shared edit seam rather than by severity, on the reasoning that severity decides membership in the gate sets and not position. Claude spot-checked every claim that could be checked against source; nine of nine held, including that `requests` and `colorama` are imported and undeclared while `ccxt` is declared and appears nowhere else in the codebase, that `close_time` arrives in the raw response and is discarded one line later so any staleness check must precede that line, that `engine_version` exists in config and is written nowhere, and that the indicator cache can never hit across runs because its key embeds the last close and the dict is per-process. Two things came out of the verification rather than the plan. The codebase contains two comments acknowledging previous instances of the same defect class — a gate comparing against a string literal no producer ever emits — which makes the dead `trend_failure` gate the third instance and argues for a class-level check rather than a third individual fix. And Claude came within one sentence of telling Viktor the plan understated that fix, on the grounds that the `STRUCTURE` column is never assigned — it is, at `structure.py` line 4502 of the audit bundle, via a `.loc` form the first search pattern did not match. That would have been the third false assertion of absence on this project. The exhaustive re-search caught it; the pattern is now a working note.

#34

The Machine Was Rebuilt, and the Baseline Survived It

August 30, 2026

REFERENCE — FILED FOR THE RECORD

“A golden baseline is a claim about an environment, and the environment was destroyed.”

Windows was reinstalled and every drive wiped. Nothing of value was lost: the repository stood at 375334a on GitHub and nothing lived only on disk except the gitignored `logs/` tree. The reason this earns an entry rather than a shrug is Item 5. The golden snapshot is not a claim about the code alone; it is a claim about the code *in an environment*, and that environment was destroyed and rebuilt from nothing.

`docs/audit_package/environment_before_reinstall.txt` records the library versions and the Python 3.12.0 that produced the original baseline. The rebuild resolved different ones — `numpy` moved from 1.26.4 to 2.2.6, `pandas` and `pandas-ta` unchanged — and the snapshot did not move. That was checked rather than assumed, and it is a real result. It is also the strongest available argument for what became audit Finding 15: **nothing in the repository required those versions**. The baseline held across a major version bump of the numerical library underneath every indicator, and it held by luck that nobody had arranged. Viktor's machine is now on Python 3.12.10 and the remediation sandbox runs 3.12.3; both produce identical suite results, verified on every delivery since.

"It tests that selected implementation details have not changed more strongly than it tests whether the engine is correct."

Sequence item 16 executed. GPT-5.6 Luna Pro received the engine source, the test suite, the frozen audit copy of the register and a written instruction (docs/audit_package/item16_review_instruction.md), and returned docs/audit_package/luna_pro_audit_report.md — a full 44-rule verdict table, fifteen findings, a test-suite assessment and a release-gate determination. **Release gate: not met.** Five findings rated Critical: Item 3, abnormal volume reaching analysis unchecked; Items 8 and 13 together, a failed macro input rendered as an ordinary neutral reading; Item 13, partially invalid indicator columns passing without degradation; Item 11, bias and confidence reusing derived evidence as if it were independent; and Item 14, AGGRESSIVE selected from conviction and entry quality with no independent risk decision. Ten further findings at Major or Moderate, and twelve rules graded Not verifiable. Two things about this run belong in the record rather than in a summary of it. First, the sentence quoted above — the auditor's judgement of the suite — is the sharpest single observation any reviewer has produced about this project, and it was aimed at a harness built four days earlier specifically to be better than that. Second, see Entry #41: the model that produced this report was not on the clean list the Remediation Plan had drawn up for this step.

"Under pytest, a return is a pass."

Two batches, one commit — dba1b63. **Batch 1, the skip mechanism.** The suite contained 46 occurrences of `if not _engine_available(): return`. A returning test is a passing test, so on a machine without `pandas_ta` the suite reported success for work it had not done — audit Finding 13, and precisely the defect Entry #29's harness had been built to prevent. All 46 became real `pytest.skip()` calls. Running the suite without `pandas_ta` then reported 46 SKIPPED instead of 46 false passes, which is how four tests that import `core.engine_core` directly, bypassing the guard, were found to fail outright in that environment. Recorded at the time and deliberately not fixed in the same batch — see Entry #38. **Batch 2, five unambiguous fixes.** The "full size" text in `decision_model.py` (Finding 8, path B); `main.py`'s bare 'Logs' (Part 6 observation 2); `trade_quality_current` and the vacuous assertion that guarded it (Findings 10 and 13); the hardcoded `Lookback 8` on the panel, now interpolated from `config.STRUCT_LOOKBACK` (Finding 8, path A); and `np.isfinite` in risk validation. **None of those five named the audit finding it closed.** Three whole findings were later discovered to be already shut for exactly that reason — Entry #40 records the cost from the other side.

“Decide items 3, 11, 14 myself.”

Viktor's instruction, delegating all three trading-judgment calls rather than ruling on each in turn. Recorded as a delegation because Roles & Authority assigns this class of decision to Viktor, and this is the second time he has knowingly handed a specific class of them over. Commit c4dfcc7. **Item 3 — abnormal volume: reject, degrade or accept, per case.** All-zero volume is now rejected at `data/validation.py`, because there is no measurement from which to build a volume-weighted read. An isolated extreme spike is still accepted at that layer: the original “deliberately not implemented” reasoning held, in that a spike is real data and rejecting a run over a busy market makes the engine least available exactly when it matters most. What changed is that a spike no longer reaches every downstream score unflagged — `indicators.py` detects one above ten times the recent rolling median and records it as a degradation, capping confidence rather than substituting a value. **Item 11 — remove the duplicated factors rather than argue for their independence.** The earlier sequence-item-11 pass had removed one duplicated term and left three siblings standing. `bias_score` is now the one place all six weighted factors are combined and confidence is exactly its magnitude; `continuation_strength` no longer carries a trend-health-derived component; and `bias_engine.py` gained the explicit dependency-graph comment the audit's required action had actually asked for. **Item 14 — the labels survive, gated.** `classify_risk_regime()` already computed a four-tier regime, and only the EXTREME-or-not boolean ever escaped `validate_risk_parameters`. The regime is now threaded through as its own contract field, and AGGRESSIVE is refused when it is HIGH VOLATILITY RISK or worse. Direction and whether a trade is permitted at all are untouched; this gates intensity only. Ten net new tests. The golden snapshot moved, in exactly the four fields predicted before the run.

“A test docstring naming what it did not fix is a debt with an address on it.”

Four commits. **Item 8/13, the macro half (5d2cbb6)**. A failed macro-timeframe fetch left `macro_bias` at its initialised "NEUTRAL" with nothing added to degradation, so a failed higher-timeframe read and a genuinely flat one rendered identically on the panel. Both the validation-failure path and the processing-exception path now record a degradation. `test_the_macro_series_is_actually_read` had documented this in its own docstring as “recorded rather than fixed... a rider on sequence item 9's degrade ruling”; that assertion was rewritten alongside the fix, which is what the docstring had asked whoever fixed it to do. **Finding 15, dependency pinning (a26c545)**. `requirements.txt` and `requirements-dev.txt` now pin exact versions, verified by building a virtualenv from empty and confirming `pip freeze` matched before the suite was run. Closes the risk Entry #34 records the project having survived by luck. **The four unguarded tests (3d0b410)**. The ones Entry #36 recorded and left. Three are entirely about `Phase7Engine` and now carry the same `_engine_available()` guard as the rest of the suite. The fourth, `test_every_module_imports`, needed a different fix: it checks all 21 engine modules for import-time defects and only three of them need `pandas_ta`, so a blanket skip would have stopped checking the other eighteen for an unrelated reason. It now excuses that one exception by name and still fails on anything else. **Stale labels (2be405f)**. Two golden-path tests still carried “EXPECTED TO FAIL until sequence item 12” in their docstrings; item 12 had fixed both defects before the Luna Pro audit began and both tests had been passing ever since. Documentation only — but a docstring that lies about the state of the code is the same class of defect as a panel line that does.

“A derived document cannot tell you what it never contained.”

Found by reading `luna_pro_audit_report.md` end to end rather than the roadmap written from it. Finding 3 — Item 13, partially invalid indicator columns passing without degradation — is one of the audit's five Criticals and had never been entered into any roadmap, so it was never scheduled, never ruled on, and never noticed missing. From 31 August to 2 September `PHASE7_NEXT.md` stated that remediation of the five Criticals was complete. Four of them were. Commit 108cc9f. **The defect.** Every indicator guard asked one question, `.isna().all()` — did the calculation return nothing at all. That cannot catch a series with 299 good values and no value at the bar the decision is made on, and it could not even in principle: `clean_series(method="forward_fill")` had already filled that gap with the previous bar's number before the guard ran. **The class was wider than the two instances the audit named.** Injecting a trailing NaN into each indicator in turn, before the fix: ATR, RSI, ADX, SuperTrend and both EMAs — every one carried a stale prior-bar value into the decision row, and not one recorded a failure. Two had no guard at all to fix, the SuperTrend *level* and the EMAs. The same trailing fill ran on the raw OHLCV columns, turning a truncated final candle into a synthetic bar repeating the previous close. **And the consumers were re-fabricating what sequence item 9a had removed.** A missing RSI fell back to `50.0`, inside the “not extended” band, scoring the full 15 of 15 — while `indicators.py`'s own failure text told the operator it scored 0 of 15. A missing HVN fell back to `close`, making the distance exactly zero and scoring 12 of 12; that is byte-for-byte the defect item 3 had fixed for VWMA six days earlier, sitting forty lines above it. A missing ATR fell back to `close * 0.02`, the flat constant item 9a had deleted — on the pinned fixture, 0.016035 against a real 0.010554, a 52% overstatement of the distance that sets the stop. **Ruled:** a value not measured at the decision bar is absent, and absent means that indicator failed — which hands it to the degradation machinery that already exists. No new policy was invented; the existing one was applied one row over. Verified end to end: before the fix, a trailing-NaN ATR produced `degraded: False, missing_inputs: []`, a full stop and three targets, and moved entry quality from 45.18 to 45.25 — a different answer, reported clean. Eighteen new tests, fourteen of which fail against the pre-fix code; the other four are controls that must pass on both sides. The golden snapshot did not move, predicted in advance and for the right reason: real pinned data has zero trailing NaNs, which is why this went unseen. **Recorded against Claude:** the first draft of that test file reintroduced the defect fixed in 3d0b410 the same morning — five of its tests imported `indicators.indicators` with no `_engine_available()` guard and errored instead of skipping. Caught by running the suite in a `pandas_ta`-free virtualenv before packaging. Same class, same day, inside the fix for the class.

#40

The Status Sweep, and a Count That Was Wrong by Three

September 2, 2026

OBSERVATION — FILED FOR
THE RECORD

“Neither document knows the state of the code.”

Claude told Viktor on 1 September that nine Major and Moderate findings remained open. That number came from the audit report rather than from the code, and it was wrong by three. Checked by opening each location the audit quoted rather than trusting the summary: **Finding 8** (inaccurate user-facing claims) was fully closed — Lookback interpolated from config and the “full size” text gone, both in batch 2, with the macro path closed by 5d2cbbbe. **Finding 10** (the unconsumed trade_quality_current) was closed in batch 2. **Finding 15** was closed at a26c545 the previous day. None of those three commits mentioned a finding number, which is why nobody knew they had shut one; Entry #36 records the same omission from the authoring side. Genuinely open: Findings 6, 7, 9, 11 and 12 in full, plus the remainders of 13 (the plotting test still asserts only the absence of ERROR records) and 14 (the required-shape assertion still omits provenance, degradation and the log path). Seven, not nine. The status table now in PHASE7_NEXT.md cites a file and a line for every verdict, so the next reader can re-check it rather than trust it. This is the mirror image of Entry #39: that one was the plan missing what the report had, and this one is the report not knowing what had been fixed since.

#41

The Step 8 Auditor Was Not on the Clean List

September 2, 2026

OBSERVATION — FILED FOR
THE RECORD

“The party that made the error should not be the party that decides the evidence of it disappears.”

Phase7_Remediation_Plan.pdf, 29 August, records that Luna Pro was considered and rejected for Step 5 because it had already seen the Constitution during the hostile review, and names the six families still clean on both lists for the Step 8 re-audit: Meta, Mistral, Qwen, Cohere, Amazon and MiniMax. **Step 8 was run on Luna Pro the following day.** Entry #31 sets out why the ledger exists — independence is tracked at the lab, not the checkpoint — and Entries #18 and #20 record Grok being removed from the auditor plan for precisely this reason, having read the Constitution in an earlier conversation. The same disqualifier applied here and the ledger was not consulted. **What this does not mean.** The findings are sound. Every one was checked against source before any of it was fixed, and Finding 3 was verified by running the engine rather than by trusting the report. Nothing here argues for discarding the report or redoing the work it produced. **What it costs** is the one thing the ledger buys: the next re-audit cannot treat agreement with this one as independent corroboration, because it would be the second reading by a model that had already read the standard. The next re-audit should go to one of the six families named above, and its agreement with Luna Pro should be read as agreement between two readings, one of them contaminated — not as confirmation. Recorded rather than quietly noted, on Entry #21's reasoning.

“A document you have sampled is a document you have not read.”

Viktor asked whether the Constitution needed upgrading. Answering that honestly required reading all nine governing documents rather than sampling them — Entry #17's lesson, and rule 17 of PHASE7_NEXT.md's earned list, which exists because Claude once declared this exact PDF clean after searching four guessed phrases. **The answer is no, and the project had already ruled so.** Entry #26, written the hour the freeze lifted: the useful next act is fixing the Non-compliances rather than reopening the rulebook that found them. Four things the read surfaced that were in no working document. **One: the release gate is in force and blocks on Item 6.** Viktor's 29 August ruling raised Traceability from Major to Critical; Item 6 is the audit's Finding 7, still open. Until it lands *and is re-audited*, no output may be relied on for a real trading decision, and the backtest architecture is not built. Nothing in PHASE7_NEXT.md had said so. **Two: the adjudications were ruled on 29 August**, in Roadmap Revision 4, whose title says exactly that. Claude reported four still open, having read the Engineering Notes and the Constitution — both of which stop before that ruling — and not the Roadmap. Only the Item 20 amendment remains. **Three: Item 20's gap is already served operationally.**

Phase7_Credential_Security_Protocol.pdf §6.2 covers telemetry, crash reporting, error aggregation and usage analytics, and is tagged “Enforces: Tier 1, Items 20 and 21” — written the same day as the invariant whose channel list omits it. The pending amendment tidies the register to match a practice that already exists; it does not close a live hole, and the engine holds no credentials at all. Unblocking it needs one uninvolved model given a text extract, not the PDF that Gemini's and Copilot's classifiers refused. **Four: the documentation is a week in arrears** — this log stopping at #33 being the largest part of it, alongside no Change Impact Records since the standard took effect on 25 August, no Version History row for Step 8, and two companion documents still pointing at Rev 6. **One check nobody has run.** The item-16 instruction requires the re-audit to grade against the frozen audit copy rather than the live version annotated with the outcomes of a previous audit.

Phase7_Constitution_v1.0_RATIFIED_AUDITCOPY.pdf is 69,656 bytes against the live document's 112,290; the smaller size is consistent with a frozen copy and there is no reason for alarm. It is still worth opening once to confirm it carries no AUDITED, AMENDED or DEFECT row, because it is a one-minute check on whether the re-audit graded the exam or the answer key. *(An earlier draft of this entry compared that byte count against the 101,831 recorded in Phase7_Audit_Execution_Instructions.pdf and called it a mismatch. That was wrong — those instructions describe the Step 3 package for Kimi K3, a different audit with a different material set. The claim was made to Viktor before it was checked. Corrected here rather than deleted.)*

“A hash proves the input changed. Only an archive lets you rebuild it.”

Item 6, Traceability, which Viktor's 29 August ruling raised from Major to Critical — and therefore the finding holding both the release gate and the backtest gate shut. Commit 302db8b. **The gap.** Sequence item 12 built a decision log: what the engine concluded, plus a five-field fingerprint of what it saw. The fingerprint was a last-candle timestamp and a row count, and two different frames can share both. Nothing stored told them apart, so “reconstructable” was a word in the Constitution rather than a property of the engine. **Viktor's ruling: hash AND archive, pruned at ninety days.** The Constitution says under Item 5 that the retention decision is one “the audit should force explicitly rather than leave implicit”; this is that decision made rather than assumed. The two halves do different jobs and the difference is the whole design. **The hash detects** — it costs nothing, it lives in the log, and the log is never pruned, so a decision from two years ago can still be checked against data fetched today. **The archive reconstructs** — the actual candles, the only thing that can rebuild a run whose source has since changed, and the only part with a cost, which is why it is the only part with a limit. About 31 KB a run; roughly 2.8 MB steady-state at the cap. Past the window a run does not become unverifiable — it becomes *verifiable but not rebuildable*, and the record says which by whether the file is still there. **Two design points worth keeping.** The hash is taken on the OHLCV as validated and *before* any indicator column exists: hashing the final frame would fingerprint this engine's own arithmetic together with the market data, so changing an indicator length would present as the exchange having changed its candles. And the canonical form is written out explicitly rather than leaning on `to_csv`, because pandas has changed its float repr and NaN spelling across versions this project has already run on — Entry #34 records numpy moving 1.26 to 2.2 across the machine rebuild alone. **Also closed:** the six bias weights live in `bias_engine.py`, not `config`, so changing 0.30 to 0.35 changed every decision the engine makes while leaving two runs byte-identical in the log. They are now fingerprinted, read live from the module rather than restated. **Verification.** Fifteen tests, and the central one does not inspect fields — a test asserting “lineage” in `decision` would pass just as happily over a section full of nulls, which is precisely the shape Luna Pro's assessment of this suite warned about. Instead it takes the archive, rebuilds the candles from it, runs the engine again against the rebuilt data and nothing else, and requires the identical decision. The golden snapshot moved exactly as predicted before the run — lineage added, provenance gained seven keys, and not one existing value changed. This commit adds a record and alters no output.

“Every hash matched. Every number matched. The whole difference was one separator.”

The Findings 6 and 7 work built the archive path with `os.path.join`, which returns a forward slash on Linux and a backslash on Windows. That path is written into the decision log, so the same run archived on the two platforms recorded two different strings for one file — and the golden snapshot, baselined in a Linux sandbox, could only ever match on Linux. It failed on the first run on Viktor's machine. **Two things about it are worth more than the fix.** The first is the near miss: `decision_log.log_path()` has the identical shape and escapes the bug by luck rather than design — `config.LOG_DIR` already ends in a separator, so its single join never inserts one. The archive path joins a directory level deeper, which is where the backslash appeared. Existing code working was not evidence that the pattern was safe, and it was read as if it were. The second is why nothing caught it earlier: **the entire verification for that commit ran on Linux, where the bug is invisible**, and a golden snapshot only ever compares a value against itself on whatever platform baselined it. The patch had been called “verified against a fresh checkout” — true, and on one platform. Fixed by normalising to forward slashes for both the file I/O and the recorded value, since Python accepts them on Windows and one spelling then serves both purposes. Pinned by a test that fails on Linux too if it regresses. Recorded as its own entry rather than folded into #43, because a defect that reached the operator's machine inside a commit whose message claimed thorough verification is not a footnote to that commit.

“Reading finds what is written. Running finds what happens.”

Viktor ran `python main.py` against live MEXC data the evening Findings 6 and 7 landed. AEROUSDT 4h came back bearish on every measure the engine has — bearish bias, bearish regime, bearish structure, an LH-LL sequence, strong bearish distribution, SuperTrend freshly flipped — with one dissenting reading, `MACRO TREND: BULLISH`. It printed **DECISION: CONSERVATIVE LONG**, with the stop at \$0.4889 above a price of \$0.4725 and three targets descending to \$0.4233. **A long label on a short plan.** Every number in that plan was correctly computed from a correct bearish analysis; the word attached to them was not. An operator following the `DECISION` line would have bought an instrument the engine had just analysed, in detail and correctly, as a short. It also printed “Bias is bullish and the broader macro trend agrees” four lines above its own Validation Note reading “The higher timeframe disagrees with this bias.” **Two causes.** `decision_model.py` opened a direction from any of three independent sources — `raw_bias` or `long_signal` or `macro_bias` — so the macro clause alone was enough; and `trend_health >= 50` then passed because trend health is an `UNSIGNED` magnitude, so a strong bearish trend scores 69 and no bearish evidence anywhere in the run could block it. The bearish block below never ran, because the bullish one returned first. Meanwhile `risk_model.py` builds stop and targets from `detailed_bias` alone. Two direction sources, never reconciled, in two modules neither of which knew the other existed. **Viktor ruled: bias is the sole direction source.** Macro keeps its existing 10% vote inside `bias_score` and gets no second, overriding one — letting it override the blend counts one piece of evidence twice, which is Item 11 in the module that picks the side. Entry-zone signals lose the same privilege for the same reason. Two further instances surfaced in testing that nobody had seen: the mirror case returned `CONSERVATIVE SHORT`, and a `NEUTRAL` bias with a long entry signal returned `AGGRESSIVE LONG` — the most confident label the engine has, from no directional view at all. **And a guard, because narrowing a source is not the same as checking.** `_refuse_incoherent_plan()` reads direction off the targets themselves and refuses any action whose label contradicts its own levels. Deliberately a refusal rather than a relabelling: one of the two sources is wrong and nothing inside that function can tell which, so `NO-TRADE` is the only answer available that is certainly not the wrong one. **What this says about the audit.** Luna Pro read the source and the tests and did not find this. Neither did four earlier passes across three models. It needs bias and macro to actually disagree on live data, and no pinned fixture makes them.

“A disk problem must neither destroy an analysis nor veto one.”

Found while writing the halt-safety test for the raw-input archive, and deliberately left unfixed in that commit: fixing it there would have let that test pass for a reason unrelated to what it was written to prove. **The cause was smaller than “the engine cannot log.”** `route()` opened with two unguarded `os.makedirs` calls that wrote nothing. Every writer in this engine already creates its own directory on demand inside its own error handling — `decision_log.write` returns `None`, `_save_state` warns, `plot_engine_chart` warns, `lineage.write_archive` returns `None`. The two calls duplicated all four and added a failure mode at the worst point in the run: the top of `route()`, before anything had been computed. An unwritable log directory raised there, the broad handler reported “Router execution failed”, and the operator lost the entire analysis rather than merely the log of it. Four independently recoverable conditions collapsed into one total failure. Same class as sequence item 14's own `REQUIRED_DIRS` finding, which removed the list and left these two calls standing — earned rule 22 collecting on the item that wrote it. **Viktor ruled: a run whose decision log cannot be written still authorizes a trade.** It warns, the panel makes no claim that anything was logged, and the operator decides. His 29 August degrade-not-halt ruling applied literally. **Claude recommended the opposite and was overruled.** The argument made and rejected: Item 6 is Critical, and a trade taken on a decision that left no trace is unauditably by construction — which differs from a failed *archive*, where the hash still lands in the log and the run stays verifiable even though it is no longer rebuildable. When the log itself fails there is no record at all. Recorded here because the difference matters to whoever audits this next: a ruling with its trade-off on the table reads differently from an absence, and the cost is stated plainly — the one decision an operator acts on without a record is the one an auditor would ask about first. The tests pin the ruling and say so in their own docstrings, so a later reader who thinks the assertion looks wrong finds the reasoning instead of guessing. **Categorisation note:** a failed write is not added to degradation. That list blocks trading and is about inputs the ANALYSIS was computed without. Nothing here was missing from the analysis — only from the filing of it.

"I really should have run the engine more often, but I was too caught up in the workflow."

Viktor's observation after Entry #45. The better answer than resolving to be more diligent — which decays — is to make the suite able to reach the state, so that it does not depend on anyone remembering. **The reason nothing caught it is exact: no pinned fixture makes bias and macro disagree.** The committed series trends one way and its daily aggregate trends the same way, because the aggregate is built from it. A suite cannot reach a state its fixtures never enter, so a hundred and seventy tests passed over a defect that only exists in that state.

`tests/test_timeframe_disagreement.py` generates a series where the timeframes genuinely disagree: a long rally, then a sharp multi-day break, so the daily EMA-50 lags and still sits below the daily close while the 4h structure has decisively turned. An ordinary market condition of the kind that happens most weeks, not a contrivance built to trip a branch. Generated rather than committed, following `test_golden_path._write_pinned_set`'s discipline: one pure function, all three files derived from it, and no RNG anywhere — the wobble is `sin()`, so the series is byte-identical on every machine and every numpy, which matters on a project whose baseline has already survived a numpy major version by luck. **The parameters were chosen with margin rather than at the edge.** A 0.16 drop put the daily close within 0.0007 of its EMA-50; 0.14 leaves real daylight. A fixture that only just produces its condition stops producing it the first time an indicator length changes, and stops silently. **Two of the seven tests assert that the fixture still splits the timeframes,** before anything is asserted about behaviour — if it ever stops, every other test in the file would go on passing while testing agreement. That is section 7.3's "setup contradicts what it claims to test", and a file like this is exactly the shape that fails that way. The property pinned is deliberately not "the engine returns WAIT here" — that is today's answer to today's thresholds. What must never be true at any threshold is a LONG label above descending targets. Four of the seven fail against the pre-fix code, reproducing the live run from generated data.

"Nobody could have caught it by thinking harder. Only by reading the bill."

Viktor delegated this one — "This is your call and I trust you make the right decision" — so it is recorded with its reasoning rather than only its conclusion, and can be overturned on the argument. **The project held two positions and had not noticed they cannot both be true.** Entry #18 resolved Grok's prior exposure by running Step 3 "in a fresh conversation with no shared memory of the earlier one." The Remediation Plan, eleven days later, rejected Luna Pro for Step 5 *because* it had read the Constitution during the hostile review — the same kind of exposure, treated as permanent. Entry #41 then graded Step 8 against the stricter reading without noticing the looser one existed. Which answer you got depended on which document you opened. **Three different things were being recorded as one.** SESSION exposure is a model reading the document in a chat; TRAINING exposure is the artifact being in the weights; LINEAGE exposure is a sibling from the same lab having worked on it. Entry #31's rule — independence is tracked at the lab, not the checkpoint — was written for the second and third. Applying it to the first is a category error, and it is why the clean list emptied faster than it needed to. **Ruled: session exposure is cleared by a fresh conversation where training-on-conversations is off; training and lineage exposure never clear.** Where it cannot be established whether a session fed training, treat it as permanent. This ruling is LESS strict than the reading the project had been using, which is worth stating plainly: over-strictness had a measurable price here, in models spent for a reason a fresh conversation removes. Kimi K3 becomes available again. **Then the assumption was checked rather than assumed.** Viktor exported the OpenRouter activity log — 713 requests, every model, every route. All `variant=standard`; not one free endpoint in the entire history, so no material was ever sent anywhere that trains. It also dated the Step 3 truncation exactly: `completion=16,384`, `finish=length`, at 2^{14} output tokens — Entry #23's "default token ceiling" with a number on it. **And it showed the ledger was wrong.** Five models ran through Aider on 22–24 August while the engine was being built. Entry #31 names three of them and misses `mistral-nemo` (71 requests) and `claude-3-haiku`. Mistral's lineage worked on this codebase through exactly the mechanism used to disqualify DeepSeek, and had sat on the clean list of eligible auditors for a week. Entry #31's table was written from recollection. The correction matters less than how it was found: every provider bills per request, and the bill cannot be mistaken about what was called.

"Eight tools were enabled. One `git clone` would have handed over the answer key."

The package the round-2 auditor receives is now **generated** by docs/build/build_audit_package.py rather than assembled by hand. Entry #24 records what hand assembly cost: "a false claim found in Claude's own Step 2a package." A script that walks the repository and computes the manifest from the bytes it wrote cannot make that class of error, and it *refuses* to build if docs/ would ship, because intending to withhold the answers is not the same as withholding them. Three things Luna Pro never got: version-control history (metadata only — subject lines withheld, since "Audit Findings 6 and 7" leaks a finding number and its outcome in eleven words), the project files that decide how dependencies are pinned, and execution transcripts. Five rules came back Not verifiable in Step 8 *only* because history was withheld; that was a defect in the package, not the project. **The instruction's Rev 2 states two things the project would rather not say:** that Step 8's own auditor was not on the clean list, and that this repository is public, so any model trained since may hold the codebase in its weights — a channel the ledger cannot see and which applies to the chosen auditor as much as anyone. It also says plainly, in its own section, what this round cannot be: the fixes shipped with comments and docstrings describing the defects in detail, so this is not a blind re-discovery but a check of whether the claimed fixes are real. **And the near miss.** The run started with eight tools enabled, including a shell, web search and web fetch. Nobody chose that and nobody checked. It probed an empty sandbox and found nothing — but the repository is public, and one command would have handed over PHASE7_NEXT.md, the previous audit report, the Engineering Notes and every commit message. Three further tools were worse in a quieter way: Fusion, Advisor and Subagent would each have silently put other models inside the audit, leaving the report labelled with one model's name while being partly the work of models nobody chose. It was caught by looking at a screenshot at the right moment. That is luck, and earned rule 28 — written four hours earlier the same day — says luck is the signal to change the setup rather than to be more careful next time. A pre-flight checklist now exists, and the omission was Claude's: the setup walkthrough covered model, provider, max tokens, streaming, file parser and data retention, and never mentioned the tools.

"This file is named but its content is not there."

The round-2 re-audit has now failed three times, and not once on its own merits. Attempt 1, Qwen3.8-Max: repeated provider-side failures before grading began. Attempt 2, Kimi K3: it received the Constitution as a PDF, could see the filename and not the contents, said so, and stopped. Attempt 3, Qwen again: the same defect. **All three refusals were correct behaviour.** The common factor across two different providers on two different days was the PDF, so the Constitution now ships to auditors as text — a pdftotext -layout extraction carrying the source PDF's SHA-256 in its own header and verified byte-for-byte against a fresh extraction. Earned rule 28 applied: a format the recipient cannot open is not fixed by asking the recipient to try harder. **One thing recorded against Claude.** The first diagnosis was built from a truncated sentence visible in a screenshot — "this file is named but its content is" — and asserted as a parsing failure. The full reasoning, recovered later, said something stronger and more specific. The fix was right; the confidence behind it was not earned at the time it was stated.

#51

Eleven Claims, and the Rule About Believing Them

September 3, 2026

FINDINGS — FROM AN
UNFINISHED AUDIT

“An audit that produced no report still produced eleven defects.”

Attempt 3 could not read the standard and correctly refused to grade 44 rules it had not seen. It then ran the instruction's own Section 7 checks — which are defined in the instruction rather than in the Constitution — and returned eleven observations. Its reasoning is preserved at `qwen_reasoning_1.txt` through `_4.txt` in the repository. **Every one was treated as a CLAIM and checked against source before a line of code was written.** None evaporated. Two got worse under checking. That discipline is the entry: a reviewer's finding is a hypothesis with a file and a line attached, and the difference between an audit and a rumour is whether someone opened the file. The eleven span fabricated levels, false claims about files, a run identity that did not identify the run, a smoothing mismatch between an indicator and its own fallback, an unset HTTP timeout, and a directory created as a side effect of an import.

#52

The Run Hash That Did Not Identify the Run

September 3, 2026

DECISION — ADOPTED

“Two different trading plans on the same candles, recorded as the same run.”

`core/decision_log.py` has carried this sentence since the Finding 6 fix was written: the audit's required action asks for “all decision-affecting configuration, **including risk-model multipliers and bias weights.**” Fifteen lines below it, the fingerprint dictionary named one module — the bias engine. The weights were covered. The risk multipliers were not. Changing `ATR_STOP_MULT` from 1.2 to 1.5 moves the stop and all three targets on every run, and left the hash, the config snapshot, the module snapshot and provenance byte-identical. Item 6 is Traceability, raised to Critical by Viktor on 29 August — broken inside the fix written to satisfy it. **It was not a name missing from a list.** The multipliers were instance attributes on `RiskModel`, and the snapshot mechanism reads module attributes; there was nothing for the list to name. Viktor ruled the constants move to module level rather than the mechanism grow to reach inside objects — the smaller change, and the shape the one compliant module already had. `__init__` is gone entirely, so no instance state can drift from what the record reports. Proven behaviour-preserving across 1,728 stop and target combinations and 192 regime combinations: zero differences.

#53

A Test That Described the Defect It Could Not Catch

September 3, 2026

PROCESS — RECORDED

“...so the default never fires and the panel prints 'saved to None'.”

That sentence has been the failure message of a test in `test_traceability.py` since sequence item 12. It is correct. It describes, exactly, a defect that was live in the engine the whole time — and the assertion beneath it only checked that the OLD hardcoded default was gone from `panel_render.py`. The `None` arrived from somewhere else: the router built the decision object with `str(chart_path)`, and `str(None)` is the four-character string “None”, which is truthy, which passes the panel's own gate. The panel printed **Chart saved to None** — a file that does not exist — for the life of the engine. **A second test had the same gap from the other direction.** `test_unwritable_log_dir.py` already pointed the chart directory inside an unwritable path, so every run of that file drove the engine through this exact state, and it asserted only the log path and the archive. Nobody looked at the chart. A test whose failure message documents a defect it cannot detect is worse than no test: to anyone auditing the suite it reads as coverage. The claim is made in one file and the value built in another, so the check now follows the value.

#54

The Door Closed and the Window Left Open

September 3, 2026

DECISION — ADOPTED

“Twelve of twelve points for a volume node that was never located.”

`StructureEngine.analyze()` wrapped each of five sub-routines in its own handler and substituted a value on failure — the previous bar's regime, “NONE”, “NEUTRAL VOLUME”, and for both volume nodes and the swing level, **the current price**. None of the five recorded anything, and the function had no channel to record through, so a crashed detector produced an ordinary-looking reading and the run was never marked degraded. **The volume-node case is Finding 3 again.** Entry quality scores structure proximity as $\text{abs}(\text{close} - \text{hvn}) / \text{close}$, which with `hvn` equal to `close` is exactly zero — inside the < 0.015 band, awarding the full 12 of 12. The comment describing that arithmetic has sat in `entry_model.py` since 1 September, written when Finding 3 was fixed. That fix changed the *consumer's* fallback to `NaN` and left the *producer* handing down a finite number equal to the price. It even cites rule 18 — fixed the instance it was looking at and left its twin. This is the twin of that twin. No handler invents a measurement now: an unlocated level is `NaN`, which every consumer already guards, an undetermined label says UNKNOWN rather than borrowing NEUTRAL, and every failure reaches the run's degradation list.

#55

Found by Reading the Panel, Twice in One Morning

September 3, 2026

FINDINGS — FROM EXECUTION

“226 tests, and neither defect was found by any of them.”

The first was Claude's, shipped that hour. The patch above replaced a panel line that ended in a newline with a computed one that did not, and the engine printed `SWING STRUCT : $0.4700 (Lookback 8)STOP LOSS : $0.4636`. Thirteen new tests, a negative control, a full-suite run and a golden-snapshot check all passed — because every assertion asked whether a substring was present, and it was. Nothing in the suite rendered the panel and looked at its shape. Viktor found it in one live run, which is earned rule 30. **The second was not Claude's.** The same panel printed `MACRO TREND: BULLISH` and, four lines below, “The higher timeframe is neutral.” The branch that produced it fired when the BIAS was neutral and then made a claim about the MACRO — false on every neutral-bias run, which is a common state and was the state that morning. Same shape as the contradiction that exposed the direction Critical the day before. The logic is now a named function with a fourth branch, extracted so the four cases can be tested at all: inline, the false branch could only be reached with market data that happened to produce a neutral bias under a directional macro, which is why no fixture held it and no test could have.

#56

A Conversation Is Not a Record

September 3, 2026

PROCESS — RECORDED

“A full day of findings, rulings and sizings existed in a chat window and nowhere else.”

Late in the session Viktor asked whether Claude would still have everything the following day. It would not have. Eleven verified findings with their file and line, four patches, three rulings made, three owed, sizes for the remaining work and the pre-audit checklist were all established in one conversation and written into no document. It was caught because he asked — which is precisely the save-by-vigilance earned rule 28 exists to replace. **The fix is a handover check, run by Claude unprompted at the end of every session**, asking six questions: is the state in the entry-point document rather than only in chat; are the day's rulings recorded; does `git status` show untracked files that matter; are these notes current or the gap stated; are loose patch files unapplied; is any evidence still in a chat window. On its first run it found three files left on the floor, one of them sitting there since the previous day. **The same conversation split the project's goal in two.** Phase 7 is now both the technical portfolio project in Viktor's 2026–2028 career plan and an engine he intends to finish, and those end months apart. Portfolio-ready is reached, tagged, and only then does backtesting begin — on top of the tag, so that the phase which destroyed the previous build cannot damage the version he submits. Earned rules 30 through 34 were written the same day.

#57

The Trend the Bias Engine Could Not See

September 4–5, 2026

DECISION — ADOPTED

“Continuation was zero, so the trend was flat — whatever the trend was doing.”

Ruling 3, and patch I. `bias_engine` was never told which way the trend pointed. It inferred direction from the sign of `continuation_strength`, so whenever continuation came out at zero the engine read the trend as flat — silently, and regardless of what the trend actually was. `indicators/trend_health.py` had computed the direction all along and thrown it away. It is now exposed as `trend_direction` and `trend_direction_sign`, and the sign is a **required** parameter of the bias engine, validated against (-1, 0, 1) and raising rather than defaulting — an inferred value replaced by a declared one. **The order in which this was decided is the entry.** How often the engine silently flattened a real trend is a question with an answer in the data, and the answer was obtained across 9,800 live bars *before* the ruling was made, not afterwards to justify it. Earned rule 36 was written from this: a ruling made first and measured second is a hypothesis with a decision attached to it.

#58

A Magnitude Offered as a Direction, and a Lean Too Small to Act On

September 5, 2026

DECISION — ADOPTED

“A magnitude was being cited as evidence for a direction it cannot indicate.”

Rulings 1 and 2, both in patch K. **Ruling 1:** the reason strings the engine writes for a BULLISH or BEARISH call cited trend health as supporting evidence. Trend health is a magnitude; it says how strong, never which way. Sitting in a directional sentence it read as directional support that it was not. It now prints as `trend strength 78/100 (rising)`, with the direction word coming from the sign patch I had just exposed — ruling 1 was only buildable because ruling 3 landed first. **A correction to the record:** Claude had claimed the panel's TREND line carried the same defect. It does not. It prints two true facts side by side and asserts no relationship between them; only the reason strings made the claim. **Ruling 2:** an ACTION could be issued off a bias that barely leaned at all. `MIN_ACTION_BIAS = 30.0` now floors it, and is fingerprinted so it enters the run's identity. Claude first argued for folding it into the existing `RAW_BIAS_THRESHOLD` and reversed that while building it: one asks *does the blend lean far enough to call a side*, the other asks *is that lean strong enough to act on*, and one number tuned for both makes one of the two answers an accident of the other. Seven fixtures that built a bias with no score at all were **completed** rather than worked around — relaxing a floor to accommodate incomplete test data is how a floor stops meaning anything. Patch J, alongside, extracted `volume_agreement` into a module-level function for the same reason patch D extracted `macro_agreement`: logic reachable only through a full run cannot be tested in isolation.

#59

Three Patches Built on a Base That Had Already Moved

September 5, 2026

PROCESS — RECORDED

"The staged copy is a snapshot, not a view."

Patch J was generated against an `engine_core.py` staged before patch I landed, and failed 39 tests. Patch K was rebuilt twice against a `decision_log.py` missing an anchor added hours earlier. None of the three patches was wrong; the base each was diffed against was. **The failure mode is asymmetric, which is why it earned a rule.** A patch built on a stale base fails loudly if you are lucky and applies cleanly if you are not — and the clean application is the one that reaches the repository. Earned rule 35: re-stage every base file immediately before generating a diff. Recorded alongside it, because it is the same week and the same shape: a test written in this batch asserted on the substring "support" against a note that quotes a label reading BULLISH VOLUME SUPPORT. Earned rule 30 had been written three days earlier about exactly that, and was violated in a test written to guard its cousin. Rule 37 came from this: knowing a rule and applying it are different acts, and the rules that repeat are candidates for a mechanical check rather than a firmer intention.

#60

Explaining in the Right Place

September 5, 2026

DECISION — ADOPTED

"A long thread pays for its own history on every turn."

Viktor asked Claude to decide and implement how to reduce token usage, on one explicit condition: it must not weaken the engineering. Two changes, decided by Claude and accepted. **First, when a patch ships with a commit message, the chat reply carries only the predictions to check and the commands to run.** What the fix does, why, the weights, the verification and the mistakes made along the way are already in the message he is about to read; saying it twice is pure waste, and the chat copy is the one that does not survive the session. **Second, a thread is closed and a fresh one started once a piece of work has finished** — "Continue Phase 7" is enough to resume, and the handover check of Entry #56 is what makes that safe rather than reckless. What was explicitly **not** cut: negative controls, golden checks, equivalence runs, and naming every consequence of a change before it is made. Those are the reason the patches have held, and they were the condition. This must not become a reason to explain less. It is a reason to explain where the explanation survives.

#61

Unreachable Is Not the Same as Safe

September 5, 2026

FINDINGS — FROM EXECUTION

“Nothing could reach it — which is exactly why two audit passes left it there.”

Patch L, finding 6. Two fabrication constants survived in `engine_core.py` after item 9a removed their siblings from `indicators.py` and Finding 3 removed them from `entry_model.py`: a substitute trend dict `{"trend_health": 50.0, ...}`, and `atr_val = ... else current_price * 0.02`. Both were unreachable. `compute_trend_health` is total, so nothing could reach the first handler; section 2 halts when ATR is absent, so nothing could reach the second. **The trend one was worse than a fabrication, and only the negative control showed it.** Its substitute dict has no `trend_direction_sign`, which the next stage reads by subscript *outside* the try. Run against pre-fix code, a broken trend contract does not degrade the run — it dies with `KeyError: 'trend_direction_sign'`, reported as neither a trend failure nor a bias failure. Found by running the test, not by reading the code.

#62

A Fallback That Claimed an Equivalence It Did Not Have

September 5, 2026

DECISION — ADOPTED

“A test had asserted, on that claim, that these paths cost nothing.”

Patch L, finding 5. The RSI and ATR fallbacks smoothed with a simple moving average; `pandas_ta` smooths both with Wilder's RMA. So the sentence in `add_technical_indicators`' own docstring — “recomputes the same quantity by another route” — was false, and it was the stated grounds on which `test_degraded_state` asserts these paths are **not** degradations. On the pinned fixture, final bar: RSI 84.45 against `pandas_ta`'s 69.14, and an ATR 1.80% low. That test had been passing on a false premise for as long as it existed. **The choice was Claude's, and was made by reading the repository rather than by preference.** Viktor: *“None needs a ruling from me.”* Match the smoothing, or record which path ran? Recording the path leaves `test_degraded_state`'s assertion false and adds a flag to explain why. Matching makes the assertion true. The repository already contained the argument; the job was to find it, not to have an opinion.

#63

The Run That Reported Two Errors and Ran Nothing

September 5, 2026

FINDINGS — FROM EXECUTION

““2 errors” — and the result for the other 251 tests was unobtainable.”

Both of these were found because the verification step ran, not because anyone was looking for them. **The `pandas_ta`-free control could not run at all.** `test_macro_agreement` and `test_volume_agreement` import a pure function whose module chain reaches `pandas_ta`, so both errored at collection; pytest reported two errors and executed nothing. That had been silently true since those files were written, and the degraded-dependency result nobody could obtain was the one the fallback work depended on. `pytest.importorskip` makes it the skip it should always have been — 0 errors, 159 passed, 88 skipped. **And a module-level object defeated a deliberate lazy import.** `SignalRouter.__init__` imports `engine_core` inside the function, which is what lets `live_trading` be imported without `pandas_ta`. A module-scope `LiveTradingSimulator()` called that constructor at import time and undid it. Verified in both directions rather than argued: pre-fix, `import live_trading` raises `ModuleNotFoundError`; post-fix it succeeds.

#64

Two Series Paired by Position, Correct by Accident

September 5, 2026

DECISION — ADOPTED

"The printed label changed in 4 of 31 positions, and the evidence line gave no tell."

Patch O, finding (a). Correlation and beta between AERO and BTC were computed by pairing the two series **by position** after both timestamp indexes had been discarded. In the ordinary case both fetches return the same 450 candles, so positional pairing is timestamp pairing and the code is right by accident — which is why it survived, and why the golden snapshot did not move when it was fixed. On the pinned fixtures the two indexes share all 450 timestamps and old and new agree to the last decimal. **It stops being harmless the moment the series differ by one bar:** a candle closing between two sequential API calls, an exchange gap, a stale feed. Measured on the fixtures by dropping one BTC bar from inside the window, at all 31 positions: correlation moved by a median of 0.105, beta by 0.135, and the printed label changed in 4 of the 31. `n_observations` read 30 either way, so the panel's own evidence line could not reveal it. The join is now on the timestamp index, and the observation count is the count of the aligned pairs.

#65

A Zone Whose Lower Bound Was Above Its Upper, and a Price Nobody Read

September 5, 2026

FINDINGS — FROM EXECUTION

"ENTRY ZONE : \$0.4981 - \$0.4918"

Patch P closed finding (4) and a display defect that was not on any list. **The display defect was found by reading the panel.** `engine_core` assigned `EMA_20` to `zone_lower` and `EMA_50` to `zone_upper` unconditionally, so in an uptrend they printed inverted. `entry_model` swaps them before scoring, so the arithmetic was correct the whole time and only the display was wrong — which is why no test could have caught it. Every test asserted on the numbers. **Finding (4) was the familiar shape:** `close * 0.99 / close * 1.01` for a missing EMA pair, reachable only when an EMA is missing, which a healthy run never is. Five fabricated constants across two files, and the worst was not on the list — a close price that could not be read became `$1.00`, so a run with no data at all reported ACTIVE ENTRY ZONE, 30 of 30, against a price nobody had read. **Three of this project's defects have now been found by running the engine and looking at the output.** It remains the cheapest detector the project has. Left open and written down rather than assumed: whether any other reader of the zone is missing the same swap.

#66

Five Components Summing to 39, Under a Total of 45.18

September 5, 2026

DECISION — ADOPTED

“Not wrong, unexplained — and for a number an operator acts on, that is its own defect.”

Patch Q. The entry-quality panel had been printing a total its own component list did not add up to, recorded on 2 September as “not wrong, unexplained”. An operator meant to act on that number either works the gap out themselves or stops trusting the number, and both are failures. Three causes, none visible anywhere: sub-scores rounded for display while the total was not, three confluence multipliers applied after the sum, and a clip at 100.

And a fourth thing the clip was hiding. The five components add to **102, not 100**, while the docstring directly above the list said 100. With full confluence a perfect setup reaches 118.08 and loses the difference silently — the clip was concealing an arithmetic error in the scoring definition itself, not merely bounding an output. The panel now prints the subtotal out of 102, the multiplier broken into its three factors, and a Clipped line when the clip actually fires. The column adds up.

#67

One Fact Declared in Four Places

September 5, 2026

OBSERVATION — FILED FOR
THE RECORD

“Nine new fields vanished between the engine and the panel, with no error and no failing test.”

Recorded, not fixed. `signal_router.py` rebuilds the entry block **field by field**. Anything the engine adds that its list does not name is dropped before the panel or the decision log ever sees it — silently. Patch Q's nine new reconciliation fields disappeared exactly that way on the first attempt: the lines simply did not appear, and nothing failed. The entry block's shape is now declared in four places — `entry_model`'s return, `engine_core`'s dict, `signal_router`'s rebuild, and `decision_contract`'s TypedDict — four copies of one fact, which is the defect class this project has recorded more often than any other. **Restructuring it is larger than any single finding warranted, so it is written down here instead of being done quietly at the end of a patch about something else.** Earned rule 24 in a different costume: the fix that belongs in its own commit does not get smuggled into one that has already made its claim.

#68

Four for Four: the Fabrications That Survive Audits Are the Unreachable Ones

September 5, 2026

OBSERVATION — FILED FOR
THE RECORD

“Nothing tests a path a healthy run never takes, and nothing else notices it is wrong.”

Item 9a, Finding 3, Finding 6 and Finding 4 were four separate findings across three weeks, and all four removed fabricated constants from code paths that a healthy run never reaches. That is not a coincidence and it is now the strongest generalisation this project has produced. **A constant on a live path is checked by every run and every fixture. A constant on a dead path is checked by nothing** — not by the suite, which cannot enter the state, and not by a reader, who sees a defensive default and moves on. So it survives audit after audit while looking exactly like prudence. Unreachable is not safe: it is one edit to an invariant elsewhere from becoming the live path, and on the day that happens the engine will print a fabricated number with complete confidence and nothing in the record will say where it came from. Finding 6 proved the sharper version — its dead branch would not merely have fabricated, it would have crashed the run with a `KeyError` attributed to neither stage. **The generalisation worth carrying to other projects:** when auditing for invented values, search the unreachable branches first, not last.

#69

A Test Suite That Could Reach the Internet

September 5, 2026

OBSERVATION — FILED FOR
THE RECORD

“Two rejected `CONNECT`s to an exchange, from a run in which every route was pointed at a dead port.”

During a full-suite run the sandbox's egress proxy logged two rejected `CONNECT` attempts to `api.mexc.com:443`. Every test that routes or fetches sets `base_url` to a dead local port first, and each of them was checked; the source of the two attempts was not identified. **Recorded rather than chased, and recorded rather than dismissed.** The attempts were blocked, so nothing reached the network and no result depended on what an exchange returned — this run. A suite that *can* reach the internet is a suite whose result depends on the network on some other run, on a machine with no proxy in front of it, and the failure would arrive as an unexplained flake rather than as an error naming its cause. It is written here so that the next unexplained flake has somewhere to start, which is what this log is for.

“Every queued fix has landed. That is not the same as a clean report.”

Between 3 and 5 September the suite went from 226 tests to **319 passing**, and the “What is left to fix” table in PHASE7_NEXT.md reached zero rows — three rulings built, and all seven queued findings closed, across patches I, J, K, L, M, O, P and Q. **The release gate stays shut, and the distinction is the entry.** Unresolved in this project means fixed *and* re-audited; every one of these fixes was written by the party under audit, and a fix nobody independent has seen is a claim, not a verdict. What stands between here and the re-audit is the five-item pre-audit checklist, and the first two items on it are the two that were skipped last round: rebuild the package against current code so the auditor grades what exists, and commit everything before building it, so the packaged files match a commit rather than a working tree. **A completed audit is not an open gate either.** If the report returns a Critical, the gate stays shut and the cycle repeats. The milestone is a clean report, not a finished one.

“Three Qwen calls, then Kimi. The record has the last two attempts the wrong way round.”

Earned rule 29 says to reconstruct a factual record from the provider's billing log rather than from memory, because every provider bills per request and the bill cannot be mistaken about what was called. It was written on 2 September after the independence ledger turned out to have been written from recollection. Applied a second time, to the question of which model ran which attempt at round 2, it found three things. **One: the order is wrong.** The log for 2 September reads Qwen 10:53, Qwen 15:38, Qwen 15:42, Kimi K3 15:58. Entry #50 records attempt 2 as Kimi and attempt 3 as Qwen; Kimi ran last, after all three Qwen calls, so the eleven observations of Entry #51 came from a Qwen run that preceded it. Viktor's recollection — “we ran Qwen three times” — was exactly right. **Two: Claude's own reasoning was wrong, and in a way worth naming.** Having found that those transcripts cite a superseded revision of the reviewer instruction, Claude offered the files' modification times as evidence that the correct package existed before the run and was not used. A modification time records when a transcript was *saved*, not when a model *ran*. Ordering of writes was presented as ordering of events; the claim was withdrawn. **What the clock does establish, once both sets of times sit on one axis, is stronger:** the UPLOAD_THESE folder was written at 16:47:56 and every one of the four calls predates it. **No attempt at this round has ever received it.** Nobody uploaded the wrong folder — the folder did not yet exist, and the eleven observations came from a run holding rev 2 and a file named for the Constitution whose contents were the engine source bundle. Two distinct payloads confirm the shape: 213,581 input tokens in the morning, then 196,283 for all three afternoon calls across two providers — one upload, retried. **Three: a model named in this log was never called.** Entry #22 records Run 1 as a blind review by DeepSeek V4 *Pro*. The bill says DeepSeek V4 **Flash 0731**. Independence is tracked at the laboratory, so the ledger's verdict is unchanged and DeepSeek is spent either way — but the record names a checkpoint that was never billed, which is the same error rule 29 was written about. **Left open rather than resolved:** the log shows nine Luna Pro calls between 27 August and 3 September, considerably more than this document's account of Luna Pro accounts for, and three of them on 3 September map to nothing recorded anywhere. Kimi K3's 2 September call returned 36,085 output tokens, which is not the truncated stub the record describes when it leaves Kimi's spent status undecided. Both are noted here so the next reader starts from the discrepancy rather than rediscovering it. **The generalisation, now paid for twice:** where a project keeps a record of what happened, ask whether some system already recorded it as a side effect of doing its own job, and prefer that system. This project's account of one failure was built first from a truncated screenshot, then from file timestamps, and only the bill settled it.

“Seven of the forty-four rules carry their previous verdict in a code comment.”

The round-3 package was built and then **searched** before it was sent — the step that the previous two rounds did not take. Section 2 of the reviewer instruction told the auditor, in as many words, to stop if it found a named prior auditor, a count of Compliant and Non-compliant items, or a list of Critical findings *in either of the two code bundles*. The bundles contain all three. Items 3 and 6 described as rated Critical; Item 18 as kept Compliant; Item 16 as having gone Non-compliant; Tier 3 items 3 and 4 as currently Non-compliant; Tier 4 item 2 as rated Compliant with the previous auditor's reasoning quoted. Three separate counts of Criticals. Luna Pro and GLM named. **A reviewer obeying the instruction would have refused to grade, and that would have been the fourth attempt at this round to end without a report and the third in a row killed by the package rather than by the code.** The contradiction had been half-seen already: Rev 3 records scoping this same stop condition to resolve a clash with Section 9, and that scoping drew the line at *the instruction may tell you counts, the artifact must not tell you content*. Nobody then checked whether the artifact told content. **Viktor's ruling:** scope the condition to the Constitution file, count and disclose exactly what the bundles hold, and ask the reviewer to grade all forty-four and say in Part 6 — “I don't know” included — whether a comment moved it. The comments are not stripped: Section 4's argument stands, that removing them is the audited party editing its own evidence and destroying what Item 8 exists to test. **And separately, the structural fix, deferred on purpose:** the audit narrative moves out of code comments and into this document, but not inside the audit it would affect — doing it now would hand the auditor a codebase edited to look better for its grader. **The reason this matters more than a fixed contradiction:** a reviewer that stops tells us loudly. The failure one step over is silent — a reviewer that does not stop, reads “Item 18, kept Compliant”, and grades Item 18 Compliant. Nothing in its report would reveal that, and the report would look entirely ordinary. That is earned rule 28's test applied to the audit process itself, and it is why the same instruction now asks the reviewer to search its own package before grading.

#73

The Disclosure That Counted Wrong

September 5, 2026

PROCESS — RECORDED

“Re-reading the instruction could not have found this. The error was in the counting.”

Entry #72 records the ruling that the code bundles' prior-audit content be counted and disclosed rather than made a reason to stop. Rev 4 carried that count. The count was wrong in three places, and wrong in the direction that matters: it under-stated what the package leaks. **Five** distinct Critical-count phrases are in the bundles, not three — “the first Critical” and “the last Critical” were missed. **Three** AI parties are named, not two: Claude appears in both bundles, twice as having been overruled by Viktor. And the disclosure covered only the two code bundles, while a third file in the upload folder, `version_control_history.md`, lists every file each commit touched and so names the previous audit report and the four prior transcripts by filename. **How it was found:** by grepping the built package, which is the pre-send step Entry #72's ruling created. Re-reading the instruction could not have found it — the prose was internally consistent and the defect was arithmetic. **What was checked and came back clean in the same pass:** the Constitution audit copy is 29 pages with no Version History row recording an outcome, every Compliant / Non-compliant in it being rubric text; and all 71 SHA-256 values in the manifest were recomputed from the bundle bytes and matched, so the manifest is not a claim taken on trust. **Why this is an entry rather than a fix:** the step was created one day and earned its keep the first time it ran, against the document written by the party that created it.

#74

A Number Read Off a Generated File

September 5, 2026

OBSERVATION — FILED FOR
THE RECORD

“117 was true when it was generated and stale when it was used.”

The commit that landed Entry #73's fix predicted the repository's commit count would go 117 to 118. It went to 119. The 117 came from the `version_control_history.md` generated three hours earlier; a commit had landed in between. **The shape, for the third time in this project:** a figure taken from an artifact that recorded it correctly at the moment it was written, and treated as current afterwards. Entry #71 withdrew an inference from file modification times on the same grounds. Here the artifact was one this project generates on purpose, which makes it more trustworthy-looking and no more current. **Why it was caught:** because the prediction was stated as a number before the command was run. An unpredicted line can only stop the work if the prediction was complete enough to be contradicted.

#75

The Transcripts Were Kimi's All Along

September 5, 2026

OBSERVATION — FILED FOR
THE RECORD

"It said it was not Qwen, in the section written to make it say so, and the sentence never reached the repository."

The largest correction this project has had to make. The four files in the repository root named `qwen_reasoning_1.txt` through `_4.txt` are a **Kimi K3** transcript. Verified rather than inferred: with whitespace stripped, the repository copy and the copy saved from the Kimi room of 2 September are identical for **57,631 characters from the first byte**. Two models do not coincide on 57,000 characters. So the eleven prior observations, and the Part 8 document built from them, are Kimi's work. **Why it went unnoticed for three days:** the repository copy is lossy, roughly 60,000 characters shorter than the room copy, and one of the passages in the gap is the model identifying itself — "the instruction says I'm Qwen3.8-Max — I'm not; I'm Kimi (Moonshot AI) ... if the ledger records Qwen3.8-Max but the actual reviewer is Kimi, the ledger is wrong." Section 5 of the instruction was written to elicit exactly that sentence. It was produced, and it did not survive being saved. **What inherited the wrong name:** Entry #50, the attempts table, the Part 8 document, the correction made from the bill in Entry #71, and a ruling made and withdrawn the same day. **The instructive part is Entry #71's own reasoning.** It argued from call timestamps — Qwen at 15:42 before Kimi at 15:58 — to the conclusion that the eleven came from Qwen. The inference was valid and the conclusion was false, because the bill records *what was called* and not *which output was saved under which name*. Earned rule 29 says prefer the system that recorded the fact as a side effect of doing its own job. Applied correctly, that system here is the transcript, not the invoice — and it had been sitting in the repository under the wrong name the entire time.

#76

An Accident That Produced a Report

September 5, 2026

AUDIT RUN D — RECORDED

"The Chatroom tab was left on Auto Router, and the round-3 package went out."

The router chose `z-ai/glm-5.3-flash-20260826`, served by BaseTen: 251,148 input against 11,475 output, \$0.0434, 102 seconds. It returned a complete Parts 1-6 — **26 Compliant, 13 Partially, 1 Not verifiable, 0 Non-compliant, eleven findings all rated Minor** — so the round happened whether or not it was intended, and Viktor ruled it **is** round 3 rather than discarding it. Its model-identity check fired correctly: its first line states it is not Qwen3.8-Max. **Three findings were verified against source and all three confirmed**, and one is under-scoped by the report in a way that matters. F-7, the permissive `risk_valid` default, was filed against the simulated-order module and rated Minor on the grounds that nothing there authorises a trade. The same default sits in `decision_model._determine_final_action` — the authorization gate — where a risk block missing that key reads as risk-passed. The severity argument does not reach it. **What the run measures beyond its findings:** 11,475 output tokens with no reasoning spend, against Kimi's 36,085 tokens of reasoning on a smaller package that never reached a report at all. Depth of read is visible in the bill. The strongest part of the report is its Part 6b, which names three places where a comment led its reasoning before the code confirmed it, unprompted beyond the instruction asking.

“Second family the ledger had wrong, and the reviewer that ran was not independent.”

The full provider activity export, 723 requests. **All 723 are variant=standard** — no free endpoint, ever — so the no-training conclusion the 2 September independence ruling rests on holds unchanged. Two things it corrected. **Z.ai is not clean.** z-ai/glm-5.3 ran on 28 August: 104,394 in, 55,179 out, a full substantive session on this project eleven days before GLM 5.3 Flash audited it. Under the lab-not-checkpoint rule the reviewer that produced Entry #76's report was **not independent**. Its findings stand on their own evidence; the caveat travels with the comparison permanently. That is the second family this ledger had listed clean and wrong, after Mistral in Entry #48. **Kimi clears, on the check that mattered.** Every Aider call in the history — 681 of them — carries the Aider app key, across exactly the five models the corrected record names. All fourteen Kimi calls are Chatroom. Kimi never worked on the codebase, so its exposure is session-level, which the 2 September ruling clears. But it is far larger than the record said: eight substantive reads on 27 August of 93K-110K input each, roughly 100,000 tokens of output about this engine, plus 36,085 on 2 September. “A Step 3 attempt that truncated” described one row of thirteen. **And outputs this project paid for and never saved:** four completed Kimi runs on 27 August totalling 60,537 tokens; Qwen's 2 September call returning 11,964 tokens, of which no transcript exists under any name, so what Qwen actually said is lost; three Luna Pro calls totalling 393,663; and one row with no model recorded at all.

“A check that looks for untracked files that matter cannot see files the repository was told to ignore.”

Both evidence sets — the Kimi transcript and the GLM report — were first saved under docs/audit_package/round*/ , which .gitignore excludes except for the manifest. That rule is right for generated package bytes and wrong for reviewer responses, which are primary evidence and cannot be regenerated. **The part worth carrying to other projects:** ignored files do not appear in git status --short, so the standing end-of-session handover check could not have caught it. The check asks whether untracked files that matter are present, and the one class of file most likely to be lost is the class the check is structurally unable to see. **Structural fix, not an instruction to be careful:** reviewer responses now live in docs/audit_reports/<round>_<model>_<date>/ , which is tracked by default. No exception list to maintain and no judgment required at the moment of saving, because the safe location is the only one on offer. round*/ stays ignored for the generated packages the rule was written for.

#79

Rev 5, and Part 8 Removed

September 5, 2026

DECISION — ADOPTED

"The instruction had been telling the reviewer who it was, and getting it wrong."

Two changes and nothing else. **Section 5 no longer asserts the reviewer's identity.** Rev 4 told the reader it was Qwen3.8-Max and that one clean-list slot had been spent on an earlier attempt at this same audit — a sentence which, sent to Kimi, describes Kimi as somebody else. It now asks the reviewer to state its own identity and says what is known about who ran what. **Part 8 and Section 13 are removed entirely.** Their comparison document was the transcript of Entry #75, so asking Kimi to reconcile its findings against it would have been a model grading its own earlier reasoning, and asking anyone else to reconcile against a document this project had misattributed was not worth doing either. The comparison it existed to produce is made outside the instruction, between two reports, by Viktor. **What did not move:** what is graded, how it is graded, the severity rubric, Section 7's checks, and everything handed over besides Part 8. GLM ran against rev 4, so the revision history names the delta and the comparison carries its own caveat.

#80

Three PDFs in a Downloads Folder

September 5, 2026

PROCESS — RECORDED

"Two were preserved. One held a paragraph that existed nowhere else."

Viktor found three documents in his Downloads folder and asked whether he could delete them. Each was checked against the repository rather than recognised by name. **The Gemini review of the Constitution** hashes to exactly the SHA-256 recorded in docs/constitution_reviews/gemini_2026-09-03.md, and its extracted text matches that file word for word, 390 words each. Fully preserved; deleted. **The independence ruling** was preserved except for one section: the argument that the ruling turned on *mechanism* rather than on the value of the model it freed — that had a mechanism been found, Kimi would have stayed out regardless of rank, and that a rule reading "an exception was made because the model was valuable" is one anyone can invoke later to justify anything, where a rule about mechanism can be attacked and overturned on its merits. That paragraph was recovered into docs/PHASE7_NEXT.md before the file was deleted. **The third document** is Claude's candid criticism of Viktor's own role in the project. Viktor ruled it does not enter the repository, because the repository is public. It is kept outside version control by decision, and that is recorded here so a later reader knows it exists and why it is absent rather than assuming it was lost. **The generalisation:** the conclusion had survived in the record and the reasoning that makes it challengeable had not, which is the failure this document exists to prevent.

#81

"Could Not Run It" Was Two Claims, and Only One Was True

September 5, 2026

OBSERVATION — FILED FOR
THE RECORD

"The impossible case was stated where the merely-undone one belonged."

The patch pointing the package builder at round 4 shipped with a commit message saying the build itself had not been run, and offering a reason: the builder executes the engine and there is no engine in the sandbox. Viktor asked why, and the question found the defect. **Two claims had been folded into one.** True: it cannot be run against *this* repository from the sandbox, because the file bridge copies files and not the .git object store, so the real history never arrives and neither the version-control history nor the commit messages can be generated. Not true: that the builder could not be run at all. A scaffolded repository — two commits, stub modules, the real instruction, a placeholder Constitution — exercises every changed line, costs minutes, and was then done: it produced a round-4 directory with seven files in the upload folder, one in the withheld folder, and no Part 8 document anywhere. **The shape, and it is the same one this project keeps finding:** a conclusion carrying more confidence than the check behind it. Here it took a new costume — a genuine impossibility placed where an undone task belonged, which reads as a limit rather than as a gap and so does not invite the question that would expose it.

#82

Kimi Runs Round 4, and the Reason That Was Rejected

September 5, 2026

DECISION — ADOPTED

"Model quality is a preference. Scarcity of clean reviewers is a fact about the ledger."

Viktor's ruling, argued and then narrowed. **What stands.** Kimi K3 has graded **none of the forty-four rules:** its earlier attempt could not read the standard, refused to grade, and ran Section 7's checks instead — so on the question this round asks, it is blind, and "self-review" overstates what it is. And it is the last high-ranked model clean on both independence lists, so a round that does not spend it spends a worse reviewer or none. **What was rejected.** "Kimi K3 is a very important model to use." That is the shape the 2 September independence ruling warned against in its own words — the paragraph recovered in Entry #80 — and admitting it here would leave a precedent anyone can invoke later. Both readings pick Kimi, which is exactly why the distinction is worth making: the decision is unchanged and the rule it leaves behind is not. **What is disclosed rather than solved.** Eleven of Kimi's own observations are already fixed in this code, with comments describing them. The exposure is not that it repeats a verdict — it has none — but that it may accept a fix to its own diagnosis more readily than a stranger would. Small, real, not removable without spending a different reviewer, and accepted rather than argued away. **A hold that outlives this entry:** the four transcripts keep their wrong filenames until the package has been sent. A rename is a commit, commits appear in the version-control history, and that file ships inside the upload folder — so correcting the name now would put "kimi" in front of Kimi before Section 5 has asked it to identify itself unprompted.

“Completing on the negative control is what makes the positive case mean something.”

Kimi's round-4 report asked for exactly one run, and this is it — made on unmodified code, before any fix, so **Finding 1 is observed rather than argued**. Three offline runs against the committed pinned fixtures. **A:** fixtures exactly as committed, negative control — completed, correlation +0.2326 over 30 observations. **B:** BTCUSDT_4h with every timestamp shifted +2h, zero shared timestamps — `Decision` object construction failed: `float()` argument must be a string or a real number, not 'NoneType', the exact string Kimi named. **C:** BTC's last 60 candles flat, the optional zero-variance variant — the same message, secondary, since a flat series also moves BTC-side ATR and volatility. **What the run found that the report does not contain:** the decision-log record written on the merge failure carries none of the run's lineage or provenance — only an error, symbol, timeframe and log path — where the healthy record beside it carries both, with input hashes and a run hash. Item 6, the Critical keeping the release gate shut, is about traceability; a defect that manufactures untraceable records inside it is not a display concern, and whether it is folded into the Finding 1 fix or filed as its own item is left to Viktor. **Two wrong turns, recorded rather than cleaned up.** The harness's first version ran with `config.LOG_DIR` untouched and wrote three records — two of them errors — into the engine's real permanent log, overwriting the persisted bias state machine with a fixture-derived value; the fix repoints the harness at a temporary directory for its duration, structural rather than a warning. And the corrected harness was staged under the filename of the version it replaced, so the commit shipped the older bytes and failed with `ModuleNotFoundError` — passing a `py_compile` check that tested syntax when the question was identity. Two rules taken from it: a delivered file gets a name unique to its version, and a delivery is verified by reading it back off the device and diffing it, not by compiling the copy that was sent.

“The producer was tested, the consumer was tested, and the seam between them was not.”

Landed against unmodified code, after the Section 11 run above had already reproduced the defect. Suite 319 to 331 passing, 0 failed; the golden snapshot did not move, predicted before the run on the same reasoning as the negative control — the pinned fixtures measure a correlation, so every value the snapshot records goes through the unchanged path. **Two changes, in** `models/signal_router.py`. A new `_optional_number()` helper passes through what the producer wrote instead of defaulting to 0.0: None stays None, a non-finite number becomes None, a measured number is unchanged. **None rather than 0.0 deliberately** — a correlation of 0.00 is a measurement a real pair of independent assets can produce, so substituting it would report a finding the engine does not have, the same defect already removed from three other modules. And the record written when the assembly fails now carries the run's lineage and provenance, closing the Item 6 gap the confirmation run found and neither reviewer filed. Twelve new tests in `tests/test_router_btc_seam.py`, written at the seam itself rather than as one more assertion about either side of it — the pinned fixtures share all 450 timestamps, so nothing in the suite had ever made the producer emit its not-measured shape and handed it to this consumer. **Evidence against pre-fix code:** 8 failed, 2 passed, 2 skipped, all eight behavioural — a `TypeError` from the bare `float()` call and a `KeyError: 'lineage'` on the failed record, no import or collection errors. **Where verification stopped:** built and verified in a Linux sandbox on the pinned dependency versions; `pandas_ta` could not be installed there, so the two end-to-end tests and the golden snapshot were unverified until the Windows run — stated at delivery rather than after the fact. **Not fixed:** the merge's other invented defaults ("NEUTRAL", "NORMAL", 0.0 for trend health) stand on the same “the producer always sets it” argument that was just proven wrong; and a failure inside the merge still cannot degrade rather than halt, because `decision_model.evaluate()` runs before the BTC block is assembled — a real gap in the 29 August ruling, recorded rather than patched in passing.

"The maximum-intensity authorization this engine can issue, on a run where risk was never assessed."

Landed at 26a05dc, pushed. Suite 331 to 363 passing, 0 failed; golden snapshot unmoved, confirmed by diffing the fixture rather than only by the test passing. **The 5 September note under-rated this as a permissive default in three places; reading it found it reachable, not latent.** `signal_router._validate_engine_output` checked that the key "risk" was present and never looked at what was under it, so an empty risk block passed validation and reached a gate whose `.get("risk_valid", True)` supplied an authorization the engine never computed. On the conviction fixture: an empty risk block or `{"risk_valid": "OK"}` both produced AGGRESSIVE LONG. **A validator that checks presence plus a consumer that defaults permissively is the composition**, and neither half looks wrong alone — the same shape as the seam Finding 1 lived in. **What changed:** a new `read_risk_verdict()` in `models/risk_model.py` treats absent, present-as-None and present-as-not-a-boolean as one state, none of them a pass, deliberately strict so an unrecognised shape fails closed. The gate now refuses with its own reason, NO-TRADE (RISK NOT ASSESSED), rather than reusing NO-TRADE (RISK TOO HIGH) — a bare flip of the default to False would have printed "Risk check failed (OK)", the same fabrication with the opposite sign. The contract keeps `risk_valid: bool` rather than becoming `Optional[bool]`, a decision on the record: Optional looks like the honest declaration and is the weaker fix, since the contract test's type check passes any Union and the field would stop being checked at all. 32 new tests. **Evidence against pre-fix code:** 14 failed, 18 passed — 11 behavioural, 3 source-text, 0 import or collection errors, verified by grafting the new function alone into an otherwise untouched file so the pre-fix run cannot collapse into one collection error. 7 of the 18 pre-fix passes are named negative controls. **Verified both ways** before delivery, unlike Finding 1: 363/0 with `pandas_ta`, 252 passed / 100 skipped without it, against a 222/98 baseline. The run owed against live data was discharged two commits later, at 10:04 on 6 September — see Kimi Finding 5 item 6 below. **Not fixed:** the reason string that blames "the risk check" on every NO-TRADE action regardless of which check actually blocked it, pre-existing; the manifest-hash test that fails on any non-Windows checkout, worked around locally rather than fixed; and the same two invented defaults ("NORMAL RISK", 50.0) elsewhere in the same assembly.

“A located level sitting exactly on the current price, on precisely the runs that located nothing.”

Landed at 4678f45, pushed. Suite 363 to 377 passing, 0 failed; golden snapshot unmoved, confirmed by the fixture's absence from `git status`. **The margin was one row.** `_detect_swing_structure` returned `current_price` on both of its not-located paths. $2 * \text{lookback} + 5$ is 21 at the configured lookback of 8; the engine's own frame-length guard rejects below 20 — the two minimums disagree by one, on the wrong side, so a 20-row frame passed validation and got its own price back as a structural level. Reachable by the same presence-only validation GLM F-7 named two commits earlier, and the third time this exact field was patched at a different layer while the producer kept inventing — `entry_mode`'s fallback on 1 September, `panel_render`'s identical default on 2 September, both fixed at the consumer while this stayed open at the source. Both not-located returns are now `float("nan")`, in `structure/structure.py` and in the router's assembly default. Deliberately no `degraded_inputs` entry — this field is read by nothing but the panel and the record, not by entry scoring or the risk stop, so a display-only value moving an authorization would be a worse defect than the one being fixed. 14 new tests. **Evidence against pre-fix code:** 9 failed, 5 passed — 7 behavioural, 2 source-text, 0 import or collection errors, no harness needed this time since the new tests import no new name. **The live run owed for both this fix and GLM F-7 was made on the patched tree before the commit** — AEROUUSD 4h, 6 September 10:04: SWING STRUCT : \$0.5063 against a current price of \$0.5437, a located level that is not the current price, and NO-TRADE (RISK TOO HIGH) naming a real reason, “Risk regime classified as EXTREME RISK,” rather than the F-7 fabrication. The NaN paths themselves are covered by the suite only — a 300-bar live frame always locates a swing. **Not fixed:** the router assembly's other invented defaults (`hvn/lvn` to 0.0, regime to NEUTRAL, sequence to NONE, volume sentiment to NEUTRAL VOLUME), named here so the sweep has them; and the same field still raises `TypeError` if present as None, Kimi Finding 1's exact shape, latent only because the producer never writes it — the wrong argument for the third time in three days. **What went wrong delivering it:** the command list was reissued with the same steps renumbered when the live run was inserted, and one `git add` was lost between the two numberings — the commit went in without its test file, caught by reading “2 files changed” rather than by any check, and amended before pushing.

"The number is behaving as designed; the sentence is false."

Landed at 1045748, pushed. Suite 377 to 400 passing, 0 failed; the golden snapshot moved, predicted in exactly five fields before the run and moved in exactly those five. **The list could not be fixed by extending the list.** Kimi named seven decision-affecting settings missing from the fingerprint; six of the seven are not list omissions but things a list of that shape cannot hold at all — class attributes, local variables, bare literals in comparisons and call sites. **What changed:** a new `core/code_fingerprint.py` reads every `.py` file outside `tests/` and `docs/`, parses it, strips docstrings, and hashes the sorted per-file digests — no list of constants or modules, so a constant written tomorrow is covered the moment it exists. The parse tree rather than the text, deliberately: this repository rewrites its comments in nearly every commit. `code_hash` is deliberately **not** folded into `run_hash` — that would move the archive filename and force a re-baseline on every commit touching any engine file, spending the check that has caught more defects here than any other. Claude's call under the delegation, reversible in one line. 23 new tests. **Evidence against pre-fix code:** 21 failed, 2 passed on the raw count; the honest number is 7 behavioural failures, the rest proving only that the new module is new. Three negative controls: a comment does not move the hash, neither does a docstring, and CRLF and LF source hash identically. **Figures:** 400/0 with `pandas_ta`; 286 passed / 103 skipped without it; `run_tests.py` 333 passed, 0 failed, 29 errors unchanged. **The live run made on the patched tree, AEROUSD 4h, 6 September 10:51, matched the sandbox hash bit for bit** — 44e085cf... on Viktor's Windows machine under 3.12.10 against a CRLF checkout, and the identical value in a Linux sandbox under 3.12.3 against LF, confirming platform-, line-ending- and 3.12-minor-version-independence as observed rather than assumed. **The same run found what neither audit report contains.** The BTC block raised confidence 73.86 to 75.80, explaining it as "BTC is also neutral, agreeing with AERO's own bias" over a BULLISH CONFIRMED AERO — a neutral value cannot be "also" what a bullish one is. Verified at `decision_model.py` 704-768: agreement is decided by the sign of `btc_score`, and the sentence's words come from the label — the two disagree whenever the score is non-zero but below the labelling threshold. **Why this is not simply Kimi Finding 2:** Finding 2's exhibit is entirely about `abs(correlation)` discarding the sign, at -0.90; here the correlation was +0.51 and the sign was never in question. A second, independent route to the same class of false sentence, filed as input to decision 3 rather than fixed, because it is Viktor's ruling and it changes what he is ruling on. **Not fixed:** `engine_version` as a static string; `code_hash`'s interpreter dependence, recorded by carrying the Python version beside it in the archive; the unnamed literals the finding named, covered by the hash but absent from the readable snapshot; and a rerun on identical inputs after a code change still overwriting the earlier archive's `meta.code`, which follows directly from the `run_hash` decision.

“A pristine checkout, seeded with marker files, came out of a green suite with every marker gone.”

Landed at fc35a2f, pushed. Suite 400 to 406 passing, 0 failed; no engine source file touched, so no live run was owed and code_hash did not move. **Found by running the suite against a clean clone and watching the filesystem, not by reading** — every individual line involved is unremarkable, and the damage exists only at runtime. Two independent reviewers read these files; one filed a finding about the very test that does the worst of it, and the mechanism survived both. **Four routes into the engine's own permanent record.** Twenty-eight tests across eleven files ran the engine against the real log directory, appending nine real decision records plus one error record on every suite run — including the two files whose entire subject is what the permanent record contains. A twelfth test deleted the directory outright, in a finally block after an archive-path spelling test. A thirteenth recreated it as a side effect of purging and re-importing every module. A fourteenth wrote into it through the capitalised spelling, Logs/ — one directory on Windows, two on Linux, exactly the defect a hardcoded-path guard exists to catch, in a file that guard does not scan because it checks modules and these lines are in a test. **The first patch was wrong, and the way it was wrong is the lesson.** v1 was verified green in the Linux sandbox, applied, and failed on Viktor's machine at the first run: the new guard watched the directory named exactly Logs, so route four's Logs was invisible to it on Linux and immediate on Windows — **a guard whose result depends on which platform it ran on is not a guard.** A second failure was a negative control hardcoding a 4-byte file size for "one\n"; Windows text mode writes 5. Both recorded rather than quietly fixed, because the first one changes how every sandbox result should be read from here: evidence about Linux until something makes it evidence about Windows. **What was lost:** the live log held eleven records — ten from a single suite run at 08:51 UTC, identical run hash, identical bias score, and one error record with no code hash — plus the real 10:04 AEROUSDT run and the three synthetic Section 11 records this file had just finished describing, both gone with the directory the rmtree deleted. Nothing in the repository can rebuild them; Logs/ is gitignored, so nothing held a second copy. **What changed:** every writer redirected to a temporary directory at import time rather than as a pytest fixture, so both test runners pick it up; the path-spelling and golden-snapshot tests keep their real relative literals so the coverage that caught the Windows failure is not retired; the hardcoded-path guard now scans case-insensitively. Six new tests. **Figures:** 406/0 with pandas_ta, 292 passed / 103 skipped without it, 339 passed / 0 failed / 29 errors unchanged under run_tests.py. Reproduced on a pristine checkout seeded with four marker files: pre-fix, 400 tests passed and every marker was gone; post-fix, through all three configurations, the seeded log's SHA-256 came through unchanged. **Not fixed:** the lost records are not recoverable; Logs/ remains gitignored, part of decision 7; and the hardcoded-path guard still scans engine modules only, not test files — recorded as a candidate rather than adopted, since this repository has been burned by source-text guards before and the new behavioural guard already covers the same class on both platforms.

“The last time this document set was read end to end, it found a whole Critical that had never been entered into any roadmap.”

Prompted by nothing in particular, which is the point. **The audit-package record is intact.** All three manifests verify byte-for-byte against their recorded build commits — round 2 60/60, round 3 71/71, round 4 71/71 — against an LF checkout, because `build_audit_package.py` normalises line endings on read and write while `tests/test_pinned_source.py` hashes raw bytes without normalising, two hashing schemes over the same repository with opposite line-ending behaviour, filed with the LF-checkout test hygiene items. One suspicion, checked and wrong: two 30 August bundle files read like a stale-bundle trap when cited by name in the round-4 instruction, but the builder documents them as the deliberately kept round-1 record, with fresh bundles copied in separately. **Two structural problems.** No build script could write into this repository — nine of ten `docs/build/` scripts hardcoded `/tmp/outputs/`, which is why the Engineering Notes gap grew at every handover; this audit's own count of “all ten” was itself off by one, corrected and fixed the same day, see below. And `README.md` and `docs/build/README.md` both cite `docs/Phase7_Audit_Findings_Complete.pdf`, which does not exist anywhere in the repository and cannot be produced from here — its builder also reads source material that was never committed. Every other dangling reference found is benign: gitignored build outputs, or already recorded elsewhere in this project's own documents. **What this audit did not do:** the eight PDFs — Constitution, Roadmap, Engineering Notes, Remediation Plan, the Documentation and Change Log Standard, the Tier 0 Companion, the Credential Security Protocol, the Audit Execution Instructions — were not opened. The mechanical half is done; the semantic half, reading each against the current state of the project, is not, and is where a second Finding 3 would be.

#90

The Build Scripts Can Write Into the Repository

September 6, 2026

DECISION — ADOPTED

“The fix did not need designing; it needed extracting from the one script that already had it.”

Suite 406 to 412 passing, 0 failed. No engine source file touched: `code_hash` unchanged at 44e085cf..., computed on both trees rather than assumed, golden snapshot unmoved, no live run owed. **The count from the document audit above was off by one.** Nine of the ten `docs/build/build_*.py` scripts hardcoded `/tmp/outputs/`, not ten — `build_audit_package.py` already resolved its output from `__file__`, which is why its manifests verify byte-for-byte while every PDF beside them was carried across by hand. Caught by recounting the evidence, not by review, the same shape as a suite-log patch that once said “twelve files” when eleven wrote and the twelfth deleted. **What changed:** a new `docs/build/_output.py` resolves `REPO_ROOT` from `__file__` and rejects any output path that is not a bare filename — absolute, containing a separator, or carrying a drive letter — so a build script no longer has a path of its own to get wrong. The nine import it; seven now write their PDF into `docs/` directly. The other two, `build_findings_bundle.py` and `build_remediation_plan.py`, also read source material that was never committed — they now exit 2 and say so before doing any work, instead of a bare `FileNotFoundError` naming a directory that does not exist on Viktor's machine. Six new tests. **Evidence against pre-fix code:** 5 failed, 1 passed — two behavioural (the scanner finding eleven absolute assignments across nine files, none from a shared resolver), three proving only that the new module is new, one negative control. **Verified rather than asserted:** all nine scripts produce the same output path run from the repository root or from inside `docs/build/`; the rebuilt Engineering Notes and Roadmap extract text identical to the committed PDFs, 63 and 15 pages, same SHA-256 over the extracted text — a rebuild is the same document, not a new one. **Which platform this is evidence about:** Linux, and the filename check is platform-neutral by construction — it reads the path string itself rather than asking `os.path.isabs` alone, which answers differently for a Windows-style path on the two platforms. Not verified on Windows: that reportlab renders these documents there — reportlab was installed on Viktor's machine the same day, and the first Windows evidence will be the first Notes rebuild. **Not fixed by this patch itself:** it makes publishing the eight entries this document now carries one command; writing their content — this row among them — was a separate, later step. Both undocumentable PDFs still cannot be built from here. Both README files still cite the audit-findings PDF that cannot be produced. `reportlab` stays out of `requirements.txt` deliberately, with its own install line in `docs/build/README.md`.

#91

Fourteen-Item Latent/Minor Sweep Closed

September 7-8, 2026

MILESTONE — RECORDED

“Three classes of latent defect closed in three commits without moving the golden object through the sweep itself.”

Grok session, Viktor numbered-CMD workflow. Items 1-5 at 22afea2 (invented defaults to NaN/UNKNOWN). Items 6-8 at 4d56f2a (dead code, `exit_watch` zero, `test_live`). Items 9-14 at c3b0d43 (banner, comment, BTC degradation, snapshot shape, NO-TRADE blame, structure read path). Suite 345/0/29; golden 8/8 through the sweep. Grok is not independent for re-audit of this code. Release gate not opened.

#92 **Kimi Finding 2 Fixed**
September 9, 2026

FINDINGS — FIXED

"Sub-threshold BTC score could print neutral agreeing; abs(correlation) discarded sign."

Viktor ruled fix under decision 3. Agreement direction uses RAW_BIAS_THRESHOLD. Adjustment uses signed correlation. Commit 5e2e9f3; docs 179f4f9. Golden re-baselined on btc_context only (78.14 to 79.26). Suite 345/0/29.

#93 **code_hash Excludes Claude outputs/**
September 9, 2026

PROCESS — RECORDED

"Backup .py copies under Claude outputs/ moved code_hash with no engine change."

Excluded Claude outputs from EXCLUDED_DIR_NAMES. Commit ef00765. Suite 345/0/29.

#94 **qwen_reasoning_*.txt Transcripts Relocated Into docs/**
September 11, 2026

PROCESS — RECORDED

"Four transcripts had sat in the repo root since round 2, under the hold recorded at Entry #93."

Commit bb6d222. The four qwen_reasoning_*.txt files move from the repository root into docs/, alongside the rest of the audit material — a rename, discharging that hold. No engine module touched; code_hash unaffected.

#95 **This Document's Own Body Reconciled With 76380ec, a Second Time**
September 11, 2026

PROCESS — RECORDED

"The 7-9 September Grok session updated this file's head block and builder but left two "Open" items describing already-landed work as still open."

Commit 4705d03. "Open — decisions" #3 and "Open — work" #2 described the fourteen-item sweep and Kimi Finding 2 (Entries #91, #92) as unresolved after both had already landed. Corrected in the same shape as the first reconciliation against 76380ec. Docs-only and code_hash-neutral by directory exclusion.

#96 **Phase7_Sweep_Session_Commands.pdf Committed**
September 11, 2026

PROCESS — RECORDED

"A generated PDF had sat untracked on Viktor's machine rather than in the repository."

Commit e75f7be. docs/Phase7_Sweep_Session_Commands.pdf committed as-is. No engine module touched.

#97

**Item 14 / Decision 2 Closed: Risk Regime Now Reads ADX,
Not trend_health**

September 11, 2026

DECISION — ADOPTED

"classify_risk_regime shared its largest input (0.30 of bias_score) with the conviction number it exists to check independently."

Commit 0c7dec5. classify_risk_regime now reads ADX instead of trend_health; REGIME_LOW_TREND_HEALTH/REGIME_HIGH_TREND_HEALTH renamed to REGIME_CHOP_ADX = 20.0 / REGIME_STRONG_ADX = 25.0, values already load-bearing in trend_health.py for the same two market states. code_hash moved and was confirmed on three trees, 47bc557119bb0350... to 2741062fae070f61..., matching Viktor's own Windows value. Suite 412 to 421 passed under pytest (0 failed both sides), 345 to 354 under run_tests.py, watched 29-error count unmoved. Golden re-baselined at six sites: four predicted, two not — the archive filename, built from run_hash, should have followed the run_hash prediction mechanically and did not, recorded as an incomplete prediction rather than a defect. No decision field moved on the pinned fixture, which exits before classify_risk_regime is reached; a live run the same day (AEROUSDT 4h, 22:50) exercised neither new branch, for two independent reasons — recorded as that, not as validation. Found by reading the panel and left unfixed: RISK REGIME prints without the ADX number now deciding it, the same transparency-regression class as GLM F-5.

#98

GPT-6 Astra Ruled for Round 5

September 11-12, 2026

DECISION — ADOPTED

"Two of Luna Pro's OpenRouter sessions had been exposed to prior audit context."

Commit 4d99cdb. Viktor's billing-export check cleared both sessions as variant=standard, no training routing. send_audit_round.py and build_audit_package.py repointed to GPT-6 Astra; Part 7 folded into the single upload set per ruling 3.

#99

The 2 September Luna Pro Hedge Closed for Good

September 11-12, 2026

PROCESS — RECORDED

"Luna Pro's independence had stood on "probably clean" since the 2 September ledger correction."

Commit de7d135. All nine Luna Pro OpenRouter calls checked individually: all standard tier, "No data training" on each. Evidence upgraded from probably clean to provably.

#100

Round 5 Sent to GPT-6 Astra: Three New Criticals

September 11-12, 2026

EXTERNAL ASSESSMENT —
GPT-6 ASTRA

"None of the three findings returned had been on record before."

Commit a12640a. \$12.21636, 435,612 prompt tokens, confirmed twice over. F1 (volatility never reaching stop construction), F2 (entry quality scored for the wrong direction), F3 (a stale decision-bar VWMA presented as current) — all three independently reproduced against live code this session. Viktor RULED to fix all three; no limitation accepted for any.

#101

Round 5 F1 Fixed — Volatility Reaches Stop Construction

September 11-12, 2026

FINDINGS — FIXED

"The production call to `calculate_stop_targets()` omitted `volatility_state`, defaulting to `NORMAL` regardless of measured conditions."

Commit 83e334d. Now passes the same `volatility_mode` its sibling call already receives. Golden snapshot unmoved — the pinned fixture's own volatility state produces the same multiplier the bug's default did. `code_hash` e774b93d... to f8e4cc7d....

#102

Round 5 F2 Fixed — Entry Quality Scored for the Right Direction

September 11-12, 2026

FINDINGS — FIXED

"`eq_trade_direction` was derived from `short_signal`, which zeroes for reasons unrelated to direction, letting a fully bearish setup score as a long entry."

Commit 3f92230. Now derived from `raw_bias` directly, the same signal `decision_model.py`'s own action selection reads — the identical ruling already made for that module on 2 September, one caller downstream. `code_hash` f8e4cc7d... to 8b0e97fd....

#103

Round 5 F3 Fixed, Closing Round 5 — VWMA's Trailing-Edge Forward-Fill

September 11-12, 2026

FINDINGS — FIXED

"VWMA's own forward-fill had no trailing-edge exception, so a decision-bar miss was silently carried forward from the previous bar."

Commit 7fcf646. Fixed at the source; VWMA added to the `critical_indicators` sweep, with three new regression tests. Accepted consequence: some currently-clean runs will now show `DEGRADED`, naming VWMA. `code_hash` 8b0e97fd... to 00c8d4ea..., confirmed on independently built trees at each of the three fixes (base e774b93d... through this one). Across F1/F2/F3: suite 429 to 432 passed under `pytest`, 362 to 365 under `run_tests.py`, watched error count unmoved; golden snapshot unmoved throughout, confirmed by running it each time. Register unchanged at 21 / 7 / 10 / 6. The release gate remains shut: fixed is not re-audited, and unresolved means both.

#104

Round 6 (Meta Muse Spark 1.3) Sent and Graded: Release Gate Holds

September 13, 2026

EXTERNAL ASSESSMENT —
META MUSE SPARK 1.3*"No Critical Tier-1 finding returned."*

Report at docs/audit_reports/round6_muse-spark-1.3_2026-09-13/report.md. Second send attempt succeeded after the first hit OpenRouter's account-level 18+ gate (Viktor resolved it; no charge incurred). finish_reason=stop, provider Meta, 470,488 prompt / 10,765 completion tokens, cost \$0.63386125.

#105

Round 6 F1 Fixed — Six Bare Literals Promoted to Named Constants

September 13, 2026

FINDINGS — FIXED

"T2-4 Explicit Configuration, Moderate. Viktor ruled fix all three of Round 6's findings, narrower than the report's own F1 Location section."

Commit 447966b. Six bare literals/locals promoted to named, fingerprinted constants: CONFLUENCE_BOOST_MULT/CONFLUENCE_PENALTY_MULT, SPIKE_RATIO, CORRELATION_WINDOW, StructureEngine.REGIME_HYSTERESIS_THRESHOLD. Values unchanged. decision_model.py's trend-health bands and the panel_render.py NaN gap, both named in the report's own F1 Location section, left open — closed later this session as F4 and F5.

#106

Round 6 F2 Fixed — Duplicate confidence_score Meaning Removed

September 13, 2026

FINDINGS — FIXED

"Item 10 Consistent Semantics, Minor. engine_core.py's raw risk dict duplicated confidence_score under a different meaning than signal_router.py gives it."

Commit 3c7e9ae. Dropped rather than renamed — the value survives intact at trend.trend_health, a few lines below in the same return object.

#107

Round 6 F3 Fixed — Six Absent Measurements No Longer Default to 0.0

September 13, 2026

FINDINGS — FIXED

"Item 13 Fail Safely / Item 8 Epistemic Honesty, Minor, unreachable on the live path today."

Commit 9b34163. signal_router.py's _build_decision_object no longer substitutes a finite 0.0 for six absent measurements (zone_lower, zone_upper, distance_from_zone, atr_stop, targets, current_price) — all six now go through _finite_or_nan, the same treatment structure.hvn/lvn/swing_struct already get. Found, not fixed here: atr_stop/targets/current_price reach panel_render.py through a separate, still-not-NaN-aware safe_float() call — closed as F4, below.

#108

PHASE7_NEXT.md Head Block Brought Current for F1/F2/F3; Rule 38 Added

September 13, 2026

PROCESS — RECORDED

"The git add -A trap was rediscovered three more times this same session even after being documented."

Commit 1f1c935. Head block rewritten for Round 6's F1/F2/F3. Rule 38 added: always git reset the delivery files before git commit — a sequence missing that step is what let git add -A sweep in files it shouldn't have, each time.

#109

The Orphaned 12 September Handover Note Committed

September 13, 2026

PROCESS — RECORDED

"A complete, never-committed 12 September session handover note was found stray and uncommitted."

Commit da2e231. Claude outputs/phase7_handover_2026-09-12.md committed as-is, confirmed by diff not to be a Viktor edit.

#110

F4 Fixed — panel_render.py's atr_stop/targets/current_price No Longer Default to 0.0

September 13, 2026

FINDINGS — FIXED

"The gap F3 found but did not fix: three measurements reached panel_render.py through a separate, still-not-NaN-aware safe_float() call."

Commit 154e534. Same _finite_or_nan-style treatment as F3, one layer downstream. A real UnboundLocalError bug (color variables referenced before definition) was found and fixed while building this — caught by running the suite, not by review, and named honestly in the commit message. code_hash 36420c9f... to 3c4ad5ca....

#111

F5 Fixed — decision_model.py's Trend-Health/Entry-Score Bands Named

September 13, 2026

FINDINGS — FIXED

"The other half of Round 6 F1's own Location section, left open when F1 landed."

Commit ac0a211. AGGRESSIVE/CONSERVATIVE trend-health/entry-score bands (75/70/50) promoted to named, fingerprinted class constants: AGGRESSIVE_TREND_HEALTH_MIN, AGGRESSIVE_ENTRY_SCORE_MIN, CONSERVATIVE_TREND_HEALTH_MIN. Values unchanged. Investigated and found NOT a defect: the same finding's MIN_ACTION_BIAS naming split against RAW_BIAS_THRESHOLD — both already named, already fingerprinted, deliberately two different thresholds; nothing changed there. code_hash 3c4ad5ca... to 13908bba.... Golden snapshot moved (7 record-completeness fields, zero decision fields) — predicted this time from the F1 precedent rather than discovered by surprise. Test counts at tip (Linux sandbox): pytest+pandas_ta 466 passed; pytest without pandas_ta 338p/117s; run_tests.py 395p/0f/32e (32 pre-existing, unrelated to anything this session touched — all fixture errors on a fixture-free runner). All commits confirmed landed and byte-identical via fresh clone + md5 at the time.

#112

Encoding Defect Found and Fixed — send_audit_round.py's OpenRouter Stream

September 13, 2026

FINDINGS — FIXED

“Every non-ASCII character in round 6's report.md as originally written — every em dash, ±, → — was mojibake.”

Commits 7fbf77c, 9c5f954, 4e0110b. Root cause: `resp.iter_lines(decode_unicode=True)` was read without setting `resp.encoding` first; OpenRouter's response carries no explicit charset, so requests fell back to Latin-1 instead of the API's actual UTF-8 — corrupting every non-ASCII character it decoded, in the streamed report/reasoning text and any `http_error.txt` body alike. Found by verifying the report before acting on it, not by a failed test. Fixed at the source: `resp.encoding` is now forced to "utf-8" immediately after the POST returns, before anything reads `resp.text` or the stream — `send_audit_round.py`'s own docstring names this as the fourth failure mode the script exists to make impossible. The corruption itself is a deterministic, reversible byte transform; `report.md` was replaced with the corrected decode (same model output, not a content edit), original corrupted bytes' SHA-256 kept for the record in the new `ENCODING_CORRECTION.md` rather than silently discarded. A second pass, prompted by this project's own handover check, found three more evidence files from the same round (`reasoning.txt`, `generation.json`, `run_metadata.json`) untracked and none of them gitignored — closed the same day rather than repeating the failure mode the `.gitignore`'s own comment already named. Docs/tooling only; `code_hash` unaffected throughout (`docs/build/` excluded by directory).

#113

The git add -A Trap, Closed Structurally Instead of a Fifth Restatement

September 14, 2026

PROCESS — RECORDED

“3e93fff — this session's own PHASE7_NEXT.md head-block patch — committed two of Claude's own delivery files alongside it.”

Commit 214c5d8. Root cause named plainly as Claude's: a redelivered patch (rebuilt to fill in a commit hash) used a shortened command sequence that dropped the `git reset` steps rule 38 already exists for, and Viktor's pasted `git status` showing both files staged was called "exactly the expected status" without checking it against what the sequence should have produced. Same shape as rule 38 itself: `git add -A` does what it does regardless of whether a numbered sequence remembers the reset steps. Fixed structurally rather than with a fifth restatement of rule 38: `.gitignore` now carries `*.patch` and `*_commit_message.txt`, patterns that match nothing else currently tracked, so `git add -A` can no longer stage either file regardless of which sequence is followed. Verified against the state that actually caused the failure — a fresh clone of the true tip with both files already absent from disk, matching Viktor's own state after he had already run `del` on them — rather than a tree where they still happened to be present. `code_hash` and golden snapshot unmoved (confirmed, not assumed: neither removed file is a `.py` file inside the fingerprinted tree, and `.gitignore` is outside the walk entirely).

#114

send_fix_verification.py Built — a Narrower Tool for a Narrower Job

September 14, 2026

DECISION — ADOPTED

“Verifying that round 6's own findings were actually fixed does not need everything a fresh independent audit needs.”

Commit 2fc578b. Implements Viktor's ruling the same day: the round-6 auditor (meta/muse-spark-1.3) is asked to verify its own findings were fixed, not to re-grade the Constitution from scratch — finding something new would be a bonus, not the goal. Deliberately not called "Round 7"; that name is reserved for the next genuinely independent audit. A new script rather than a mode on send_audit_round.py: that script gets repointed at a new reviewer every round (three repoints already), so importing its MODEL constant would let a future repoint silently redirect this call too — the "one edit upstream" shape muse-spark-1.3's own F3 named for a fabricated-zero default three call frames from where it fires. MODEL/PROVIDER_SLUG/pricing pinned independently here instead. Sends a single-turn API call (no memory persists between calls): round 6's own report.md verbatim, a review note pointing at each of the five items (with explicit pointers for F4/F5, which are not separately numbered in muse-spark-1.3's own report), and the full current content of the eight touched files — the middle ground between a diff too thin to check and a full-codebase resend too expensive to justify. code_hash unaffected (docs/build/ excluded by directory, confirmed).

#115

Round-6 Fix-Verification Sent and Graded: F1 Through F5, All Fixed

September 14, 2026

EXTERNAL ASSESSMENT —
META MUSE SPARK 1.3

“Every item round 6 itself raised, verified fixed by that same auditor — not re-graded by a fresh model.”

Commits 59de938, e65a0f7. Sent via send_fix_verification.py --send; provider confirmed Meta (matches the pin, no substitution), finish_reason=stop (not truncated), 72,542 / 2,008 native prompt/completion tokens, cost \$0.0992115 — below the estimated range because the response was terse (five short verdicts, no full report) rather than the 5k-20k output tokens budgeted. All five verdicts: Fixed — F1 (four locations named and fingerprinted), F2 (confidence_score dropped from engine_core.py, sole meaning left in signal_router.py), F3 (all six fields NaN-passthrough via _finite_or_nan), F4 (panel_render.py's downstream defaults NaN-safe), F5 (decision_model.py's bands and threshold split named/fingerprinted). Independently checked against the actual landed code, not just relayed: every citation in the model's report grepped against source and confirmed to match. Bonus finding, unprompted by the review note, confirmed real: models/signal_router.py's duplicate @staticmethod above _finite_or_nan — harmless, found rather than fixed at this point (see Entry #116). Response saved to docs/audit_reports/round6_fix-verification_2026-09-14/, byte-identical between Viktor's disk and the landed repo, confirmed by md5.

#116

**Bonus Finding Fixed — Duplicate @staticmethod Above
_finite_or_nan**

September 14, 2026

FINDINGS — FIXED

“models/signal_router.py has duplicate @staticmethod stacked above _finite_or_nan — harmless, cleanup only.”

Commit a70fc1b. One of the two stacked decorator lines removed; _finite_or_nan keeps its single @staticmethod, same body, same six call sites in _build_decision_object — no behavior change. code_hash moved, predicted and confirmed: 13908bba... to 3867ba49..., because code_fingerprint.py hashes the docstring-stripped parse tree and a duplicate decorator is a second entry in the function's decorator_list. Golden snapshot unmoved — confirmed via tests/test_golden_path.py passing, not assumed from the AST theory alone. Verified in the Linux sandbox against a fresh clone in all three configurations (466 passed / 338p,117s / 395p,0f,32e) and matched exactly on Viktor's own machine before he committed.

#117

**Viktor's Request: the Setup's Direction Stated in Words, Not
Just Via Targets**

September 14, 2026

DECISION — ADOPTED

“A live run correctly produced a SHORT setup on a bearish market, under a NO-TRADE decision, with nothing on the panel saying “SHORT” — worked out from the descending targets instead.”

Commit 66f1479. A solid colored box added directly under the DECISION line in panel_render.py — green “THIS ANALYSIS IS FOR A LONG” when bias is BULLISH, red “...FOR A SHORT” when BEARISH — on every run, not only ones that authorize a trade. Rendering only; source is bias.raw, the same field decision_model.py's _determine_final_action reads for direction per Viktor's 2 September ruling, so this box cannot disagree with an authorized LONG/SHORT by construction. One case handled deliberately rather than reused: _refuse_incoherent_plan firing (“NO-TRADE (PLAN CONTRADICTS ACTION)”) means bias and the risk plan's own targets disagreed this run — see tests/test_direction_source.py, the 2 September “long label on a short plan” defect this box exists to never repeat through a new line — so that case prints “SETUP DIRECTION: CONTRADICTORY” rather than taking a side. Verified by rendering four decision objects directly, since no existing test asserts on this box. code_hash moved, predicted and confirmed: 3867ba49... to c7c50a99.... Golden snapshot unmoved. Live-data run on Viktor's machine matched the prediction exactly before he committed.

#118

NEUTRAL Bias Ruled Silent, Not a "No Lean" Line

September 14, 2026

PROCESS — RECORDED

"Claude's own addition to Entry #117's box, not requested: a NEUTRAL bias printed "SETUP DIRECTION: NEUTRAL — no directional lean this run" for consistency with this panel's other "not located"/"not measured" lines."

Commit 39e0e79. Flagged as optional; Viktor ruled: print nothing. The distinction that changes it — swing_struct, the entry zone and correlation say "not located"/"not measured" because those are measurements that failed to produce a value. A NEUTRAL bias is not a failed reading; it is bias correctly reporting no lean, which an absent box already communicates. Only the NEUTRAL branch changed — BULLISH/BEARISH still render their colored box, and PLAN CONTRADICTS ACTION (Entry #117) is a correctness guard, not a style choice, and was not touched. code_hash moved, predicted and confirmed: c7c50a99... to 38458f20.... Golden snapshot unmoved. Verified in the Linux sandbox, all three configurations unchanged (466p,2w / 338p,117s / 395p,0f,32e), and by rendering a NEUTRAL-bias decision object directly.

#119

Portfolio Document Added — Item 4 of 6, Portfolio-Ready

September 14, 2026

DECISION — ADOPTED

"Drawn from the Constitution, the engineering log, the test suite, and the actual current source — deliberately not from README.md, whose own status table was checked and found to predate round 5/6 and four now-existing core/ modules."

Commit 69f1972. docs/build/build_portfolio_document.py added, generating the fourteen-section Phase7_Portfolio_Document.pdf: objective, problem definition, architecture, data sources, modules, decision logic, testing methodology, debugging process, major problems, solutions, results, limitations, lessons learned, future development — the career plan's own section list. Item 5 of the six portfolio-ready criteria, the separate AI-attribution accounting, deliberately left out of this document per Viktor's instruction to do the two in order (see Entry #120); this document does not itself declare the release gate open, the remaining findings closed, or the project portfolio-ready, leaving all three to Viktor the way PHASE7_NEXT.md does throughout. One factual error caught in proofreading before Viktor ever saw it: an early draft's Results section attributed the sixteen-item remediation sequence to GPT-5.6 Luna Pro's Step 8 re-audit; checked against PHASE7_NEXT.md and build_roadmap.py's own text and corrected — the sequence is GLM 5.3's Step 5 output, Luna Pro's Step 8 came after it and found further findings of its own, including the decision-bar-integrity Critical no prior audit had scheduled. code_hash and the golden snapshot unaffected — docs/ excluded from core/code_fingerprint.py's file walk by directory name, confirmed before and after rather than assumed. Full three-configuration suite unaffected (466 passed / 338 passed, 117 skipped / 395 passed, 0 failed, 32 pre-existing errors), same identity as every prior patch that session.

#120

AI-Attribution Statement Added — Item 5 of 6; All Six Portfolio-Ready Criteria Now on Record

September 14, 2026

MILESTONE — RECORDED

“A first draft’s engineering-log entry-count table summed to 111 against the true total of 118, from a regex that silently missed several tag strings.”

Commit eee218e. docs/build/build_ai_attribution.py added, generating the five-page Phase7_AI_Attribution.pdf — the last unwritten piece of item 5: what Viktor designed, decided, tested and ruled on, versus what Claude produced, assembled into one place for the first time. Covers the method (every claim checkable against a code comment, a dated log entry, a commit, or PHASE7_NEXT.md, not memory alone), the Constitution’s own “Roles & Authority” text, a dated table of ten rulings Viktor made himself each cited to where it lives in current source, a table of five concrete cases where Claude was wrong and Viktor caught it (including the two-September session’s five verification errors and the unrequested NEUTRAL-panel line Viktor overruled at Entry #118), and the independent third-party auditors — including the one exposure (DeepSeek-lineage models touching the codebase via Aider before a later blind DeepSeek audit) stated rather than glossed over. The entry-count table’s error found by recounting from this file’s own AST directly — every entry_box()/highlighted_entry_box() call, its tag argument resolved as a constant — which also confirmed entries numbered 1 through 118 with no gap and no duplicate; the wrong 111-total table never reached Viktor. With this and Entry #119 landed, all six portfolio-ready criteria in docs/PHASE7_NEXT.md have their supporting material in place. Whether the release gate is open, whether the remaining findings count as closed, and whether to declare the project portfolio-ready and tag portfolio-v1 remained Viktor’s own calls — not decided by this patch (see Entry #124). code_hash and the golden snapshot unaffected, confirmed before and after. Full three-configuration suite unaffected.

#121

PHASE7_NEXT.md Head Block Brought Current for Entries

#119 and #120

September 14, 2026

PROCESS — RECORDED

“A session-close handover check found the project’s own entry-point document still reading as if the last thing that happened was Entry #116’s fix.”

Commit 59b915a (sixteenth patch). 293f310 (this document’s own v1.29 regeneration), Entry #119’s Portfolio Document, and Entry #120’s AI-Attribution Statement had all landed the same session, but none of the three touched PHASE7_NEXT.md — exactly the gap the standing end-of-session check exists to catch. New head-block entry records what landed and in what order; that all six portfolio-ready criteria now have supporting material on record, not five; and, explicitly, that criterion 1 (the release gate) has its technical condition on record — round 6 re-audited round 5’s Criticals and found no Critical Tier-1 defect, the Constitution’s own bar — but had never been declared met by Viktor, with declaring the project portfolio-ready and tagging portfolio-v1 named as his calls, not decided here. code_hash confirmed unmoved, computed independently before and after rather than assumed from the exclusion rule alone. Full three-configuration suite unaffected; tests/test_golden_path.py’s eight cases individually confirmed passing, snapshot unmoved.

#122

Five Token-Saving Workflow Rules Recorded in Working Practice

September 15, 2026

PROCESS — RECORDED

“Already implemented into the patch-delivery skill and recorded in memory earlier the same day — this is the in-repo copy Viktor asked for, not a new ruling.”

Commit 48757fe (seventeenth patch). New bullet under Working practice: scoped-files review packages (the touched files, their direct callers, and the tests that assert their contract) as the default for fix-verification sends and narrow patches, with the full ~400K+-token audit package reserved for a genuine fresh Tier-1 audit round rather than decided per-round; batched Engineering Notes / Portfolio Document regeneration at session close or once several commits have accumulated, rather than after every small patch, with any resulting gap stated explicitly rather than closed on reflex; no routine staging or re-reading of docs/audit_package/s qwen_reasoning_*.txt files or superseded audit-round folders; chat-reply discipline restated as holding regardless of patch size; the fresh-session-at-closure preference restated. Docs-only, no engine module touched. code_hash confirmed unmoved on both the pristine tip and the edited tree, not assumed from the directory-exclusion rule alone. Full three-configuration suite re-run twice — once against the working edit, once independently against a fresh clone with the patch applied — identical both times and matching the recorded baseline.

#123

Correction: the Five Round-6 Mutant-Escape Findings Were Already Closed at 044b055

September 15, 2026

OBSERVATION — FILED FOR THE RECORD

“Claude reported these five as still open in chat, from reading this file's own narrative rather than checking git.”

Commit f7f2e51 (eighteenth patch). Viktor asked to fix the five round-6 mutant-escape findings (requested-run 6, GPT-6 Astra's round-5 report), which PHASE7_NEXT.md's own narrative still described as recorded “for Viktor to rule on” with no later entry saying otherwise. Checking git before touching any code found they were already fixed, verified, and folded into round 6's audit batch at 044b055 — three days before this session existed: compute_trend_health's always-neutral escape, _determine_final_action's always-WAIT escape, plot_engine_chart's no-op escape on both affected tests, SignalRouter.route's canned-error escape, and run_tests.py's zero-match-exits-0 bug, each mutant reproduced in-process, confirmed to escape the old test and be caught by the new one. New head-block entry corrects the record rather than re-doing work already done, and names the lesson plainly: this file cited 044b055 by hash without ever saying what it fixed — the exact failure shape rule 19 already describes, work from the audit, not from the summary of it. No code or test file touched; the full three-configuration suite re-run was skipped as disproportionate to a documentation-accuracy fix no suite run could validate — only git log and git show could, and both were checked first. code_hash confirmed unmoved regardless, since that check is cheap and catches any accidental non-docs change.

#124

The Release Gate Declared Open; the Project Declared Portfolio-Ready; Tagged portfolio-v1

September 15, 2026

MILESTONE —
PORTFOLIO-READY
DECLARED

“Round 6's targeted re-audit satisfies the Constitution's literal bar and is a new pattern for this project rather than a repeat of prior practice.”

Commit 99e022e. Viktor's ruling, 15 September 2026, made after writing his own position first and having it checked rather than handed to him — the one thing Entries #121 through #123 all named as outstanding and none of them decided. New permanent section, “Declared — 15 September 2026,” under “Two goals, and the order they finish in”: round 6's targeted (not whole-codebase) re-audit satisfies the Constitution's literal bar; the stronger basis offered is the suite's positive, checked result across all three post-round-5 fix commits (466 passed, golden snapshot and code_hash as predicted) on a suite now hardened against the vacuous-pass failure mode, not merely an absence of suspicion; criterion 2 fully closed as of Entry #123's correction; criteria 3 through 6 already on record via Entries #119 through #121. This closes the two-goals-ordering question standing since 3 September: Goal B (backtesting) can now start without risk to what Viktor submits. Tagged portfolio-v1 at this commit. Docs-only; code_hash confirmed unmoved, no engine module or test touched; golden snapshot not applicable — no code path touched.

"Viktor read an unrelated agent-harness project and asked for critical feedback on two of its components, then delegated the final call: "It is your call, do what is best for the engine and the project.""

Commit 63d8620 (twentieth patch). Adopted, in modified form: the seven-item handover checklist in Working practice had been an instruction ("run these checks") rather than a structural fact of the repo — Viktor's own standing preference is a structural fix over a reminder to be careful. The new script answers four of the seven mechanically: untracked/modified files, ignored files (filtered for routine build/cache noise so a real signal isn't buried, the same blind spot this repo's own .gitignore comments already document), loose .patch/*_commit_message.txt files anywhere in the tree, and a staged-but-uncommitted index. The other three — is this session's state written into PHASE7_NEXT.md, are today's rulings recorded, are the Engineering Notes current or the gap stated — need someone to read and judge content, so the script prints them as an explicit reminder rather than pretending to answer them. Not automatic, and said so rather than overclaimed: the source project wires this class of check to persistent-session hooks; this project's actual workflow is single commands in cmd.exe with no session to hook into, so the script runs by hand at session close, computed instead of recollected. Declined: confidence-scored auto-pruning of "Lessons carried forward" — conflicts directly with this project's own standing rule that wrong turns are recorded rather than quietly corrected, and recurrence count is not a proxy for importance (the git add -A trap, Entry #113, stopped recurring specifically because it was fixed structurally; a recurrence-based prune would have deleted that lesson first, precisely because it worked). All four detectors exercised against deliberately introduced test conditions before delivery and confirmed to actually fire, then the test scaffolding removed and the tree confirmed clean again. code_hash confirmed unmoved on the pristine tree, the pre-apply working copy, and a fresh clone with the patch applied — all three identical. Full three-configuration suite run on a fresh clone with the patch applied, matching baseline exactly, including the 32 pre-existing run_tests.py errors diffed by name, not just counted.

"README.md had not been touched since 30 August and still said the release gate was shut, the day after Entry #124 declared it open — the first thing anyone visiting the repository sees."

Commits ebb0a46, d5bb329. Found while starting to scope Goal B (backtesting): the status table, the four-Critical framing, and the test counts were all stale relative to the 15 September portfolio-ready declaration. Fixed by reading current source directly rather than carrying the old table forward — `decision_log.py` and `lineage.py`'s hash-and-archive mechanism for the old Item 6, `bias_score/confidence` de-duplication for the old Item 11, three further Criticals a later round found and closed, current test counts, the release-gate declaration and `portfolio-v1` tag. The repository-layout note on backtesting deliberately softened from a first draft that would have flatly repeated "blocked by the release gate": the Constitution's own backtest-start condition (Items 2, 3, 6, 18) is a separate clause, evidence points to it already being satisfied but nobody has formally re-checked it since 15 September, stated as not-yet-confirmed rather than assumed. Viktor's own correction, recorded plainly: this was not a case of him failing to ask for document checks — the standing handover check had never covered README.md at all, only PHASE7_NEXT.md, the Engineering Notes, and git status; that checklist gap is named, not fixed, here (closed two commits later, Entry #127). Separately, a live instance of a known blind spot, found by hand-checking what the new handover script checks: `phase7_mutant_escape_correction_commit_message.txt` sitting at the repo root, gitignored and therefore invisible to `git status --short` — very likely Entry #123's own delivery file, never deleted per its numbered sequence's own step 11. Not deleted by Claude; flagged for Viktor to remove himself. Also recorded: today's backtesting-scoping research — the backtest-gate's separate Constitution condition, Item 17's actual blast-radius requirement, a real precedent already fixed for what it guards against, and two governance questions (what "known-good checkpoint" and "fixed evaluation dataset" concretely mean) put to Viktor and not yet answered — written down so a new session can resume it rather than only a closed chat window remembering it. `code_hash` confirmed unmoved on both commits, README.md and PHASE7_NEXT.md both outside `core/code_fingerprint.py`'s file walk. Full suite re-run on the README fix specifically (466 passed, matching baseline exactly — "obviously safe" has been wrong on this project before); the pure head-block addition's own re-run skipped as disproportionate, per this document's own precedent for that class of change.

“Viktor named the split himself — the checkpoint/dataset questions are his to write, this one is Claude's engineering call — then said “Yes. Fix this.””

Commit a441023 (twenty-second patch). New `check_readme_currency()` function, printed as “## 5.” in the handover script's output: README.md's last-touched commit (hash, date, subject), HEAD's own, and the commit count between them. Informational only — it does not feed the script's exit code, the same way the four “still manual” reminders don't, because whether README.md's content still matches what PHASE7_NEXT.md's head block declares is a judgment call, not a git fact; what the script removes is the reason that judgment call went unmade for sixteen days (Entry #126) — nobody ran `git log` on README.md to see how old it was. The other half of the open question needed no code change: whether the ignored-file sweep should run “every time” was already answered by the script's own existence — `check_ignored_files()` runs on every invocation, unconditionally, since Entry #125 created it; the stray commit-message file Entry #126 found slipped through because the script itself didn't run that session (no shell access to Viktor's machine), not because its sweep logic was incomplete. Working practice's checklist gets item 8 to match. Three scenarios run against a throwaway git repo — a clean tree with README.md one commit behind HEAD; a stray commit-message file plus an ignored file plus a staged index together, confirming the new check doesn't disturb the existing flagged/clean logic; and a repo where README.md has no commit history at all, confirming graceful degradation — all three produced the expected output and exit code. `code_hash` confirmed unmoved: docs/ excluded from core/code_fingerprint.py's file walk by directory name, checked directly in `EXCLUDED_DIR_NAMES` rather than assumed from precedent. Test suite not re-run, stated as a scoping call rather than a skip: this script has no pytest coverage and nothing under test imports it. First real Windows run confirmed working, closing the platform question every prior entry in this row could only leave unconfirmed.

#128

Goal B Specified and Ratified Before Any Backtesting Code Exists; Two of Claude's Own Errors Recorded Inside It

September 15, 2026

DECISION — ADOPTED

"After four rounds of writing his own position first, Viktor delegated the remaining open points: "make the changes you want and are content with... it is your call."

Commit dff7d00 (twenty-fourth patch), with f9e5127 (twenty-third patch) — the latter only the head-block note recording this document's own v1.30 regeneration. New permanent section, "Goal B — the backtesting phase, specified before it starts," written into PHASE7_NEXT.md and since moved, verbatim, to docs/PHASE7_DECISIONS.md by the 18 September split (Entry #130). Same shape as the portfolio-ready declaration (Entry #124): Viktor wrote his position on each governance question first, Claude critiqued it, corrections were folded in, Viktor ratified — four rounds covering checkpoint, rollback and isolation rules; the fixed evaluation dataset; the verdict thresholds; and the completion boundary. Closes the two questions open since Entry #126 (what "known-good checkpoint" and "fixed evaluation dataset" concretely mean) and settles three that were never on the list. A verdict procedure fixed before any result exists: INCONCLUSIVE below 30 out-of-sample trades, then POSITIVE or NEGATIVE on a symmetric band of 0.30 absolute Sharpe units with drawdown guards, NEUTRAL otherwise, every verdict printed beside absolute return and time-in-market. A completion boundary that is outcome-shaped and accepts a negative verdict as completion. And a re-run clause: the first out-of-sample evaluation is the verdict of record, and any later re-run is a new phase with its own pre-registration, because pre-registration only works once. Three findings came from reading the code rather than the documents: PHASE7_PINNED_DATA silently falling back to live data (named a precondition here, fixed at Entry #131); the evaluation dataset being three series per run rather than one file, which makes truncating each series by candle open time the highest-risk line of the phase; and the pinned path already serving df.iloc[-limit:], so the harness is a time cursor on existing code rather than a new pipeline. Two of Claude's own errors, recorded in the section rather than dropped, both the shape Entry #123 exists to warn about: Claude said no lookahead sweep had ever run, and git showed 793e863 (30 August) had done it, with tests/test_no_lookahead.py pinning the result; and Claude called fc35a2f a precedent covering most of the Item 17 guard without reading it — it is a test-time path redirect, not a runtime guard. Both were caught only because Viktor asked for a critical review. Standing mitigation adopted: every claim Claude makes about the codebase's state says whether it came from the code or from a document. No Goal B code written. Docs-only; code_hash unmoved at 38458f20... on both commits.

“Six audit rounds have checked whether this engine is honestly built. Not one has checked whether it is analytically sound.”

Commit 4a97c32. Decided in conversation at the close of 15 September and kept until then only in a handover document outside git; written into docs/PHASE7_DECISIONS.md as a subsection of Goal B so the reasoning survives in the repository, per the standing rule that a conversation is not a record. Two rulings by Viktor. First, the course correction: before any Goal B implementation, the engine itself is reviewed for logical soundness — is the decision logic coherent, are the constants justified, are the signals actually independent, and is the engine's thesis written down anywhere at all. Everything Entry #128 ratified stands; nothing is cancelled, only deferred. The reasoning: a backtest grades a decision process against outcomes, but it cannot show that the process's own inputs are non-redundant or that its constants are more than convenient — and backtesting an arbitrary threshold tests the threshold, not the idea. Of Goal B's preconditions only the PHASE7_PINNED_DATA fix was judged urgent on its own merits (Entry #131). Two scope conditions were agreed the same day and have not yet been acted on: the review needs its own completion boundary, written before it starts, since “review the whole engine” is unbounded by construction; and a boundary on what it can settle — internal coherence is reviewable, but whether the market thesis is correct is not something a review, or an AI, can establish; only data can. Second, the pace ruling, made by Viktor on his own judgment: 226 commits across 22 days, active on 20 of them, had become hasty, and the project slows down considerably. A session no longer opens by proposing work from the open items as a queue to clear; it asks what Viktor wants to do first. Split into an entry of its own, rather than folded into Entry #130, at Viktor's choice. The review's first two findings are Entries #132 and #134.

“Reassembling the eight extracted blocks in their original order reproduced the 5,831-line file byte for byte.”

Commit 4a97c32. docs/PHASE7_NEXT.md had grown to 57,833 words and was what every session read first. Split three ways: PHASE7_NEXT.md becomes a current-state entry point only, rewritten each session rather than appended to (1,059 words, a 98% reduction in what a session reads first); docs/PHASE7_DECISIONS.md holds standing material — the two goals, Goal B, the rulings, Working practice, the earned rules; docs/PHASE7_HISTORY.md holds the head-block archive and every dated session narrative, append-only. Proven programmatically rather than eyeballed: every one of the original's 5,831 lines was extracted by exact line range into one of eight blocks, and reassembling the blocks in original order matched the pre-split file by string equality. Four sections titled “Rulings” needed a judgment call — standing rulings went to DECISIONS, rulings inseparable from build narrative went to HISTORY — each named in the commit message so Viktor can override it. Resolved in the same commit by reading the code, rather than re-running the malformed grep that had left it unverified: the six bias weights are not in FINGERPRINTED_CONFIG, which holds config.py-level names only, but they are fingerprinted through the separate FINGERPRINTED_MODULES / module_snapshot() mechanism, and both feed run_hash — already pinned by tests/test_lineage.py::test_the_run_hash_moves_when_a_bias_weight_moves. No code change was needed. Named and deliberately not fixed: seven test files and three docs/build/ scripts, this document's own build script among them, still cite PHASE7_NEXT.md for facts that now live in HISTORY; none of them reads the file, so the citations are stale rather than broken. Earlier entries of this document that cite PHASE7_NEXT.md are left as written, under this document's no-silent-edits rule — they were true when written. One consequence surfaced a day later and was repaired at 5c73e24 (Entry #135): HISTORY had not been appended to after the split, so the 18 September session's account existed only in NEXT until that commit moved it across verbatim. code_hash unmoved at 38458f20...; full three-configuration suite unchanged (466 passed; 338 passed, 117 skipped; 395 passed, 0 failed, the same 32 errors by name).

"Viktor: "Let us do Tier 1 and then start a new chat.""

Commit 5be5d82, with its landing update 2ebbd6b. Three defects, each confirmed real against source before the patch touched anything. (1) data/data_fetcher.py: with PHASE7_PINNED_DATA set to a path that did not resolve to a directory — a typo, another machine's path, a dropped mount — pinned_source() returned None, which get_tf() reads as "go live." A broken pin request served live data with nothing in the output saying so, reintroducing the exact nondeterminism the pinned mechanism exists to remove. It now raises ValueError naming the resolved path, the choice set_pinned_source() already made for the sibling case; engine_core.py's existing top-level handler turns that into a clean decision-object error, traced before anything was written. (2) core/engine_core.py: _save_state opened the state file with open(path, "w"), which truncates before a byte of the new state is written, so an interrupted write left an empty or partial file — and _load_state read that silently as a first run, losing Exit Watch's comparison baseline. It now writes a temp file in the same directory and os.replace()s it over the real path; the degrade-not-halt behaviour of 29 August is unchanged. (3) docs/build/session_handover_check.py flagged logs/ and the generated docs/audit_package/round*/ content as unexplained ignored files every session an engine run or package rebuild had happened, burying a real surprise among expected clutter; both joined the routine-noise filter, with a negative control proving an unrelated ignored file still surfaces. Seven new tests; each fix's test fails against the pre-fix code, reproducing its defect. One wrong turn recorded: the atomic-write test was first written with pytest's monkeypatch fixture, which passed under pytest but moved run_tests.py's error count from 32 to 33 — caught only by running all three configurations, and rewritten fixture-free. The first engine-code change since the portfolio-v1 tag: code_hash moved from 38458f20... to 5080ccf0..., in exactly the two touched engine files' per-file digests, and had not moved on any commit between the tag and this one. Golden snapshot unmoved, as predicted, since neither changed path is reachable on the golden fixture. Suite 466 to 473 on the Linux sandbox. The landing update, 2ebbd6b, also removed a duplicate open-item bullet its own first editing pass had missed, found by rereading the whole file rather than trusting the diff of the block it had touched.

#132

**Bias-Weight Independence: ADX's Second, Near-Duplicate
Read Removed From continuation_strength**

September 19, 2026

FINDINGS — FIXED

“Asked to write a position on whether the six bias weights are defensible, Claude found 40% of bias_score's weight reading the same raw ADX value on two different curves.”

Commit f24a6e9, with its landing update 2fee78f. The engine review's first finding (Entry #129). trend_health (weight 0.30) and reversal_continuation (0.10) both read adx_val from indicators/trend_health.py and scored it on two different monotonic curves. Not Item 11's defect — that was a reused computed value; this was a reused raw input reaching bias_score through two formulas, which is why it passed Item 11's audit and six later rounds that checked for shared values, not shared raw indicators. models/risk_model.py's Item 14 comment had already named and accepted this exact fact on 11 September, under the narrower standard that mattered there. Ruling: offered (a) independence only, (b) a reasoned rationale for the weight magnitudes, or (c) both, Viktor chose (a). Removing ADX rather than RSI from continuation_strength was Claude's engineering call within that scope, not separately ruled on. The fix deletes continuation's adx_component; its ceiling honestly drops from 60 to 35 rather than being rescaled back up, the same standard as the 31 August health_component removal. RSI was deliberately left in both factors, on the grounds that one reads it as a symmetric question and the other as a directional one — an exemption Entry #134 later measured. Golden snapshot predicted to move in bias.score and what derives from it, and nothing else; the prediction held. Re-baselined with PHASE7_UPDATE_SNAPSHOT=1 and diffed at leaf-field level: exactly 16 fields, all downstream of bias_score or confidence (bias.score 78.70 to 77.10 on the TESTUSDT fixture); entry, exit, structure, degradation and provenance untouched. code_hash moved from 5080ccf0... to bb47ab53..., in indicators/trend_health.py's digest alone — the comment-only edits to bias_engine.py and risk_model.py cannot move it, because the fingerprint hashes the parse tree and # comments are not part of it. run_hash unchanged on the golden run, as designed: it captures input data and fingerprinted configuration, not formula code, and code_hash is what tells the two code versions apart. One new regression test. Not fixed, by Viktor's choice: the weight magnitudes remain unreviewed hand-picked judgment calls, recorded as open; the other four factors were not checked for hidden shared inputs in this pass.

“Caught only by reading git status --short before committing — a carefulness step, not a structural one.”

Recorded in 2fee78f, the landing update for Entry #132. Delivering f24a6e9, the command sequence's git add -A swept an untracked file, Claude outputs/Phase7_Session_Handover_2026-09-15.pdf, toward staging. It did not reach the commit, because git status --short was read first. The project had paid for this exact trap before — at least four earlier occurrences are recorded in docs/PHASE7_HISTORY.md, among them commits bd79800 and 77d822f and a commit-message file swept into a commit — and each time the answer had been to be more careful. Viktor's standing preference for anything that cannot be undone is a structural fix over an instruction to be careful, and a commit to a public repository is that kind of thing. So the delivery procedure changed: on Viktor's repository every file is added by name, one git add per file, so an unrelated untracked file cannot be staged whatever else is lying in the tree. git add -A survives only inside the throwaway repository used to build a patch, where nothing else exists. The change lives in the patch-delivery skill, outside the repository; 2fee78f records it as proposed for Viktor's review, not yet confirmed saved. The PDF itself was left where it was, neither committed nor deleted, with its fate recorded as an open item for Viktor. Split into an entry of its own, rather than folded into Entry #132, at Viktor's choice.

"In a sustained trend the four agree because they must, and bias_score reports that agreement as four independent confirmations."

Commit e2c6637, with its landing update 742b514. The engine review's second pass, tracing the four factors Entry #132 left unchecked. Trend health (0.30), structure regime (0.20), SuperTrend direction (0.15) and macro bias (0.10) — three quarters of the blend — are four transforms of one measurement, the recent direction of close: an EMA slope, the gap between two close means, an ATR-banded flip, and price against its 50-EMA one timeframe up. Different windows and arithmetic, so they genuinely disagree at turning points and this is not the ADX defect repeated; but all four are monotonic in the same quantity and enter the blend with the same sign convention. Only volume sentiment (0.15) and continuation's acceleration term read something other than price direction. Entry #132's RSI exemption was measured across RSI 0 to 100: it holds in downtrends ($r = 0.37$) and largely fails in uptrends ($r = 0.83$), worth at most 1.5 points of bias_score against the 2.5 the ADX fix removed — a note, not an alarm. Why nothing is fixed: acting means reweighting or dropping factors, a trading judgment about what corroboration between correlated trend reads is worth, which this project cannot yet evaluate — the golden baseline proves a change is attributable, never that it is correct. The decision is Viktor's and has not been made. Two errors in the project's own record corrected with it. models/bias_engine.py's dependency graph described structure_regime as a swing-based label; it never was — it is a close-mean gap with hysteresis, and swing structure reaches only the panel — so the first pass had cleared the factor as independent against a mechanism it does not use. And PHASE7_NEXT.md had claimed the session's test counts matched on Linux and Windows; that was false, and Claude's error: on any LF checkout test_manifest_hashes_match_the_files fails, because the pinned CSVs are normalised by .gitattributes while MANIFEST.json pins raw-byte hashes, so the test is green only where the manifest was generated. Recorded as a standing platform difference, not fixed here. Two new tests pin facts, not judgments — structure_regime responds to close alone, and both RSI correlations stay inside bands — each with a negative control. code_hash unmoved at bb47ab53..., the bias_engine.py edit being comment-only; golden snapshot untouched.

"Viktor's first answer was that a finding should 'only warn, then I consult with you and we push or not.'"

Commit 2f5fdaa, with its landing update 5c73e24. githooks/pre-push runs docs/build/session_handover_check.py on every git push: exit 0 lets the push through, anything else stops it with the findings printed. Installed once per clone with git config core.hooksPath githooks, since git does not version .git/hooks. Proposed the session before and never built; Viktor raised it himself. Ruling: three options were put to him — (A) stop the push, consult, override with git push --no-verify; (B) a y/N prompt at the terminal; (C) warn and let the push through. Claude pointed out that warn-only could not deliver the second half of his first answer: the push would already have left before the consult. He chose (A), and --no-verify is his override after the findings have been read together, never a routine step. (B) was set aside not on merit but because reading the keyboard from a hook under cmd.exe through Git for Windows' sh could not be verified from the sandbox. The hook fails closed — no python on PATH, a missing script or a crash all stop the push too. The handover check gained a section 6 that flags a clone without the hook installed. Two findings from verification. Applying the patch in one step wrote the new hook with CRLF despite the same patch adding githooks/* eol=lf, because git apply works under the attributes in force before the patch — so a patch that changes .gitattributes and adds files now applies in two steps, .gitattributes first. And on Linux git 2.43 a hook that cannot be executed makes git push hang rather than fail; the first negative control hung the sandbox for five minutes, and the push tests now carry a timeout. Nine new tests, no fixtures; against the pre-change code 8 fail with FileNotFoundError, proving only that the file is new, and 1 fails behaviourally. The hook's first run on Viktor's Windows machine was the push of 2f5fdaa itself, clean. code_hash unmoved at bb47ab53...; golden snapshot untouched. Suite on the Linux sandbox, clone made with core.autocrlf=true: 485 passed, 0 failed with pandas_ta. The same count on Viktor's machine is recorded as his confirmation by proceeding, not a count Claude read. The landing update, 5c73e24, also found PHASE7_HISTORY.md had not been appended to since the 18 September split, and moved the previous account there verbatim before rewriting NEXT — proven by un-demoting the moved block and comparing it byte for byte.

#136

PHASE7_NEXT.md Brought Current at c7ced36; a Wrong Prediction About git status Recorded

September 19, 2026

PROCESS — RECORDED

"The prediction was made from memory, not checked; the patch-delivery skill itself says to confirm the delivery files are absent, which was right."

Commits e431714 and e115272, docs-only. PHASE7_NEXT.md had been written at 2f5fdaa/5c73e24 and did not know about c7ced36 (this document's v1.31 regeneration), so two sessions had to open from a resume prompt that corrected it. e431714 made the entry-point file true on its own again: the previous version moved verbatim into PHASE7_HISTORY.md under a dated entry, headings demoted one level, and proven a move rather than an edit by un-demoting the block and comparing it byte for byte, with a one-character negative control. README.md, fifteen commits behind, had its test counts and layout tree brought current; two older omissions it found (Claude outputs/ and decision_log_backups/ missing from the tree, no pointer to the three PHASE7_* files) were recorded as an open item rather than fixed out of scope, and closed at Entry #140. e115272 recorded a wrong prediction from that delivery: Claude said git status --short would list the two delivery files as untracked, and they did not appear at all, because .gitignore ignores *.patch and *_commit_message.txt. No harm followed, but the error existed only in chat until Viktor asked for it to be recorded. Test counts unchanged at 485 / 0 with pandas_ta; on Viktor's machine recorded as his confirmation by proceeding. code_hash unmoved at bb47ab53..., recomputed on both commits' trees for this regeneration.

#137

The Pinned-Data Manifest Check Made Portable: One Hash Per Content, on Every Checkout

September 20, 2026

FINDINGS — FIXED

"One committed file, two byte forms — and a test that hashed the bytes."

Commit b68de08, tests and tooling only. tests/test_pinned_source.py compared the raw bytes of the three pinned CSVs with MANIFEST.json. With *text=auto in .gitattributes, git checks those CSVs out CRLF on Windows and LF elsewhere; the manifest had been generated on a CRLF tree, so the test passed on Windows and failed on every default LF checkout (484 / 1). Recorded since Entry #134 as a standing platform difference and never quietly subtracted from a count. The fix hashes content with CRLF normalised to LF, in both places: make_pinned.py's sha256(), which writes the manifest, and a copy of the same function in the test, which cannot import make_pinned.py because that module writes files when imported. The manifest was regenerated by running the script itself, not by editing hashes in, and its CSVs came out byte-identical. A new test, test_manifest_hash_ignores_line_endings, builds both byte forms of each CSV itself and requires both to match, so a revert to raw-byte hashing fails on every platform — the existing test alone could not catch that on an LF checkout. Against the pre-fix manifest, 2 behavioural failures and no import errors; four negative controls, each restored and confirmed with cmp. Counts +1 everywhere: 486 / 0 with pandas_ta, 355 passed / 120 skipped without it, run_tests.py 415 / 0 / 32; a default LF clone now 486 / 0. make_pinned.py sits under docs/, excluded by directory, so code_hash stayed at bb47ab53... — computed on three trees, not assumed. Golden snapshot untouched.

"We are gonna scrap that entire list. 1-7."

Commit 16d3c1f, docs-only. Viktor's ruling, recorded as its own section in docs/PHASE7_DECISIONS.md. The seven judgment and review items open at e115272 — his written position on the bias-weight magnitudes and thesis, five points from the 15 September PDF, retiring the stale to-do PDF, the four-factor finding of Entry #134, RSI reaching bias_score through two factors, the engine review's missing scope and completion boundary, and the backtest-start condition never formally declared met — were scrapped as open items in favour of "a new audit by an independent model," with three mechanical items first: the manifest test (Entry #137), the stale citations (Entry #139) and README's omissions (Entry #140). How it came about: asked whether the weights had been changed without a record, Claude checked git log rather than answering from memory — the six weights are unchanged since 83a9425 (26 August), which replaced the first commits' four-term blend wholesale; no commit ever changed a value, and the original choice predates the repository, which is why no reason for it exists. Claude's recommendation on items 4 and 5 — that fixing either means choosing new weights, item 1 by another name — was superseded by the full scrap rather than answered. What the ruling does not change: the findings stay on record in DECISIONS and HISTORY; the Constitution's backtest-start condition is neither declared met nor waived; the 15 September course correction (Entry #129) is not amended. Still undecided, and Viktor's: which model, which package, and whether the auditor sees the scrapped findings. The same commit moved the previous NEXT to HISTORY verbatim and found it headed a day ahead of the commits that wrote it — recorded in the new entry's heading note, not corrected in the moved text, since HISTORY is append-only. code_hash unmoved (bb47ab53...).

“Viktor's first position was to update the dated records too: “the aim was to keep a real record.””

Commit 9a35f1c, comments, docstrings and docs only. After Entry #130's split, citations of PHASE7_NEXT.md written before 18 September sent readers to a file that no longer held what they cited. The item was recorded as seven test files and three build scripts; a repository-wide search found twenty-four files, including comments in six engine modules. Two rulings. (A) The item covers every stale pointer the search found — yes. (B) Whether dated records — this document's entries, session handovers, audit reports, constitution reviews — are edited too: Viktor's first answer was yes. Claude argued against: those citations were true when written, since DECISIONS and HISTORY did not exist before 18 September; this document is declared append-only in the project's own AI-Attribution Statement; and some of the records are other models' output. Viktor accepted a forward note instead, and Claude read his reply as that ruling and said so before building. Rule: a dated record is not edited because a file it cites was later split or moved; one sentence in PHASE7_NEXT.md's head block, plus the ruling in DECISIONS, tells a reader where such content now lives. The citations inside this document's own earlier entries are therefore left as written, as v1.31 had already done. Each live pointer now names the file and section that hold the content, each target checked to exist. Two citations were wrong from the day they were written, not merely stale: “rule 18” in models/entry_model.py and tests/test_decision_bar_integrity.py — the same commit that wrote them, 108cc9f, inserted a rule and moved that one from 18 to 20; it is now rule 22. Found and not fixed: DECISIONS' rules list has no rule 18, and is not renumbered because rule numbers are cited throughout. The rendered text of the Portfolio Document and the AI-Attribution Statement changed with it; both PDFs were regenerated with this document's v1.32. code_hash predicted and confirmed unmoved at bb47ab53... — comments and docstrings are stripped by the fingerprint's AST pass — and all 33 per-file fingerprints equal; a negative control (one real expression changed) moved the hash, and restoring the file restored it. Suite unchanged, 486 / 0; Windows confirmed by proceeding.

"Sixteen modules at ratification, twenty-two now — and the README still said sixteen."

Commit 077d120, docs-only. Closes the omissions Entry #136 recorded: Claude outputs/ and decision_log_backups/ are in the layout tree, and a new paragraph describes PHASE7_NEXT.md (entry point), PHASE7_DECISIONS.md (standing rules and rulings) and PHASE7_HISTORY.md (dated record), noting that older documents cite the first for content now in the other two. Because the item's recorded list had come from a narrower search at Entry #139, README.md was read end to end and each factual claim checked against the tree and the code. Four more stale statements, fixed in the same commit: the test counts, which predated the test b68de08 added; Logs/ in the tree where the directory is logs/ — one directory on Windows, two on Linux, the defect class of 6 September; the module lists and "sixteen files", which was the count at ratification (83a9425, checked with git ls-tree) and is now twenty-two; and the claim that run_tests.py needs "nothing but a Python interpreter," which the runner's own docstring had withdrawn on 6 September — most test files import pytest at module scope. Left alone, because it is not Claude's call: README.md still cites docs/Phase7_Audit_Findings_Complete.pdf, which does not exist; whether the reference goes or the source material is restored has been Viktor's open decision since 6 September, and is now an open item in PHASE7_NEXT.md rather than only in HISTORY. code_hash unmoved at bb47ab53...; suite unchanged; Windows confirmed by proceeding.

"The pre-reinstall check looked at what the repository tracked and ignored, and could not see material that had never entered it."

Commit 4629002, docs-only. README.md said all raw auditor output was published verbatim in docs/Phase7_Audit_Findings_Complete.pdf. Neither that PDF nor its source — the four round-1 auditor outputs of 27 August (Run 1, DeepSeek V4 Pro, blind; Runs A, B and C, Kimi K3), run through the OpenRouter chat room — was ever in the repository, and the same passage's claim that the later rounds' reports sat in docs/audit_reports/ "the same as the original four" was wrong for the same reason. The item had been Viktor's call since 6 September. Claude recommended restoring the source if a copy existed and correcting README.md if not; Viktor chose to search first. The repository on disk, git history on every branch, D:\Phase_7_Engine_Random_Files, the OpenRouter chat room (it keeps history only in the browser), OpenRouter's request logs (I/O logging was off: token counts, no text) and a Claude data export for 26–31 August were searched, and none held a copy. Claude's first guess, that an icon on some request-log rows meant the text was stored, was wrong; opening a row showed it. Viktor then ended the search and approved Claude's three recommendations as proposed. README.md states the loss and what survives — Phase7_Remediation_Plan.pdf, GLM 5.3's prioritisation of the findings, not the findings themselves — and says no copy "has been found" rather than that none survives, because Downloads, the Cowork task list and one unnamed project were not opened. The error in PHASE7_HISTORY.md's 30 August entry, "nothing lived only on disk except gitignored logs", is corrected by a new dated entry, not an edit, per the ruling Entry #139 records. build_findings_bundle.py is kept, as the record of how the PDF was built. Found and deliberately not fixed: README.md's "Every raw report from every round is in docs/audit_reports/" — round 2's folder holds a reasoning trace that ends before any report — recorded as an open item, since it lay outside this one. code_hash unremoved at bb47ab53..., computed on both trees. Suite unchanged, 486 / 0 with pandas_ta on the Linux sandbox; on Viktor's Windows machine recorded as his confirmation by proceeding, and the pre-push hook reported SUMMARY: clean on the push.

"The reason that holds is authorship: that session produced the Remediation Plan, and round 3 reviewed work done under it."

Commit 53394ff, docs-only. Viktor chose four loose ends from PHASE7_NEXT.md's Open items. The first was README.md's "Every raw report from every round is in docs/audit_reports/". That was narrowly overstated: round 2 produced no report, since every attempt on 2 September stopped before writing one, and its folder holds a reasoning trace. The same paragraph said every round ran "on a model reporting no prior exposure". Checking that turned up a contradiction in the record about round 3 (GLM 5.3 Flash, 5 September). HISTORY's 5 September entry called it not independent under the lab-level rule. The 2 September ruling clears chat-session exposure, and calls applying the lab-level rule to it a category error. GLM 5.3's 28 August session was chat-session exposure, read in the provider export. The reason that does hold is authorship: that session produced the Remediation Plan, and round 3 reviewed work done under it. Viktor first proposed counting round 3. He remembered the 28 August run as minor, and as a deliberate compromise. Claude checked the record, which showed a Step 5 run with the full source, the Constitution and all four round-1 outputs, and no compromise ruling anywhere. Viktor then ruled that round 3 is not counted as independent. It is recorded as its own section of PHASE7_DECISIONS.md, including what it weakens. The independent rounds after the original audit go from five as published to three that are both independent and complete (rounds 4, 5 and 6). README's status row, release-gate paragraph and audit paragraph were corrected to match. The wording was delegated to Claude. code_hash unmoved at bb47ab53..., computed on both trees.

"One rule ever changed number, and it changed three times."

Commit 53394ff, the same four-item session as Entry #142. The `api.mexc.com` connection seen during test runs was traced by hooking every outbound connection (a `sitecustomize` hook on `getaddrinfo` and on HTTP tunnel setup). The traced call is `tests/test_clean_checkout.py::test_main_runs_without_a_logs_directory`, which runs `python main.py` in a temporary copy of the repository. It is the suite's only outbound connection, with `pandas_ta`, under both `pytest` and `run_tests.py`, and it writes only inside its temporary directory. Claude recommended pointing that subprocess at `tests/fixtures/pinned/`, which ran `main.py` to exit 0 with no network in the sandbox. Viktor ruled to leave it, and Claude was overruled; no test changed. Rule numbers: `PHASE7_DECISIONS.md`'s rules list was rebuilt at all 98 commits that touched it. One rule ever changed number: "Fixing the instance you found does not close the item", 18 → 20 → 21 → 22, all on 1–2 September. Every citation of rules 18–21 in tracked text files and in every commit message was checked against the list as it stood when written. All were correct except "rule 18" meaning today's rule 22, which 9a35f1c had already recorded; PDFs were not searched. At Viktor's approval, item 18 now says the number is empty and where the rule went. Found while adding it: Markdown numbers an ordered list by position, so from rule 19 on, every rule rendered one lower than its cited number. This was checked with a CommonMark renderer, not on GitHub itself; after the note, all 38 match. A wrong turn, sandbox only: while testing the pinned-data option, a failed `cd` let `main.py` run inside the scratch clone and create `logs/` there. Nothing on Viktor's machine was involved. Found and not fixed here: the Portfolio Document's "no prior involvement" claim (fixed at Entry #146), and the Qwen-or-Kimi attribution of round 2's observations (settled at Entry #158).

“A commit that regenerates the Notes cannot cover itself, so one commit behind is the lowest the gap can go.”

Commits 92775ea, 982e70f and 65a0aef, docs-only. They follow 3a899b5, which regenerated this document at v1.33 and so has no entry of its own, as c7ced36 and c745a67 had none. 92775ea filed 3a899b5's landing facts: Windows counts confirmed by proceeding, and SUMMARY: clean from the pre-push hook. Claude laid out three options for filing them and pointed out that the third, a Notes regeneration, would not close the gap. Viktor chose to leave them for the next docs commit, then changed to a commit of their own. Claude's wording for that option said the Notes would “stay one commit behind”. That was wrong: the commit made it two. 982e70f filed 53394ff's landing. Its three recorded points were Viktor's fresh clone confirming the tip; Claude's expected git status output listing ?? lines for the two delivery files a second time, although .gitignore ignores them (lines 106–107, checked); and the to-do PDF deleted (Viktor's report, not checked by Claude). Claude noted that landing these ahead of the audit-planning commit departed from Viktor's own earlier instruction; the choice was his. 65a0aef did the once-per-session rewrite of PHASE7_NEXT.md. The previous version moved to HISTORY verbatim, with the move proven by un-demotion and a negative control. A first run of that proof reported a mismatch. The cause was the proof's own offset (+2 instead of +4 for the two separating blank lines), not the file. Ruled in 65a0aef: the Notes were not regenerated. Claude recommended waiting, because a regeneration then would have to be repeated after the audit-planning commit. Viktor agreed to regenerate once, after the audit rulings, and only if the audit package included the Notes. This v1.34 regeneration is Viktor's later call (Document History, v1.34). code_hash unmoved at bb47ab53... through all three commits.

"A refusal summarised as worth taking."

Commit a9d4b1f. `models/decision_model.py` built the decision's one-line summary from `reasons[-1]`, and `_compute_ev` appends to `reasons` last on every path through `evaluate()`. So every run's summary was the illustrative expected-value sentence, whatever the run had decided. The golden fixture recorded a NO-TRADE (RISK TOO HIGH) summarised as "positive -- worth taking on average". The summary reaches the decision object and the decision log, but `panel_render.py` does not print it, which is why six audit rounds driven by the panel did not reach it. It was found, not by an audit, while answering the five questions from Viktor's 15 September PDF against the code. None of the five turned out to be a calculation error. Claude's prediction in chat that they were "mostly checkable arithmetic" was wrong, and is recorded as wrong. The fix makes the summary name the reason of whichever stage last set `final_action`, tracked in `evaluate()`. Claude proposed that over the simpler `reasons[0]`, because two later stages can replace the action after the first reason is written. Viktor: "the summary fix can go into the patch also, whatever." The EV sentence now says it is confidence restated (a linear rescaling that cannot disagree with it). It quotes the breakeven and positive levels, 33/100 and 43/100, derived from `AVG_REWARD_R` and the new named `EV_BREAKEVEN_BAND_R`. The `VALIDATION` line gained its "/100". A `bias_engine.py` comment records where volume and macro each reach an output, and says whether three penalties is the right total weight is Viktor's judgment, unmade. Not fixed: whether the EV line belongs in the reasoning at all, the unreviewed bias weights, and the BTC adjustment scaling up with correlation. `code_hash` `bb47ab53...` → `b0213766....` Golden re-baselined in exactly the two predicted fields. A mechanical consequence was predicted only after it happened: the Linux re-baseline wrote the fixture all-LF, so it was converted back to CRLF by hand. Tests 486 → 500 with `pandas_ta`, 14 new; against the pre-fix code, 9 behavioural failures and 1 `AttributeError` for the new constant (which proves only that it is new). Confirmed on Windows by Viktor's live run: `VALIDATION 35.00/100` with its arithmetic checked, and the logged summary naming the risk reason.

#146

README and the Portfolio Document Stopped Calling the Four Original Runs Independent of the Build

September 20, 2026

FINDINGS — FIXED

"We fix this"

Commit 6e1baba, docs-only. The pre-push hook's section 5 flagged README.md three commits behind. Read end to end against the tree, it held four stale or wrong claims. First, its two test-count lines predated a9d4b1f. Second, "four independent runs ... on models with no prior involvement in the build" was contradicted by README itself seventy lines later: three models worked on the build through Aider, two of them DeepSeek's (V3 and R1), and Run 1 was DeepSeek V4 Pro. Independence is tracked at the lab level here, so the lineage counts; Runs A, B and C (Kimi K3) are unaffected. Third, "Every raw report from every round is in docs/audit_reports/" was wrong: round 2 produced no report. Fourth, docs/build/README.md listed build_findings_bundle.py's PDF with no note that neither the PDF nor its source ever reached the repository. The Portfolio Document carried the second claim in its own words. That had been a standing open item, and Viktor ruled that it rides along now ("We fix this"). The script is corrected and Phase7_Portfolio_Document.pdf regenerated. Each correction says in the text that the document used to say otherwise, and since when. A wrong turn, recorded: the first check that the PDF's text had changed extracted zero characters (a hand-rolled zlib walk that matched nothing) and reported "changed text lines: 0". That reads as "nothing changed" when in fact nothing was read, the same vacuous-check defect this project has recorded before. It was redone with pdftotext: exactly one paragraph differs. A no-change rebuild differs in 74 bytes, all timestamps. Found and not fixed here: the TREND line's missing denominator (Entry #148). code_hash unmoved at b0213766....

#147

code_hash Is Comparable Only Within One Python Minor Version

September 20, 2026

OBSERVATION — FILED FOR THE RECORD

"A hash that appears to have moved is a question about the interpreter before it is a question about the code."

Commit b869a30, docs-only: the once-per-session rewrite of PHASE7_NEXT.md at 6e1baba. The previous version moved to HISTORY verbatim, proven four ways: un-demotion reproduces the old file byte for byte; a one-character negative control fails; exactly seven lines differ, each by one leading #; and old HISTORY is a byte prefix of the new. The new fact: the sandbox's default python3 was 3.11.15, and it reported code_hash 938720cd... on an unmodified 6e1baba tree. Under Python 3.12.3 the same tree gives b0213766..., matching Viktor's Windows decision-log record, and under 3.13.13 a third value again. core/code_fingerprint.py hashes ast.dump output, a CPython implementation detail, and already said so under "WHAT IT DOES NOT SURVIVE". It is now a standing rule in PHASE7_NEXT.md: every code_hash claim about this project is computed under Python 3.12 (Viktor runs 3.12.10). The same thing happened again the next morning (a45366b3... under 3.11 at 119c8a3; Entry #149). The v1.34 row's code_hash values were all recomputed under Python 3.12.3. Also recorded in the rewrite: Viktor's position that the independent audit might be paused, "since we fixed all the issues now", recorded as a position and not a ruling. Claude's objection sits beside it: an audit exists to find what nobody has found, and the TREND-line defect, found by Viktor on a live panel, is a small instance of exactly that. The hook's results on the pushes of 65a0aef, a9d4b1f and 6e1baba were not reported, and the file says so rather than carrying earlier clean runs forward. The ruling came the next day (Entry #156).

#148

The TREND Line Printed a Score With No Denominator — Rule 22 Collecting on a9d4b1f

September 21, 2026

FINDINGS — FIXED

“Fixing the instance you found does not close the item.”

Commit 119c8a3. core/panel_render.py printed “TREND : BULLISH / HEALTHY (Score: 100.00)” directly above “VALIDATION : WEAK (Score: 35.00/100)”. A reader could not tell the top of a 0–100 scale from an unbounded total that happened to reach three figures. Viktor found it reading a live panel about two hours after a9d4b1f fixed the identical defect one line below. It was not found by an audit, nor by the work that produced a9d4b1f, which had gone looking for this class of defect and stopped at the instance it was asked about. That is rule 22 of PHASE7_DECISIONS.md, collecting. Every line render_panel() prints was then read, not sampled: after this fix, no score on the panel lacks its scale. The denominator was checked in indicators/trend_health.py, not assumed from the 100.00 on screen: the score is clamped with max(0.0, min(100.0, ...)). The new test holds the scale from two independent directions, because either alone can pass vacuously. It reads the clamp bounds from the function's parse tree, and it runs the function on four generated frames plus the degraded path. Three negative controls: revert the panel line (3 fail); change the clamp to 10 (1 fails, the parse-tree half); inflate the value after the clamp so the parse tree still reads 0–100 (1 fails, the behavioural half). The third proves the two halves separately load-bearing. code_hash b0213766... → 35718f6b...; of 33 per-file digests exactly one moved. Golden unmoved: the rendered panel is discarded by every caller and never reaches the decision object. Tests +4 in all three configurations, run_tests.py's 32 errors unmoved. Confirmed on Windows by Viktor's live run at 00:02 on 21 September: “Score: 100.00/100” above “Score: 35.00/100”.

#149

The Entry-Point Document Carried a Wrong code_hash — Eleven Statements Corrected in Place

September 21, 2026

PROCESS — RECORDED

“A wrong code_hash in the entry-point document is the one fact a later session would trust without rechecking.”

Commit 635a94e, docs-only. PHASE7_NEXT.md was written by b869a30, and 119c8a3 landed after it in the same session. A commit cannot name its own hash, and the once-per-session rewrite rule meant nothing corrected the file afterwards. So the document every fresh session reads first carried five false statements and six incomplete ones. They were found by the handover check at the session's close and by reading the file against 119c8a3 at the next opening. The worst was code_hash: the file gave b0213766... as unmoved, and 119c8a3 had moved it to 35718f6b.... The rest were the tip, the test counts, the Notes gap, and the TREND line still listed as open; plus 119c8a3's Windows live run (only in chat until then), the golden-snapshot line, and the pre-push hook. The hook reported SUMMARY: clean after a first attempt failed on DNS, but whether that push carried both b869a30 and 119c8a3 or only the latter was not established. It is written as unknown. Viktor was offered three scopes and ruled all eleven, as Claude recommended. The Notes count disagreed, nine against ten. Both were right: the file counted itself, and Viktor did not count the amendment. The file now states the convention. Tests against pre-change code were shown to be no evidence here: with both files replaced by the single word GARBAGE the suite still passed. So equal counts before and after are vacuous for this change, and the commit says so instead of offering them. The evidence was the code_hash on both trees, the byte match of the applied files, a clean reverse-apply, and a patch that refuses a drifted base. The 3.11 interpreter gave a45366b3... on an unchanged tree, a third interpreter and a third value (Entry #147). code_hash unmoved at 35718f6b....

“Organize a to do list and we start working, It is up to you.”

Commit ebb4e5c, docs-only. Viktor asked what to put to the engine before backtesting: whether its logic is correct, whether the panel shows the correct information, and whether it carries dead code. Claude read the decision path and the input side, and built an import graph of every engine module. The decision path was `risk_model`, `entry_model`, `exit_model`, `decision_model`, `panel_render`, and what they use from `engine_core`, `signal_router` and `bias_engine`. The input side was `indicators`, `trend_health`, `structure` and `volume_profile`. Thirteen findings resulted. Every one comes from reading the code, none was reproduced by running the engine, and each says how far its reachability was checked. Four are questions about what the engine should do, and are Viktor's to rule on before any backtest is designed. They are: the entry zone against a plan measured from the last close (4); a NEUTRAL bias still printing a plan (5); the stop pulled to the 75-day volume point of control (6); and indicator values beyond 5 sigma silently replaced by the previous bar's (7). The rest became Claude's work orders, B to F. Viktor delegated the order of work to Claude. He deferred the read of `data_fetcher`, `validation`, `decision_log` and `lineage` until after the work list (“we check it after”); that read is Entry #158. Two items were reported as Claude's claims rather than findings: every engine module is reachable from `main.py` except two by design (14), and `_refuse_incoherent_plan` cannot fire today, correctly (15). Also filed: SUMMARY: clean on 635a94e's push (Viktor's paste). The previous NEXT moved to HISTORY verbatim, proven by un-demotion. `code_hash` unmoved at 35718f6b....

“One function rather than six guards, so the next score line added cannot forget.”

Commit a530006, review findings 1–3. (1) Two sentences in the panel's BTC section named AERO whatever the symbol. On BTCUSDT the engine always skips BTC context, so every BTCUSDT run printed that the “AERO analysis” was unaffected. Sequence item 12 had fixed the same string in `decision_model.py` and missed these two. Both now use `asset_name()`, moved to module level so the panel and the model cannot drift apart. (2) TREND and VALIDATION could print “Score: nan/100”. The router sends NaN for a value it could not measure, and both lines formatted it directly. The price lines were guarded at Round 6 F3; these two were not, including in a9d4b1f and 119c8a3, which edited exactly these lines. Decision Reasoning's “trend strength nan/100” was the same defect. (3) Entry quality, confidence, trade quality and BTC-adjusted confidence printed 0.00/100 when absent, a score nobody computed; latent, since the router always sets the first three. The fix is one helper, `_score_text()`, used by every score line: the score with its scale, or “not computed”. Found while scoping and fixed in the same function: `asset_name` tried “USD” before “BUSD”, so ETHBUSD read as “ETHB”; longest suffix first now, with no recorded run affected. The new test renders a panel with every numeric field missing and asserts “nan” appears nowhere, so a later score line fails too. Against the pre-fix code 8 fail: 6 behavioural, 1 parse-tree, and 1 ImportError that proves only that `asset_name` is new. Nine negative controls each failed exactly their tests. A wrong turn: two controls were first contaminated by stale bytecode, because the broken suffix order was the same length and was restored within the same second. They were rerun with `PYTHONDONTWRITEBYTECODE=1`. `code_hash` 35718f6b... → ec88cf24...; golden unmoved, and the pinned run's whole panel is byte-identical. Confirmed on Windows by Viktor's live run: every score with its scale, no “nan”, AERO in the BTC section, and the record carrying ec88cf24....

#152

**Work Order C: the Risk Plan's Direction Parameter Decided
Nothing; a Wrong-Side Stop Now Refuses**

September 21, 2026

FINDINGS — FIXED

“A check at the only producer is the structural form of what D asked for.”

Commit e3f3d51, review findings 8, 9 and 10. (8) RiskModel.calculate_stop_targets tested its detailed_bias parameter for “LONG” or “SHORT”. Its only caller passes the bias state machine's “BULLISH CONFIRMED”, “BEARISH CONFIRMED” or “NEUTRAL”, so the test never matched, and the plan's direction has always been the sign of bias_score. It is the same wrong-vocabulary shape as entry_model's A1/A2 gates, in a second file. It gave no wrong answer, because raw bias comes from the same score. But the parameter read as the direction source, so it is removed, and passing it raises TypeError. (9) A fallback for a stop on the wrong side of price replaced the stop distance and left the stop where it was. It is unreachable on the engine's path, because bias_score is clipped to ± 100 . It is wrong if reached, shown on the pre-fix code with a bias_score of 400: a LONG with its stop at 101.04 above a price of 100, which validate_risk_parameters passed, since it measures distance with abs(). A NaN bias_score gave a NaN stop with a SHORT's targets. Both are refused now. (10) Nothing at runtime checked the stop's side. It was planned as work order D and folded in as Claude's call: calculate_stop_targets is the only producer of a stop, and it now refuses a wrong-side one. detailed_bias is also removed from lineage.risk_inputs, which it never fed. The new tests include a grid over every bias score, trend health, volatility state and structural level the engine can produce. It asserts the stop's side, so the unreachability claim is checked rather than argued. Against the pre-fix code, 5 fail through the signature change alone, which proves only that the signature changed, and 2 fail behaviourally. A wrong turn: the NaN control first failed to fail, because the new wrong-side refusal caught it with a different message. The test now asserts the message names bias_score. Golden re-baselined in exactly the one predicted line; no decision field moved. code_hash ec88cf24... → ac02a155.... Finding 6 was evidenced twice from the record: the live run's stop was the HVN, 18.2% below price, where the ATR stop sits 6.2% below; the golden fixture's was 19.9% below. Confirmed on Windows: the next live record carries ac02a155... with no detailed_bias in risk_inputs.

#153

Work Order E: the Last Fabricated Defaults, and the Forming Candle Found From a Wrong Prediction

September 21, 2026

FINDINGS — FIXED

"The simulated order log is an obvious thing to build a backtester on, which is why it is fixed before that work starts."

Commit afd8460, review findings 12 and 13. (12) `live_trading._build_simulated_order` recorded an absent entry zone, stop, target set or price as 0.0, and an absent risk reason as "OK". It was the third copy of the shape Round 6 F3 removed from the router and the panel. Absent is now None, which is JSON null. Found while scoping: the order's timestamp called `datetime.utcnow()`, deprecated since Python 3.12. A 5 September note in the same file records fixing the same call twice elsewhere. This one was missed, and it was the source of the suite's two DeprecationWarnings. A test now sweeps every engine module's parse tree for `utcnow`. (13) `structure.calculate_structure` copied its analysis with `.get()` defaults of "NEUTRAL STRUCTURE" and 0.0 on keys that are always present. Had they fired, the HVN column would have held a price of zero. They are indexed directly now. Six negative controls each failed exactly their tests. The first draft of the counts said 527 and 392: the file has six tests and the arithmetic added five. The suite run caught it before delivery. `code_hash` ac02a155... → 3e76c1c5...; a docstring counts, since the hash is over the parse tree. The suite now prints no warnings line. Finding 16 was found from a wrong prediction. Claude said Viktor's next live run would repeat the 05:28 numbers, because it was the same candle. It did not: same stop, a new price, each target exactly 1R, 2R and 3R from it. The engine decides on the candle still forming. That was confirmed in `data/data_fetcher.py`, which keeps every kline and reads `close_time` only to discard it. Which candle counts is a rule, not a defect, so it is Viktor's to rule on with findings 4–7. Confirmed on Windows (filed at 49de810): the 06:18 run's record carries 3e76c1c5..., and pytest printed 528 passed with no warnings line.

#154

The Constitution Checked Against the Session's Work; Two Claims Reheaded as the Builder's Own

September 21, 2026

PROCESS — RECORDED

"Pausing the audit is consistent with the Constitution; pausing it and then starting backtesting is not."

Commit 49de810, docs-only. After afd8460 the end-of-session check found five things that existed only in chat, and this commit files them in PHASE7_NEXT.md. First, work order E's Windows confirmation (Entry #153). Second, SUMMARY: clean on afd8460's push. Third, a third live instance of finding 6: the 03:28, 03:51 and 04:18 UTC records all carry a stop equal to the HVN, 0.549596, 18.2%, 18.1% and 17.6% below price, and all three are NO-TRADE (RISK TOO HIGH). They were read from the records, not the panels. Fourth, Viktor's plan, stated that morning: about four more weeks of our own review, fixes and improvements before the next audit, with no time pressure. It is recorded as a plan in his words, not a ruling. Fifth, the Constitution check he asked for. Claude read the ratified PDF's rules and its Next Steps sequence, not its version history or glossary: nothing the session did breaks it. One condition bears on the audit decision: step 8 requires an independent re-audit of what changed before step 9, backtesting, and none of the session's fixes has had one. Where the session is thinnest is recorded too: two riders found mid-work (the BUSD suffix in B, `utcnow` in E), and finding 16 as a possible Item 3 or Item 8 gap, Unknown until ruled on. Following from the check, findings 14 and 15 were reheaded. Under "Checked and found sound" they read as the builder certifying his own compliance, which the Constitution does not allow. They are now "Claude's claims, open to the independent auditor". `code_hash` unmoved at 3e76c1c5....

"Make the adjustments you want and do what is best for the engine and the project"

Commit 3f263c2, review finding 11. long_signal and short_signal were computed on every run and carried into the decision object, the log and the simulated order, but read by nothing that decides. The golden fixture recorded long_signal: true beside NO-TRADE (RISK TOO HIGH). Claude set out three options (remove, confirm, relabel) and named the choice as Viktor's, because it changes which trades are taken. He chose confirm: a trade the ladder chooses is taken only if that side's signal confirms it, else NO-TRADE (SIGNAL UNCONFIRMED). He wrote the conditions after two rounds of Claude's critique. Macro is removed, as counted twice. The reversal check is ruled by its components: HVN proximity alone no longer blocks, divergence and exhaustion do, and only a reversal pointing against the trade blocks. The CONFIRMED requirement is dropped, so the signal follows raw_bias as decision_model does. When risk and confirmation both fail, RISK TOO HIGH stays the label and the confirmation failure is appended. He then delegated the rest, in the words above. Claude's calls under that delegation: one direction-aware divergence rule on every tier, replacing the upper tiers' direction-blind veto (CONSERVATIVE had none); trend health of at least 50 dropped from the signal, since every trading branch already requires it; and macro inside CONSERVATIVE left for its own commit, finding 17, work order G. Degradation overrides only actions naming a side. Claude had told Viktor during scoping that DEGRADED INPUT would override SIGNAL UNCONFIRMED; reading the code corrected that. The live log was checked first: of 29 records, no action would change. The one SHORT passes where the old signal refused it. Golden: three fields added, none changed, as predicted. code_hash 3e76c1c5... → 44f7296b.... 21 new tests; against the pre-fix code, only an ImportError, which proves only that the file is new. The behavioural evidence is seven negative controls. Found, not fixed: finding 17; finding 18 (the bias state machine now gates no trade); and the dependency that, if finding 6 stops pulling the stop to the HVN, nothing checks HVN proximity. Not backtested: nothing yet says the gate improves decisions.

#156

The Independent Audit Paused for Our Own Work — No Backtesting Before an Independent Re-Audit

September 21, 2026

DECISION — ADOPTED

"We pause the audit. We work on the engine another four weeks."

Ruled by Viktor and recorded in PHASE7_DECISIONS.md with work order F (commit 3f263c2). Asked to write his position first, he set out about four weeks of reviewing everything, analysing, fixing and evaluating our own work, with integrity and truth kept. After Claude's critique he added that we should not audit our own work, but we have fixes to make and should prepare a new audit as well as we can. What the ruling does not change: the Constitution's step 8 still requires an independent re-audit of what changed before step 9, backtesting. So no backtesting before an independent re-audit, and the open questions of the 20 September ruling (which model, which package, whether the auditor sees the scrapped findings) stay open until the audit is planned. Agreed with Viktor: nothing done in those weeks is written up as verification. Our own review, tests and negative controls are recorded as unaudited work awaiting the auditor, never as "found sound". Preparing the audit means three things: a running list of every change since the last audit, which becomes the auditor's scope (Entry #158); Viktor's open rulings answered before the package is sent; and the deferred read done first. Claude made two points, and neither was adopted. First, the pause had no end condition: "about four weeks" is a plan, not a trigger, and Claude suggested a date around 19 October to decide again. Second, the no-backtest rule exists only as text, where a structural form would be a backtest entry point that refuses to run without a recorded re-audit. An end condition was ruled on 22 September. It is recorded in PHASE7_DECISIONS.md in the same commit that brings this document to v1.34, and enters this document at its next regeneration. The second point remains open.

#157

F's Live Run Came After Its Commit — From Then On the Command List Stops at the Live Run

September 21, 2026

PROCESS — RECORDED

"Never predict live numbers; check the record."

Commits aafded0 and cb659f1, docs-only. F's command sequence asked for a live run before git commit. It did not happen: at the push the decision log still held 29 records, the newest on the old code_hash. Claude caught it by reading the log, not from a paste. The record of Viktor's 19:24 run carries 44f7296b..., so F was confirmed on Windows after its commit, not before, and is recorded that way. The record shows F's fields as designed. The hook's output on 3f263c2's push was not pasted, so a clean result was inferred from the push landing, not seen. A wrong prediction: the expected git status put the new test file last, but git sorts by path. The once-per-session rewrite was deferred to the next opening, agreed at 25% of Viktor's weekly usage. cb659f1 did it. It condensed the fixed findings to one line each while keeping their numbers, because test docstrings cite them by number. It wrote the carried lesson: on any change that moves code_hash, the live run and the panel read happen before git commit, and Claude checks the newest decision-log record for the new hash before giving the commit steps. It also corrected "working tree clean" to "last seen clean", since the hook's summary is not a git status. README's audit and backtesting rows were brought current in Viktor's first person and are his to reword. cb659f1's HISTORY move was proven by un-demotion but without the one-character negative control. That control was run afterwards (Entry #158): the proof holds, and the mutation is detected. The rule was followed at 2c7a7d1, 4e2b1c8, 486f1a5 and 3bfa6b7: each command list stopped at the live run, Claude read the record, and only then gave the commit steps. code_hash unmoved at 44f7296b... across both commits.

#158

The Deferred Read (Findings 19–27) and a Running Change List for the Next Auditor

September 21, 2026

PROCESS — RECORDED

“we check it after”

Commit 1cc2142, docs-only. The read Viktor had deferred (Entry #150) covered data/data_fetcher.py, data/validation.py, core/decision_log.py and core/lineage.py, read in full at cb659f1 with their call sites. It produced nine findings, 19 to 27, all Claude's, none changing a decision, none fixed in this commit. They are from reading the code, except finding 19, checked against Viktor's live decision log: 1 of 30 records (6 September) carries a bare NaN token. Finding 16, the forming candle, was confirmed again. New file docs/audit_change_list.md: the running list of every change since the last independent audit, which becomes the next auditor's scope. Its baseline is e65a0f7. The nine files round 6's fix-verification was sent match every commit from ac0a211 to e65a0f7 by sha256, and code_hash is 13908bba... at both (Python 3.12, computed). Each line gives what changed, what it closes, which tests guard it, and whether output moved. Every test file named was checked to exist. Two engine changes had no test at all, the SETUP DIRECTION box and its NEUTRAL branch (Entry #159). From here on, each commit changing engine code, tests or tooling adds its line in the same commit. A wrong claim was caught before delivery: the first draft said F was the only listed change that alters which trades are taken. But f24a6e9 moves bias_score on every run (78.70 → 77.10 on the fixture), so it can change actions near a threshold too; corrected. Two record questions were settled. The 119c8a3 counts came from an autocrlf clone, and the note the question leaned on was the patch-delivery skill's, stale since b68de08. Round 2's eleven observations are Kimi's, since HISTORY's later same-day section supersedes the earlier Qwen attribution. SUMMARY: clean on cb659f1's push (Viktor's paste). code_hash unmoved at 44f7296b....

#159

The SETUP DIRECTION Box Gets Its First Tests

September 21, 2026

PROCESS — RECORDED

“There is no pre-fix code — the box already behaves this way; these tests hold it there.”

Commit 9c4917c, tests only. It is the first of the six commits Viktor chose from the third session's list. The audit change list (Entry #158) found 66f1479 (the SETUP DIRECTION box, 14 September) and 39e0e79 (its NEUTRAL branch, Viktor's ruling the same day) to be the only engine changes since the last independent audit with no test. tests/test_setup_direction_box.py, six fixture-free tests, pins them. A BULLISH bias prints the LONG box and nothing else, and a BEARISH bias the SHORT box. A NEUTRAL bias prints no box and no SETUP DIRECTION line. NO-TRADE (PLAN CONTRADICTS ACTION) prints CONTRADICTORY and neither side, for both biases: that case keeps the 2 September “long label on a short plan” defect from coming back through this line. The box follows the bias on a run that takes no trade, which is why it was added, and it sits directly under the DECISION line. Against the current code all six pass, so the evidence is four negative controls on core/panel_render.py, each restored and confirmed with cmp: the CONTRADICTORY guard disabled (1 fails), a NEUTRAL line printed (1 fails), the box shown only when a trade is authorised (2 fail), and the box moved below Decision Reasoning (1 fails). code_hash unmoved at 44f7296b..., since tests/ is excluded by directory. Tests 549 → 555 with pandas_ta, run_tests.py's 32 errors unmoved. Confirmed on Windows by Viktor proceeding past the command sequence's stop conditions.

“Python reads it; jq and JSON.parse reject the line.”

Commit 05a12c7, deferred-read findings 19, 20 and 21, core/decision_log.py only among engine files. No decision changes. (19) write() used json.dumps with its default allow_nan=True, so the router's deliberate “not measured” NaN was written as the token NaN. It was checked against the live log: 1 of 30 records. Now _json_safe() copies the record with every non-finite float as null, and json.dumps runs with allow_nan=False, without modifying the caller's object. Records already in the log keep their NaN, and read() still accepts them. A numpy.float32 keeps its old path, the string “nan”, which is valid JSON; that is out of scope and said so in the docstring. The archive's half of the same defect was left for the lineage commit (Entry #162). (20) read() skipped any unparseable line anywhere, silently, though its comment named only a torn final line. read_with_report() now returns the records and the skipped line numbers. read() keeps skipping, so one damaged line cannot make the history unreadable, but now logs a warning naming the file and the lines. (21) The module docstring named “the pinned directory” as the recorded source; the engine records the literal “pinned”. Against the pre-fix code, 6 of 8 new tests fail: 4 behavioural, and 2 AttributeError that prove only that read_with_report() is new. Five negative controls on the module each failed exactly their tests. code_hash 44f7296b... → f691c4d8.... It was not confirmed on Windows by itself: the one live run was taken on the next commit's tree, which carries this code (Entry #161). There the record was strict JSON, but held no non-finite value to convert, so the null conversion is evidenced on Linux only.

"This list has fallen behind the code three times before, each fixed by adding names."

Commits 2c7a7d1 and 5d3a4b4, deferred-read finding 22, on its own because it moves `run_hash` and so the golden snapshot. `FINGERPRINTED_MODULES` listed `DecisionModel.BTC_ADJUSTMENT_CAP`, but not its sibling `BTC_STRESS_PENALTY`, nor `AVG_REWARD_R` and `EV_BREAKEVEN_BAND_R`. All three were inside `code_hash`, so a change was detected, but the record could not say which value a run used. They are added. The structural part is `tests/test_fingerprint_names_every_constant.py`. It does not name the three: it reads every module the list covers, and fails on any `UPPER_CASE`, numeric, finite constant at module or class level that is not listed. Its exclusions are stated in the test. Without `pandas_ta` the scan is skipped whole rather than narrowed, because a pass that silently left a module out would look like full coverage. Negative controls: one of the three names removed (2 fail), and a new constant added to `risk_model.py` (1 fails, naming it). The first attempt at the second control appended to a file with no final newline and broke its syntax, failing for the wrong reason. It was redone, and is recorded because a control that fails for the wrong reason proves nothing. The golden re-baseline changed exactly seven leaves: `run_hash` in two places, the archive name in two, and three added `module_constants` leaves. No decision field moved. An incomplete prediction is recorded: Claude had named only `run_hash` and the archive name, missing the three leaves that are the change itself, and that `run_hash` sits in two places. `code_hash` f691c4d8... → 82c14ef9.... Confirmed on Windows before the commit, by Viktor's 22:09 run, read by Claude from his disk: 31 records, the first 30 byte-identical to the earlier copy, the new one strict JSON on 82c14ef9... with the three constants at 15.0, 2.0 and 0.3. 5d3a4b4 filed that check and the hook's SUMMARY: clean, section 1 "none", on the push of 2c7a7d1, and brought README's counts current.

“An edit that rewrote both the frame and its digest passed.”

Commit 4e2b1c8, deferred-read findings 23, 24 and the archive's half of 19, core/lineage.py only among engine files. No decision changes. (24) `verify_archive()` compared each stored frame with the digest stored beside it in the same file. So an edit that rewrote both passed, and its docstring overstated what it checked. The check that means something, the archive against the decision log's input hashes, existed only in the tests. New `verify_against_record()` re-hashes each archived frame, trusting no stored digest, and compares the result with the record's input hashes and `run_hash`. Run on Viktor's live log on Linux, 30 of 31 records matched in every frame and in `run_hash`. The remaining one, from 6 September, carries no input hashes. Not built: a command-line wrapper, and a comparison against re-fetched data. (23) The archive is named by `run_hash`, which excludes `code_hash` by design, so a rerun on the same candles under changed code overwrites the file. It is documented and pinned by a test, not renamed, because the name is pinned by the golden snapshot and the record keeps its own `code_hash`. (19, archive half) `write_archive` now uses the decision log's sanitiser with `allow_nan=False`. This is latent, since only meta can hold a float and none is non-finite today. All 22 of Viktor's archives re-serialise byte-identically under the old and the new call (Linux). Not changed: `run_hash()` still uses the default `allow_nan`, because changing it could move the pinned `run_hash`, and no non-finite value reaches it. Against the pre-fix code, 10 of 12 new tests fail: 8 `AttributeErrors` that prove only that the function is new, and 2 behavioural. Six negative controls each failed exactly their tests. `code_hash` 82c14ef9... → 4fe5084e...; golden unmoved. Confirmed on Windows before the commit, by Viktor's 22:54 run, the 32nd record: the earlier 31 byte-identical, and the new record strict JSON on 4fe5084e.... Its archive had the same hash, and both verify functions returned all True. SUMMARY: clean, section 1 “none”, on the push.

“One bar, not zero: the exchange's clock and the local clock are two clocks.”

Commit 486f1a5, deferred-read findings 25 and 26, and 28, found while fixing 25; data/validation.py only among engine files. No decision changes. The commit also did the once-per-session rewrite of PHASE7_NEXT.md, with the move to HISTORY proven by un-demotion and a one-character negative control. (25) The staleness check bounded a series' age from above only, so a last candle dated after the reference time passed as current. A last candle more than FUTURE_TOLERANCE_BARS = 1 bar after now is now rejected as future-dated. Reachability on MEXC was not measured: the sandbox cannot reach api.mexc.com. (26) The timeframe table lower-cased its input, so MEXC's month “1M” became one minute, and MEXC's “60m” was unlisted, which switched the spacing and staleness checks off. The lookup is exact now, with “60m” and “1W” listed and “1M” deliberately unlisted, since a month is not a fixed number of minutes. The stated consequence: an upper-case spelling the table does not list, such as “4H”, now skips the checks where before it was lower-cased and checked. Nothing in the engine uses one. The evidence for MEXC's spellings is weak, and is said to be. (28) An aware now was stripped with tz_localize(None), which keeps the wall-clock reading, so a Stockholm now was read two hours late. It is converted to UTC first now. The engine's own caller passes naive UTC, but after 25 a relabelled now west of UTC would make a current series look future-dated, so the two land together. Recorded, not changed: an unlisted timeframe still skips the checks silently. Against the pre-fix code, 7 of 11 new tests fail, all behavioural. Five negative controls each failed exactly their tests. code_hash 4fe5084e... → 2c8ebe32...; golden unmoved. Confirmed on Windows before the commit, by Viktor's 23:17 run, the 33rd record: no fetch error, so the forming candle passed 25's check as predicted. The earlier records were byte-identical, and the new record and its archive strict JSON on 2c8ebe32.... SUMMARY: clean on the push.

"Fixing only the non-list branch would have left the usual case as it was."

Commits 3bfa6b7 and 5e55eb8, deferred-read finding 27, data/data_fetcher.py only among engine files. No decision changes. fetch_ohlc answered every reply that was not a non-empty list with "Empty or invalid API response.", so MEXC's own code and message were read and thrown away. Now a reply that is not a list is quoted: MEXC's code and message as that pair, anything else as its repr. An empty list is reported apart. The fix was widened past the finding as written, as Claude's call: a 4XX or 5XX reply's body is quoted too. MEXC documents 4XX for a malformed request, so that is where its message usually arrives, and requests' HTTPError text holds only the status line and the URL. Quoted text is cut at 300 characters, and the unused import time is removed. The text reaches the panel's "Data fetch failed" line only, because a failed fetch writes no decision-log record. What MEXC actually returns for a bad symbol was not observed, since the sandbox cannot reach it. Against the pre-fix code, 6 of 8 new tests fail, all behavioural; one fails on any unused import in the module. Five negative controls each failed exactly their tests. code_hash 2c8ebe32... → b3c2308f...; golden unmoved. Confirmed on Windows before the commit, by Viktor's 23:29 run, the 34th record: strict JSON on b3c2308f..., the earlier records byte-identical, and both verify functions all True. The fetch succeeded, so the new error text is evidenced on Linux only. 5e55eb8 filed that check, the hook's SUMMARY: clean on the push, and the session's close: the six commits for findings 19–28 and the direction-box tests are done. Filed with this regeneration, from 22 September: the push of 5e55eb8 printed SUMMARY: clean, section 1 "none", section 6 "installed". Section 5 reported README.md one commit behind, which was expected. The previous session's sandbox fetched the tip afterwards, and the verified close patch reverse-applied to it cleanly (from Viktor's message opening 22 September). This regeneration's own clone of GitHub's tip is 5e55eb8.

“We do the audit after we have done G and my findings.”

Ruled by Viktor on 22 September and recorded in PHASE7_DECISIONS.md, “Ruling, 22 September 2026 — when the audit resumes”, in commit e804b64. That commit is v1.34’s own regeneration, which by precedent has no entry; this entry is for the ruling it carried. It answers the first of Claude’s two points in the 21 September pause ruling (Entry #156): the pause had no end condition. Viktor wrote his position first, in the words above. Claude’s critique named three gaps. An open list could grow without end, since every new finding would push the audit back, which is the no-end-condition point again by another route. “Done” was undefined. And it was unclear whether the trigger replaced the four weeks. His answers closed all three. The list is closed: work order G, and findings 4, 5, 6, 7, 16 and 18. Anything found afterwards goes on a separate list for after the audit, kept in PHASE7_NEXT.md, and does not move it. His argument was that a fixed list is itself the end condition, and Claude accepted it and withdrew the objection. A finding is done when he has ruled it and any code his ruling calls for has landed. A ruling to leave something as it is counts as done. Claude proposed that definition, and he accepted it. The four weeks was his estimate with a buffer, not a rule, and no date is attached. Recorded but not ruled: a tension Claude named between “until ... it actually generates profit” and Goal B’s verdict-of-record and re-run clauses. Working toward a profitable engine fits Goal B one pre-registered phase at a time; re-running until a result passes does not. What the ruling weakens is stated in DECISIONS: a defect found during the work waits for after the audit, whatever its severity, because the ruling makes no exception. Claude raised that consequence before the patch was applied, and Viktor did not add an exception. Claude’s second point, that the no-backtest rule exists only as text, is untouched and open. The Constitution’s order is unchanged: an independent re-audit comes before backtesting.

"A commit cannot file its own push."

Commit c5dc4cd, docs-only, at Viktor's request for the session's finishing touches, nothing major. It filed the pre-push hook's result on the push of e804b64: SUMMARY: clean, section 1 "none", section 6 "installed", and section 5 README.md two commits behind, as predicted (Viktor's paste). The tip was fetched into the sandbox afterwards, and its six changed files matched the verified build byte for byte. README.md's Independent-audit row still said the next round was paused "for about four weeks of fixes and audit preparation", stale since the ruling in Entry #165. It now says the round is paused until work order G and Viktor's rulings on six open findings are done, with anything found meanwhile waiting for after the audit. The row is in his first person, like the rest of README, and his to reword. PHASE7_NEXT.md was amended in place, because the session's rewrite and HISTORY move had already happened at e804b64. Filed with this regeneration: the push of c5dc4cd printed SUMMARY: clean, section 1 "none", section 6 "installed", and section 5 README.md 0 commits behind, as predicted (Viktor's paste). The tip fetched afterwards reverse-applies the verified patch cleanly. Each such filing leaves the next push's result owed to the commit after it. That regress cannot be closed, only kept one step long. code_hash unmoved at b3c2308f..., computed on the applied tree under Python 3.12.3. The suite was unchanged at 597, 460 with 126 skipped, and 526 / 0 / 32.

Document History

Version	Date	Notes
v1.0	August 25, 2026	Converted from the standalone Engineering Note #1 into this standing log, at Viktor's request. Entries 1–3 carried over unchanged, split from one combined note into three individually tagged entries.
v1.1	August 26, 2026	Entries 4 through 9 added, tracking the Constitution's Rev 3 through Rev 5 and the new Tier0 Companion document. No earlier entry rewritten — Entry #9 corrects scope drift introduced in Entry #8's phrasing without altering Entry #8 itself, per this document's own no-silent-edits rule.
v1.2	August 26, 2026	Entry #10 added, closing Entries #1 and #6 — the credential gap is adopted into Constitution Rev 6 as Tier 1 Items 18 through 21, with operational detail in the new Phase7_Credential_Security_Protocol.pdf, and the audit now has a named independent auditor. New DECISION — ADOPTED status tag introduced for entries that record a decision actually made rather than a candidate awaiting one. Entries #1 and #6 left unedited; Entry #10 is where their closure is recorded.
v1.3	August 26, 2026	Entry #11 added: Viktor named Grok as the independent auditor for Steps 3, 4, and 8, once released, recorded in Constitution Rev 8. Answers the authorship-vs-judgment-independence question Rev 7 had just sharpened via Reviewer 4's identity.
v1.4	August 26, 2026	Entry #12 added, and format changed to hold it: a new gold-framed highlighted_entry_box treatment, used only when Viktor asks for an entry to be marked especially important. Documented in "How This Document Works" as a rare, deliberately non-standard marker rather than a new status tag. First use: a reflection on the audit loop's design, and the two places its soundness still depends on discipline rather than mechanism.
v1.5	August 26, 2026	Entry #13 added: criteria for evaluating any future auditor substitution — independence as a gate, power and cost as the actual trade-off, cost mattering because Step 8 makes the audit a recurring expense rather than a one-time one. Extends Entry #11 without editing it; Grok remains the named plan.
v1.6	August 26, 2026	Entry #14 added: Viktor ratified Constitution Rev. 8. Recorded in the Constitution's own Version History as a new RATIFIED row, not a Revision 9 — no rule text changed, only status. Scope freeze now in force for real; Step 2a of the audit sequence may begin.
v1.7	August 26, 2026	Entry #15 added: Gemini's assessment of the finished, ratified Constitution, logged and gold-highlighted at Viktor's request — second use of the highlighted_entry_box treatment, still rare. Not a structured finding and not the audit; flagged as such in the entry itself.
v1.8	August 26, 2026	Entry #16 added, recording how Entry #15's assessment was obtained — solicited by the party being evaluated — and the critical counter-review Viktor then commissioned. Entry #15 left unedited. Deliberately not gold-framed: the highlight marks Viktor's request for emphasis, not epistemic weight, and a counterweight entry that competed for attention would defeat its own point.
v1.9	August 26, 2026	Entry #17 added: Step 2a complete. Claude assembled the audit package from the actual working repository (19 source files, 4 evidence files not present in the public repo) at Viktor's instruction. Zero findings written, per Step 2a's own rule. Notes that Entry #3's file list is now superseded by the package's manifest, without editing Entry #3 itself.

Version	Date	Notes
v1.10	August 26, 2026	Entry #18 added: before running Step 3, Viktor noted Grok had already seen the Constitution in an earlier, separate conversation. Named as an independence risk on the same grounds as Entry #16, resolved by running Step 3 in a fresh conversation with no shared memory of the earlier one. The public GitHub repository was considered and deliberately left published.
v1.11	August 27, 2026	Entries #19, #20 and #21 added. #19 logs Grok's review of the Constitution as an external assessment, not Step 3, with its factual claims verified and its weaknesses named. #20 records Step 3 being paused and the auditor plan changing from Grok to a panel of models with no prior involvement, run pay-per-use; no revision of the Constitution was needed, since Rev. 8 named Grok as a plan rather than a requirement. #21 records an internal contradiction found in Claude's own drafting of the Minimum Viable Audit gate, deliberately recorded rather than repaired. Two new Constitution Version History rows (AUDITOR, DEFECT); no rule text touched and the register still frozen at 21 / 7 / 10 / 6.
v1.12	August 27, 2026	Entries #22, #23 and #24 added, covering the audit actually being run. #22 records Run 1, the blind source-only review by DeepSeek V4 Pro, whose ten findings were each checked against the real source and held. #23 records the first Step 3 attempt truncating inside its own reasoning at a default token ceiling, and the resulting split of the register across three scoped runs; the corresponding wording change is recorded as a new Section 6 in Phase7_Audit_Execution_Instructions.pdf rather than made silently. #24 records Run A, the Minimum Viable Audit gate: Items 2 and 18 Compliant, Items 3 and 6 Non-compliant, the Entry #21 DEFECT contradiction resolved by the auditor in favour of the four-item gate, and a false claim found in Claude's own Step 2a package. #24 also records, and withdraws, a false accusation Claude made against the auditor — a citation Claude called invented turned out to be accurate and located in the other evidence file, which Claude had not searched. Two genuine disagreements between Claude and the auditor are recorded unresolved and go to Viktor. No rule changed; register still frozen at 21 / 7 / 10 / 6, and the freeze stays in force until every Tier 1 item has a finding.
v1.13	August 27, 2026	Entries #25 and #26 added. #25 records Run B — the seventeen Tier 1 invariants outside the gate — with every substantive claim verified against the source, two precision errors in the auditor's evidence noted, two Compliant ratings disputed by Claude over the indicator cache, and the observation that two material Item 14 defects appear in Kimi's truncated first attempt but in neither run that completed. #26 records the scope freeze lifting: all 21 Tier 1 items now have findings, meeting the mechanical trigger Reviewer 4 forced into Next Steps at Revision 7. A new AUDITED row was added to the Constitution's Version History recording the same fact; no rule text was touched and the register stands at 21 / 7 / 10 / 6, no longer frozen. Four adjudications remain open for Viktor.

Version	Date	Notes
v1.14	August 27, 2026	Entry #27 added, recording Run C — Tiers 2, 3 and 4 — which completes the register: all 44 rules now carry a finding. Twenty-three items graded; four Tier 2 Non-compliances including a module that rewrites its caller's DataFrame in place, and a dependency manifest that omits two packages the code imports while declaring one nothing imports. Recorded as a method difference: Run C had web access and fetched the public repository, which Runs 1, A and B did not. Claude repeated that fetch independently and closed two Unknowns — no tags or releases (T3-7 becomes Non-compliant) and twenty-three commits on master (T3-9's depth question resolved). One contradiction between Run B and Run C is recorded with Run B's reading preferred, on source evidence.
v1.15	August 27, 2026	Entry #28 added, recording a hostile external review of the Constitution commissioned by Viktor from an uninvolved model. Two of its findings adopted into the Constitution the same day — a release gate, since the freeze-lifting definition established that the audit ran but attached no consequence to what it found; and amendment control, since the now-live amendment process would permit a Tier 1 invariant to be weakened by a short note reviewed only by Claude and Viktor. A front-matter statement of what the engine is was added alongside them. A third finding, Item 20 not covering crash-reporter capture of the process environment, is recorded and left unfixed: it requires amending a Tier 1 invariant, which the newly adopted rule says Claude may not do without external review. The review's remaining nineteen proposed invariants are declined as calibrated to a production trading system this engine is forbidden to become, and go to Future Amendment Candidates. Register unchanged at 21 / 7 / 10 / 6.
v1.16	August 29, 2026	Entries #29 through #33 added, covering the work between the audit closing and the remediation sequence arriving. #29 records Phase A — eighteen tests built and run, five passing, thirteen failing as written acceptance criteria, and the golden-path test that was shipped unverified and turned out to be wrong. #30 records the three defects that first run found, which four audit passes across three models had missed: a severity escalation on the clone failure, and two findings nobody had. #31 records the model-independence table in one place for the first time, and the rule that independence is tracked at the lab rather than the checkpoint. #32 records machine learning put on ICE with five written conditions for revisiting. #33 records Step 5 running on GLM 5.3, nine of nine source claims verifying, and Claude coming one sentence from a third false assertion of absence. Register unchanged at 21 / 7 / 10 / 6; the adjudications Step 5 asks for remain open.

Version	Date	Notes
v1.17	September 2, 2026	Entries #34 through #42 added, covering the six days between the audit closing and this row — a period this log had not recorded at all. #34 the machine rebuild and the golden baseline surviving it across a numpy major version. #35 Step 8, the independent re-audit: five Criticals, ten further findings, release gate not met. #36 remediation batches 1 and 2, including 46 tests that returned where they should have skipped. #37 items 3, 11 and 14 delegated by Viktor and ruled. #38 the macro degradation, dependency pinning, the four unguarded tests and two stale docstring labels. #39 Finding 3 — a Critical that had never been entered into any roadmap, found by reading the audit report rather than the summary derived from it, with a defect class wider than the two instances the report named. #40 the status sweep: three findings already closed by commits that never named them, and a count wrong by three. #41 the Step 8 auditor not being on the clean list, recorded rather than quietly noted. #42 the full document set read end to end, correcting four things no working document held. Register unchanged at 21 / 7 / 10 / 6, no longer frozen. Two amendments owed, neither adopted; the release gate stands closed on Item 6.
v1.18	September 2, 2026	Entries #43 through #49 added, covering the day the last Critical closed and a new one opened. #43 Findings 6 and 7 — input hashing, a raw-candle archive pruned at ninety days on Viktor's ruling, and the Item 6 lineage chain persisted. #44 a defect that shipped: a path separator that made the golden snapshot match only on the platform it was baselined on, caught on Viktor's machine on the first run. #45 the direction-source Critical — a CONSERVATIVE LONG printed over a stop above price and three descending targets, found by running the engine on live data, and missed by an independent 44-rule audit plus four passes across three models. #46 an unwritable log directory destroying a whole run, with Claude recommending refusal and being overruled. #47 a generated fixture that lets the suite reach the state that hid #45. #48 the session / training / lineage distinction ruled, and the independence ledger corrected from the billing log — Mistral had been wrongly listed clean for a week. #49 the round-2 audit package generated rather than assembled, and eight tools left enabled on the first attempt at it. Register unchanged at 21 / 7 / 10 / 6. Every Critical the Step 8 audit raised now has a fix that has landed; the release gate stays shut until a re-audit has seen them.
v1.19	September 3, 2026	Entries #50 through #56 added. #50 the re-audit's third failed attempt and the Constitution shipped as text, because two providers on two days could see its filename and not its contents. #51 eleven defects returned by an audit that produced no report, every one checked against source before it was believed. #52 the run hash that did not cover the risk multipliers — Item 6 broken inside the fix written to satisfy it. #53 a test whose failure message described a live defect it could not detect. #54 the volume-node fabrication reached through the producer after Finding 3 closed it at the consumer. #55 two defects found by reading the panel in one morning, one of them shipped by Claude an hour earlier and passed by all 226 tests. #56 the handover check, and the project goal split into a tagged portfolio milestone and a later finished engine. Suite 196 to 226. Register unchanged at 21 / 7 / 10 / 6. Earned rules 30 through 34 written the same day.

Version	Date	Notes
v1.20	September 5, 2026	Entries #57 through #70 added, covering 4–5 September — the three open rulings built and the whole of the “What is left to fix” table closed, across eight patches. #57 ruling 3, the bias engine inferring trend direction from the sign of continuation strength and reading a real trend as flat, measured across 9,800 live bars before the ruling rather than after. #58 rulings 1 and 2 — trend health, a magnitude, offered in reason strings as directional support, and a minimum bias strength kept as its own fingerprinted constant rather than folded into an existing threshold. #59 three patches diffed against a base that had already moved, and earned rule 30 violated in a test written to guard its own class. #60 the usage discipline Viktor delegated: reasoning goes where it survives, verification is not cut. #61 two fabrication constants that survived two passes because they were unreachable, one of which would have killed the run rather than degraded it. #62 a fallback whose docstring claimed an equivalence it did not have, and the test that had asserted on that claim for as long as it existed. #63 a degraded-dependency control that reported two errors and executed nothing, and a module-level object that defeated a deliberate lazy import. #64 correlation and beta paired by position after both timestamp indexes were discarded — correct by accident, and the printed label moving in 4 of 31 measured positions once the series differ by one bar. #65 an entry zone printed with its lower bound above its upper, found by reading the panel, and a close price that could not be read becoming \$1.00. #66 five entry sub-scores summing to 39 under a printed total of 45.18, and the clip that was hiding a component list adding to 102 against a docstring saying 100. #67 the entry block’s shape declared in four places, recorded rather than restructured inside a patch about something else. #68 the four-for-four generalisation: the fabrications that survive audits are the ones on paths a healthy run never takes. #69 two rejected CONNECTs to an exchange from a suite in which every route was pointed at a dead port, source unidentified, recorded rather than chased. #70 suite 226 to 319 passing and the fix table empty, with the release gate still shut because unresolved means fixed and re-audited. Register unchanged at 21 / 7 / 10 / 6. Earned rules 35, 36 and 37 written across the same two days.
v1.21	September 5, 2026	Entries #71 and #72 added, both corrections to this project’s own record rather than to its code. #71 applies earned rule 29 a second time: the provider’s billing log shows three Qwen calls followed by Kimi K3 on 2 September, so Entry #50 has the last two attempts of round 2 in the wrong order and the eleven observations of Entry #51 came from a Qwen run that preceded Kimi rather than followed it. The same log establishes that no attempt at round 2 ever received the UPLOAD_THESE folder — it was written after the last of the four calls — and corrects Entry #22, which names a DeepSeek checkpoint the bill says was never called. Claude’s own inference from file modification times is withdrawn in the entry. Luna Pro’s call count and Kimi K3’s spent status are recorded as open questions, not resolved. #72 records a defect found by searching the round-3 package before sending it: Section 2 of the reviewer instruction told the auditor to stop if the code bundles contained prior-audit outcomes, and they contain a stated verdict for seven of the forty-four rules, three counts of Criticals, and two named prior reviewers. Viktor ruled the condition scoped to the Constitution file and the bundles’ contents counted and disclosed, with the audit narrative to move out of code comments into this document later — not inside the audit it would affect. Register unchanged at 21 / 7 / 10 / 6. The release gate remains shut.

Version	Date	Notes
v1.22	September 5, 2026	Entries #73 through #82 added, covering the afternoon and evening of 5 September — the day's second half, in which nothing in the engine changed and most of what this project believed about its own audit did. #73 the package-search step created by Entry #72 finding, on its first run, that Entry #72's own disclosure had under-counted what the bundles leak, in three separate ways. #74 a commit count read off a generated file three hours after it was generated: the third instance of a figure treated as current because the artifact holding it looked trustworthy. #75 the largest correction this project has made — the four transcripts named for Qwen are Kimi's, identical for 57,631 characters, and the passage where the model says so never reached the repository, so Entry #71's own correction was valid reasoning to a false conclusion. #76 an Auto Router left on by accident sending the round-3 package to GLM 5.3 Flash, which returned a complete report; ruled to be round 3 rather than discarded, with one finding verified as under-scoped in a way that reaches the trade-authorization gate. #77 the ledger rebuilt from all 723 provider rows: Z.ai not clean and so the reviewer that ran was not independent, Kimi cleared on the check that mattered with its exposure thirteen times larger than recorded, and several paid-for outputs that were never saved. #78 the handover check's structural blind spot — ignored files do not appear in git status — and reviewer responses moved to a tracked directory where the safe location is the only one on offer. #79 rev 5: the instruction stops telling the reviewer who it is, and Part 8 is removed because its comparison document was the reviewer's own reasoning. #80 three PDFs checked against the repository before deletion, one holding a paragraph of ruling reasoning that existed nowhere else. #81 a commit message stating an impossibility where an undone task belonged, found because Viktor asked why. #82 Kimi ruled for round 4 on two grounds, with a third — that it is an important model — argued out and recorded as rejected. Register unchanged at 21 / 7 / 10 / 6. The release gate remains shut, and the round-4 package is built and waiting to be sent.

Version	Date	Notes
v1.23	September 6, 2026	Entries #83 through #90 added, covering the sixth session of 6 September — the round-4 confirmation run and the four fixes it and the two reviewers' reports led to, plus the two defects found by running the suite and reading the document set rather than by review. #83 the Section 11 confirmation run, made before any fix, confirming Finding 1 end to end and finding a second Item 6 gap neither reviewer filed. #84 Finding 1 fixed: a producer's not-measured value now passes through instead of being defaulted, and a failed merge's record keeps its lineage. #85 GLM F-7 fixed as extended: a risk block that was never assessed could authorize the engine's most aggressive action, closed by a validator that reads the verdict rather than assumes one. #86 Kimi Finding 5 item 6 fixed, and the live run it and F-7 were both owed: a swing-structure field invented its own current price as a location, the third time this class of defect was patched at a consumer while the producer kept inventing. #87 Kimi Finding 3 fixed — a source hash over the parse tree of every engine file, replacing a list of names nothing of that shape could hold — and a live run at 10:51 that found a second, independent route to the false-sentence class Kimi Finding 2 describes, filed as input to a decision rather than fixed. #88 the suite was destroying its own decision record: four routes, one of them a directory delete, found by running a clean clone rather than by reading, with the 10:04 live run's raw record among what could not be recovered. #89 the document set audited end to end: the audit-package record verifies byte-for-byte, no build script could write into the repository, and the semantic half — the eight PDFs read against current state — remains undone. #90 the build scripts fixed: seven of nine now write into the repository directly, the other two fail early and say why, and publishing the entries in this row became one command. Suite 319 to 412 passing across the batch; the golden snapshot moved once, in exactly the five fields predicted. Register unchanged at 21 / 7 / 10 / 6. The release gate remains shut — nothing that landed this batch has been re-audited, and eight decisions remain open, all Viktor's.
v1.24	September 9, 2026	Entries #91-#93: fourteen-item sweep, Kimi Finding 2 fix, Claude outputs fingerprint exclusion.

Version	Date	Notes
v1.25	September 11, 2026	Entries #94 through #97 added, covering the fourth commit batch of 11 September. #94 the four qwen_reasoning_*.txt transcripts moved from the repo root into docs/ (bb6d222), discharging the hold noted at Entry #93. #95 this document's own body reconciled with 76380ec a second time, in a new shape (4705d03): the 7-9 September Grok session had updated the head block and this builder but left "Open - decisions" #3 and "Open - work" #2 describing the sweep and Kimi Finding 2 as open after both had landed, docs-only and code_hash-neutral by directory exclusion. #96 docs/Phase7_Sweep_Session_Commands.pdf committed after sitting untracked on Viktor's machine (e75f7be). #97 Item 14 / decision 2 ruled and closed (0c7dec5): classify_risk_regime now reads ADX instead of trend_health, so the risk gate no longer shares its largest input (0.30 of bias_score) with the conviction number it is meant to check independently; REGIME_LOW_TREND_HEALTH/REGIME_HIGH_TREND_HEALTH renamed to REGIME_CHOP_ADX = 20.0 / REGIME_STRONG_ADX = 25.0, values already load-bearing in trend_health.py for the same two market states. code_hash moved and was confirmed on three trees, 47bc557119bb0350... to 2741062fae070f61..., matching Viktor's own Windows value. Suite 412 to 421 passed under pytest (0 failed both sides), 345 to 354 passed under run_tests.py with the watched 29-error count unmoved. Golden re-baselined at six sites: four predicted (module constants, risk_inputs, run_hash twice) and two not (the archive filename, built from run_hash, which should have followed the run_hash prediction mechanically and did not) - recorded as an incomplete prediction rather than a defect. No decision field moved on the pinned fixture, which exits before classify_risk_regime is reached; a live run the same day (AEROUSDT 4h, 22:50) did not exercise the new branches either, for two independent reasons, and is recorded as that rather than as validation. Found by reading the panel and left unfixed: RISK REGIME prints without the ADX number now deciding it, a transparency regression of the same class as GLM F-5. Register unchanged at 21 / 7 / 10 / 6. The release gate remains shut - nothing landed since round 4 has been re-audited.

Version	Date	Notes
v1.26	September 11-12, 2026	Entries #98 through #103 added, covering the round-5 send through all three Criticals landing as fixes. #98 (4d99cdb) GPT-6 Astra ruled for round 5 after Viktor's billing-export check cleared both of Luna Pro's exposed sessions as variant=standard, no training routing; send_audit_round.py and build_audit_package.py repointed, Part 7 folded into the single upload set per ruling 3. #99 (de7d135) the 2 September hedge closed for good: all nine Luna Pro OpenRouter calls checked individually, all standard tier, "No data training" -- evidence upgraded from "probably clean" to provably. #100 (a12640a) round 5 sent to GPT-6 Astra (\$12.21636, 435,612 prompt tokens, confirmed twice over); three new Criticals returned, none previously on record -- F1 (volatility never reaching stop construction), F2 (entry quality scored for the wrong direction), F3 (a stale decision-bar VWMA presented as current) -- all three independently reproduced against live code this session and RULED to fix, no limitation accepted for any. #101 (83e334d) F1 fixed: the production call to calculate_stop_targets() omitted volatility_state, defaulting to NORMAL regardless of measured conditions; now passes the same volatility_mode its sibling call already receives. Golden snapshot unmoved -- the pinned fixture's own volatility state produces the same multiplier the bug's default did. #102 (3f92230) F2 fixed: eq_trade_direction was derived from short_signal, which zeroes for reasons unrelated to direction, letting a fully bearish setup score as a long entry; now derived from raw_bias directly, the same signal decision_model.py's own action selection reads -- the identical ruling already made for that module on 2 September, one caller downstream. #103 (7fcf646) F3 fixed, closing round 5: VWMA's own forward-fill had no trailing-edge exception, so a decision-bar miss was silently carried forward from the previous bar; fixed at the source and VWMA added to the critical_indicators sweep, with three new regression tests. Accepted consequence: some currently-clean runs will now show DEGRADED, naming VWMA. code_hash moved once per fix, confirmed on independently built trees at each step: e774b93d... to f8e4cc7d... to 8b0e97fd... to 00c8d4ea.... Suite 429 to 432 passed under pytest, 362 to 365 under run_tests.py, watched error count unmoved. Golden snapshot unmoved across all three fixes, confirmed by running it each time. Register unchanged at 21 / 7 / 10 / 6. The release gate remains shut: fixed is not re-audited, and unresolved means both.
v1.27	September 13, 2026	Entries #94 through #103 reconstructed from the v1.25/v1.26 Document History rows' own text into entry_box() form — a document-integrity gap found where those two rows claimed the entries existed as numbered log entries and they did not.
v1.28	September 13, 2026	Entries #104 through #111 added, covering 13 September: Round 6 (Meta Muse Spark 1.3) sent and graded clean, Round 6's F1/F2/F3 fixed, the docs and handover-file commits, F4 and F5 fixed. code_hash and the golden snapshot addressed per fix within the entries above; nothing in this row's own change touches either.

Version	Date	Notes
v1.29	September 13-14, 2026	Entries #112 through #118 added, closing the gap this document had carried since Entry #111 -- PHASE7_NEXT.md's own ninth-patch note had flagged this file as stale, last regenerated before Round 6. #112 (7fbf77c, 9c5f954, 4e0110b) the encoding defect that corrupted round 6's report.md found and fixed at the source -- OpenRouter's response carries no explicit charset and requests silently fell back to Latin-1; report.md corrected, three more evidence files caught untracked the same day. Appended out of strict date order, deliberately: it precedes Entries #104-#111 chronologically, matching this document's own Entry #71/#72 precedent for corrections filed when found rather than renumbering. #113 (214c5d8) the git add -A trap staging two of Claude's own delivery files a third time, closed structurally -- *.patch and *_commit_message.txt gitignored -- rather than a fifth restatement of rule 38. #114 (2fc578b) send_fix_verification.py built: a narrower tool for verifying round 6's own findings were fixed, deliberately not sharing send_audit_round.py's repointed-every-round MODEL constant. #115 (59de938, e65a0f7) round-6 fix-verification sent and graded: F1 through F5 all Fixed, independently checked against code, plus one bonus finding (a duplicate @staticmethod) found and not yet fixed at that point. #116 (a70fc1b) the bonus finding fixed. #117 (66f1479) Viktor's request: the setup's LONG/SHORT direction stated in words under the DECISION line, sourced from the same field the actual decision logic uses so it cannot disagree with an authorized trade, with a deliberate non-answer for the one case (PLAN CONTRADICTS ACTION) where stating a side would repeat the 2 September wrong-direction-label defect. #118 (39e0e79) Claude's own unrequested NEUTRAL-case line overruled by Viktor -- silence instead, since a neutral bias is not a failed measurement. Of the ten commits in this row, seven -- #112's three, #113, #114, and #115's two -- are docs/tooling-only and left code_hash and the golden snapshot untouched, confirmed rather than assumed in each case; the other three (#116, #117, #118) each moved code_hash by exactly the amount predicted in their own entries, and none moved the golden snapshot.

Version	Date	Notes
v1.30	September 14-15, 2026	Entries #119 through #127 added, closing the gap since Entry #118 -- ten commits landed without a fresh regeneration, batched per the rule Entry #122 itself records. #119 (69f1972) the Portfolio Document built, the fourteen-section deliverable drawn from current source rather than the (found-stale) README.md, with a Luna-Pro/GLM-5.3 misattribution caught and fixed before Viktor ever saw it. #120 (eee218e) the AI-Attribution Statement built, closing item 5 and -- with #119 -- all six portfolio-ready criteria; its own entry-count table caught and corrected (111 vs the true 118) by parsing this file's AST directly rather than trusting a regex. #121 (59b915a) this document's own entry-point brought current for #119/#120, a handover-check gap otherwise. #122 (48757fe) five token-saving workflow rules recorded in Working practice, already adopted in the patch-delivery skill and memory the same day. #123 (f7f2e51) a correction, not new work: the five round-6 mutant-escape findings were already closed at 044b055 three days before this session existed; Claude had reported them open from this file's own narrative rather than from git. #124 (99e022e) Viktor's ruling: the release gate declared open, the project declared portfolio-ready, tagged portfolio-v1 -- closing the question Entries #121 through #123 all left open. #125 (63d8620) docs/build/session_handover_check.py built, answering four of the standing checklist's seven items mechanically; confidence-scored auto-pruning of the Lessons list considered and declined, on the project's own record-don't-hide-errors rule. #126 (ebb0a46, d5bb329) README.md brought current after sixteen days stale -- it still said the gate was shut the day after Entry #124 declared it open -- plus a stray gitignored delivery file found and left for Viktor to remove, and Goal B's backtesting-scoping findings written down rather than left only in chat. #127 (a441023) item 8 added to the handover checklist closing the README-currency gap #126 exposed, and the ignored-file-sweep question closed with no code change since #125's own script already ran that sweep every invocation. All ten commits in this row are docs/tooling-only or pure PHASE7_NEXT.md prose; code_hash confirmed unmoved throughout (38458f20779d2709ed6402c69407339f797c69089ace3b18922ca5c8bb289fbb, computed on the tip this document was regenerated from, matching every commit's own individual confirmation) and the golden snapshot was never in play. The release gate itself moved this row -- shut at the start of it, open and tagged by #124 -- the first row in this document's own history to record that.

Version	Date	Notes
v1.31	September 15-19, 2026	Entries #128 through #135 added, closing the gap since v1.30 -- eleven commits landed without a regeneration, batched per the rule Entry #122 records and regenerated at Viktor's call. #128 (dff7d00, f9e5127) Goal B specified and ratified before any of its code exists, with two of Claude's own errors of asserting from the documents instead of the code recorded inside it. #129 (4a97c32) the course correction -- an engine review for analytical soundness before any Goal B work -- and Viktor's pace ruling. #130 (4a97c32) PHASE7_NEXT.md split into NEXT, DECISIONS and HISTORY, proven a move by byte-for-byte reassembly, and the bias weights confirmed fingerprinted by reading the code. #131 (5be5d82, 2ebbd6b) three Tier-1 fixes, PHASE7_PINNED_DATA's silent live fallback first among them. #132 (f24a6e9, 2fee78f) ADX's second read removed from continuation_strength on Viktor's independence-only ruling. #133 (2fee78f) git add -A nearly sweeping an untracked PDF, and files added by name from then on. #134 (e2c6637, 742b514) four of the six bias factors found to measure one thing, recorded and not fixed, with two errors in the project's own record corrected. #135 (2f5fdaa, 5c73e24) the pre-push hook, on Viktor's ruling to stop and consult. #129 and #133 were split into entries of their own at Viktor's choice. The first row in this document's history in which code_hash moved, and it moved twice: 38458f20779d2709ed6402c69407339f797c69089ace3b18922ca5c8bb289fbb, unchanged on every commit from portfolio-v1 through 4a97c32, to 5080ccf0490502ad72dadaf6c3c525d19792139ea52d7fb15ee28a923ab5d189 at 5be5d82, to bb47ab537314953e16b2db2fcf24003ead64eb5538a30221a6578d518bb7d34d at f24a6e9, unchanged since -- each recomputed on that commit's own tree for this regeneration rather than taken from the commit messages. The golden snapshot was re-baselined once, at f24a6e9, 16 leaf fields as predicted. Entries before the 18 September split that cite PHASE7_NEXT.md for material now in PHASE7_DECISIONS.md or PHASE7_HISTORY.md are left as written, per this document's own no-silent-edits rule.
v1.32	September 19-20, 2026	Entries #136 through #140 added, closing the gap since v1.31 -- six commits landed without a regeneration, batched per the rule Entry #122 records and regenerated once the three mechanical items before the independent audit had all landed. #136 (e431714, e115272) PHASE7_NEXT.md made current at c7ced36, and a wrong prediction about git status recorded. #137 (b68de08) the pinned-data manifest check made portable, closing the standing LF/CRLF platform difference. #138 (16d3c1f) Viktor's ruling scrapping the seven open items in favour of an independent audit. #139 (9a35f1c) stale citations of PHASE7_NEXT.md fixed in live files, and the ruling that dated records -- this document included -- are not edited to follow a move. #140 (077d120) README.md brought current. code_hash bb47ab537314953e16b2db2fcf24003ead64eb5538a30221a6578d518bb7d34d on every one of the six commits, each recomputed on its own tree for this regeneration. Golden snapshot unchanged since f24a6e9. Phase7_Portfolio_Document.pdf and Phase7_AI_Attribution.pdf regenerated in the same commit, their text differing from the previous build only by 9a35f1c's pointer changes (compared word by word).

Version	Date	Notes
v1.33	September 20, 2026	Entry #141 added for 4629002, the one commit since v1.32, regenerated at Viktor's call rather than left for the next batch. #141 (4629002) the round-1 audit outputs recorded as lost, and README.md and PHASE7_HISTORY.md's 30 August reinstall entry corrected. code_hash bb47ab537314953e16b2db2fcf24003ead64eb5538a30221a6578d518bb7d34d on 4629002's tree, recomputed for this regeneration. Golden snapshot unchanged since f24a6e9. Phase7_Portfolio_Document.pdf and Phase7_AI_Attribution.pdf were rebuilt from 4629002 and their extracted text compared with the committed PDFs: identical, because 4629002 changed neither script. They were therefore not regenerated, and PHASE7_NEXT.md's statement at 4629002 that they were one commit behind is corrected there.
v1.34	September 22, 2026	Entries #142 through #164 added, closing the gap since v1.33 -- twenty-seven commits landed without a regeneration, by Viktor's choice under the batching rule Entry #122 records, and regenerated at his call on 22 September, before work order G and his rulings on findings 4, 5, 6, 7, 16 and 18. 65a0aef had planned the regeneration for after the audit rulings and only if the audit package included the Notes (Entry #144); the audit is not yet planned, so this regeneration is Viktor's later call and not that condition met. 3a899b5, v1.33's own regeneration, has no entry, as c7ced36 and c745a67 had none. Twenty-three entries for twenty-six commits: #142 and #143 (53394ff) round 3 ruled not independent; the suite's live call left in place and the empty rule number marked. #144 (92775ea, 982e70f, 65a0aef) landing facts, and why the Notes gap cannot fall below one. #145 (a9d4b1f) the run summary named the EV line. #146 (6e1baba) the four-runs independence claim corrected. #147 (b869a30) code_hash comparable only within one Python minor version. #148 (119c8a3) the TREND denominator. #149 (635a94e) the entry-point document's eleven corrections. #150 (ebb4e5c) the pre-backtest review and the delegated order of work. #151 to #153 (a530006, e3f3d51, afd8460) work orders B, C with D folded in, and E, with finding 16 found from a wrong prediction. #154 (49de810) the Constitution check. #155 and #156 (3f263c2) work order F, and the audit paused. #157 (aafded0, cb659f1) the live run moved before the commit. #158 (1cc2142) the deferred read and the audit change list. #159 (9c4917c) the SETUP DIRECTION tests. #160 to #164 (05a12c7, 2c7a7d1 with 5d3a4b4, 4e2b1c8, 486f1a5, 3bfa6b7 with 5e55eb8) findings 19 to 28. Each written from its commit's own message, and #155 and #156 also from PHASE7_DECISIONS.md's two 21 September rulings; none re-checked against the code or the running engine. code_hash recomputed under Python 3.12.3 on every one of the twenty-seven trees for this regeneration, and every value matched its commit message: bb47ab53... through 65a0aef; b0213766... at a9d4b1f; 35718f6b... at 119c8a3; ec88cf24... at a530006; ac02a155... at e3f3d51; 3e76c1c5... at afd8460; 44f7296b... at 3f263c2; f691c4d8... at 05a12c7; 82c14ef9... at 2c7a7d1; 4fe5084e... at 4e2b1c8; 2c8ebe32... at 486f1a5; b3c2308f8f3e05981af25ee82468c071f7bf0b9d49519b6e75bd78c37b5fb365 at 3bfa6b7, unchanged at 5e55eb8 -- eleven moves. The golden snapshot was re-baselined four times (a9d4b1f, e3f3d51, 3f263c2, 2c7a7d1), read from git log of the fixture. Phase7_Portfolio_Document.pdf and Phase7_AI_Attribution.pdf not regenerated: their scripts last changed at 6e1baba, which regenerated the Portfolio Document, and neither script reads another file.

Version	Date	Notes
v1.35	September 22, 2026	Entries #165 and #166 added, at Viktor's call the same day as v1.34. #165 is the 22 September ruling on when the audit resumes, carried by e804b64; e804b64 itself is v1.34's regeneration and, by the precedent #144 states, has no entry of its own, so the entry is for the ruling it carried, not the commit. #166 (c5dc4cd) e804b64's push filed and README's audit row brought current. The gap goes from two commits to one, the floor. Both entries were written from their commits' own messages and PHASE7_DECISIONS.md's 22 September section, not re-checked against the code or the running engine. code_hash b3c2308f8f3e05981af25ee82468c071f7bf0b9d49519b6e75bd78c37b5fb365 at both commits, recomputed under Python 3.12.3 on c5dc4cd's tree; e804b64's was computed on its applied tree before delivery. Golden snapshot unchanged since 2c7a7d1. The portfolio PDFs were not rebuilt, as at v1.34.

Future entries get appended below this line, in order, each with its own number, date, and status tag — this document grows with the project rather than being rebuilt each time.