Phase-7 Engineering Notes

A Running Log of Observations, Ideas, and Notes for the Phase-7 Project

---

Status: **Living document — updated as new entries are added**

Started: **August 25, 2026**

Origin: **Began as a single standalone note (credentials gap + pacing observation); converted into this standing log at Viktor's request, so future notes and ideas get appended here instead of each needing a new document.**

Relationship to the Constitution: **entries here are never automatically part of it.** The frozen 21 / 7 / 10 / 6 scope (ratified August 26, 2026) only changes through the Constitution's own process — see “How This Document Works.”

> *“Evidence gets the final vote, on both sides of that relationship” — including when the evidence is a note neither of us would have thought to write down on its own.*

Contents

How This Document Works

Entry Log

Document History

How This Document Works

A place for anything worth writing down that doesn't belong in the Constitution and isn't substantial enough — yet — to deserve its own standalone PDF. New entries get appended at the end, in order, as they come up.

Every entry gets the same shape: a number, a date, a short title, a status tag, and a body. Entries are never rewritten after the fact — if something changes or gets resolved, that becomes a new entry that references the old one by number, the same no-silent-edits discipline the Constitution holds itself to.

Status tag	What it means
CANDIDATE — NOT ADOPTED	A possible future addition to the Constitution. Logged for the record, not adopted — the scope freeze means it waits for the audit, the same as every other future candidate.
OBSERVATION — FILED FOR THE RECORD	A reflection, insight, or pacing note. Not asking for any action — just worth having written down rather than left in chat history.
REFERENCE — FILED FOR THE RECORD	Practical, factual information worth keeping — a file list, a readiness confirmation, something a future entry or the audit might need to point back to.
RESOLVED — SEE ENTRY #N	Marks an entry whose open question has since been settled elsewhere, pointing to the entry that resolved it, rather than editing the original.
IDEA — UNEVALUATED	A raw feature idea for the engine itself. Not a Constitution candidate, not an observation about the project, not a settled fact — just a concept written down before it has been scoped, designed, or checked against the Constitution.
DECISION — ADOPTED	A decision that has actually been made and acted on, usually closing one or more earlier entries. Records what was decided, by whom, and what it changed — including any judgment call Claude made rather than Viktor, so it can be reversed knowingly rather than discovered later.
★ Gold-framed entries	Not a status tag — a small number of entries Viktor has asked to be marked especially important get a gold frame and a kicker line instead of, or on top of, their normal status tag. Rare on purpose. If it stops being rare, it stops meaning anything.

None of this bypasses the Constitution's own process. A CANDIDATE entry becomes part of the Constitution only the way any future amendment does — proposed once the audit has run and the scope freeze is deliberately revisited, not by accumulating enough notes here.

Entry Log

In order added. Entries 1 through 3 below were the original standalone Engineering Note #1, split into individual entries when this became a standing log.

#1	Credentials & Secrets Handling Gap August 25, 2026	CANDIDATE — NOT ADOPTED

“API keys, exchange access tokens, and any other credentials the engine or its operator relies on must never be logged, hardcoded, committed to version control, or surfaced in decision-support output.”

Why it's being logged: Nothing in the Constitution, and nothing in any of the three outside assessments, addresses credentials handling. This project is explicitly headed toward being run by other people, potentially with their own exchange credentials in play. A leaked key is a categorically different kind of failure than a wrong prediction — it's irreversible and harms the customer directly, not just the engine's credibility. Sits close in spirit to Tier 1, Item 1 (“Analysis ≠ authority”): the engine already isn't trusted to act on markets on its own; it certainly shouldn't be trusted to mishandle the keys that would let it, or an attacker, do so. Proposed tier, whenever this is actually considered: Tier 1, alongside “Tool, Not Autonomous Actor” — though exact placement is a judgment call for the audit, not a decision made here.

#2	Project Pacing, Mid-Review August 25, 2026	OBSERVATION — FILED FOR THE RECORD

Said plainly, without alarm: this is a pacing observation, not a criticism. Nearly everything built in this project's governance thread so far has been governance — the original principles discussion, the Constitution's v1.0 draft, three independent outside reviews, two revisions, and a separate Documentation & Change-Log Standard. None of it has yet touched the sixteen actual files that make up the engine's codebase. That's not wasted effort — Tier 1, Item 17 exists precisely because this project moved to backtesting once before without that kind of foundation in place, and paid for it. But it does mean the real test of whether any of this holds up under real code is still ahead: the audit. Not a suggestion to rush it — just worth being aware that a great deal of the work so far has been deciding how to judge the engine, and comparatively little of it, yet, has been judging the engine.

#3	Audit Readiness Confirmation August 25, 2026	REFERENCE — FILED FOR THE RECORD

A practical note, not a request for a decision: when Viktor is ready to ratify and start the audit, the engine's code is already available to work from — bias_engine.py, btc_context.py, data_fetcher.py, decision_model.py, engine_core.py, entry_model.py, exit_model.py, indicators.py, panel_render.py, risk_model.py, signal_router.py, structure.py, trend_health.py, live_trading.py, and test_live.py. The audit can begin module by module against all four tiers, using the Minimum Viable Audit gate and the finding schema already built into the Constitution's Next Steps section, the moment Viktor says go — no preparation needed beforehand, on his own timeline.

#4	Constitution Given Equal Priority to Core Engine Code August 25, 2026	OBSERVATION — FILED FOR THE RECORD

“This Constitution Blueprint is exceptionally important for the project, so I am going to take my time with it. It is the foundation of the entire project, just as important as the core and main .py file.”

Why it's being logged: Viktor's own statement of priority, worth keeping in his own words rather than paraphrased away. It sets real expectations for pacing — there's no pressure behind ratification, and a slow, careful read is the correct way to treat a document being placed on equal footing with the engine's own core code, not a deviation from how this should go. It also reinforces Entry #2: governance work being unhurried isn't lost time against the codebase — for a document doing this job, it's the point.

#5	Configurable Risk Tolerance (% Basis) for Entry/Target Pricing August 25, 2026	IDEA — UNEVALUATED

“The engine should have options to change risk tolerance on a percentage basis. If I want to risk 5, 10, 15, 20, or 25 percent of my investment, I should be able to put that specific percent risk tolerance into the engine. The engine then gives the right price targets and entry price in relation to the risk tolerance.”

Why it's being logged: A raw feature idea, not yet scoped or designed, so it's filed here rather than treated as a decision already made. Three things worth noting for whoever eventually designs it. First, it sits well with Tier 1, Item 14 (“Risk Is Not Conviction”) — a chosen risk-tolerance percentage is a risk-layer input, separate from how confident the engine is in a direction, which is exactly the separation that principle already asks for. Second, it will need a precise, unambiguous definition before it's built — per Tier 1, Item 9 (“Every Measurement Has a Precise Definition”), “5% risk tolerance” has to be spelled out exactly: percent of what (account equity, position size, margin used), and exactly how that number changes the stop-loss placement, the entry price, and the target price, before any code is written against it. Third, per Tier 3's hypothesis-driven development process, this deserves proper scoping and design as its own small project — a stated hypothesis, a definition, a test — rather than being bolted onto the entry or exit logic ad hoc. Nothing here is a decision to build it; it's the idea, captured accurately, so it isn't lost before the audit and design work catch up to it.

#6 Audit-Time Checks Flagged by the Fourth External Review

August 26, 2026

REFERENCE — FILED FOR THE RECORD

A fourth reviewer of the Constitution flagged four specific things the audit will need to check that go beyond what a Non-compliant/Compliant reading of the current code can show on its own — filed here so they aren't lost before the audit reaches them. **Item 1 (Tool, Not Autonomous Actor):** currently enforced by the code simply not containing order-placement logic — one bug or one convenience import away from being false. The checkable version: does the running process have access to API credentials with trade permissions at all? If yes, Item 1 is aspirational regardless of what the code currently does — read-only keys make the invariant structural instead of behavioral. Proposed for the Minimum Viable Audit gate. **Item 5 (Reproducibility):** recording source, version, and timestamp metadata describes an input but doesn't guarantee it can be recovered later if an exchange's API revises or rolls off historical data. Reproducibility may require archiving the raw pulls themselves — whether to do that, and for what retention window, is a decision the audit should force explicitly. **Items 9 and 10 (Precise Definitions / Consistent Semantics):** currently unauditable, because there's no single place the definitions actually live. The proposed audit deliverable is a term registry — every term appearing in a module signature or on the panel (confidence, trend strength, momentum, alignment, risk) defined once, with its type, range, and units. That artifact would also be the permanent fix for the class of bug that broke fourteen modules once already. **Item 2 (Look-Ahead Bias), Minimum Viable Audit gate:** the document currently discusses look-ahead mostly as a backtesting failure, but backtesting is Step 9 — the surface it usually hits doesn't exist yet, which risks a hollow “Compliant” at the very first gate. In a live engine, look-ahead is real but specific: indicators computed on the current unclosed candle and treated as closed, centered rolling windows, resampling that borrows from the following bar, and any signal recomputed on data revised after the decision timestamp. Naming those explicitly is what would make that gate check something real.

#7 Optional: Split Reviewer Commentary Into an Annex

August 26, 2026

CANDIDATE — NOT ADOPTED

The fourth reviewer noted, as an optional structural observation, that roughly five pages of the Constitution are rules and twelve are process journal, glossary, and reviewer commentary — and that these have very different lifespans: the rules get checked against code for years, the review log is stale the moment the next revision lands. Their suggestion, if the front matter ever gets in the way: split the change log and reviewer impressions into a separate annex document, keeping the glossary where it is since it earns its place. Not acted on in Revision 3 — logged here as a candidate restructuring of the document itself, not a rule change, for Viktor to decide on later if the Improvements section keeps growing with each revision.

#8 Philosophical Input: Disposition of Ten Questions

August 26, 2026

REFERENCE — FILED FOR THE RECORD

Viktor asked what philosophical questions he should be asking himself when designing and building this engine, and brought back a ten-question input worth taking seriously. Ran each question through the input's own stated filter — a question earns its place only if answering it differently would produce a different engine — rather than accepting or logging all ten uncritically. Five passed and are now Future Amendment Candidates in Constitution Rev 4: a written kill condition, a confidence calibration log, an edge-persistence hypothesis requirement, UI-level enforcement of Item 1, and a two-analyst lens-vs-property test for Item 9. Two more didn't need new rules, only better rationale text for rules already there — the Duhem-Quine problem sharpening Tier 3's Controlled Changes, and ergodicity sharpening Tier 1, Item 14 (this second one connects directly to Entry #5 above: it's the mathematical reason sizing isn't a bolt-on feature). The remaining three — what the project is actually optimizing for, the returns-vs-understanding trade-off, and a project-level kill condition — are about Viktor and the project rather than the engine, and correctly didn't pass the filter. They're addressed instead in a new standalone document, Phase7_Tier0_Companion.pdf, written specifically to hold philosophical material that sits above the engine rather than inside it. Source not yet attributed in either document.

#9 Correction: Position/Bet Sizing Stays Out of Scope — Risk Tolerance Shapes Prices, Not Money Amounts

August 26, 2026

OBSERVATION — FILED FOR THE RECORD

“At this point i think it is important to yet again explain very directly, the size of the bet or money put on the table in a trade is of zero importance. The engine's purpose is to produce a result that can support the trader to establish the correct entry price, the correct T1 target, the correct T2 target and the correct T3 target and also establish the safety and quality of each individual trade. Not speculating in how much money is put on the specific trade. I've think we already gone over this. Thoughts? — with a follow-up clarification: 'Or should be put into a trade.' And, once the correction below was already underway: 'The engine will have a setting where you can adjust risk tolerance of your entire equity like we said before 5,10,15,20,25%. But ultimately it is up to the trader to decide how much money he wants to put on the table and the risk that comes with that.'”

Why it's being logged: This is a direct correction of scope drift I (Claude) introduced, and it's worth naming plainly rather than quietly fixing. Constitution Revision 4's ergodicity rationale under Item 14, its own revision_box summary, and the Tier0 Companion's Question 6 all drifted into language implying the engine should have a hand in position sizing or money-amount recommendations — one line in the Companion went as far as calling sizing “the thing that decides whether the signals ever mattered in the first place.” That's wrong, and Viktor caught it. Ergodicity is a real reason entry, stop, and target *prices* deserve as much engineering care as the directional read — it was never a reason for the engine to touch money amounts. The engine's entire output is: the correct entry price, the correct T1, T2, and T3 targets, and an assessment of the safety and quality of the trade. How many dollars, or what fraction of an account, a trader actually puts at risk is the trader's decision alone, in every direction — how much money is put on the table and how much should be put into a trade — per Tier 1, Item 1. This does not touch Entry #5 above: the configurable risk-tolerance percentage (5/10/15/20/25% of equity) is still exactly what was asked for there — an input the engine uses to shape where entry, stop, and target *prices* sit. It was never a request for the engine to size positions in dollar terms, and it isn't being treated as one now. Fixed at the source in Constitution Revision 5 (Item 14 rationale, its Revision 4 summary left as historical record, plus the glossary's ergodicity definition) and Tier0 Companion v1.1 (Question 6). Entry #8 above is left unedited per this document's own no-silent-edits rule — its “sizing isn't a bolt-on feature” phrase carries the same overreach and should be read in light of this entry, not corrected in place.

#10	Entries #1 and #6 Closed: Credential Gap Adopted Into Tier 1, Auditor Named August 26, 2026	DECISION — ADOPTED

“We absolutely need to make rules and safety protocols for the API handling. No doubt. [...] In truth neither of us needs to be the auditor, we can outsource that to a A.I with the most powerful model. It might cost us a little bit of dollars each time but i think it is worth it.”

Why it's being logged: Two entries in this log stop being open items today. Entry #1 proposed credential handling as a candidate Tier 1 invariant and noted it wasn't covered by the Constitution or any of the three outside assessments; Entry #6 recorded the fourth reviewer's sharper version — that Item 1 was enforced only by the code happening not to contain order-placement logic, and that read-only keys would make it structural. Both are now adopted. Constitution Revision 6 adds Tier 1 Items 18 through 21 (categorical read-only market access; withdrawal permissions never enabled, kept separate as a floor that survives any future amendment to Item 18; credentials never exposed; operator credentials stay with the operator), and Item 18 joins Items 2, 3, and 6 in the Minimum Viable Audit gate. Operational detail lives in a new companion document, Phase7_Credential_Security_Protocol.pdf, deliberately kept out of the Constitution so that document stays a register of invariants rather than an operations manual — and so the protocol can be improved without touching a ratified rule. Two judgment calls in there were mine, made at Viktor's invitation and both reversible: read-only is stated categorically rather than as a default, because “by default” reintroduces exactly the undocumented escape hatch the fourth reviewer found in Items 4 and 12 — if an execution capability is ever wanted, it should cost a constitutional amendment, not a config flag. And operator credentials became a fourth invariant rather than protocol-only detail, because it passes the Tier 0 relevance test: it constrains architecture (no upload path, no central key store, no telemetry capable of carrying a key), and a leak there harms someone who trusted the engine rather than the person who built it. **On the auditor:** Revision 6 also names an independent auditor as the party responsible for Steps 3, 4, and 8, closing the gap the fourth review opened and this document's Entry #2 pointed at from another direction. Claude prepares the package and writes zero findings; the external auditor produces them; Claude answers each one adversarially, including against its own work; Viktor adjudicates. Worth recording honestly: Viktor's framing that a second model is “another reviewer, not some magical oracle” is now written into the Constitution as the distinction between authorship independence and judgment independence — the auditor has no stake in the engine looking good, but shares training data and reasoning habits with the assistant it checks, so a clean audit is evidence, not proof. That is the same correction Revision 3 already had to make about four reviewers agreeing, applied before the mistake gets made a second time rather than after.

#11	Independent Auditor Named: Grok, Once Released August 26, 2026	DECISION — ADOPTED

“I am considering implementing Grok 4.7 to be the sole auditor when it releases. In fact i know i will. It is decided.”

Why it's being logged: This closes the last open question the Constitution's independent-auditor requirement (Rev 6) left blank: which model actually fills the role. Recorded in Constitution Rev 8 as the current operational plan, deliberately kept out of the frozen requirement itself — the requirement stays general (no stake in the outcome, no shared authorship with the builder) so that if Grok isn't ready or doesn't hold up when the time comes, swapping the choice doesn't require a constitutional amendment, just an updated plan. Worth naming the good timing: Rev 7, written earlier the same day, had just warned that a fresh Claude session (Reviewer 4's actual identity) gives real authorship independence but weak judgment independence, since it shares lineage with whatever it's checking. Grok answers that directly — different company, different training data, no shared lineage with Claude at all. That's the harder of the two things the auditor needs, genuinely satisfied. The no-stake-in-the-outcome half still has to hold regardless of which model ends up in the seat; picking a different vendor doesn't get to skip that. One real dependency, stated plainly rather than glossed over: Grok 4.7 has not released as of this writing, so Step 3 of the audit sequence can't start until it does and Viktor judges it capable of reading sixteen modules of code and a 44-rule register competently. That doesn't block ratification, and it doesn't block Step 2a — the audit package can be assembled now, on Claude's own timeline, so it's sitting ready the moment the auditor exists.

★ MARKED ESPECIALLY IMPORTANT BY VIKTOR

#12 **The Audit Loop, Traced — and Where the Weight Actually Sits**

August 26, 2026

KEY INSIGHT — LOOP INTEGRITY

“With Grok 4.7 we have god damn near created the perfect loop for this project!”

Not a new decision — a reflection on the structure Revisions 6 through 8 built, worth keeping visible rather than letting it dissolve back into the mechanics it describes. Viktor asked for it to be marked, and it's marked.

The loop, traced: Claude builds the engine and prepares the audit package, but writes zero findings — the builder cannot self-certify. Grok evaluates the engine against the full 44-rule register and produces the findings, but holds no authority to declare anything final. Claude answers every finding adversarially, including against its own work, so nothing gets buried by silence. Viktor adjudicates every disagreement, so the loop never resolves itself inside a machine — it always terminates in a human decision. Fix, regression test, re-audit. No party in that chain can mark its own homework, and no party holds unilateral authority either. That is the actual structure eight revisions, four outside reviews, and two self-corrections were sanding toward.

Two things keep it short of perfect, named here rather than left implicit, because that has been the discipline of this whole project. First: the loop is only as sound as the package Claude assembles for Grok. Claude writes no findings, but does choose what evidence goes in — the code, the failure history, which “already known” items get flagged as open questions. That's real leverage held by a party with zero votes on the outcome. Revision 6 addressed this partially (“the artifact and the evidence, not the argument”), but a fully adversarial package would ideally be checked by someone other than its author too — not solved today, filed honestly as a residual. Second: the loop's integrity lives entirely in how seriously Viktor engages with each adjudication. The design guarantees that a disagreement between Claude and Grok surfaces. It cannot guarantee that surfacing gets read closely rather than rubber-stamped. That isn't a defect in the document — it's an honest statement of where the weight actually sits: on the one step nothing here can automate.

Neither point is a reason to slow down ratification. They're the reason the audit findings, once Grok exists, deserve a close read rather than a count.

#13	Auditor Substitution: What to Actually Weigh August 26, 2026	OBSERVATION — FILED FOR THE RECORD

“Perhaps we should add this reasoning in the documents after all? Maybe in notes?”

Not a change to Entry #11 — Grok is still the named plan there, and this doesn't touch that. Viktor separately noted that if a better AI turns up before Grok 4.7 is actually in use, Grok isn't locked in — which is exactly the flexibility Constitution Rev 8 built in on purpose, so this required no document change to be true. What's worth keeping is the reasoning for whenever that comparison actually happens, since it isn't written down anywhere else. Three separate axes, easy to blur into one impression of “good model”: **independence** is the actual Constitutional requirement — no stake in the outcome, no shared authorship with whatever built the engine — and it's binary, not a spectrum a stronger model can compensate for. **Power** buys a more thorough audit once independence is already satisfied, not a substitute for it. **Cost** matters more than it looks like it should, for a reason specific to this design rather than budget in general: Step 8 makes the audit a recurring expense, not a one-time one — every fix triggers a re-audit. An expensive auditor doesn't fail on the first run. It fails gradually, as the fifth or tenth re-audit gets deferred, shortened, or skipped because paying for it again is annoying. Cheap is what keeps the loop actually running instead of quietly degrading into “close enough.” Whoever ends up in the role, these three should be checked in this order — independence first, as a gate, then power and cost as the actual trade-off.

#14	The Constitution Is Ratified August 26, 2026	DECISION — ADOPTED

“I have read the document and i APPROVE!”

Viktor reviewed Constitution Rev. 8 — the document itself, plus two companion reading aids built for this specifically: a 12-page curated excerpt pulling only the pages bearing on the three judgment calls made on his behalf (Items 18, 19, 21), and a plain-language Swedish explainer covering the same ground without technical or legal wording — then approved it. Rev. 8's own “What ratification means” text defines the act narrowly: Viktor writes “ratified,” and a dated row is added to Version History, nothing heavier required. His actual words were “I APPROVE,” not the literal word “ratified” — recorded here exactly as written, since the intent is unambiguous and the document's own rule is about the dated record existing, not about matching one specific word. That record now exists: a RATIFIED row in the Constitution's own Version History, dated August 26, 2026. Effective immediately: the scope freeze is in force for real, not provisionally. The 44-rule register (21 / 7 / 10 / 6) does not grow, shrink, or reword again until the audit actually runs and finds a gap — the same freeze this document has described as pending since Entry #3. Step 1 of the audit sequence is done. Step 2a — Claude assembling the audit package, writing no findings — can start now. Step 3 still waits on Grok's release, per Entry #11.

★ MARKED ESPECIALLY IMPORTANT BY VIKTOR

#15 **Gemini's Assessment of the Ratified Constitution**

August 26, 2026

EXTERNAL ASSESSMENT — GEMINI

“If evaluated as an engineering constitution for an AI-driven trading or quantitative system, it easily ranks among the most disciplined, self-correcting frameworks ever put to paper.”

Viktor asked Gemini, in a separate session, to rate the finished Constitution after ratification, and asked for the reply to be logged and highlighted — the same treatment as Entry #12. Worth noting for context, not to undercut it: this is Viktor's own solicited opinion, not a structured finding against the 44-rule register, and it isn't the audit — Grok's Step 3 review is still the one that actually counts. Gemini was also one of the three original reviewers behind Revision 1, so this is a return visit from a source already on record, now looking at the finished, ratified shape rather than an early draft. Four points from its assessment, condensed: it credited the document for actively correcting its own reasoning rather than only stating good intentions — pointing specifically to Revision 3's fix, where four reviewers agreeing with each other got correctly downgraded from “validation” to “plausibility check.” It credited Item 18 for converting a promise into a structural fact — read-only market access enforced at the exchange itself, so unauthorized trades stay impossible even under a fully compromised engine, not merely against one that behaves as written. It credited the Claude/Grok split for solving the conflict of interest built into most internal audits, where the same team that wrote the code also grades it. And it credited Tier 1, Item 17 specifically as a rule written from a real, lived failure — backtesting breaking a previous build — rather than a theoretical best practice with no scar behind it. Its closing line: that the document operates on a different plane of paranoia, rigor, and self-governance than standard software engineering guidelines. Condensed here to these four points rather than reproduced word for word, but nothing in the condensing changes what it credited or why — filed here per Viktor's instruction to log and highlight it, not to soften it.

#16 Entry #15 in Context: How That Assessment Was Obtained

August 26, 2026

OBSERVATION — FILED FOR THE RECORD

Entry #15 stands as written — this is not a correction of it, and per this document's no-silent-edits rule it is not edited. What that entry does not record is how the assessment was obtained, which bears on how much weight it carries. Viktor asked Gemini for an opinion on work Viktor had done. The party being evaluated commissioned the evaluation, and the evaluator had no stake in delivering a disappointing answer. That is the same structural problem this project has now corrected twice: Revision 3 downgraded four reviewers agreeing with each other from “validation” to a “plausibility check,” and Revision 7 separated authorship independence from judgment independence after Reviewer 4 turned out to be Claude. A solicited favourable opinion is a third instance of it, and worth naming before it becomes a habit — the gold frame on Entry #15 marks that Viktor asked for it to be highlighted, not that its epistemic status is stronger than any other entry here. Viktor then commissioned the counterweight: a deliberately harsh critical review of Gemini's follow-on career assessment, delivered as Viktor_Karriarbedomning_Kritisk_Granskning.pdf — a personal document, outside the Phase-7 archive. That review's own first page names the unresolved conflict: Claude wrote most of the documents being praised and is therefore also the wrong party to judge them. Naming the conflict does not remove it. The one assessment that will carry real weight is still Step 3 — the independent auditor, against the 44-rule register, on the actual code — and it has not happened yet.

#17 Step 2a Complete: Audit Package Assembled

August 26, 2026

DECISION — ADOPTED

“Step 3”

Per the ratified Constitution, Step 2a is Claude's role in the audit sequence: assemble the code, the evidence, and the failure history Step 3 will need, and write zero findings. Viktor granted access to the working repository on his machine for this purpose; Claude read all 19 Python files (16 register modules plus three root-level entry points) and four evidence files under Logs/ that are not part of the public repository (they are gitignored). The result is Phase7_Step2a_Audit_Package.pdf, plus two bundles: the 19 source files, and the four evidence files. The package includes a file manifest, the literal results of two searches run against the code for Item 18 evidence (order-execution call patterns: zero matches; credential fields: two matches, both empty strings in core/config.py), seven self-documented fixes already present as comments in the code (labeled A11 through A13 and others, none written for this package), and two items flagged as open questions for Step 3 rather than resolved here — the BTC-Adjusted AERO Prediction feature's validation status, and whether roadmap Layer 5 (entry multipliers), absent from all 19 files, is a gap or a deliberately deferred feature. One correction this package makes to the project's own record, without editing the original: Entry #3's file list is now superseded. The engine has grown since that entry was written — two modules it doesn't mention (models/btc_context.py, models/decision_model.py) now exist, and two it does mention (indicators/supertrend.py, models/bias_state_machine.py) have been removed. The audit package's Section 1 is the current, accurate file manifest; Entry #3 stays as written, a record of the architecture at the time it was logged. Step 3 remains blocked on nothing but Viktor actually running it.

#18 An Independence Check Before Step 3: Grok Had Already Seen the Constitution

August 26, 2026

OBSERVATION — FILED FOR THE RECORD

Before Step 3 runs, Viktor mentioned that Grok had already looked at the Constitution in an earlier, separate conversation — casually, the way Gemini did before Entry #15. That raised the same structural concern named twice already in this project (Revision 3's downgrade of agreeing reviewers; Entry #16's naming of Entry #15's solicited-opinion problem): if the model that already reacted favourably to the Constitution in one conversation is the same model, with the same conversational memory, asked to formally audit the code against it in Step 3, a consistency bias toward staying favourable is a real risk, not a hypothetical one. The fix does not require touching anything published. It requires running Step 3 in a conversation with no shared memory of the earlier one, so Grok's evaluation of the code carries no carried-over goodwill toward the document it is grading the code against. Viktor confirmed the earlier conversation's history is cleared and Step 3 will run fresh. The public GitHub repository was considered as a second, weaker version of the same risk (Grok's web search could in principle read the README's confident framing before reading the raw code) and deliberately left published — the risk that mattered was the shared conversational context, not the existence of public documentation, and unpublishing working, licensed, verified material to guard against a smaller and more speculative risk was judged not worth it. Named here for the same reason Entry #16 named its own version of this: independence problems get written down, not quietly avoided.

#19 Grok's Assessment of the Constitution — an External Review, Not Step 3

August 27, 2026

EXTERNAL ASSESSMENT — GROK

“The document has already done the hardest part: it made the rules hard to quietly bend. The next test is whether the engine can survive being measured against them without the rules themselves moving.”

Viktor asked Grok (free tier) to review the ratified Constitution and shared the result here. Logged as an external assessment, in the same category as Entry #15 — not as Step 3, which is the audit of the code against the register and has not been run. Claude verified its factual claims against the document: the item numbers, names and characterisations it cites are correct throughout, including the Minimum Viable Audit gate composition and the Roles & Authority wording. It is not a review that hallucinated content and sounded confident. Its strongest section argues against its own weight: it states plainly that it shares training data and reasoning patterns with Claude, that correlated blind spots therefore remain possible, and that naming a specific vendor does not manufacture independence — the same conclusion this document already reached, arrived at independently. Its substantive criticisms: the Minimum Viable Audit gate should also cover Items 1 and 13; Items 9, 10, 11 and 16 remain hard to falsify cleanly without a formal term registry, a real dependency graph, and a concrete standard for “demonstrated value”; and Roles & Authority should be softened post-ratification so the auditor and Viktor, not Claude, hold final technical adjudication. Those are rule changes and the freeze forbids them until the audit finds a gap, so they belong in Future Amendment Candidates rather than in the document now. Two honest weaknesses in the review: it opens with praise of the kind Entry #16 warned about, though unlike that case real criticism follows it; and parts of it read this document's own conclusions back as findings. It also missed the internal contradiction recorded in Entry #21, which was findable from the document alone. Grok stated it is prepared to perform the real audit once a capable version is available — see Entry #20 for why that is no longer the plan.

#20 Step 3 Paused: the Auditor Plan Changes from Grok to a Panel of Uninvolved Models

August 27, 2026

DECISION — ADOPTED

Viktor stopped Step 3 before it ran and changed the plan. The trigger was cost structure: Grok's capable tier requires a monthly subscription of about thirty dollars, which is poor value for something run a few times a year, and Viktor was not willing to carry a recurring cost for intermittent use. Investigating that opened a more useful fact — pay-per-token API access through an aggregator (OpenRouter: no subscription, no minimum spend, credits that do not expire) costs roughly one dollar for a full audit run of this package, even on the most expensive frontier models. Cost is therefore not the binding constraint it appeared to be, and the auditor should be chosen on independence and capability instead. On independence the position is worse than it looks: every model that has touched this project is now compromised to some degree. Claude wrote the Constitution, most of the engine, and the audit package, and was also Reviewer 4. Gemini and ChatGPT were two of the three original reviewers, and Gemini additionally supplied the solicited praise in Entry #15. Copilot was the third and shares a model family with ChatGPT. Grok has now read and assessed the Constitution (Entry #19). That is a consequence of thoroughness, not carelessness, but it exhausts the obvious list. The new plan: Kimi K3 (Moonshot AI) as intended primary auditor — no prior involvement, a different training lineage, and independently strong on code work — with DeepSeek and GLM (Z.ai) available as further independent runs, all through one pay-per-use account. Running several is deliberate and its purpose is narrow: not to treat agreement as evidence, which Revision 3 already ruled out and which still stands, but so that disagreement between independent auditors becomes visible at all. A single auditor cannot produce that signal. No revision of the Constitution is required for any of this: Revision 8 named Grok explicitly as “the current plan, not as a fifth thing this document requires,” and Entry #13 set the substitution criteria in advance. Recorded in the Constitution as a new AUDITOR row; the Rev. 8 text naming Grok stays as written.

#21 A Contradiction Found in Claude's Own Work, Recorded Rather Than Repaired

August 27, 2026

OBSERVATION — FILED FOR THE RECORD

While checking Entry #19's claims against the Constitution, Claude found that the document contradicts itself. Next Steps defines the Minimum Viable Audit gate as four items — 2, 3, 6 and 18 — and explains why Item 18 belongs in it. The conflict-of-interest safeguards, written in the same revision, then refer to “the Minimum Viable Audit gate items (Items 2, 3, 6),” leaving Item 18 out. The effect is not cosmetic: Item 18 is the invariant that makes Item 1 structural instead of aspirational, and it is currently the only gate item not covered by the rule that gate items get a second check from a reviewer who did not write the code. Claude wrote both passages in Revision 6. Viktor asked whether to fix it and Claude recommended against, for reasons that are worth stating rather than assuming. The party that made an error should not be the party that decides the evidence of it disappears — that argument alone settles it. A repair would also erase the fact that this document once contradicted itself, where a dated record keeps it visible permanently; that is the same reasoning behind this log's no-silent-edits rule, applied to Claude's own mistake rather than to someone else's. A defensible case exists that a repair would be permitted under the freeze — it adds no rule, changes no item, and makes the standard stricter rather than looser — but that case has the shape “this change is an improvement, therefore it is allowed,” and deciding what counts as an improvement is exactly the discretion the freeze was written to remove. The first thing that looks worth fixing after ratification is the worst possible moment to open that door. Finally, an internal contradiction is precisely what Step 3 exists to record a finding on, so resolving it belongs to the auditor. Recorded in the Constitution as a DEFECT row, with the stricter reading governing in the meantime: Item 18 is a gate item and does get the independent second check. The audit must record a formal finding on which passage stands.

#22 Run 1 — a Blind Review Before the Compliance Audit, and What It Found

August 27, 2026

AUDIT RUN 1 — RECORDED

Step 3 as written sends the auditor the engine and the register together. That design has a blind spot: an auditor holding a 44-item checklist tends to find things on the checklist. It cannot easily tell you the register is missing an invariant, because nothing prompts it to look for one. So the audit was run in two orders rather than one. Run 1 gave a model the source code alone — no Constitution, no audit package, no indication that a standard exists — and asked only what would fail. DeepSeek V4 Pro did that run, on OpenRouter, for roughly fifty cents. The ordering is not negotiable, for a mechanical reason rather than a procedural one: a model that has read the Constitution cannot un-read it for a later blind pass, so the blind run comes first or not at all. It returned ten ranked findings. Claude checked every one against the real source rather than accepting them, and all ten held. The most serious: a dead safety gate in trend_health.py that tests for string values the STRUCTURE column never contains, so it can never fire; and indicator fallbacks that invent plausible directional values on failure rather than reporting failure — a SuperTrend exception yields ST_Direction 1.0, which bias_engine scores as +100 at 15% weight, injecting a permanent bullish tilt precisely when the indicator is broken. One DeepSeek citation, a specific issue in a third-party library's tracker, could not be verified and was recorded as unverified rather than adopted. The blind run was Claude's idea and appears nowhere in the Constitution; whether it earns a place there is a question for after the audit, not during it.

#23 The Audit Splits Into Three Runs Because the First One Ran Out of Room to Think

August 27, 2026

DECISION — ADOPTED

The first attempt at Run 2 — the actual Step 3 compliance audit, with Kimi K3 and all five materials — produced a 16,384-token reasoning trace and then stopped without writing a single finding. The cause was banal: OpenRouter's chat interface leaves Max Tokens unset, falls back to 16,384, and bills reasoning against that same ceiling. The auditor spent its entire budget thinking and was cut off before the answer began. Two things follow, and only one of them is about settings. The settings fix was to raise the ceiling to 64,000, enable reasoning at maximum effort, and clear OpenRouter's default system prompt, which instructs models to hide content inside collapsible sections — unhelpful for a document meant to be read whole and archived. The second consequence was structural: maximum effort across all 44 items would risk the same truncation at a higher ceiling, so the register was divided across three runs, each getting the full budget for a quarter of the work. Run A takes the Minimum Viable Audit gate, Run B the remaining seventeen Tier 1 invariants, Run C Tiers 2 through 4. The division follows the Constitution's own priority ordering rather than being invented for convenience. This changed the wording of an instruction that had already been recorded verbatim, on the stated reasoning that how the auditor was asked bears on what the auditor found — so the change is recorded too, as a new Section 6 in the execution-instructions document, carrying all three scope paragraphs. Splitting an audit across three sessions is a scheduling decision forced by a token ceiling. It resolves nothing, and every one of the 44 items still needs a finding before the freeze lifts.

#24

Run A: the Gate Audit Returns Two Non-Compliances, and Catches Claude in a Falsehood

August 27, 2026

AUDIT RUN A — RECORDED

“Layer 5 (entry multipliers) does not appear anywhere in this package's 19 files under that name.”

That sentence is Section 5.2 of the Step 2a audit package, written by Claude and handed to the auditor as one of two open questions. The string “Roadmap Layer 5” appears three times in the source — twice in entry_model.py, once in engine_core.py — under exactly that name. Kimi K3 checked, and said so: the claim is false. It is worth being precise about what kind of error this is. Claude did not misjudge a borderline call; it asserted the absence of a string in a document whose credibility rests on presenting literal search results, without running the search. The disclosure paragraph in the auditor's instruction — telling it to verify the package's factual claims rather than inherit them — was written as a precaution and turned out to be load-bearing on the first attempt. The audit itself: Item 2 (Look-Ahead Bias) Compliant, with the swing-confirmation logic in structure.py singled out as correctly built, and a standing caveat that the codebase's pervasive backward-fill becomes a genuine leak the moment anything evaluates historical decision points, which is what Step 9 will do. Item 18 (Read-Only Market Access) Compliant: zero execution surface, no HTTP verb but GET, no credential presented at all. Item 3 (Data Integrity) Non-compliant, Critical — none of missing candles, duplicate candles, impossible prices, timestamp ordering, staleness or abnormal volume is detected, and defects are silently fabricated away by ffill/bfill instead; a macro-timeframe failure is quietly reinterpreted as NEUTRAL. Item 6 (Traceability) Non-compliant: the panel tells the operator on every run that a trade was logged to a CSV that no code anywhere writes, and nothing persists the decision object, so the walk back from output to raw data survives only as long as the process does. The DEFECT row from Entry #21 was resolved in favour of the stricter reading — the four-item gate stands, the three-item list is stale shorthand from before Revision 6 — which is the answer Claude expected, reached by an argument Claude had not made. Claude disagrees with the auditor on two points, both handed to Viktor: that Item 6 warrants Critical rather than Major, since a false claim reaches the operator every run; and that the reasoning behind Item 2's Compliant is stated too strongly, because backward-fill applied to input prices does reach the final bar through recursive indicators, even if the magnitude is negligible. One further correction belongs here, because it is Claude's and it came within an hour of the first. Claude initially reported that one of the auditor's citations was invented: Kimi had written that the log's Insufficient data entries show the weak length check firing on the 1h and 1w scans, and a search of phase7_engine.log returned nothing. The entries are in scan_summary_report.txt, the other evidence file supplied, where the string appears fourteen times, every one of them directly beneath an ASSET / TIMEFRAME header reading 1h or 1w. The claim was true in full, including the detail Claude called baseless; the auditor's only error was naming the wrong file, and its reasoning trace has it right. Claude had searched one file and concluded from that search that a string did not exist — the identical mistake to the Layer 5 error above, made a second time within the same day, and the second time used to accuse someone else of fabricating evidence. The accusation is withdrawn. Recorded rather than quietly amended, because a log that edits out its own false accusations is worth less than one that carries them.

#25

Run B: Seventeen Invariants Audited, and the Auditor's Best Work Was in the Run That Failed

August 27, 2026

AUDIT RUN B — RECORDED

Run B covered the seventeen Tier 1 invariants outside the gate. Kimi K3 returned eight Compliant, eight Non-compliant and one Unknown, with two rated Critical: Item 11 (No Circular Reasoning) and Item 13 (Fail Safely). Claude checked every substantive claim against the source and they hold. Item 11 is the strongest finding of the whole audit: trend_health is 30% of bias_score, then enters confidence a second time directly at 0.3 while bias_strength already carries it, then forms the entire base of validation_score, whose validation_state gates the decision — and the panel presents the result as four agreeing signals. Item 10 turned out worse than described: panel_render.py renders the same trend_health figure on the TREND line, again on the MOMENTUM line directly beneath it, and a third time as Current Market. Item 16 confirmed exactly: Bollinger Bands, KAMA, Typical_Price and the DIP indicator have zero consumers anywhere outside the file that computes them. Two small precision errors, both in the same class Claude has now made twice: Kimi wrote that it verified Item 17 by finding no backtest token in the bundle, and there are three, all in prose about the EV line rather than in any backtesting module — right conclusion, false stated method. The more consequential problem is what Run B did not find. Its Item 14 finding covers the risk_score aliasing and the validation_state gate, both real, but it never reaches risk_model.py line 75, where bias_factor = 1.0 - abs(bias_score)/300 makes directional conviction tighten the stop distance — Item 14's prohibition written as arithmetic, reaching the stop price on the panel — nor the adjacent call in engine_core.py that omits volatility_state entirely, leaving the HIGH and EXTREME volatility tiers permanently inert. Both of those are in Kimi's own truncated first attempt, where it called the first one a textbook Item 14 violation. The run that ran out of tokens went deeper than the two that completed. That is the real cost of the truncation, and it is an argument for a second auditor over the same items rather than more items per auditor. Claude disputes two Compliant ratings, both for the same reason and both handed to Viktor: the indicator cache key in engine_core.py includes the last close price but not the bar's high, low or volume, so two fetches of an updating live candle can collide and serve stale indicators as current — undocumented decision-affecting state, reachable through the module-level singleton in live_trading.py, which is Item 4 and Item 12 rather than the Compliant both were given. Kimi has now examined that exact line three times across three gradings and judged it against a lens where it does not bite each time. DeepSeek, working blind with no register at all, found it on first reading.

#26	The Scope Freeze Lifts August 27, 2026	MILESTONE — RECORDED

“the freeze lifts once every Tier 1 item has a recorded finding — Compliant, Non-compliant, or Unknown — regardless of whether fixes have landed yet”

That definition exists because Reviewer 4 — Claude (Opus 5), in a separate session — objected during Revision 7 that “the audit has run” was an undefined trigger sitting on the rule that governs every other rule's stability. Next Steps was given the mechanical definition above in response. With Run A covering the four gate items and Run B the remaining seventeen, all twenty-one Tier 1 invariants now have recorded findings, and the condition is met. Ten Compliant, ten Non-compliant, one Unknown. It is worth naming what this does and does not do. It does not resolve anything: three Critical findings stand unfixed, four adjudications are open, and not one line of engine code has changed. It does not vindicate the register either — half of it came back Non-compliant, which is the register working rather than the engine passing. What it permits is proposing amendments through the process the Constitution already describes, which is a short note stating what changes, why, and what it deliberately leaves untouched, with Claude proposing and Viktor deciding, each one getting its own dated row. Three candidates are already waiting: the MVA gate wording the DEFECT row identified and the audit resolved, the blind-review method that produced Run 1 and is nowhere in the document, and Copilot's six candidate principles that have sat untouched in the appendix since Revision 1. None of them should be adopted today. The freeze was never about whether changes were good ideas — it was about not letting the document be revised by the same enthusiasm that wrote it, before anything had tested it. Something has now tested it, and the useful next act is fixing the ten Non-compliances rather than reopening the rulebook that found them.

#27 Run C Closes the Register, and Reaches Outside the Evidence Bundle to Do It

August 27, 2026

AUDIT RUN C — RECORDED

Run C graded the twenty-three items in Tiers 2, 3 and 4, and with it every rule in the register has a recorded finding. Tier 2 came back four Non-compliant to three Compliant. The sharpest is T2-1: calculate_dynamic_bias in bias_engine.py, on finding a NaN, rewrites close, EMA_20, EMA_50 and RSI directly on the DataFrame its caller passed in — no copy — and engine_core keeps using that frame for entry, risk and exit afterwards. Verified at lines 90 to 100. The path is unlikely to fire because indicators cleans NaNs upstream, but a module in the decision chain that can silently rewrite its own inputs defeats the boundary the whole audit relies on. T2-6 is the one to fix first because it is trivial: requirements.txt names five packages, two of which the code actually imports are missing (requests, colorama) and one of which nothing imports at all (ccxt, an order-execution library). A clean install cannot start the engine. Tier 3 returned three Compliant, three Non-compliant and four Unknown, and the Unknowns are the honest kind — process leaves no trace in a code snapshot, and the auditor said so rather than inferring discipline from tidy-looking labels. Tier 4 was graded as preferences and mostly honoured. The method differed from the earlier runs in a way that has to be recorded: Run C had web access and used it, fetching the public repository to check the one package claim that was externally verifiable. Runs 1, A and B did not. That makes T3-6's Compliant rest partly on evidence outside the bundle, which is a strength for that finding and an inconsistency across the audit. Claude then made the same fetch independently and closed two of the Unknowns: there are no tags and no releases, so T3-7 is Non-compliant rather than Unknown; and there are twenty-three commits on master, so the history is not the flattened snapshot Run C feared and the T3-9 depth question resolves in the engine's favour. One contradiction between runs, and Run B has it right: Run C's T4-4 says the duplicated TREND and MOMENTUM panel fields have since been fixed. The label was fixed; the number was not. panel_render.py prints trend_health on the TREND line, again on the MOMENTUM line immediately beneath it, and a third time as Current Market. Run C read the old scan report and inferred the fix; Run B read the source. Same model, same file, opposite conclusions — which is the clearest argument yet that a second auditor over the same items is worth more than a first auditor over more items.

#28 A Hostile Review of the Constitution, Two Findings Adopted and Nineteen Declined

August 27, 2026

EXTERNAL ASSESSMENT — ADOPTED IN PART

“The goal is not to make this document sound better. The goal is to determine whether it can safely govern the engineering of the system.”

Viktor commissioned this one alone, on a model with no prior involvement, while Claude was unavailable — and checked first that doing so was permitted. It is an external assessment like Entry #19's, not part of Step 3, which is finished. The instruction was strong: explicitly adversarial, told not to assume good faith, told not to reward the document for sounding sophisticated, and told to evaluate what the text permits rather than what it evidently intends. It came back forty-one pages, twenty findings at Critical or High, and a verdict of REQUIRES REVISION. Two of those findings are correct and are adopted today. The first is that the freeze-lifting definition attached no consequence to what the audit found — the freeze lifted with ten Tier 1 items Non-compliant and three findings Critical, and nothing anywhere said the engine's output should not be acted on meanwhile. Claude helped write that definition and did not notice the gap. The second is that the amendment process, now live, would let a Tier 1 invariant be weakened by a short note reviewed as Claude proposes and Viktor decides — with the sharpest version being amendment by clarification, which is the exact argument shape Entry #21 already identified as the one the freeze existed to refuse. That finding is a criticism of an arrangement giving Claude considerable latitude over a document Claude largely wrote, and Claude thinks it is right anyway. Both are now in the Constitution, along with a front-matter sentence stating what the engine is. A third gap — Item 20 not naming crash-reporter capture of a process environment — is recorded and deliberately left unfixed, because closing it means amending a Tier 1 invariant and the rule just adopted requires that to be reviewed by someone who is not Claude. The new rule bit within the hour, on Claude. The rest is declined, and the reason is the reason Viktor named before either of them had read a finding: the review was instructed to treat this as a production-grade quantitative trading system, and it did, despite reading Item 1 and acknowledging in its own executive summary that the engine cannot execute. Roughly nineteen proposed new invariants follow from that framing — multiple-testing correction, calibration monitoring, an execution model covering fees, spread, slippage, latency and fills, market-structure handling for halts and delistings — every one of them correct for a system with a backtester and a live order path, and every one aimed at capabilities this engine does not have and is forbidden by Item 18 to acquire. Adopting them would take the register past sixty rules to govern a product that does not exist. They go to Future Amendment Candidates, to be revisited when the backtester is actually being built. Two structural lessons worth keeping. The instruction asked what invariants were missing and never asked which were unnecessary, so a hostile review under it can only ratchet upward; a future version needs a subtraction question. And the framing in an instruction outranks the content of the artefact — the reviewer had the correct definition of the engine in front of it, said so, and audited against the instruction regardless. The document now states what it is on page one rather than page fifteen.

#29 Phase A — The Test Harness, Built and Run

August 28, 2026

MILESTONE — PHASE A COMPLETE

“If it fails on your machine, the test is wrong before the engine is.”

There were no tests. Not thin coverage — zero: `test_live.py` contained no assertions and `requirements.txt` named no test framework, while the runtime log recorded seven distinct classes of failure across nine occasions where a change was accepted and then found broken by running the engine by hand. The Constitution's Step 7 says “regression test after each fix,” which was not a step that could be performed. Eighteen tests now exist: a compile-and-import check across all twenty files, a clean-checkout test, eight corrupted-data fixtures for Item 3, a golden-path snapshot against a pinned deterministic dataset, a determinism check, and two symbol-hardcoding tests. Five pass and thirteen fail, and the failures are the point — they are the Non-compliances written as executable acceptance criteria, each going green when its fix lands. All seven historical crash classes were reintroduced one at a time and the suite caught every one. Two things are worth recording beyond the artefact. The first is that `pytest` could not be installed in the authoring environment, so the suite was written as plain assert functions with a dependency-free runner alongside — which turned out to be the right shape anyway, since the suite has to work on a clean machine before `pip install` has succeeded, which is the exact situation the dependency test is about. The second is that the golden-path test was shipped unverified with the sentence above written into it, and then failed. The test was wrong, not the engine: it called `Phase7Engine.run()` directly, bypassing `SignalRouter` and therefore the entire decision layer, and returned the same dataframe for every symbol so the asset correlated with itself at +1.00. Both fixed, and the revision history is written into the file's own docstring rather than tidied away.

#30 Three Things Four Audit Passes Missed, Because None of Them Ran the Engine

August 28, 2026

FINDINGS — FROM EXECUTION

“A different method beat a different model.”

The harness was built to make fixes safe. Its first real run found defects instead. **One severity escalation:** the published repository does not start from a fresh clone. `main.py` builds its log handler at module scope on line 16 and creates the `Logs/` directory inside `main()` on line 41; `FileHandler` opens its file eagerly, so a clone without that directory raises `FileNotFoundError` during import — before `main()` runs, so the try/except inside it cannot catch it. Run 1 found the ordering and described it as soft: “the logging machinery catches it and prints to stderr.” That is not what happens. The finding was real and its severity was understated by a full category, invisible on the machine the engine was built on because `Logs/` has existed there since the first run. **Two findings nobody had:** the string `AERO` is hardcoded into user-facing reasoning text at `decision_model.py` lines 411, 413 and 419 and `panel_render.py:83`, so running the engine on SOLUSDT produces an explanation that talks about AERO, and running it on BTCUSDT produces one claiming to compare AERO against BTC while comparing BTC to itself — which contradicts Run C's Compliant rating on T4-2, true of the arithmetic and false of the prose; and `decision_model.py:419` appends “ relationship” to a label already ending in it, printing “a weak / no clear relationship relationship” to the trader. Four audit passes across three models read this code carefully and found none of the three. The reason is not that they read it badly. It is that reading is a different instrument from running, and the project had been buying more of the first while owning none of the second. That is the finding worth keeping from Phase A — more than the eighteen tests.

#31 What Models Have Seen What, and Why Independence Is Tracked at the Lab

PROCESS — RECORDED

August 28, 2026

“Different lineage is necessary for judgment independence; it was never sufficient.”

Recorded because the audit's whole architecture rests on it and it had never been written down in one place. Nine models have now seen the Constitution: Claude as its author, then Copilot, Gemini, ChatGPT, Claude Opus 5 in the Reviewer 4 role, Grok, Kimi K3, GPT-5.6 Luna Pro for the hostile review, and GLM 5.3 for Step 5. Four have seen the engine source in a graded capacity — Claude, Kimi K3, DeepSeek V4 Pro and GLM 5.3 — and three more saw it during the build itself: Claude Sonnet 4, DeepSeek V3 and DeepSeek R1, through Aider. That last group is the one that matters, because it means DeepSeek's lineage had prior exposure to this codebase, and Run 1's claim of no prior involvement is therefore weaker than it was stated to be. Independence is tracked at the lab, not the checkpoint: a newer model from a family that has already worked on the artefact is not an independent reviewer of it, and treating a version bump as a reset would make the safeguard ceremonial. Still clean on both lists: Meta, Mistral, Qwen, Cohere, Amazon, MiniMax. One consequence already acted on — Luna Pro was considered for Step 5 and rejected, because it had seen the Constitution during the hostile review, and GLM was chosen on the basis of being clean on both lists at the time. It no longer is. Every run spends independence that cannot be recovered, which is an argument for deciding what a reviewer is for before opening the room.

#32 Machine Learning Put On ICE, With Conditions Written Down

August 29, 2026

DECISION — ON ICE

“ML feels like the milestone that makes the project serious. That instinct is what this document exists to catch.”

Viktor asked what could be counterproductive about adding machine learning — the obvious use being to fit the six bias weights against real outcomes rather than choosing them by hand — and then parked it. Not cancelled; parked, with conditions, which is the difference between a decision and a mood. The argument, verified against source before it was given: **training weaponises a dormant bug.** Item 2 is Compliant, but input-side `bfill` reaches the final bar through recursive indicators. Live mode only ever evaluates the current bar, so the leak has nothing future to pull from and stays inert. Training evaluates thousands of historical decision points, and at each one the backward fill pulls from bars that had not yet happened. The result is an excellent fit with no live edge, and the failure is silent — a good-looking equity curve, which is the most expensive kind of wrong this project could manufacture. **The training data does not exist:** `live_trading.py:210` does persist a decision, but nothing anywhere records what price did afterward and `engine_core.py:466` overwrites its state file each run rather than accumulating. Decisions without outcomes are unlabelled examples. **Item 11 makes fitted weights worse, not better:** fitting over features that are secretly the same feature produces large offsetting coefficients that swing on tiny data changes — crude weights wearing a lab coat. **Item 13's sentinels become learned signal:** `RSI=50` means the indicator broke, and a model would learn to predict from it. Two costs that are not bugs: reproducibility gains a whole new surface, since a decision would then depend on a weights file depending on a training run depending on a data snapshot, all needing version-pinning to the decisions they influenced; and explainability drops, which matters for a tool whose entire job is helping a human judge. Revisiting requires all five of: Items 2, 3, 11 and 13 fixed and verified by the harness; a working decision *and* outcome log with accumulated history; a backtester validated against known-answer cases; a kill condition written before the first fit; and an out-of-sample test that beats the current hand weights. If the fit does not beat 0.30 / 0.20 / 0.15 / 0.15 / 0.10 / 0.10 on data it never saw, that is a real result — it means the hand weights were fine.

#33 Step 5 — The Remediation Sequence, and What Verifying It Found

August 29, 2026

AUDIT STEP 5 — VERIFIED

“A fix made before its apparatus exists can be made, never resolved.”

Step 5 ran on GLM 5.3 — the pinned slug, not the floating `glm-latest` alias, which silently repoints when a successor ships and would have quietly poisoned this document's record of which model produced what. The instruction was Viktor's, not Claude's: it asked for position, dependencies, effort, risk-of-doing and risk-of-not per item, five direct questions, and an explicit flag wherever the reviewer's reasoning might be correlated with the model family that built the engine. Two edits before it went out — the build-timeline phrase removed, and a paragraph added about Phase A, because the instruction predated the harness and would otherwise have hidden three real findings from the reviewer. What came back organises sixteen items by verifiability prerequisite and shared edit seam rather than by severity, on the reasoning that severity decides membership in the gate sets and not position. Claude spot-checked every claim that could be checked against source; nine of nine held, including that `requests` and `colorama` are imported and undeclared while `ccxt` is declared and appears nowhere else in the codebase, that `close_time` arrives in the raw response and is discarded one line later so any staleness check must precede that line, that `engine_version` exists in config and is written nowhere, and that the indicator cache can never hit across runs because its key embeds the last close and the dict is per-process. Two things came out of the verification rather than the plan. The codebase contains two comments acknowledging previous instances of the same defect class — a gate comparing against a string literal no producer ever emits — which makes the dead `trend_failure` gate the third instance and argues for a class-level check rather than a third individual fix. And Claude came within one sentence of telling Viktor the plan understated that fix, on the grounds that the STRUCTURE column is never assigned — it is, at `structure.py` line 4502 of the audit bundle, via a `.loc` form the first search pattern did not match. That would have been the third false assertion of absence on this project. The exhaustive re-search caught it; the pattern is now a working note.

#34 The Machine Was Rebuilt, and the Baseline Survived It

August 30, 2026

REFERENCE — FILED FOR THE RECORD

“A golden baseline is a claim about an environment, and the environment was destroyed.”

Windows was reinstalled and every drive wiped. Nothing of value was lost: the repository stood at `375334a` on GitHub and nothing lived only on disk except the gitignored `logs/` tree. The reason this earns an entry rather than a shrug is Item 5. The golden snapshot is not a claim about the code alone; it is a claim about the code *in an environment*, and that environment was destroyed and rebuilt from nothing. `docs/audit_package/environment_before_reinstall.txt` records the library versions and the Python 3.12.0 that produced the original baseline. The rebuild resolved different ones — `numpy` moved from 1.26.4 to 2.2.6, `pandas` and `pandas-ta` unchanged — and the snapshot did not move. That was checked rather than assumed, and it is a real result. It is also the strongest available argument for what became audit Finding 15: **nothing in the repository required those versions.** The baseline held across a major version bump of the numerical library underneath every indicator, and it held by luck that nobody had arranged. Viktor's machine is now on Python 3.12.10 and the remediation sandbox runs 3.12.3; both produce identical suite results, verified on every delivery since.

#35 Step 8 — The Independent Re-Audit, and Five Criticals

August 30, 2026

AUDIT STEP 8 — RECORDED

“It tests that selected implementation details have not changed more strongly than it tests whether the engine is correct.”

Sequence item 16 executed. GPT-5.6 Luna Pro received the engine source, the test suite, the frozen audit copy of the register and a written instruction (`docs/audit_package/item16_review_instruction.md`), and returned `docs/audit_package/luna_pro_audit_report.md` — a full 44-rule verdict table, fifteen findings, a test-suite assessment and a release-gate determination. **Release gate: not met.** Five findings rated Critical: Item 3, abnormal volume reaching analysis unchecked; Items 8 and 13 together, a failed macro input rendered as an ordinary neutral reading; Item 13, partially invalid indicator columns passing without degradation; Item 11, bias and confidence reusing derived evidence as if it were independent; and Item 14, `AGGRESSIVE` selected from conviction and entry quality with no independent risk decision. Ten further findings at Major or Moderate, and twelve rules graded Not verifiable. Two things about this run belong in the record rather than in a summary of it. First, the sentence quoted above — the auditor's judgement of the suite — is the sharpest single observation any reviewer has produced about this project, and it was aimed at a harness built four days earlier specifically to be better than that. Second, see Entry #41: the model that produced this report was not on the clean list the Remediation Plan had drawn up for this step.

#36 Remediation Batches 1 and 2

August 31, 2026

MILESTONE — RECORDED

“Under pytest, a return is a pass.”

Two batches, one commit — `dba1b63`. **Batch 1, the skip mechanism.** The suite contained 46 occurrences of `if not _engine_available(): return`. A returning test is a passing test, so on a machine without `pandas_ta` the suite reported success for work it had not done — audit Finding 13, and precisely the defect Entry #29's harness had been built to prevent. All 46 became real `pytest.skip()` calls. Running the suite without `pandas_ta` then reported 46 SKIPPED instead of 46 false passes, which is how four tests that import `core.engine_core` directly, bypassing the guard, were found to fail outright in that environment. Recorded at the time and deliberately not fixed in the same batch — see Entry #38. **Batch 2, five unambiguous fixes.** The “full size” text in `decision_model.py` (Finding 8, path B); `main.py`'s bare `'Logs'` (Part 6 observation 2); `trade_quality_current` and the vacuous assertion that guarded it (Findings 10 and 13); the hardcoded `Lookback 8` on the panel, now interpolated from `config.STRUCT_LOOKBACK` (Finding 8, path A); and `np.isfinite` in risk validation. **None of those five named the audit finding it closed.** Three whole findings were later discovered to be already shut for exactly that reason — Entry #40 records the cost from the other side.

#37	Items 3, 11 and 14 — Delegated, Ruled, Implemented August 31, 2026	DECISION — ADOPTED

“Decide items 3, 11, 14 myself.”

Viktor's instruction, delegating all three trading-judgment calls rather than ruling on each in turn. Recorded as a delegation because Roles & Authority assigns this class of decision to Viktor, and this is the second time he has knowingly handed a specific class of them over. Commit `c4dfcc7`. **Item 3 — abnormal volume: reject, degrade or accept, per case.** All-zero volume is now rejected at `data/validation.py`, because there is no measurement from which to build a volume-weighted read. An isolated extreme spike is still accepted at that layer: the original “deliberately not implemented” reasoning held, in that a spike is real data and rejecting a run over a busy market makes the engine least available exactly when it matters most. What changed is that a spike no longer reaches every downstream score unflagged — `indicators.py` detects one above ten times the recent rolling median and records it as a degradation, capping confidence rather than substituting a value. **Item 11 — remove the duplicated factors rather than argue for their independence.** The earlier sequence-item-11 pass had removed one duplicated term and left three siblings standing. `bias_score` is now the one place all six weighted factors are combined and `confidence` is exactly its magnitude; `continuation_strength` no longer carries a trend-health-derived component; and `bias_engine.py` gained the explicit dependency-graph comment the audit's required action had actually asked for. **Item 14 — the labels survive, gated.** `classify_risk_regime()` already computed a four-tier regime, and only the EXTREME-or-not boolean ever escaped `validate_risk_parameters`. The regime is now threaded through as its own contract field, and AGGRESSIVE is refused when it is HIGH VOLATILITY RISK or worse. Direction and whether a trade is permitted at all are untouched; this gates intensity only. Ten net new tests. The golden snapshot moved, in exactly the four fields predicted before the run.

#38	Three Small Items, and the Four Tests From Entry #36 September 1–2, 2026	MILESTONE — RECORDED

“A test docstring naming what it did not fix is a debt with an address on it.”

Four commits. **Item 8/13, the macro half (`5d2cbbe`).** A failed macro-timeframe fetch left `macro_bias` at its initialised `"NEUTRAL"` with nothing added to `degradation`, so a failed higher-timeframe read and a genuinely flat one rendered identically on the panel. Both the validation-failure path and the processing-exception path now record a degradation. `test_the_macro_series_is_actually_read` had documented this in its own docstring as “recorded rather than fixed… a rider on sequence item 9's degrade ruling”; that assertion was rewritten alongside the fix, which is what the docstring had asked whoever fixed it to do. **Finding 15, dependency pinning (`a26c545`).** `requirements.txt` and `requirements-dev.txt` now pin exact versions, verified by building a virtualenv from empty and confirming `pip freeze` matched before the suite was run. Closes the risk Entry #34 records the project having survived by luck. **The four unguarded tests (`3d0b410`).** The ones Entry #36 recorded and left. Three are entirely about `Phase7Engine` and now carry the same `_engine_available()` guard as the rest of the suite. The fourth, `test_every_module_imports`, needed a different fix: it checks all 21 engine modules for import-time defects and only three of them need `pandas_ta`, so a blanket skip would have stopped checking the other eighteen for an unrelated reason. It now excuses that one exception by name and still fails on anything else. **Stale labels (`2be405f`).** Two golden-path tests still carried “EXPECTED TO FAIL until sequence item 12” in their docstrings; item 12 had fixed both defects before the Luna Pro audit began and both tests had been passing ever since. Documentation only — but a docstring that lies about the state of the code is the same class of defect as a panel line that does.

#39 Finding 3 — the Critical That Was Never Scheduled

September 2, 2026

FINDINGS — FROM EXECUTION

“A derived document cannot tell you what it never contained.”

Found by reading `luna_pro_audit_report.md` end to end rather than the roadmap written from it. Finding 3 — Item 13, partially invalid indicator columns passing without degradation — is one of the audit's five Criticals and had never been entered into any roadmap, so it was never scheduled, never ruled on, and never noticed missing. From 31 August to 2 September `PHASE7_NEXT.md` stated that remediation of the five Criticals was complete. Four of them were. Commit `108cc9f`. **The defect.** Every indicator guard asked one question, `.isna().all()` — did the calculation return nothing at all. That cannot catch a series with 299 good values and no value at the bar the decision is made on, and it could not even in principle: `clean_series(method="forward_fill")` had already filled that gap with the previous bar's number before the guard ran. **The class was wider than the two instances the audit named.** Injecting a trailing NaN into each indicator in turn, before the fix: ATR, RSI, ADX, SuperTrend and both EMAs — every one carried a stale prior-bar value into the decision row, and not one recorded a failure. Two had no guard at all to fix, the SuperTrend *level* and the EMAs. The same trailing fill ran on the raw OHLCV columns, turning a truncated final candle into a synthetic bar repeating the previous close. **And the consumers were re-fabricating what sequence item 9a had removed.** A missing RSI fell back to `50.0`, inside the “not extended” band, scoring the full 15 of 15 — while `indicators.py`'s own failure text told the operator it scored 0 of 15. A missing HVN fell back to `close`, making the distance exactly zero and scoring 12 of 12; that is byte-for-byte the defect item 3 had fixed for VWMA six days earlier, sitting forty lines above it. A missing ATR fell back to `close * 0.02`, the flat constant item 9a had deleted — on the pinned fixture, 0.016035 against a real 0.010554, a 52% overstatement of the distance that sets the stop. **Ruled:** a value not measured at the decision bar is absent, and absent means that indicator failed — which hands it to the degradation machinery that already exists. No new policy was invented; the existing one was applied one row over. Verified end to end: before the fix, a trailing-NaN ATR produced `degraded: False`, `missing_inputs: []`, a full stop and three targets, and moved entry quality from 45.18 to 45.25 — a different answer, reported clean. Eighteen new tests, fourteen of which fail against the pre-fix code; the other four are controls that must pass on both sides. The golden snapshot did not move, predicted in advance and for the right reason: real pinned data has zero trailing NaNs, which is why this went unseen. **Recorded against Claude:** the first draft of that test file reintroduced the defect fixed in `3d0b410` the same morning — five of its tests imported `indicators.indicators` with no `_engine_available()` guard and errored instead of skipping. Caught by running the suite in a `pandas_ta`-free virtualenv before packaging. Same class, same day, inside the fix for the class.

#40	The Status Sweep, and a Count That Was Wrong by Three September 2, 2026	OBSERVATION — FILED FOR THE RECORD

“Neither document knows the state of the code.”

Claude told Viktor on 1 September that nine Major and Moderate findings remained open. That number came from the audit report rather than from the code, and it was wrong by three. Checked by opening each location the audit quoted rather than trusting the summary: **Finding 8** (inaccurate user-facing claims) was fully closed — `Lookback` interpolated from config and the “full size” text gone, both in batch 2, with the macro path closed by `5d2cbbe`. **Finding 10** (the unconsumed `trade_quality_current`) was closed in batch 2. **Finding 15** was closed at `a26c545` the previous day. None of those three commits mentioned a finding number, which is why nobody knew they had shut one; Entry #36 records the same omission from the authoring side. Genuinely open: Findings 6, 7, 9, 11 and 12 in full, plus the remainders of 13 (the plotting test still asserts only the absence of ERROR records) and 14 (the required-shape assertion still omits `provenance`, `degradation` and the log path). Seven, not nine. The status table now in `PHASE7_NEXT.md` cites a file and a line for every verdict, so the next reader can re-check it rather than trust it. This is the mirror image of Entry #39: that one was the plan missing what the report had, and this one is the report not knowing what had been fixed since.

#41	The Step 8 Auditor Was Not on the Clean List September 2, 2026	OBSERVATION — FILED FOR THE RECORD

“The party that made the error should not be the party that decides the evidence of it disappears.”

`Phase7_Remediation_Plan.pdf`, 29 August, records that Luna Pro was considered and rejected for Step 5 because it had already seen the Constitution during the hostile review, and names the six families still clean on both lists for the Step 8 re-audit: Meta, Mistral, Qwen, Cohere, Amazon and MiniMax. **Step 8 was run on Luna Pro the following day.** Entry #31 sets out why the ledger exists — independence is tracked at the lab, not the checkpoint — and Entries #18 and #20 record Grok being removed from the auditor plan for precisely this reason, having read the Constitution in an earlier conversation. The same disqualifier applied here and the ledger was not consulted. **What this does not mean.** The findings are sound. Every one was checked against source before any of it was fixed, and Finding 3 was verified by running the engine rather than by trusting the report. Nothing here argues for discarding the report or redoing the work it produced. **What it costs** is the one thing the ledger buys: the next re-audit cannot treat agreement with this one as independent corroboration, because it would be the second reading by a model that had already read the standard. The next re-audit should go to one of the six families named above, and its agreement with Luna Pro should be read as agreement between two readings, one of them contaminated — not as confirmation. Recorded rather than quietly noted, on Entry #21's reasoning.

#42 **The Document Set Read End to End**

September 2, 2026

REFERENCE — FILED FOR THE RECORD

“A document you have sampled is a document you have not read.”

Viktor asked whether the Constitution needed upgrading. Answering that honestly required reading all nine governing documents rather than sampling them — Entry #17's lesson, and rule 17 of `PHASE7_NEXT.md`'s earned list, which exists because Claude once declared this exact PDF clean after searching four guessed phrases. **The answer is no, and the project had already ruled so.** Entry #26, written the hour the freeze lifted: the useful next act is fixing the Non-compliances rather than reopening the rulebook that found them. Four things the read surfaced that were in no working document. **One: the release gate is in force and blocks on Item 6.** Viktor's 29 August ruling raised Traceability from Major to Critical; Item 6 is the audit's Finding 7, still open. Until it lands *and is re-audited*, no output may be relied on for a real trading decision, and the backtest architecture is not built. Nothing in `PHASE7_NEXT.md` had said so. **Two: the adjudications were ruled on 29 August**, in Roadmap Revision 4, whose title says exactly that. Claude reported four still open, having read the Engineering Notes and the Constitution — both of which stop before that ruling — and not the Roadmap. Only the Item 20 amendment remains. **Three: Item 20's gap is already served operationally.** `Phase7_Credential_Security_Protocol.pdf` §6.2 covers telemetry, crash reporting, error aggregation and usage analytics, and is tagged “Enforces: Tier 1, Items 20 and 21” — written the same day as the invariant whose channel list omits it. The pending amendment tidies the register to match a practice that already exists; it does not close a live hole, and the engine holds no credentials at all. Unblocking it needs one uninvolved model given a text extract, not the PDF that Gemini's and Copilot's classifiers refused. **Four: the documentation is a week in arrears** — this log stopping at #33 being the largest part of it, alongside no Change Impact Records since the standard took effect on 25 August, no Version History row for Step 8, and two companion documents still pointing at Rev 6. **One check nobody has run.** The item-16 instruction requires the re-audit to grade against the frozen audit copy rather than the live version annotated with the outcomes of a previous audit. `Phase7_Constitution_v1.0_RATIFIED_AUDITCOPY.pdf` is 69,656 bytes against the live document's 112,290; the smaller size is consistent with a frozen copy and there is no reason for alarm. It is still worth opening once to confirm it carries no AUDITED, AMENDED or DEFECT row, because it is a one-minute check on whether the re-audit graded the exam or the answer key. *(An earlier draft of this entry compared that byte count against the 101,831 recorded in* `Phase7_Audit_Execution_Instructions.pdf` *and called it a mismatch. That was wrong — those instructions describe the Step 3 package for Kimi K3, a different audit with a different material set. The claim was made to Viktor before it was checked. Corrected here rather than deleted.)*

#43	Findings 6 and 7 — the Last Critical, and What Reconstructable Means September 2, 2026	DECISION — ADOPTED

“A hash proves the input changed. Only an archive lets you rebuild it.”

Item 6, Traceability, which Viktor's 29 August ruling raised from Major to Critical — and therefore the finding holding both the release gate and the backtest gate shut. Commit `302db8b`. **The gap.** Sequence item 12 built a decision log: what the engine concluded, plus a five-field fingerprint of what it saw. The fingerprint was a last-candle timestamp and a row count, and two different frames can share both. Nothing stored told them apart, so “reconstructable” was a word in the Constitution rather than a property of the engine. **Viktor's ruling: hash AND archive, pruned at ninety days.** The Constitution says under Item 5 that the retention decision is one “the audit should force explicitly rather than leave implicit”; this is that decision made rather than assumed. The two halves do different jobs and the difference is the whole design. **The hash detects** — it costs nothing, it lives in the log, and the log is never pruned, so a decision from two years ago can still be checked against data fetched today. **The archive reconstructs** — the actual candles, the only thing that can rebuild a run whose source has since changed, and the only part with a cost, which is why it is the only part with a limit. About 31 KB a run; roughly 2.8 MB steady-state at the cap. Past the window a run does not become unverifiable — it becomes *verifiable but not rebuildable*, and the record says which by whether the file is still there. **Two design points worth keeping.** The hash is taken on the OHLCV as validated and *before* any indicator column exists: hashing the final frame would fingerprint this engine's own arithmetic together with the market data, so changing an indicator length would present as the exchange having changed its candles. And the canonical form is written out explicitly rather than leaning on `to_csv`, because pandas has changed its float repr and NaN spelling across versions this project has already run on — Entry #34 records numpy moving 1.26 to 2.2 across the machine rebuild alone. **Also closed:** the six bias weights live in `bias_engine.py`, not config, so changing 0.30 to 0.35 changed every decision the engine makes while leaving two runs byte-identical in the log. They are now fingerprinted, read live from the module rather than restated. **Verification.** Fifteen tests, and the central one does not inspect fields — a test asserting `"lineage" in decision` would pass just as happily over a section full of nulls, which is precisely the shape Luna Pro's assessment of this suite warned about. Instead it takes the archive, rebuilds the candles from it, runs the engine again against the rebuilt data and nothing else, and requires the identical decision. The golden snapshot moved exactly as predicted before the run — `lineage` added, `provenance` gained seven keys, and not one existing value changed. This commit adds a record and alters no output.

#44	A Defect That Shipped, and the One Machine That Could See It September 2, 2026	OBSERVATION — FILED FOR THE RECORD

“Every hash matched. Every number matched. The whole difference was one separator.”

The Findings 6 and 7 work built the archive path with `os.path.join`, which returns a forward slash on Linux and a backslash on Windows. That path is written into the decision log, so the same run archived on the two platforms recorded two different strings for one file — and the golden snapshot, baselined in a Linux sandbox, could only ever match on Linux. It failed on the first run on Viktor's machine. **Two things about it are worth more than the fix.** The first is the near miss: `decision_log.log_path()` has the identical shape and escapes the bug by luck rather than design — `config.LOG_DIR` already ends in a separator, so its single join never inserts one. The archive path joins a directory level deeper, which is where the backslash appeared. Existing code working was not evidence that the pattern was safe, and it was read as if it were. The second is why nothing caught it earlier: **the entire verification for that commit ran on Linux, where the bug is invisible**, and a golden snapshot only ever compares a value against itself on whatever platform baselined it. The patch had been called “verified against a fresh checkout” — true, and on one platform. Fixed by normalising to forward slashes for both the file I/O and the recorded value, since Python accepts them on Windows and one spelling then serves both purposes. Pinned by a test that fails on Linux too if it regresses. Recorded as its own entry rather than folded into #43, because a defect that reached the operator's machine inside a commit whose message claimed thorough verification is not a footnote to that commit.

#45 The Critical the Audit Did Not Have, Found by Running the Engine

September 2, 2026

FINDINGS — FROM EXECUTION

“Reading finds what is written. Running finds what happens.”

Viktor ran `python main.py` against live MEXC data the evening Findings 6 and 7 landed. AEROUSDT 4h came back bearish on every measure the engine has — bearish bias, bearish regime, bearish structure, an LH-LL sequence, strong bearish distribution, SuperTrend freshly flipped — with one dissenting reading, `MACRO TREND: BULLISH`. It printed **DECISION: CONSERVATIVE LONG**, with the stop at $0.4889 *above* a price of $0.4725 and three targets descending to $0.4233. **A long label on a short plan.** Every number in that plan was correctly computed from a correct bearish analysis; the word attached to them was not. An operator following the DECISION line would have bought an instrument the engine had just analysed, in detail and correctly, as a short. It also printed “Bias is bullish and the broader macro trend agrees” four lines above its own Validation Note reading “The higher timeframe disagrees with this bias.” **Two causes.** `decision_model.py` opened a direction from any of three independent sources — `raw_bias or long_signal or macro_bias` — so the macro clause alone was enough; and `trend_health >= 50` then passed because trend health is an UNSIGNED magnitude, so a strong bearish trend scores 69 and no bearish evidence anywhere in the run could block it. The bearish block below never ran, because the bullish one returned first. Meanwhile `risk_model.py` builds stop and targets from `detailed_bias` alone. Two direction sources, never reconciled, in two modules neither of which knew the other existed. **Viktor ruled: bias is the sole direction source.** Macro keeps its existing 10% vote inside `bias_score` and gets no second, overriding one — letting it override the blend counts one piece of evidence twice, which is Item 11 in the module that picks the side. Entry-zone signals lose the same privilege for the same reason. Two further instances surfaced in testing that nobody had seen: the mirror case returned CONSERVATIVE SHORT, and a NEUTRAL bias with a long entry signal returned AGGRESSIVE LONG — the most confident label the engine has, from no directional view at all. **And a guard, because narrowing a source is not the same as checking.** `_refuse_incoherent_plan()` reads direction off the targets themselves and refuses any action whose label contradicts its own levels. Deliberately a refusal rather than a relabelling: one of the two sources is wrong and nothing inside that function can tell which, so NO-TRADE is the only answer available that is certainly not the wrong one. **What this says about the audit.** Luna Pro read the source and the tests and did not find this. Neither did four earlier passes across three models. It needs bias and macro to actually disagree on live data, and no pinned fixture makes them.

#46	An Unwritable Log Directory Killed the Whole Run — and Claude Was Overruled on It September 2, 2026	DECISION — ADOPTED

“A disk problem must neither destroy an analysis nor veto one.”

Found while writing the halt-safety test for the raw-input archive, and deliberately left unfixed in that commit: fixing it there would have let that test pass for a reason unrelated to what it was written to prove. **The cause was smaller than “the engine cannot log.”** `route()` opened with two unguarded `os.makedirs` calls that wrote nothing. Every writer in this engine already creates its own directory on demand inside its own error handling — `decision_log.write` returns None, `_save_state` warns, `plot_engine_chart` warns, `lineage.write_archive` returns None. The two calls duplicated all four and added a failure mode at the worst point in the run: the top of `route()`, before anything had been computed. An unwritable log directory raised there, the broad handler reported “Router execution failed”, and the operator lost the entire analysis rather than merely the log of it. Four independently recoverable conditions collapsed into one total failure. Same class as sequence item 14's own `REQUIRED_DIRS` finding, which removed the list and left these two calls standing — earned rule 22 collecting on the item that wrote it. **Viktor ruled: a run whose decision log cannot be written still authorizes a trade.** It warns, the panel makes no claim that anything was logged, and the operator decides. His 29 August degrade-not-halt ruling applied literally. **Claude recommended the opposite and was overruled.** The argument made and rejected: Item 6 is Critical, and a trade taken on a decision that left no trace is unauditable by construction — which differs from a failed *archive*, where the hash still lands in the log and the run stays verifiable even though it is no longer rebuildable. When the log itself fails there is no record at all. Recorded here because the difference matters to whoever audits this next: a ruling with its trade-off on the table reads differently from an absence, and the cost is stated plainly — the one decision an operator acts on without a record is the one an auditor would ask about first. The tests pin the ruling and say so in their own docstrings, so a later reader who thinks the assertion looks wrong finds the reasoning instead of guessing. **Categorisation note:** a failed write is not added to `degradation`. That list blocks trading and is about inputs the ANALYSIS was computed without. Nothing here was missing from the analysis — only from the filing of it.

#47 Letting the Suite Reach the State That Hid the Defect

September 2, 2026

MILESTONE — RECORDED

“I really should have run the engine more often, but I was too caught up in the workflow.”

Viktor's observation after Entry #45. The better answer than resolving to be more diligent — which decays — is to make the suite able to reach the state, so that it does not depend on anyone remembering. **The reason nothing caught it is exact: no pinned fixture makes bias and macro disagree.** The committed series trends one way and its daily aggregate trends the same way, because the aggregate is built from it. A suite cannot reach a state its fixtures never enter, so a hundred and seventy tests passed over a defect that only exists in that state. `tests/test_timeframe_disagreement.py` generates a series where the timeframes genuinely disagree: a long rally, then a sharp multi-day break, so the daily EMA-50 lags and still sits below the daily close while the 4h structure has decisively turned. An ordinary market condition of the kind that happens most weeks, not a contrivance built to trip a branch. Generated rather than committed, following `test_golden_path._write_pinned_set`'s discipline: one pure function, all three files derived from it, and no RNG anywhere — the wobble is `sin()`, so the series is byte-identical on every machine and every numpy, which matters on a project whose baseline has already survived a numpy major version by luck. **The parameters were chosen with margin rather than at the edge.** A 0.16 drop put the daily close within 0.0007 of its EMA-50; 0.14 leaves real daylight. A fixture that only just produces its condition stops producing it the first time an indicator length changes, and stops silently. **Two of the seven tests assert that the fixture still splits the timeframes**, before anything is asserted about behaviour — if it ever stops, every other test in the file would go on passing while testing agreement. That is section 7.3's “setup contradicts what it claims to test”, and a file like this is exactly the shape that fails that way. The property pinned is deliberately not “the engine returns WAIT here” — that is today's answer to today's thresholds. What must never be true at any threshold is a LONG label above descending targets. Four of the seven fail against the pre-fix code, reproducing the live run from generated data.

#48	Three Kinds of Exposure, and a Ledger That Was Wrong for a Week September 2, 2026	DECISION — ADOPTED

“Nobody could have caught it by thinking harder. Only by reading the bill.”

Viktor delegated this one — “This is your call and I trust you make the right decision” — so it is recorded with its reasoning rather than only its conclusion, and can be overturned on the argument. **The project held two positions and had not noticed they cannot both be true.** Entry #18 resolved Grok's prior exposure by running Step 3 “in a fresh conversation with no shared memory of the earlier one.” The Remediation Plan, eleven days later, rejected Luna Pro for Step 5 *because* it had read the Constitution during the hostile review — the same kind of exposure, treated as permanent. Entry #41 then graded Step 8 against the stricter reading without noticing the looser one existed. Which answer you got depended on which document you opened. **Three different things were being recorded as one.** SESSION exposure is a model reading the document in a chat; TRAINING exposure is the artifact being in the weights; LINEAGE exposure is a sibling from the same lab having worked on it. Entry #31's rule — independence is tracked at the lab, not the checkpoint — was written for the second and third. Applying it to the first is a category error, and it is why the clean list emptied faster than it needed to. **Ruled: session exposure is cleared by a fresh conversation where training-on-conversations is off; training and lineage exposure never clear.** Where it cannot be established whether a session fed training, treat it as permanent. This ruling is LESS strict than the reading the project had been using, which is worth stating plainly: over-strictness had a measurable price here, in models spent for a reason a fresh conversation removes. Kimi K3 becomes available again. **Then the assumption was checked rather than assumed.** Viktor exported the OpenRouter activity log — 713 requests, every model, every route. All `variant=standard`; not one free endpoint in the entire history, so no material was ever sent anywhere that trains. It also dated the Step 3 truncation exactly: `completion=16,384, finish=length`, at 2^14 output tokens — Entry #23's “default token ceiling” with a number on it. **And it showed the ledger was wrong.** Five models ran through Aider on 22–24 August while the engine was being built. Entry #31 names three of them and misses `mistral-nemo` (71 requests) and `claude-3-haiku`. Mistral's lineage worked on this codebase through exactly the mechanism used to disqualify DeepSeek, and had sat on the clean list of eligible auditors for a week. Entry #31's table was written from recollection. The correction matters less than how it was found: every provider bills per request, and the bill cannot be mistaken about what was called.

#49 The Round-2 Package, and the Tools Nobody Turned Off

September 2, 2026 — **PROCESS — RECORDED**

“Eight tools were enabled. One `git clone` would have handed over the answer key.”

The package the round-2 auditor receives is now **generated** by `docs/build/build_audit_package.py` rather than assembled by hand. Entry #24 records what hand assembly cost: “a false claim found in Claude's own Step 2a package.” A script that walks the repository and computes the manifest from the bytes it wrote cannot make that class of error, and it *refuses* to build if `docs/` would ship, because intending to withhold the answers is not the same as withholding them. Three things Luna Pro never got: version-control history (metadata only — subject lines withheld, since “Audit Findings 6 and 7” leaks a finding number and its outcome in eleven words), the project files that decide how dependencies are pinned, and execution transcripts. Five rules came back Not verifiable in Step 8 *only* because history was withheld; that was a defect in the package, not the project. **The instruction's Rev 2 states two things the project would rather not say:** that Step 8's own auditor was not on the clean list, and that this repository is public, so any model trained since may hold the codebase in its weights — a channel the ledger cannot see and which applies to the chosen auditor as much as anyone. It also says plainly, in its own section, what this round cannot be: the fixes shipped with comments and docstrings describing the defects in detail, so this is not a blind re-discovery but a check of whether the claimed fixes are real. **And the near miss.** The run started with eight tools enabled, including a shell, web search and web fetch. Nobody chose that and nobody checked. It probed an empty sandbox and found nothing — but the repository is public, and one command would have handed over PHASE7_NEXT.md, the previous audit report, the Engineering Notes and every commit message. Three further tools were worse in a quieter way: Fusion, Advisor and Subagent would each have silently put other models inside the audit, leaving the report labelled with one model's name while being partly the work of models nobody chose. It was caught by looking at a screenshot at the right moment. That is luck, and earned rule 28 — written four hours earlier the same day — says luck is the signal to change the setup rather than to be more careful next time. A pre-flight checklist now exists, and the omission was Claude's: the setup walkthrough covered model, provider, max tokens, streaming, file parser and data retention, and never mentioned the tools.

#50 The Standard the Auditor Could Not Read

September 3, 2026 — **PROCESS — RECORDED**

“This file is named but its content is not there.”

The round-2 re-audit has now failed three times, and not once on its own merits. Attempt 1, Qwen3.8-Max: repeated provider-side failures before grading began. Attempt 2, Kimi K3: it received the Constitution as a PDF, could see the filename and not the contents, said so, and stopped. Attempt 3, Qwen again: the same defect. **All three refusals were correct behaviour.** The common factor across two different providers on two different days was the PDF, so the Constitution now ships to auditors as text — a `pdftotext -layout` extraction carrying the source PDF's SHA-256 in its own header and verified byte-for-byte against a fresh extraction. Earned rule 28 applied: a format the recipient cannot open is not fixed by asking the recipient to try harder. **One thing recorded against Claude.** The first diagnosis was built from a truncated sentence visible in a screenshot — “this file is named but its content is” — and asserted as a parsing failure. The full reasoning, recovered later, said something stronger and more specific. The fix was right; the confidence behind it was not earned at the time it was stated.

#51 Eleven Claims, and the Rule About Believing Them

September 3, 2026

FINDINGS — FROM AN UNFINISHED AUDIT

“An audit that produced no report still produced eleven defects.”

Attempt 3 could not read the standard and correctly refused to grade 44 rules it had not seen. It then ran the instruction's own Section 7 checks — which are defined in the instruction rather than in the Constitution — and returned eleven observations. Its reasoning is preserved at `qwen_reasoning_1.txt` through `_4.txt` in the repository. **Every one was treated as a CLAIM and checked against source before a line of code was written.** None evaporated. Two got worse under checking. That discipline is the entry: a reviewer's finding is a hypothesis with a file and a line attached, and the difference between an audit and a rumour is whether someone opened the file. The eleven span fabricated levels, false claims about files, a run identity that did not identify the run, a smoothing mismatch between an indicator and its own fallback, an unset HTTP timeout, and a directory created as a side effect of an import.

#52 The Run Hash That Did Not Identify the Run

September 3, 2026

DECISION — ADOPTED

“Two different trading plans on the same candles, recorded as the same run.”

`core/decision_log.py` has carried this sentence since the Finding 6 fix was written: the audit's required action asks for “all decision-affecting configuration, **including risk-model multipliers and bias weights**.” Fifteen lines below it, the fingerprint dictionary named one module — the bias engine. The weights were covered. The risk multipliers were not. Changing `ATR_STOP_MULT` from 1.2 to 1.5 moves the stop and all three targets on every run, and left the hash, the config snapshot, the module snapshot and provenance byte-identical. Item 6 is Traceability, raised to Critical by Viktor on 29 August — broken inside the fix written to satisfy it. **It was not a name missing from a list.** The multipliers were instance attributes on RiskModel, and the snapshot mechanism reads module attributes; there was nothing for the list to name. Viktor ruled the constants move to module level rather than the mechanism grow to reach inside objects — the smaller change, and the shape the one compliant module already had. `__init__` is gone entirely, so no instance state can drift from what the record reports. Proven behaviour-preserving across 1,728 stop and target combinations and 192 regime combinations: zero differences.

#53 A Test That Described the Defect It Could Not Catch

September 3, 2026 — **PROCESS — RECORDED**

“”…so the default never fires and the panel prints ‘saved to None’.””

That sentence has been the failure message of a test in `test_traceability.py` since sequence item 12. It is correct. It describes, exactly, a defect that was live in the engine the whole time — and the assertion beneath it only checked that the OLD hardcoded default was gone from `panel_render.py`. The None arrived from somewhere else: the router built the decision object with `str(chart_path)`, and `str(None)` is the four-character string “None”, which is truthy, which passes the panel's own gate. The panel printed **Chart saved to None** — a file that does not exist — for the life of the engine. **A second test had the same gap from the other direction.** `test_unwritable_log_dir.py` already pointed the chart directory inside an unwritable path, so every run of that file drove the engine through this exact state, and it asserted only the log path and the archive. Nobody looked at the chart. A test whose failure message documents a defect it cannot detect is worse than no test: to anyone auditing the suite it reads as coverage. The claim is made in one file and the value built in another, so the check now follows the value.

#54 The Door Closed and the Window Left Open

September 3, 2026 — **DECISION — ADOPTED**

“Twelve of twelve points for a volume node that was never located.”

`StructureEngine.analyze()` wrapped each of five sub-routines in its own handler and substituted a value on failure — the previous bar's regime, “NONE”, “NEUTRAL VOLUME”, and for both volume nodes and the swing level, **the current price**. None of the five recorded anything, and the function had no channel to record through, so a crashed detector produced an ordinary-looking reading and the run was never marked degraded. **The volume-node case is Finding 3 again.** Entry quality scores structure proximity as `abs(close - hvn) / close`, which with hvn equal to close is exactly zero — inside the < 0.015 band, awarding the full 12 of 12. The comment describing that arithmetic has sat in `entry_model.py` since 1 September, written when Finding 3 was fixed. That fix changed the *consumer's* fallback to NaN and left the *producer* handing down a finite number equal to the price. It even cites rule 18 — fixed the instance it was looking at and left its twin. This is the twin of that twin. No handler invents a measurement now: an unlocated level is NaN, which every consumer already guards, an undetermined label says UNKNOWN rather than borrowing NEUTRAL, and every failure reaches the run's degradation list.

#55 Found by Reading the Panel, Twice in One Morning

September 3, 2026 — FINDINGS — FROM EXECUTION

“226 tests, and neither defect was found by any of them.”

The first was Claude's, shipped that hour. The patch above replaced a panel line that ended in a newline with a computed one that did not, and the engine printed `SWING STRUCT : $0.4700 (Lookback 8)STOP LOSS : $0.4636`. Thirteen new tests, a negative control, a full-suite run and a golden-snapshot check all passed — because every assertion asked whether a substring was present, and it was. Nothing in the suite rendered the panel and looked at its shape. Viktor found it in one live run, which is earned rule 30. **The second was not Claude's.** The same panel printed `MACRO TREND: BULLISH` and, four lines below, “The higher timeframe is neutral.” The branch that produced it fired when the BIAS was neutral and then made a claim about the MACRO — false on every neutral-bias run, which is a common state and was the state that morning. Same shape as the contradiction that exposed the direction Critical the day before. The logic is now a named function with a fourth branch, extracted so the four cases can be tested at all: inline, the false branch could only be reached with market data that happened to produce a neutral bias under a directional macro, which is why no fixture held it and no test could have.

#56 A Conversation Is Not a Record

September 3, 2026 — PROCESS — RECORDED

“A full day of findings, rulings and sizings existed in a chat window and nowhere else.”

Late in the session Viktor asked whether Claude would still have everything the following day. It would not have. Eleven verified findings with their file and line, four patches, three rulings made, three owed, sizes for the remaining work and the pre-audit checklist were all established in one conversation and written into no document. It was caught because he asked — which is precisely the save-by-vigilance earned rule 28 exists to replace. **The fix is a handover check, run by Claude unprompted at the end of every session**, asking six questions: is the state in the entry-point document rather than only in chat; are the day's rulings recorded; does `git status` show untracked files that matter; are these notes current or the gap stated; are loose patch files unapplied; is any evidence still in a chat window. On its first run it found three files left on the floor, one of them sitting there since the previous day. **The same conversation split the project's goal in two.** Phase 7 is now both the technical portfolio project in Viktor's 2026–2028 career plan and an engine he intends to finish, and those end months apart. Portfolio-ready is reached, tagged, and only then does backtesting begin — on top of the tag, so that the phase which destroyed the previous build cannot damage the version he submits. Earned rules 30 through 34 were written the same day.

#57	The Trend the Bias Engine Could Not See September 4–5, 2026	DECISION — ADOPTED

“Continuation was zero, so the trend was flat — whatever the trend was doing.”

Ruling 3, and patch I. `bias_engine` was never told which way the trend pointed. It inferred direction from the sign of `continuation_strength`, so whenever continuation came out at zero the engine read the trend as flat — silently, and regardless of what the trend actually was. `indicators/trend_health.py` had computed the direction all along and thrown it away. It is now exposed as `trend_direction` and `trend_direction_sign`, and the sign is a **required** parameter of the bias engine, validated against (-1, 0, 1) and raising rather than defaulting — an inferred value replaced by a declared one. **The order in which this was decided is the entry.** How often the engine silently flattened a real trend is a question with an answer in the data, and the answer was obtained across 9,800 live bars *before* the ruling was made, not afterwards to justify it. Earned rule 36 was written from this: a ruling made first and measured second is a hypothesis with a decision attached to it.

#58	A Magnitude Offered as a Direction, and a Lean Too Small to Act On September 5, 2026	DECISION — ADOPTED

“A magnitude was being cited as evidence for a direction it cannot indicate.”

Rulings 1 and 2, both in patch K. **Ruling 1:** the reason strings the engine writes for a BULLISH or BEARISH call cited trend health as supporting evidence. Trend health is a magnitude; it says how strong, never which way. Sitting in a directional sentence it read as directional support that it was not. It now prints as `trend strength 78/100 (rising)`, with the direction word coming from the sign patch I had just exposed — ruling 1 was only buildable because ruling 3 landed first. **A correction to the record:** Claude had claimed the panel's TREND line carried the same defect. It does not. It prints two true facts side by side and asserts no relationship between them; only the reason strings made the claim. **Ruling 2:** an ACTION could be issued off a bias that barely leaned at all. `MIN_ACTION_BIAS = 30.0` now floors it, and is fingerprinted so it enters the run's identity. Claude first argued for folding it into the existing `RAW_BIAS_THRESHOLD` and reversed that while building it: one asks *does the blend lean far enough to call a side*, the other asks *is that lean strong enough to act on*, and one number tuned for both makes one of the two answers an accident of the other. Seven fixtures that built a bias with no score at all were **completed** rather than worked around — relaxing a floor to accommodate incomplete test data is how a floor stops meaning anything. Patch J, alongside, extracted `volume_agreement` into a module-level function for the same reason patch D extracted `macro_agreement`: logic reachable only through a full run cannot be tested in isolation.

#59 Three Patches Built on a Base That Had Already Moved

September 5, 2026 — **PROCESS — RECORDED**

“The staged copy is a snapshot, not a view.”

Patch J was generated against an `engine_core.py` staged before patch I landed, and failed 39 tests. Patch K was rebuilt twice against a `decision_log.py` missing an anchor added hours earlier. None of the three patches was wrong; the base each was diffed against was. **The failure mode is asymmetric, which is why it earned a rule.** A patch built on a stale base fails loudly if you are lucky and applies cleanly if you are not — and the clean application is the one that reaches the repository. Earned rule 35: re-stage every base file immediately before generating a diff. Recorded alongside it, because it is the same week and the same shape: a test written in this batch asserted on the substring `"support"` against a note that quotes a label reading BULLISH VOLUME SUPPORT. Earned rule 30 had been written three days earlier about exactly that, and was violated in a test written to guard its cousin. Rule 37 came from this: knowing a rule and applying it are different acts, and the rules that repeat are candidates for a mechanical check rather than a firmer intention.

#60 Explaining in the Right Place

September 5, 2026 — **DECISION — ADOPTED**

“A long thread pays for its own history on every turn.”

Viktor asked Claude to decide and implement how to reduce token usage, on one explicit condition: it must not weaken the engineering. Two changes, decided by Claude and accepted. **First, when a patch ships with a commit message, the chat reply carries only the predictions to check and the commands to run.** What the fix does, why, the weights, the verification and the mistakes made along the way are already in the message he is about to read; saying it twice is pure waste, and the chat copy is the one that does not survive the session. **Second, a thread is closed and a fresh one started once a piece of work has finished** — “Continue Phase 7” is enough to resume, and the handover check of Entry #56 is what makes that safe rather than reckless. What was explicitly **not** cut: negative controls, golden checks, equivalence runs, and naming every consequence of a change before it is made. Those are the reason the patches have held, and they were the condition. This must not become a reason to explain less. It is a reason to explain where the explanation survives.

#61 Unreachable Is Not the Same as Safe

September 5, 2026 — **FINDINGS — FROM EXECUTION**

"Nothing could reach it — which is exactly why two audit passes left it there."

Patch L, finding 6. Two fabrication constants survived in `engine_core.py` after item 9a removed their siblings from `indicators.py` and Finding 3 removed them from `entry_model.py`: a substitute trend dict `{"trend_health": 50.0, …}`, and `atr_val = … else current_price * 0.02`. Both were unreachable. `compute_trend_health` is total, so nothing could reach the first handler; section 2 halts when ATR is absent, so nothing could reach the second. **The trend one was worse than a fabrication, and only the negative control showed it.** Its substitute dict has no `trend_direction_sign`, which the next stage reads by subscript *outside* the try. Run against pre-fix code, a broken trend contract does not degrade the run — it dies with `KeyError: 'trend_direction_sign'`, reported as neither a trend failure nor a bias failure. Found by running the test, not by reading the code.

#62 A Fallback That Claimed an Equivalence It Did Not Have

September 5, 2026 — **DECISION — ADOPTED**

"A test had asserted, on that claim, that these paths cost nothing."

Patch L, finding 5. The RSI and ATR fallbacks smoothed with a simple moving average; `pandas_ta` smooths both with Wilder's RMA. So the sentence in `add_technical_indicators`' own docstring — "recomputes the same quantity by another route" — was false, and it was the stated grounds on which `test_degraded_state` asserts these paths are **not** degradations. On the pinned fixture, final bar: RSI 84.45 against `pandas_ta`'s 69.14, and an ATR 1.80% low. That test had been passing on a false premise for as long as it existed. **The choice was Claude's, and was made by reading the repository rather than by preference.** Viktor: *"None needs a ruling from me."* Match the smoothing, or record which path ran? Recording the path leaves `test_degraded_state`'s assertion false and adds a flag to explain why. Matching makes the assertion true. The repository already contained the argument; the job was to find it, not to have an opinion.

#63 The Run That Reported Two Errors and Ran Nothing

September 5, 2026 — **FINDINGS — FROM EXECUTION**

""2 errors" — and the result for the other 251 tests was unobtainable."

Both of these were found because the verification step ran, not because anyone was looking for them. **The `pandas_ta`-free control could not run at all.** `test_macro_agreement` and `test_volume_agreement` import a pure function whose module chain reaches `pandas_ta`, so both errored at collection; pytest reported two errors and executed nothing. That had been silently true since those files were written, and the degraded-dependency result nobody could obtain was the one the fallback work depended on. `pytest.importorskip` makes it the skip it should always have been — 0 errors, 159 passed, 88 skipped. **And a module-level object defeated a deliberate lazy import.** `SignalRouter.__init__` imports `engine_core` inside the function, which is what lets `live_trading` be imported without `pandas_ta`. A module-scope `LiveTradingSimulator()` called that constructor at import time and undid it. Verified in both directions rather than argued: pre-fix, `import live_trading` raises `ModuleNotFoundError`; post-fix it succeeds.

#64	Two Series Paired by Position, Correct by Accident September 5, 2026	DECISION — ADOPTED

“The printed label changed in 4 of 31 positions, and the evidence line gave no tell.”

Patch O, finding (a). Correlation and beta between AERO and BTC were computed by pairing the two series **by position** after both timestamp indexes had been discarded. In the ordinary case both fetches return the same 450 candles, so positional pairing is timestamp pairing and the code is right by accident — which is why it survived, and why the golden snapshot did not move when it was fixed. On the pinned fixtures the two indexes share all 450 timestamps and old and new agree to the last decimal. **It stops being harmless the moment the series differ by one bar**: a candle closing between two sequential API calls, an exchange gap, a stale feed. Measured on the fixtures by dropping one BTC bar from inside the window, at all 31 positions: correlation moved by a median of 0.105, beta by 0.135, and the printed label changed in 4 of the 31. `n_observations` read 30 either way, so the panel's own evidence line could not reveal it. The join is now on the timestamp index, and the observation count is the count of the aligned pairs.

#65	A Zone Whose Lower Bound Was Above Its Upper, and a Price Nobody Read September 5, 2026	FINDINGS — FROM EXECUTION

“ENTRY ZONE : $0.4981 - $0.4918”

Patch P closed finding (4) and a display defect that was not on any list. **The display defect was found by reading the panel.** `engine_core` assigned EMA_20 to `zone_lower` and EMA_50 to `zone_upper` unconditionally, so in an uptrend they printed inverted. `entry_model` swaps them before scoring, so the arithmetic was correct the whole time and only the display was wrong — which is why no test could have caught it. Every test asserted on the numbers. **Finding (4) was the familiar shape**: `close * 0.99 / close * 1.01` for a missing EMA pair, reachable only when an EMA is missing, which a healthy run never is. Five fabricated constants across two files, and the worst was not on the list — a close price that could not be read became `$1.00`, so a run with no data at all reported ACTIVE ENTRY ZONE, 30 of 30, against a price nobody had read. **Three of this project's defects have now been found by running the engine and looking at the output.** It remains the cheapest detector the project has. Left open and written down rather than assumed: whether any other reader of the zone is missing the same swap.

#66 Five Components Summing to 39, Under a Total of 45.18

September 5, 2026

DECISION — ADOPTED

“Not wrong, unexplained — and for a number an operator acts on, that is its own defect.”

Patch Q. The entry-quality panel had been printing a total its own component list did not add up to, recorded on 2 September as “not wrong, unexplained”. An operator meant to act on that number either works the gap out themselves or stops trusting the number, and both are failures. Three causes, none visible anywhere: sub-scores rounded for display while the total was not, three confluence multipliers applied after the sum, and a clip at 100. **And a fourth thing the clip was hiding.** The five components add to **102, not 100**, while the docstring directly above the list said 100. With full confluence a perfect setup reaches 118.08 and loses the difference silently — the clip was concealing an arithmetic error in the scoring definition itself, not merely bounding an output. The panel now prints the subtotal out of 102, the multiplier broken into its three factors, and a Clipped line when the clip actually fires. The column adds up.

#67 One Fact Declared in Four Places

September 5, 2026

OBSERVATION — FILED FOR THE RECORD

“Nine new fields vanished between the engine and the panel, with no error and no failing test.”

Recorded, not fixed. `signal_router.py` rebuilds the entry block **field by field**. Anything the engine adds that its list does not name is dropped before the panel or the decision log ever sees it — silently. Patch Q's nine new reconciliation fields disappeared exactly that way on the first attempt: the lines simply did not appear, and nothing failed. The entry block's shape is now declared in four places — `entry_model`'s return, `engine_core`'s dict, `signal_router`'s rebuild, and `decision_contract`'s TypedDict — four copies of one fact, which is the defect class this project has recorded more often than any other. **Restructuring it is larger than any single finding warranted, so it is written down here instead of being done quietly at the end of a patch about something else.** Earned rule 24 in a different costume: the fix that belongs in its own commit does not get smuggled into one that has already made its claim.

#68 Four for Four: the Fabrications That Survive Audits Are the Unreachable Ones

September 5, 2026

OBSERVATION — FILED FOR THE RECORD

“Nothing tests a path a healthy run never takes, and nothing else notices it is wrong.”

Item 9a, Finding 3, Finding 6 and Finding 4 were four separate findings across three weeks, and all four removed fabricated constants from code paths that a healthy run never reaches. That is not a coincidence and it is now the strongest generalisation this project has produced. **A constant on a live path is checked by every run and every fixture. A constant on a dead path is checked by nothing** — not by the suite, which cannot enter the state, and not by a reader, who sees a defensive default and moves on. So it survives audit after audit while looking exactly like prudence. Unreachable is not safe: it is one edit to an invariant elsewhere from becoming the live path, and on the day that happens the engine will print a fabricated number with complete confidence and nothing in the record will say where it came from. Finding 6 proved the sharper version — its dead branch would not merely have fabricated, it would have crashed the run with a `KeyError` attributed to neither stage. **The generalisation worth carrying to other projects:** when auditing for invented values, search the unreachable branches first, not last.

#69 A Test Suite That Could Reach the Internet

September 5, 2026

OBSERVATION — FILED FOR THE RECORD

“Two rejected CONNECTs to an exchange, from a run in which every route was pointed at a dead port.”

During a full-suite run the sandbox's egress proxy logged two rejected CONNECT attempts to `api.mexc.com:443`. Every test that routes or fetches sets `base_url` to a dead local port first, and each of them was checked; the source of the two attempts was not identified. **Recorded rather than chased, and recorded rather than dismissed.** The attempts were blocked, so nothing reached the network and no result depended on what an exchange returned — this run. A suite that *can* reach the internet is a suite whose result depends on the network on some other run, on a machine with no proxy in front of it, and the failure would arrive as an unexplained flake rather than as an error naming its cause. It is written here so that the next unexplained flake has somewhere to start, which is what this log is for.

#70	The Table Is Empty and the Gate Is Still Shut September 5, 2026	MILESTONE — RECORDED

“Every queued fix has landed. That is not the same as a clean report.”

Between 3 and 5 September the suite went from 226 tests to **319 passing**, and the “What is left to fix” table in `PHASE7_NEXT.md` reached zero rows — three rulings built, and all seven queued findings closed, across patches I, J, K, L, M, O, P and Q. **The release gate stays shut, and the distinction is the entry.** Unresolved in this project means fixed *and* re-audited; every one of these fixes was written by the party under audit, and a fix nobody independent has seen is a claim, not a verdict. What stands between here and the re-audit is the five-item pre-audit checklist, and the first two items on it are the two that were skipped last round: rebuild the package against current code so the auditor grades what exists, and commit everything before building it, so the packaged files match a commit rather than a working tree. **A completed audit is not an open gate either.** If the report returns a Critical, the gate stays shut and the cycle repeats. The milestone is a clean report, not a finished one.

#71 The Bill Knew, Again

September 5, 2026 — **PROCESS — RECORDED**

“Three Qwen calls, then Kimi. The record has the last two attempts the wrong way round.”

Earned rule 29 says to reconstruct a factual record from the provider's billing log rather than from memory, because every provider bills per request and the bill cannot be mistaken about what was called. It was written on 2 September after the independence ledger turned out to have been written from recollection. Applied a second time, to the question of which model ran which attempt at round 2, it found three things. **One: the order is wrong.** The log for 2 September reads Qwen 10:53, Qwen 15:38, Qwen 15:42, Kimi K3 15:58. Entry #50 records attempt 2 as Kimi and attempt 3 as Qwen; Kimi ran last, after all three Qwen calls, so the eleven observations of Entry #51 came from a Qwen run that preceded it. Viktor's recollection — “we ran Qwen three times” — was exactly right. **Two: Claude's own reasoning was wrong, and in a way worth naming.** Having found that those transcripts cite a superseded revision of the reviewer instruction, Claude offered the files' modification times as evidence that the correct package existed before the run and was not used. A modification time records when a transcript was *saved*, not when a model *ran*. Ordering of writes was presented as ordering of events; the claim was withdrawn. **What the clock does establish, once both sets of times sit on one axis, is stronger:** the `UPLOAD_THESE` folder was written at 16:47:56 and every one of the four calls predates it. **No attempt at this round has ever received it.** Nobody uploaded the wrong folder — the folder did not yet exist, and the eleven observations came from a run holding rev 2 and a file named for the Constitution whose contents were the engine source bundle. Two distinct payloads confirm the shape: 213,581 input tokens in the morning, then 196,283 for all three afternoon calls across two providers — one upload, retried. **Three: a model named in this log was never called.** Entry #22 records Run 1 as a blind review by DeepSeek V4 *Pro*. The bill says DeepSeek V4 **Flash 0731**. Independence is tracked at the laboratory, so the ledger's verdict is unchanged and DeepSeek is spent either way — but the record names a checkpoint that was never billed, which is the same error rule 29 was written about. **Left open rather than resolved:** the log shows nine Luna Pro calls between 27 August and 3 September, considerably more than this document's account of Luna Pro accounts for, and three of them on 3 September map to nothing recorded anywhere. Kimi K3's 2 September call returned 36,085 output tokens, which is not the truncated stub the record describes when it leaves Kimi's spent status undecided. Both are noted here so the next reader starts from the discrepancy rather than rediscovering it. **The generalisation, now paid for twice:** where a project keeps a record of what happened, ask whether some system already recorded it as a side effect of doing its own job, and prefer that system. This project's account of one failure was built first from a truncated screenshot, then from file timestamps, and only the bill settled it.

#72	The Instruction Told the Reviewer to Stop, and It Would Have Been Right September 5, 2026	DECISION — ADOPTED

“Seven of the forty-four rules carry their previous verdict in a code comment.”

The round-3 package was built and then **searched** before it was sent — the step that the previous two rounds did not take. Section 2 of the reviewer instruction told the auditor, in as many words, to stop if it found a named prior auditor, a count of Compliant and Non-compliant items, or a list of Critical findings *in either of the two code bundles*. The bundles contain all three. Items 3 and 6 described as rated Critical; Item 18 as kept Compliant; Item 16 as having gone Non-compliant; Tier 3 items 3 and 4 as currently Non-compliant; Tier 4 item 2 as rated Compliant with the previous auditor's reasoning quoted. Three separate counts of Criticals. Luna Pro and GLM named. **A reviewer obeying the instruction would have refused to grade, and that would have been the fourth attempt at this round to end without a report and the third in a row killed by the package rather than by the code.** The contradiction had been half-seen already: Rev 3 records scoping this same stop condition to resolve a clash with Section 9, and that scoping drew the line at *the instruction may tell you counts, the artifact must not tell you content*. Nobody then checked whether the artifact told content. **Viktor's ruling:** scope the condition to the Constitution file, count and disclose exactly what the bundles hold, and ask the reviewer to grade all forty-four and say in Part 6 — “I don't know” included — whether a comment moved it. The comments are not stripped: Section 4's argument stands, that removing them is the audited party editing its own evidence and destroying what Item 8 exists to test. **And separately, the structural fix, deferred on purpose:** the audit narrative moves out of code comments and into this document, but not inside the audit it would affect — doing it now would hand the auditor a codebase edited to look better for its grader. **The reason this matters more than a fixed contradiction:** a reviewer that stops tells us loudly. The failure one step over is silent — a reviewer that does not stop, reads “Item 18, kept Compliant”, and grades Item 18 Compliant. Nothing in its report would reveal that, and the report would look entirely ordinary. That is earned rule 28's test applied to the audit process itself, and it is why the same instruction now asks the reviewer to search its own package before grading.

#73 The Disclosure That Counted Wrong

September 5, 2026 — **PROCESS — RECORDED**

“Re-reading the instruction could not have found this. The error was in the counting.”

Entry #72 records the ruling that the code bundles' prior-audit content be counted and disclosed rather than made a reason to stop. Rev 4 carried that count. The count was wrong in three places, and wrong in the direction that matters: it under-stated what the package leaks. **Five** distinct Critical-count phrases are in the bundles, not three — “the first Critical” and “the last Critical” were missed. **Three** AI parties are named, not two: Claude appears in both bundles, twice as having been overruled by Viktor. And the disclosure covered only the two code bundles, while a third file in the upload folder, `version_control_history.md`, lists every file each commit touched and so names the previous audit report and the four prior transcripts by filename. **How it was found:** by grepping the built package, which is the pre-send step Entry #72's ruling created. Re-reading the instruction could not have found it — the prose was internally consistent and the defect was arithmetic. **What was checked and came back clean in the same pass:** the Constitution audit copy is 29 pages with no Version History row recording an outcome, every Compliant / Non-compliant in it being rubric text; and all 71 SHA-256 values in the manifest were recomputed from the bundle bytes and matched, so the manifest is not a claim taken on trust. **Why this is an entry rather than a fix:** the step was created one day and earned its keep the first time it ran, against the document written by the party that created it.

#74 A Number Read Off a Generated File

September 5, 2026 — **OBSERVATION — FILED FOR THE RECORD**

“117 was true when it was generated and stale when it was used.”

The commit that landed Entry #73's fix predicted the repository's commit count would go 117 to 118. It went to 119. The 117 came from the `version_control_history.md` generated three hours earlier; a commit had landed in between. **The shape, for the third time in this project:** a figure taken from an artifact that recorded it correctly at the moment it was written, and treated as current afterwards. Entry #71 withdrew an inference from file modification times on the same grounds. Here the artifact was one this project generates on purpose, which makes it more trustworthy-looking and no more current. **Why it was caught:** because the prediction was stated as a number before the command was run. An unpredicted line can only stop the work if the prediction was complete enough to be contradicted.

#75 The Transcripts Were Kimi's All Along

September 5, 2026

OBSERVATION — FILED FOR THE RECORD

“It said it was not Qwen, in the section written to make it say so, and the sentence never reached the repository.”

The largest correction this project has had to make. The four files in the repository root named `qwen_reasoning_1.txt` through `_4.txt` are a **Kimi K3** transcript. Verified rather than inferred: with whitespace stripped, the repository copy and the copy saved from the Kimi room of 2 September are identical for **57,631 characters from the first byte**. Two models do not coincide on 57,000 characters. So the eleven prior observations, and the Part 8 document built from them, are Kimi's work. **Why it went unnoticed for three days:** the repository copy is lossy, roughly 60,000 characters shorter than the room copy, and one of the passages in the gap is the model identifying itself — “the instruction says I'm Qwen3.8-Max — I'm not; I'm Kimi (Moonshot AI) … if the ledger records Qwen3.8-Max but the actual reviewer is Kimi, the ledger is wrong.” Section 5 of the instruction was written to elicit exactly that sentence. It was produced, and it did not survive being saved. **What inherited the wrong name:** Entry #50, the attempts table, the Part 8 document, the correction made from the bill in Entry #71, and a ruling made and withdrawn the same day. **The instructive part is Entry #71's own reasoning.** It argued from call timestamps — Qwen at 15:42 before Kimi at 15:58 — to the conclusion that the eleven came from Qwen. The inference was valid and the conclusion was false, because the bill records *what was called* and not *which output was saved under which name*. Earned rule 29 says prefer the system that recorded the fact as a side effect of doing its own job. Applied correctly, that system here is the transcript, not the invoice — and it had been sitting in the repository under the wrong name the entire time.

#76 An Accident That Produced a Report

September 5, 2026

AUDIT RUN D — RECORDED

“The Chatroom tab was left on Auto Router, and the round-3 package went out.”

The router chose `z-ai/glm-5.3-flash-20260826`, served by BaseTen: 251,148 input against 11,475 output, $0.0434, 102 seconds. It returned a complete Parts 1-6 — **26 Compliant, 13 Partially, 1 Not verifiable, 0 Non-compliant, eleven findings all rated Minor** — so the round happened whether or not it was intended, and Viktor ruled it **is** round 3 rather than discarding it. Its model-identity check fired correctly: its first line states it is not Qwen3.8-Max. **Three findings were verified against source and all three confirmed**, and one is under-scoped by the report in a way that matters. F-7, the permissive `risk_valid` default, was filed against the simulated-order module and rated Minor on the grounds that nothing there authorises a trade. The same default sits in `decision_model._determine_final_action` — the authorization gate — where a risk block missing that key reads as risk-passed. The severity argument does not reach it. **What the run measures beyond its findings:** 11,475 output tokens with no reasoning spend, against Kimi's 36,085 tokens of reasoning on a smaller package that never reached a report at all. Depth of read is visible in the bill. The strongest part of the report is its Part 6b, which names three places where a comment led its reasoning before the code confirmed it, unprompted beyond the instruction asking.

#77 The Ledger Rebuilt From 723 Rows

September 5, 2026 — **PROCESS — RECORDED**

“Second family the ledger had wrong, and the reviewer that ran was not independent.”

The full provider activity export, 723 requests. **All 723 are** `variant=standard` — no free endpoint, ever — so the no-training conclusion the 2 September independence ruling rests on holds unchanged. Two things it corrected. **Z.ai is not clean.** `z-ai/glm-5.3` ran on 28 August: 104,394 in, 55,179 out, a full substantive session on this project eleven days before GLM 5.3 Flash audited it. Under the lab-not-checkpoint rule the reviewer that produced Entry #76's report was **not independent**. Its findings stand on their own evidence; the caveat travels with the comparison permanently. That is the second family this ledger had listed clean and wrong, after Mistral in Entry #48. **Kimi clears, on the check that mattered.** Every Aider call in the history — 681 of them — carries the Aider app key, across exactly the five models the corrected record names. All fourteen Kimi calls are Chatroom. Kimi never worked on the codebase, so its exposure is session-level, which the 2 September ruling clears. But it is far larger than the record said: eight substantive reads on 27 August of 93K-110K input each, roughly 100,000 tokens of output about this engine, plus 36,085 on 2 September. “A Step 3 attempt that truncated” described one row of thirteen. **And outputs this project paid for and never saved:** four completed Kimi runs on 27 August totalling 60,537 tokens; Qwen's 2 September call returning 11,964 tokens, of which no transcript exists under any name, so what Qwen actually said is lost; three Luna Pro calls totalling 393,663; and one row with no model recorded at all.

#78 The Handover Check Was Blind to Ignored Files

September 5, 2026 — **PROCESS — RECORDED**

“A check that looks for untracked files that matter cannot see files the repository was told to ignore.”

Both evidence sets — the Kimi transcript and the GLM report — were first saved under `docs/audit_package/round*/`, which `.gitignore` excludes except for the manifest. That rule is right for generated package bytes and wrong for reviewer responses, which are primary evidence and cannot be regenerated. **The part worth carrying to other projects:** ignored files do not appear in `git status --short`, so the standing end-of-session handover check could not have caught it. The check asks whether untracked files that matter are present, and the one class of file most likely to be lost is the class the check is structurally unable to see. **Structural fix, not an instruction to be careful:** reviewer responses now live in `docs/audit_reports/<round>_<model>_<date>/`, which is tracked by default. No exception list to maintain and no judgment required at the moment of saving, because the safe location is the only one on offer. `round*/` stays ignored for the generated packages the rule was written for.

#79	Rev 5, and Part 8 Removed September 5, 2026	DECISION — ADOPTED

“The instruction had been telling the reviewer who it was, and getting it wrong.”

Two changes and nothing else. **Section 5 no longer asserts the reviewer's identity.** Rev 4 told the reader it was Qwen3.8-Max and that one clean-list slot had been spent on an earlier attempt at this same audit — a sentence which, sent to Kimi, describes Kimi as somebody else. It now asks the reviewer to state its own identity and says what is known about who ran what. **Part 8 and Section 13 are removed entirely.** Their comparison document was the transcript of Entry #75, so asking Kimi to reconcile its findings against it would have been a model grading its own earlier reasoning, and asking anyone else to reconcile against a document this project had misattributed was not worth doing either. The comparison it existed to produce is made outside the instruction, between two reports, by Viktor. **What did not move:** what is graded, how it is graded, the severity rubric, Section 7's checks, and everything handed over besides Part 8. GLM ran against rev 4, so the revision history names the delta and the comparison carries its own caveat.

#80	Three PDFs in a Downloads Folder September 5, 2026	PROCESS — RECORDED

“Two were preserved. One held a paragraph that existed nowhere else.”

Viktor found three documents in his Downloads folder and asked whether he could delete them. Each was checked against the repository rather than recognised by name. **The Gemini review of the Constitution** hashes to exactly the SHA-256 recorded in `docs/constitution_reviews/gemini_2026-09-03.md`, and its extracted text matches that file word for word, 390 words each. Fully preserved; deleted. **The independence ruling** was preserved except for one section: the argument that the ruling turned on *mechanism* rather than on the value of the model it freed — that had a mechanism been found, Kimi would have stayed out regardless of rank, and that a rule reading “an exception was made because the model was valuable” is one anyone can invoke later to justify anything, where a rule about mechanism can be attacked and overturned on its merits. That paragraph was recovered into `docs/PHASE7_NEXT.md` before the file was deleted. **The third document** is Claude's candid criticism of Viktor's own role in the project. Viktor ruled it does not enter the repository, because the repository is public. It is kept outside version control by decision, and that is recorded here so a later reader knows it exists and why it is absent rather than assuming it was lost. **The generalisation:** the conclusion had survived in the record and the reasoning that makes it challengeable had not, which is the failure this document exists to prevent.

#81 “Could Not Run It” Was Two Claims, and Only One Was True

September 5, 2026

OBSERVATION — FILED FOR THE RECORD

“The impossible case was stated where the merely-undone one belonged.”

The patch pointing the package builder at round 4 shipped with a commit message saying the build itself had not been run, and offering a reason: the builder executes the engine and there is no engine in the sandbox. Viktor asked why, and the question found the defect. **Two claims had been folded into one.** True: it cannot be run against *this* repository from the sandbox, because the file bridge copies files and not the `.git` object store, so the real history never arrives and neither the version-control history nor the commit messages can be generated. Not true: that the builder could not be run at all. A scaffolded repository — two commits, stub modules, the real instruction, a placeholder Constitution — exercises every changed line, costs minutes, and was then done: it produced a round-4 directory with seven files in the upload folder, one in the withheld folder, and no Part 8 document anywhere. **The shape, and it is the same one this project keeps finding:** a conclusion carrying more confidence than the check behind it. Here it took a new costume — a genuine impossibility placed where an undone task belonged, which reads as a limit rather than as a gap and so does not invite the question that would expose it.

#82 Kimi Runs Round 4, and the Reason That Was Rejected

September 5, 2026

DECISION — ADOPTED

“Model quality is a preference. Scarcity of clean reviewers is a fact about the ledger.”

Viktor's ruling, argued and then narrowed. **What stands.** Kimi K3 has graded **none of the forty-four rules**: its earlier attempt could not read the standard, refused to grade, and ran Section 7's checks instead — so on the question this round asks, it is blind, and “self-review” overstates what it is. And it is the last high-ranked model clean on both independence lists, so a round that does not spend it spends a worse reviewer or none. **What was rejected.** “Kimi K3 is a very important model to use.” That is the shape the 2 September independence ruling warned against in its own words — the paragraph recovered in Entry #80 — and admitting it here would leave a precedent anyone can invoke later. Both readings pick Kimi, which is exactly why the distinction is worth making: the decision is unchanged and the rule it leaves behind is not. **What is disclosed rather than solved.** Eleven of Kimi's own observations are already fixed in this code, with comments describing them. The exposure is not that it repeats a verdict — it has none — but that it may accept a fix to its own diagnosis more readily than a stranger would. Small, real, not removable without spending a different reviewer, and accepted rather than argued away. **A hold that outlives this entry:** the four transcripts keep their wrong filenames until the package has been sent. A rename is a commit, commits appear in the version-control history, and that file ships inside the upload folder — so correcting the name now would put “kimi” in front of Kimi before Section 5 has asked it to identify itself unprompted.

#83 **The Section 11 Confirmation Run, Made Before Any Fix**

September 6, 2026

FINDINGS — FROM EXECUTION

“Completing on the negative control is what makes the positive case mean something.”

Kimi's round-4 report asked for exactly one run, and this is it — made on unmodified code, before any fix, so **Finding 1 is observed rather than argued**. Three offline runs against the committed pinned fixtures. **A:** fixtures exactly as committed, negative control — completed, correlation +0.2326 over 30 observations. **B:** `BTCUSDT_4h` with every timestamp shifted +2h, zero shared timestamps — `Decision object construction failed: float() argument must be a string or a real number, not 'NoneType'`, the exact string Kimi named. **C:** BTC's last 60 candles flat, the optional zero-variance variant — the same message, secondary, since a flat series also moves BTC-side ATR and volatility. **What the run found that the report does not contain:** the decision-log record written on the merge failure carries none of the run's lineage or provenance — only an error, symbol, timeframe and log path — where the healthy record beside it carries both, with input hashes and a run hash. Item 6, the Critical keeping the release gate shut, is about traceability; a defect that manufactures untraceable records inside it is not a display concern, and whether it is folded into the Finding 1 fix or filed as its own item is left to Viktor. **Two wrong turns, recorded rather than cleaned up.** The harness's first version ran with `config.LOG_DIR` untouched and wrote three records — two of them errors — into the engine's real permanent log, overwriting the persisted bias state machine with a fixture-derived value; the fix repoints the harness at a temporary directory for its duration, structural rather than a warning. And the corrected harness was staged under the filename of the version it replaced, so the commit shipped the older bytes and failed with `ModuleNotFoundError` — passing a `py_compile` check that tested syntax when the question was identity. Two rules taken from it: a delivered file gets a name unique to its version, and a delivery is verified by reading it back off the device and diffing it, not by compiling the copy that was sent.

#84	Finding 1 Fixed, and the Seam It Was Hiding In September 6, 2026	DECISION — ADOPTED

“The producer was tested, the consumer was tested, and the seam between them was not.”

Landed against unmodified code, after the Section 11 run above had already reproduced the defect. Suite 319 to 331 passing, 0 failed; the golden snapshot did not move, predicted before the run on the same reasoning as the negative control — the pinned fixtures measure a correlation, so every value the snapshot records goes through the unchanged path. **Two changes, in** `models/signal_router.py`**.** A new `_optional_number()` helper passes through what the producer wrote instead of defaulting to 0.0: None stays None, a non-finite number becomes None, a measured number is unchanged. **None rather than 0.0 deliberately** — a correlation of 0.00 is a measurement a real pair of independent assets can produce, so substituting it would report a finding the engine does not have, the same defect already removed from three other modules. And the record written when the assembly fails now carries the run's lineage and provenance, closing the Item 6 gap the confirmation run found and neither reviewer filed. Twelve new tests in `tests/test_router_btc_seam.py`, written at the seam itself rather than as one more assertion about either side of it — the pinned fixtures share all 450 timestamps, so nothing in the suite had ever made the producer emit its not-measured shape and handed it to this consumer. **Evidence against pre-fix code:** 8 failed, 2 passed, 2 skipped, all eight behavioural — a `TypeError` from the bare `float()` call and a `KeyError: 'lineage'` on the failed record, no import or collection errors. **Where verification stopped:** built and verified in a Linux sandbox on the pinned dependency versions; `pandas_ta` could not be installed there, so the two end-to-end tests and the golden snapshot were unverified until the Windows run — stated at delivery rather than after the fact. **Not fixed:** the merge's other invented defaults (`"NEUTRAL"`, `"NORMAL"`, 0.0 for trend health) stand on the same “the producer always sets it” argument that was just proven wrong; and a failure inside the merge still cannot degrade rather than halt, because `decision_model.evaluate()` runs before the BTC block is assembled — a real gap in the 29 August ruling, recorded rather than patched in passing.

#85	GLM F-7 as Extended, and It Was Reachable September 6, 2026	DECISION — ADOPTED

“The maximum-intensity authorization this engine can issue, on a run where risk was never assessed.”

Landed at `26a05dc`, pushed. Suite 331 to 363 passing, 0 failed; golden snapshot unmoved, confirmed by diffing the fixture rather than only by the test passing. **The 5 September note under-rated this as a permissive default in three places; reading it found it reachable, not latent.** `signal_router._validate_engine_output` checked that the key `"risk"` was present and never looked at what was under it, so an empty risk block passed validation and reached a gate whose `.get("risk_valid", True)` supplied an authorization the engine never computed. On the conviction fixture: an empty risk block or `{"risk_valid": "OK"}` both produced AGGRESSIVE LONG. **A validator that checks presence plus a consumer that defaults permissively is the composition**, and neither half looks wrong alone — the same shape as the seam Finding 1 lived in. **What changed:** a new `read_risk_verdict()` in `models/risk_model.py` treats absent, present-as-None and present-as-not-a-boolean as one state, none of them a pass, deliberately strict so an unrecognised shape fails closed. The gate now refuses with its own reason, `NO-TRADE (RISK NOT ASSESSED)`, rather than reusing `NO-TRADE (RISK TOO HIGH)` — a bare flip of the default to False would have printed “Risk check failed (OK)”, the same fabrication with the opposite sign. The contract keeps `risk_valid: bool` rather than becoming `Optional[bool]`, a decision on the record: Optional looks like the honest declaration and is the weaker fix, since the contract test's type check passes any Union and the field would stop being checked at all. 32 new tests. **Evidence against pre-fix code:** 14 failed, 18 passed — 11 behavioural, 3 source-text, 0 import or collection errors, verified by grafting the new function alone into an otherwise untouched file so the pre-fix run cannot collapse into one collection error. 7 of the 18 pre-fix passes are named negative controls. **Verified both ways** before delivery, unlike Finding 1: 363/0 with `pandas_ta`, 252 passed / 100 skipped without it, against a 222/98 baseline. The run owed against live data was discharged two commits later, at 10:04 on 6 September — see Kimi Finding 5 item 6 below. **Not fixed:** the reason string that blames “the risk check” on every NO-TRADE action regardless of which check actually blocked it, pre-existing; the manifest-hash test that fails on any non-Windows checkout, worked around locally rather than fixed; and the same two invented defaults (`"NORMAL RISK"`, 50.0) elsewhere in the same assembly.

#86 **Kimi Finding 5 Item 6, and the Live Run Two Fixes Were Owed**

September 6, 2026

FINDINGS — FROM EXECUTION

“A located level sitting exactly on the current price, on precisely the runs that located nothing.”

Landed at `4678f45`, pushed. Suite 363 to 377 passing, 0 failed; golden snapshot unmoved, confirmed by the fixture's absence from `git status`. **The margin was one row.** `_detect_swing_structure` returned `current_price` on both of its not-located paths. `2 * lookback + 5` is 21 at the configured lookback of 8; the engine's own frame-length guard rejects below 20 — the two minimums disagree by one, on the wrong side, so a 20-row frame passed validation and got its own price back as a structural level. Reachable by the same presence-only validation GLM F-7 named two commits earlier, and the third time this exact field was patched at a different layer while the producer kept inventing — `entry_model`'s fallback on 1 September, `panel_render`'s identical default on 2 September, both fixed at the consumer while this stayed open at the source. Both not-located returns are now `float("nan")`, in `structure/structure.py` and in the router's assembly default. Deliberately no `degraded_inputs` entry — this field is read by nothing but the panel and the record, not by entry scoring or the risk stop, so a display-only value moving an authorization would be a worse defect than the one being fixed. 14 new tests. **Evidence against pre-fix code:** 9 failed, 5 passed — 7 behavioural, 2 source-text, 0 import or collection errors, no harness needed this time since the new tests import no new name. **The live run owed for both this fix and GLM F-7 was made on the patched tree before the commit** — AEROUSDT 4h, 6 September 10:04: `SWING STRUCT : $0.5063` against a current price of $0.5437, a located level that is not the current price, and `NO-TRADE (RISK TOO HIGH)` naming a real reason, “Risk regime classified as EXTREME RISK,” rather than the F-7 fabrication. The NaN paths themselves are covered by the suite only — a 300-bar live frame always locates a swing. **Not fixed:** the router assembly's other invented defaults (`hvn`/`lvn` to 0.0, regime to NEUTRAL, sequence to NONE, volume sentiment to NEUTRAL VOLUME), named here so the sweep has them; and the same field still raises `TypeError` if present as None, Kimi Finding 1's exact shape, latent only because the producer never writes it — the wrong argument for the third time in three days. **What went wrong delivering it:** the command list was reissued with the same steps renumbered when the live run was inserted, and one `git add` was lost between the two numberings — the commit went in without its test file, caught by reading “2 files changed” rather than by any check, and amended before pushing.

#87	Kimi Finding 3, and a Sentence the Live Run Found That Neither Reviewer Did September 6, 2026	DECISION — ADOPTED

“The number is behaving as designed; the sentence is false.”

Landed at `1045748`, pushed. Suite 377 to 400 passing, 0 failed; the golden snapshot moved, predicted in exactly five fields before the run and moved in exactly those five. **The list could not be fixed by extending the list.** Kimi named seven decision-affecting settings missing from the fingerprint; six of the seven are not list omissions but things a list of that shape cannot hold at all — class attributes, local variables, bare literals in comparisons and call sites. **What changed:** a new `core/code_fingerprint.py` reads every `.py` file outside `tests/` and `docs/`, parses it, strips docstrings, and hashes the sorted per-file digests — no list of constants or modules, so a constant written tomorrow is covered the moment it exists. The parse tree rather than the text, deliberately: this repository rewrites its comments in nearly every commit. `code_hash` is deliberately **not** folded into `run_hash` — that would move the archive filename and force a re-baseline on every commit touching any engine file, spending the check that has caught more defects here than any other. Claude's call under the delegation, reversible in one line. 23 new tests. **Evidence against pre-fix code:** 21 failed, 2 passed on the raw count; the honest number is 7 behavioural failures, the rest proving only that the new module is new. Three negative controls: a comment does not move the hash, neither does a docstring, and CRLF and LF source hash identically. **Figures:** 400/0 with `pandas_ta`; 286 passed / 103 skipped without it; `run_tests.py` 333 passed, 0 failed, 29 errors unchanged. **The live run made on the patched tree, AEROUSDT 4h, 6 September 10:51, matched the sandbox hash bit for bit** — `44e085cf`… on Viktor's Windows machine under 3.12.10 against a CRLF checkout, and the identical value in a Linux sandbox under 3.12.3 against LF, confirming platform-, line-ending- and 3.12-minor-version-independence as observed rather than assumed. **The same run found what neither audit report contains.** The BTC block raised confidence 73.86 to 75.80, explaining it as “BTC is also neutral, agreeing with AERO's own bias” over a BULLISH CONFIRMED AERO — a neutral value cannot be “also” what a bullish one is. Verified at `decision_model.py` 704-768: agreement is decided by the sign of `btc_score`, and the sentence's words come from the label — the two disagree whenever the score is non-zero but below the labelling threshold. **Why this is not simply Kimi Finding 2:** Finding 2's exhibit is entirely about `abs(correlation)` discarding the sign, at -0.90; here the correlation was +0.51 and the sign was never in question. A second, independent route to the same class of false sentence, filed as input to decision 3 rather than fixed, because it is Viktor's ruling and it changes what he is ruling on. **Not fixed:** `engine_version` as a static string; `code_hash`'s interpreter dependence, recorded by carrying the Python version beside it in the archive; the unnamed literals the finding named, covered by the hash but absent from the readable snapshot; and a rerun on identical inputs after a code change still overwriting the earlier archive's `meta.code`, which follows directly from the `run_hash` decision.

#88	The Suite Was Destroying the Decision Record, and Had Already Done It September 6, 2026	FINDINGS — FROM EXECUTION

“A pristine checkout, seeded with marker files, came out of a green suite with every marker gone.”

Landed at `fc35a2f`, pushed. Suite 400 to 406 passing, 0 failed; no engine source file touched, so no live run was owed and `code_hash` did not move. **Found by running the suite against a clean clone and watching the filesystem, not by reading** — every individual line involved is unremarkable, and the damage exists only at runtime. Two independent reviewers read these files; one filed a finding about the very test that does the worst of it, and the mechanism survived both. **Four routes into the engine's own permanent record.** Twenty-eight tests across eleven files ran the engine against the real log directory, appending nine real decision records plus one error record on every suite run — including the two files whose entire subject is what the permanent record contains. A twelfth test deleted the directory outright, in a `finally` block after an archive-path spelling test. A thirteenth recreated it as a side effect of purging and re-importing every module. A fourteenth wrote into it through the capitalised spelling, `Logs/` — one directory on Windows, two on Linux, exactly the defect a hardcoded-path guard exists to catch, in a file that guard does not scan because it checks modules and these lines are in a test. **The first patch was wrong, and the way it was wrong is the lesson.** v1 was verified green in the Linux sandbox, applied, and failed on Viktor's machine at the first run: the new guard watched the directory named exactly `logs`, so route four's `Logs` was invisible to it on Linux and immediate on Windows — **a guard whose result depends on which platform it ran on is not a guard.** A second failure was a negative control hardcoding a 4-byte file size for `"one\n"`; Windows text mode writes 5. Both recorded rather than quietly fixed, because the first one changes how every sandbox result should be read from here: evidence about Linux until something makes it evidence about Windows. **What was lost:** the live log held eleven records — ten from a single suite run at 08:51 UTC, identical run hash, identical bias score, and one error record with no code hash — plus the real 10:04 AEROUSDT run and the three synthetic Section 11 records this file had just finished describing, both gone with the directory the `rmtree` deleted. Nothing in the repository can rebuild them; `logs/` is gitignored, so nothing held a second copy. **What changed:** every writer redirected to a temporary directory at import time rather than as a pytest fixture, so both test runners pick it up; the path-spelling and golden-snapshot tests keep their real relative literals so the coverage that caught the Windows failure is not retired; the hardcoded-path guard now scans case-insensitively. Six new tests. **Figures:** 406/0 with `pandas_ta`, 292 passed / 103 skipped without it, 339 passed / 0 failed / 29 errors unchanged under `run_tests.py`. Reproduced on a pristine checkout seeded with four marker files: pre-fix, 400 tests passed and every marker was gone; post-fix, through all three configurations, the seeded log's SHA-256 came through unchanged. **Not fixed:** the lost records are not recoverable; `logs/` remains gitignored, part of decision 7; and the hardcoded-path guard still scans engine modules only, not test files — recorded as a candidate rather than adopted, since this repository has been burned by source-text guards before and the new behavioural guard already covers the same class on both platforms.

#89	The Document Set Audited End to End, and What It Found September 6, 2026	OBSERVATION — FILED FOR THE RECORD

“The last time this document set was read end to end, it found a whole Critical that had never been entered into any roadmap.”

Prompted by nothing in particular, which is the point. **The audit-package record is intact.** All three manifests verify byte-for-byte against their recorded build commits — round 2 60/60, round 3 71/71, round 4 71/71 — against an LF checkout, because `build_audit_package.py` normalises line endings on read and write while `tests/test_pinned_source.py` hashes raw bytes without normalising, two hashing schemes over the same repository with opposite line-ending behaviour, filed with the LF-checkout test hygiene items. One suspicion, checked and wrong: two 30 August bundle files read like a stale-bundle trap when cited by name in the round-4 instruction, but the builder documents them as the deliberately kept round-1 record, with fresh bundles copied in separately. **Two structural problems.** No build script could write into this repository — nine of ten `docs/build/` scripts hardcoded `/tmp/outputs/`, which is why the Engineering Notes gap grew at every handover; this audit's own count of “all ten” was itself off by one, corrected and fixed the same day, see below. And `README.md` and `docs/build/README.md` both cite `docs/Phase7_Audit_Findings_Complete.pdf`, which does not exist anywhere in the repository and cannot be produced from here — its builder also reads source material that was never committed. Every other dangling reference found is benign: gitignored build outputs, or already recorded elsewhere in this project's own documents. **What this audit did not do:** the eight PDFs — Constitution, Roadmap, Engineering Notes, Remediation Plan, the Documentation and Change Log Standard, the Tier 0 Companion, the Credential Security Protocol, the Audit Execution Instructions — were not opened. The mechanical half is done; the semantic half, reading each against the current state of the project, is not, and is where a second Finding 3 would be.

#90	The Build Scripts Can Write Into the Repository September 6, 2026	DECISION — ADOPTED

“The fix did not need designing; it needed extracting from the one script that already had it.”

Suite 406 to 412 passing, 0 failed. No engine source file touched: `code_hash` unchanged at `44e085cf`…, computed on both trees rather than assumed, golden snapshot unmoved, no live run owed. **The count from the document audit above was off by one.** Nine of the ten `docs/build/build_*.py` scripts hardcoded `/tmp/outputs/`, not ten — `build_audit_package.py` already resolved its output from `__file__`, which is why its manifests verify byte-for-byte while every PDF beside them was carried across by hand. Caught by recounting the evidence, not by review, the same shape as a suite-log patch that once said “twelve files” when eleven wrote and the twelfth deleted. **What changed:** a new `docs/build/_output.py` resolves `REPO_ROOT` from `__file__` and rejects any output path that is not a bare filename — absolute, containing a separator, or carrying a drive letter — so a build script no longer has a path of its own to get wrong. The nine import it; seven now write their PDF into `docs/` directly. The other two, `build_findings_bundle.py` and `build_remediation_plan.py`, also read source material that was never committed — they now exit 2 and say so before doing any work, instead of a bare `FileNotFoundError` naming a directory that does not exist on Viktor's machine. Six new tests. **Evidence against pre-fix code:** 5 failed, 1 passed — two behavioural (the scanner finding eleven absolute assignments across nine files, none from a shared resolver), three proving only that the new module is new, one negative control. **Verified rather than asserted:** all nine scripts produce the same output path run from the repository root or from inside `docs/build/`; the rebuilt Engineering Notes and Roadmap extract text identical to the committed PDFs, 63 and 15 pages, same SHA-256 over the extracted text — a rebuild is the same document, not a new one. **Which platform this is evidence about: Linux**, and the filename check is platform-neutral by construction — it reads the path string itself rather than asking `os.path.isabs` alone, which answers differently for a Windows-style path on the two platforms. Not verified on Windows: that reportlab renders these documents there — reportlab was installed on Viktor's machine the same day, and the first Windows evidence will be the first Notes rebuild. **Not fixed by this patch itself:** it makes publishing the eight entries this document now carries one command; writing their content — this row among them — was a separate, later step. Both undocumentable PDFs still cannot be built from here. Both README files still cite the audit-findings PDF that cannot be produced. reportlab stays out of `requirements.txt` deliberately, with its own install line in `docs/build/README.md`.

#91	Fourteen-Item Latent/Minor Sweep Closed September 7-8, 2026	MILESTONE — RECORDED

“Three classes of latent defect closed in three commits without moving the golden object through the sweep itself.”

Grok session, Viktor numbered-CMD workflow. Items 1-5 at 22afea2 (invented defaults to NaN/UNKNOWN). Items 6-8 at 4d56f2a (dead code, exit_watch zero, test_live). Items 9-14 at c3b0d43 (banner, comment, BTC degradation, snapshot shape, NO-TRADE blame, structure read path). Suite 345/0/29; golden 8/8 through the sweep. Grok is not independent for re-audit of this code. Release gate not opened.

#92	Kimi Finding 2 Fixed September 9, 2026	FINDINGS — FIXED

“Sub-threshold BTC score could print neutral agreeing; abs(correlation) discarded sign.”

Viktor ruled fix under decision 3. Agreement direction uses RAW_BIAS_THRESHOLD. Adjustment uses signed correlation. Commit 5e2e9f3; docs 179f4f9. Golden re-baselined on btc_context only (78.14 to 79.26). Suite 345/0/29.

#93	code_hash Excludes Claude outputs/ September 9, 2026	PROCESS — RECORDED

“Backup .py copies under Claude outputs/ moved code_hash with no engine change.”

Excluded Claude outputs from EXCLUDED_DIR_NAMES. Commit ef00765. Suite 345/0/29.

#94	qwen_reasoning_*.txt Transcripts Relocated Into docs/ September 11, 2026	PROCESS — RECORDED

“Four transcripts had sat in the repo root since round 2, under the hold recorded at Entry #93.”

Commit bb6d222. The four qwen_reasoning_*.txt files move from the repository root into docs/, alongside the rest of the audit material — a rename, discharging that hold. No engine module touched; code_hash unaffected.

#95	This Document's Own Body Reconciled With 76380ec, a Second Time September 11, 2026	PROCESS — RECORDED

“The 7-9 September Grok session updated this file's head block and builder but left two “Open” items describing already-landed work as still open.”

Commit 4705d03. “Open — decisions” #3 and “Open — work” #2 described the fourteen-item sweep and Kimi Finding 2 (Entries #91, #92) as unresolved after both had already landed. Corrected in the same shape as the first reconciliation against 76380ec. Docs-only and code_hash-neutral by directory exclusion.

#96	Phase7_Sweep_Session_Commands.pdf Committed September 11, 2026	PROCESS — RECORDED

“A generated PDF had sat untracked on Viktor's machine rather than in the repository.”

Commit e75f7be. docs/Phase7_Sweep_Session_Commands.pdf committed as-is. No engine module touched.

#97	Item 14 / Decision 2 Closed: Risk Regime Now Reads ADX, Not trend_health September 11, 2026	DECISION — ADOPTED

“classify_risk_regime shared its largest input (0.30 of bias_score) with the conviction number it exists to check independently.”

Commit 0c7dec5. classify_risk_regime now reads ADX instead of trend_health; REGIME_LOW_TREND_HEALTH/REGIME_HIGH_TREND_HEALTH renamed to REGIME_CHOP_ADX = 20.0 / REGIME_STRONG_ADX = 25.0, values already load-bearing in trend_health.py for the same two market states. code_hash moved and was confirmed on three trees, 47bc557119bb0350… to 2741062fae070f61…, matching Viktor's own Windows value. Suite 412 to 421 passed under pytest (0 failed both sides), 345 to 354 under run_tests.py, watched 29-error count unmoved. Golden re-baselined at six sites: four predicted, two not — the archive filename, built from run_hash, should have followed the run_hash prediction mechanically and did not, recorded as an incomplete prediction rather than a defect. No decision field moved on the pinned fixture, which exits before classify_risk_regime is reached; a live run the same day (AEROUSDT 4h, 22:50) exercised neither new branch, for two independent reasons — recorded as that, not as validation. Found by reading the panel and left unfixed: RISK REGIME prints without the ADX number now deciding it, the same transparency-regression class as GLM F-5.

#98	GPT-6 Astra Ruled for Round 5 September 11-12, 2026	DECISION — ADOPTED

“Two of Luna Pro's OpenRouter sessions had been exposed to prior audit context.”

Commit 4d99cdb. Viktor's billing-export check cleared both sessions as variant=standard, no training routing. send_audit_round.py and build_audit_package.py repointed to GPT-6 Astra; Part 7 folded into the single upload set per ruling 3.

#99	The 2 September Luna Pro Hedge Closed for Good September 11-12, 2026	PROCESS — RECORDED

“Luna Pro's independence had stood on “probably clean” since the 2 September ledger correction.”

Commit de7d135. All nine Luna Pro OpenRouter calls checked individually: all standard tier, “No data training” on each. Evidence upgraded from probably clean to provably.

#100	Round 5 Sent to GPT-6 Astra: Three New Criticals September 11-12, 2026	EXTERNAL ASSESSMENT — GPT-6 ASTRA

“None of the three findings returned had been on record before.”

Commit a12640a. $12.21636, 435,612 prompt tokens, confirmed twice over. F1 (volatility never reaching stop construction), F2 (entry quality scored for the wrong direction), F3 (a stale decision-bar VWMA presented as current) — all three independently reproduced against live code this session. Viktor RULED to fix all three; no limitation accepted for any.

#101	Round 5 F1 Fixed — Volatility Reaches Stop Construction September 11-12, 2026	FINDINGS — FIXED

“The production call to calculate_stop_targets() omitted volatility_state, defaulting to NORMAL regardless of measured conditions.”

Commit 83e334d. Now passes the same volatility_mode its sibling call already receives. Golden snapshot unmoved — the pinned fixture's own volatility state produces the same multiplier the bug's default did. code_hash e774b93d… to f8e4cc7d….

#102	Round 5 F2 Fixed — Entry Quality Scored for the Right Direction September 11-12, 2026	FINDINGS — FIXED

“eq_trade_direction was derived from short_signal, which zeroes for reasons unrelated to direction, letting a fully bearish setup score as a long entry.”

Commit 3f92230. Now derived from raw_bias directly, the same signal decision_model.py's own action selection reads — the identical ruling already made for that module on 2 September, one caller downstream. code_hash f8e4cc7d… to 8b0e97fd….

#103	Round 5 F3 Fixed, Closing Round 5 — VWMA's Trailing-Edge Forward-Fill September 11-12, 2026	FINDINGS — FIXED

“VWMA's own forward-fill had no trailing-edge exception, so a decision-bar miss was silently carried forward from the previous bar.”

Commit 7fcf646. Fixed at the source; VWMA added to the critical_indicators sweep, with three new regression tests. Accepted consequence: some currently-clean runs will now show DEGRADED, naming VWMA. code_hash 8b0e97fd… to 00c8d4ea…, confirmed on independently built trees at each of the three fixes (base e774b93d… through this one). Across F1/F2/F3: suite 429 to 432 passed under pytest, 362 to 365 under run_tests.py, watched error count unmoved; golden snapshot unmoved throughout, confirmed by running it each time. Register unchanged at 21 / 7 / 10 / 6. The release gate remains shut: fixed is not re-audited, and unresolved means both.

#104	Round 6 (Meta Muse Spark 1.3) Sent and Graded: Release Gate Holds September 13, 2026	EXTERNAL ASSESSMENT — META MUSE SPARK 1.3

“No Critical Tier-1 finding returned.”

Report at docs/audit_reports/round6_muse-spark-1.3_2026-09-13/report.md. Second send attempt succeeded after the first hit OpenRouter's account-level 18+ gate (Viktor resolved it; no charge incurred). finish_reason=stop, provider Meta, 470,488 prompt / 10,765 completion tokens, cost $0.63386125.

#105	Round 6 F1 Fixed — Six Bare Literals Promoted to Named Constants September 13, 2026	FINDINGS — FIXED

“T2-4 Explicit Configuration, Moderate. Viktor ruled fix all three of Round 6's findings, narrower than the report's own F1 Location section.”

Commit 447966b. Six bare literals/locals promoted to named, fingerprinted constants: CONFLUENCE_BOOST_MULT/CONFLUENCE_PENALTY_MULT, SPIKE_RATIO, CORRELATION_WINDOW, StructureEngine.REGIME_HYSTERESIS_THRESHOLD. Values unchanged. decision_model.py's trend-health bands and the panel_render.py NaN gap, both named in the report's own F1 Location section, left open — closed later this session as F4 and F5.

#106	Round 6 F2 Fixed — Duplicate confidence_score Meaning Removed September 13, 2026	FINDINGS — FIXED

“Item 10 Consistent Semantics, Minor. engine_core.py's raw risk dict duplicated confidence_score under a different meaning than signal_router.py gives it.”

Commit 3c7e9ae. Dropped rather than renamed — the value survives intact at trend.trend_health, a few lines below in the same return object.

#107	Round 6 F3 Fixed — Six Absent Measurements No Longer Default to 0.0 September 13, 2026	FINDINGS — FIXED

“Item 13 Fail Safely / Item 8 Epistemic Honesty, Minor, unreachable on the live path today.”

Commit 9b34163. signal_router.py's _build_decision_object no longer substitutes a finite 0.0 for six absent measurements (zone_lower, zone_upper, distance_from_zone, atr_stop, targets, current_price) — all six now go through _finite_or_nan, the same treatment structure.hvn/lvn/swing_struct already get. Found, not fixed here: atr_stop/targets/current_price reach panel_render.py through a separate, still-not-NaN-aware safe_float() call — closed as F4, below.

#108	PHASE7_NEXT.md Head Block Brought Current for F1/F2/F3; Rule 38 Added September 13, 2026	PROCESS — RECORDED

“The git add -A trap was rediscovered three more times this same session even after being documented.”

Commit 1f1c935. Head block rewritten for Round 6's F1/F2/F3. Rule 38 added: always git reset the delivery files before git commit — a sequence missing that step is what let git add -A sweep in files it shouldn't have, each time.

#109	The Orphaned 12 September Handover Note Committed September 13, 2026	PROCESS — RECORDED

“A complete, never-committed 12 September session handover note was found stray and uncommitted.”

Commit da2e231. Claude outputs/phase7_handover_2026-09-12.md committed as-is, confirmed by diff not to be a Viktor edit.

#110	F4 Fixed — panel_render.py's atr_stop/targets/current_price No Longer Default to 0.0 September 13, 2026	FINDINGS — FIXED

“The gap F3 found but did not fix: three measurements reached panel_render.py through a separate, still-not-NaN-aware safe_float() call.”

Commit 154e534. Same _finite_or_nan-style treatment as F3, one layer downstream. A real UnboundLocalError bug (color variables referenced before definition) was found and fixed while building this — caught by running the suite, not by review, and named honestly in the commit message. code_hash 36420c9f… to 3c4ad5ca….

#111	F5 Fixed — decision_model.py's Trend-Health/Entry-Score Bands Named September 13, 2026	FINDINGS — FIXED

“The other half of Round 6 F1's own Location section, left open when F1 landed.”

Commit ac0a211. AGGRESSIVE/CONSERVATIVE trend-health/entry-score bands (75/70/50) promoted to named, fingerprinted class constants: AGGRESSIVE_TREND_HEALTH_MIN, AGGRESSIVE_ENTRY_SCORE_MIN, CONSERVATIVE_TREND_HEALTH_MIN. Values unchanged. Investigated and found NOT a defect: the same finding's MIN_ACTION_BIAS naming split against RAW_BIAS_THRESHOLD — both already named, already fingerprinted, deliberately two different thresholds; nothing changed there. code_hash 3c4ad5ca… to 13908bba…. Golden snapshot moved (7 record-completeness fields, zero decision fields) — predicted this time from the F1 precedent rather than discovered by surprise. Test counts at tip (Linux sandbox): pytest+pandas_ta 466 passed; pytest without pandas_ta 338p/117s; run_tests.py 395p/0f/32e (32 pre-existing, unrelated to anything this session touched — all fixture errors on a fixture-free runner). All commits confirmed landed and byte-identical via fresh clone + md5 at the time.

#112 Encoding Defect Found and Fixed — send_audit_round.py's OpenRouter Stream

FINDINGS — FIXED

September 13, 2026

“Every non-ASCII character in round 6's report.md as originally written — every em dash, ±, → — was mojibake.”

Commits 7fbf77c, 9c5f954, 4e0110b. Root cause: resp.iter_lines(decode_unicode=True) was read without setting resp.encoding first; OpenRouter's response carries no explicit charset, so requests fell back to Latin-1 instead of the API's actual UTF-8 — corrupting every non-ASCII character it decoded, in the streamed report/reasoning text and any http_error.txt body alike. Found by verifying the report before acting on it, not by a failed test. Fixed at the source: resp.encoding is now forced to "utf-8" immediately after the POST returns, before anything reads resp.text or the stream — send_audit_round.py's own docstring names this as the fourth failure mode the script exists to make impossible. The corruption itself is a deterministic, reversible byte transform; report.md was replaced with the corrected decode (same model output, not a content edit), original corrupted bytes' SHA-256 kept for the record in the new ENCODING_CORRECTION.md rather than silently discarded. A second pass, prompted by this project's own handover check, found three more evidence files from the same round (reasoning.txt, generation.json, run_metadata.json) untracked and none of them gitignored — closed the same day rather than repeating the failure mode the .gitignore's own comment already named. Docs/tooling only; code_hash unaffected throughout (docs/build/ excluded by directory).

#113 The git add -A Trap, Closed Structurally Instead of a Fifth Restatement

PROCESS — RECORDED

September 14, 2026

“3e93fff — this session's own PHASE7_NEXT.md head-block patch — committed two of Claude's own delivery files alongside it.”

Commit 214c5d8. Root cause named plainly as Claude's: a redelivered patch (rebuilt to fill in a commit hash) used a shortened command sequence that dropped the git reset steps rule 38 already exists for, and Viktor's pasted git status showing both files staged was called "exactly the expected status" without checking it against what the sequence should have produced. Same shape as rule 38 itself: git add -A does what it does regardless of whether a numbered sequence remembers the reset steps. Fixed structurally rather than with a fifth restatement of rule 38: .gitignore now carries *.patch and *_commit_message.txt, patterns that match nothing else currently tracked, so git add -A can no longer stage either file regardless of which sequence is followed. Verified against the state that actually caused the failure — a fresh clone of the true tip with both files already absent from disk, matching Viktor's own state after he had already run del on them — rather than a tree where they still happened to be present. code_hash and golden snapshot unmoved (confirmed, not assumed: neither removed file is a .py file inside the fingerprinted tree, and .gitignore is outside the walk entirely).

#114 send_fix_verification.py Built — a Narrower Tool for a Narrower Job

DECISION — ADOPTED

September 14, 2026

“Verifying that round 6's own findings were actually fixed does not need everything a fresh independent audit needs.”

Commit 2fc578b. Implements Viktor's ruling the same day: the round-6 auditor (meta/muse-spark-1.3) is asked to verify its own findings were fixed, not to re-grade the Constitution from scratch — finding something new would be a bonus, not the goal. Deliberately not called "Round 7"; that name is reserved for the next genuinely independent audit. A new script rather than a mode on send_audit_round.py: that script gets repointed at a new reviewer every round (three repoints already), so importing its MODEL constant would let a future repoint silently redirect this call too — the "one edit upstream" shape muse-spark-1.3's own F3 named for a fabricated-zero default three call frames from where it fires. MODEL/PROVIDER_SLUG/pricing pinned independently here instead. Sends a single-turn API call (no memory persists between calls): round 6's own report.md verbatim, a review note pointing at each of the five items (with explicit pointers for F4/F5, which are not separately numbered in muse-spark-1.3's own report), and the full current content of the eight touched files — the middle ground between a diff too thin to check and a full-codebase resend too expensive to justify. code_hash unaffected (docs/build/ excluded by directory, confirmed).

#115 Round-6 Fix-Verification Sent and Graded: F1 Through F5, All Fixed

EXTERNAL ASSESSMENT — META MUSE SPARK 1.3

September 14, 2026

“Every item round 6 itself raised, verified fixed by that same auditor — not re-graded by a fresh model.”

Commits 59de938, e65a0f7. Sent via send_fix_verification.py --send; provider confirmed Meta (matches the pin, no substitution), finish_reason=stop (not truncated), 72,542 / 2,008 native prompt/completion tokens, cost $0.0992115 — below the estimated range because the response was terse (five short verdicts, no full report) rather than the 5k-20k output tokens budgeted. All five verdicts: Fixed — F1 (four locations named and fingerprinted), F2 (confidence_score dropped from engine_core.py, sole meaning left in signal_router.py), F3 (all six fields NaN-passthrough via _finite_or_nan), F4 (panel_render.py's downstream defaults NaN-safe), F5 (decision_model.py's bands and threshold split named/fingerprinted). Independently checked against the actual landed code, not just relayed: every citation in the model's report grepped against source and confirmed to match. Bonus finding, unprompted by the review note, confirmed real: models/signal_router.py's duplicate @staticmethod above _finite_or_nan — harmless, found rather than fixed at this point (see Entry #116). Response saved to docs/audit_reports/round6_fix-verification_2026-09-14/, byte-identical between Viktor's disk and the landed repo, confirmed by md5.

#116	Bonus Finding Fixed — Duplicate @staticmethod Above _finite_or_nan September 14, 2026	FINDINGS — FIXED

“"models/signal_router.py has duplicate @staticmethod stacked above _finite_or_nan — harmless, cleanup only.""

Commit a70fc1b. One of the two stacked decorator lines removed; _finite_or_nan keeps its single @staticmethod, same body, same six call sites in _build_decision_object — no behavior change. code_hash moved, predicted and confirmed: 13908bba… to 3867ba49…, because code_fingerprint.py hashes the docstring-stripped parse tree and a duplicate decorator is a second entry in the function's decorator_list. Golden snapshot unmoved — confirmed via tests/test_golden_path.py passing, not assumed from the AST theory alone. Verified in the Linux sandbox against a fresh clone in all three configurations (466 passed / 338p,117s / 395p,0f,32e) and matched exactly on Viktor's own machine before he committed.

#117	Viktor's Request: the Setup's Direction Stated in Words, Not Just Via Targets September 14, 2026	DECISION — ADOPTED

“A live run correctly produced a SHORT setup on a bearish market, under a NO-TRADE decision, with nothing on the panel saying "SHORT" — worked out from the descending targets instead.”

Commit 66f1479. A solid colored box added directly under the DECISION line in panel_render.py — green "THIS ANALYSIS IS FOR A LONG" when bias is BULLISH, red "...FOR A SHORT" when BEARISH — on every run, not only ones that authorize a trade. Rendering only; source is bias.raw, the same field decision_model.py's _determine_final_action reads for direction per Viktor's 2 September ruling, so this box cannot disagree with an authorized LONG/SHORT by construction. One case handled deliberately rather than reused: _refuse_incoherent_plan firing ("NO-TRADE (PLAN CONTRADICTS ACTION)") means bias and the risk plan's own targets disagreed this run — see tests/test_direction_source.py, the 2 September "long label on a short plan" defect this box exists to never repeat through a new line — so that case prints "SETUP DIRECTION: CONTRADICTORY" rather than taking a side. Verified by rendering four decision objects directly, since no existing test asserts on this box. code_hash moved, predicted and confirmed: 3867ba49… to c7c50a99…. Golden snapshot unmoved. Live-data run on Viktor's machine matched the prediction exactly before he committed.

#118	NEUTRAL Bias Ruled Silent, Not a "No Lean" Line September 14, 2026	PROCESS — RECORDED

“Claude's own addition to Entry #117's box, not requested: a NEUTRAL bias printed "SETUP DIRECTION: NEUTRAL — no directional lean this run" for consistency with this panel's other "not located"/"not measured" lines.”

Commit 39e0e79. Flagged as optional; Viktor ruled: print nothing. The distinction that changes it — swing_struct, the entry zone and correlation say "not located"/"not measured" because those are measurements that failed to produce a value. A NEUTRAL bias is not a failed reading; it is bias correctly reporting no lean, which an absent box already communicates. Only the NEUTRAL branch changed — BULLISH/BEARISH still render their colored box, and PLAN CONTRADICTS ACTION (Entry #117) is a correctness guard, not a style choice, and was not touched. code_hash moved, predicted and confirmed: c7c50a99… to 38458f20…. Golden snapshot unmoved. Verified in the Linux sandbox, all three configurations unchanged (466p,2w / 338p,117s / 395p,0f,32e), and by rendering a NEUTRAL-bias decision object directly.

#119	Portfolio Document Added — Item 4 of 6, Portfolio-Ready September 14, 2026	DECISION — ADOPTED

“Drawn from the Constitution, the engineering log, the test suite, and the actual current source — deliberately not from README.md, whose own status table was checked and found to predate round 5/6 and four now-existing core/ modules.”

Commit 69f1972. docs/build/build_portfolio_document.py added, generating the fourteen-section Phase7_Portfolio_Document.pdf: objective, problem definition, architecture, data sources, modules, decision logic, testing methodology, debugging process, major problems, solutions, results, limitations, lessons learned, future development — the career plan's own section list. Item 5 of the six portfolio-ready criteria, the separate AI-attribution accounting, deliberately left out of this document per Viktor's instruction to do the two in order (see Entry #120); this document does not itself declare the release gate open, the remaining findings closed, or the project portfolio-ready, leaving all three to Viktor the way PHASE7_NEXT.md does throughout. One factual error caught in proofreading before Viktor ever saw it: an early draft's Results section attributed the sixteen-item remediation sequence to GPT-5.6 Luna Pro's Step 8 re-audit; checked against PHASE7_NEXT.md and build_roadmap.py's own text and corrected — the sequence is GLM 5.3's Step 5 output, Luna Pro's Step 8 came after it and found further findings of its own, including the decision-bar-integrity Critical no prior audit had scheduled. code_hash and the golden snapshot unaffected — docs/ excluded from core/code_fingerprint.py's file walk by directory name, confirmed before and after rather than assumed. Full three-configuration suite unaffected (466 passed / 338 passed, 117 skipped / 395 passed, 0 failed, 32 pre-existing errors), same identity as every prior patch that session.

#120 AI-Attribution Statement Added — Item 5 of 6; All Six Portfolio-Ready Criteria Now on Record

September 14, 2026

MILESTONE — RECORDED

“A first draft's engineering-log entry-count table summed to 111 against the true total of 118, from a regex that silently missed several tag strings.”

Commit eee218e. docs/build/build_ai_attribution.py added, generating the five-page Phase7_AI_Attribution.pdf — the last unwritten piece of item 5: what Viktor designed, decided, tested and ruled on, versus what Claude produced, assembled into one place for the first time. Covers the method (every claim checkable against a code comment, a dated log entry, a commit, or PHASE7_NEXT.md, not memory alone), the Constitution's own “Roles & Authority” text, a dated table of ten rulings Viktor made himself each cited to where it lives in current source, a table of five concrete cases where Claude was wrong and Viktor caught it (including the two-September session's five verification errors and the unrequested NEUTRAL-panel line Viktor overruled at Entry #118), and the independent third-party auditors — including the one exposure (DeepSeek-lineage models touching the codebase via Aider before a later blind DeepSeek audit) stated rather than glossed over. The entry-count table's error found by recounting from this file's own AST directly — every entry_box()/highlighted_entry_box() call, its tag argument resolved as a constant — which also confirmed entries numbered 1 through 118 with no gap and no duplicate; the wrong 111-total table never reached Viktor. With this and Entry #119 landed, all six portfolio-ready criteria in docs/PHASE7_NEXT.md have their supporting material in place. Whether the release gate is open, whether the remaining findings count as closed, and whether to declare the project portfolio-ready and tag portfolio-v1 remained Viktor's own calls — not decided by this patch (see Entry #124). code_hash and the golden snapshot unaffected, confirmed before and after. Full three-configuration suite unaffected.

#121 PHASE7_NEXT.md Head Block Brought Current for Entries #119 and #120

September 14, 2026

PROCESS — RECORDED

“A session-close handover check found the project's own entry-point document still reading as if the last thing that happened was Entry #116's fix.”

Commit 59b915a (sixteenth patch). 293f310 (this document's own v1.29 regeneration), Entry #119's Portfolio Document, and Entry #120's AI-Attribution Statement had all landed the same session, but none of the three touched PHASE7_NEXT.md — exactly the gap the standing end-of-session check exists to catch. New head-block entry records what landed and in what order; that all six portfolio-ready criteria now have supporting material on record, not five; and, explicitly, that criterion 1 (the release gate) has its technical condition on record — round 6 re-audited round 5's Criticals and found no Critical Tier-1 defect, the Constitution's own bar — but had never been declared met by Viktor, with declaring the project portfolio-ready and tagging portfolio-v1 named as his calls, not decided here. code_hash confirmed unmoved, computed independently before and after rather than assumed from the exclusion rule alone. Full three-configuration suite unaffected; tests/test_golden_path.py's eight cases individually confirmed passing, snapshot unmoved.

#122 Five Token-Saving Workflow Rules Recorded in Working Practice

PROCESS — RECORDED

September 15, 2026

“Already implemented into the patch-delivery skill and recorded in memory earlier the same day — this is the in-repo copy Viktor asked for, not a new ruling.”

Commit 48757fe (seventeenth patch). New bullet under Working practice: scoped-files review packages (the touched files, their direct callers, and the tests that assert their contract) as the default for fix-verification sends and narrow patches, with the full ~400K+-token audit package reserved for a genuine fresh Tier-1 audit round rather than decided per-round; batched Engineering Notes / Portfolio Document regeneration at session close or once several commits have accumulated, rather than after every small patch, with any resulting gap stated explicitly rather than closed on reflex; no routine staging or re-reading of docs/audit_package/'s qwen_reasoning_*.txt files or superseded audit-round folders; chat-reply discipline restated as holding regardless of patch size; the fresh-session-at-closure preference restated. Docs-only, no engine module touched. code_hash confirmed unmoved on both the pristine tip and the edited tree, not assumed from the directory-exclusion rule alone. Full three-configuration suite re-run twice — once against the working edit, once independently against a fresh clone with the patch applied — identical both times and matching the recorded baseline.

#123 Correction: the Five Round-6 Mutant-Escape Findings Were Already Closed at 044b055

OBSERVATION — FILED FOR THE RECORD

September 15, 2026

“Claude reported these five as still open in chat, from reading this file's own narrative rather than checking git.”

Commit f7f2e51 (eighteenth patch). Viktor asked to fix the five round-6 mutant-escape findings (requested-run 6, GPT-6 Astra's round-5 report), which PHASE7_NEXT.md's own narrative still described as recorded “for Viktor to rule on” with no later entry saying otherwise. Checking git before touching any code found they were already fixed, verified, and folded into round 6's audit batch at 044b055 — three days before this session existed: compute_trend_health's always-neutral escape, _determine_final_action's always-WAIT escape, plot_engine_chart's no-op escape on both affected tests, SignalRouter.route's canned-error escape, and run_tests.py's zero-match-exits-0 bug, each mutant reproduced in-process, confirmed to escape the old test and be caught by the new one. New head-block entry corrects the record rather than re-doing work already done, and names the lesson plainly: this file cited 044b055 by hash without ever saying what it fixed — the exact failure shape rule 19 already describes, work from the audit, not from the summary of it. No code or test file touched; the full three-configuration suite re-run was skipped as disproportionate to a documentation-accuracy fix no suite run could validate — only git log and git show could, and both were checked first. code_hash confirmed unmoved regardless, since that check is cheap and catches any accidental non-docs change.

#124 The Release Gate Declared Open; the Project Declared Portfolio-Ready; Tagged portfolio-v1

September 15, 2026

MILESTONE — PORTFOLIO-READY DECLARED

“Round 6's targeted re-audit satisfies the Constitution's literal bar and is a new pattern for this project rather than a repeat of prior practice.”

Commit 99e022e. Viktor's ruling, 15 September 2026, made after writing his own position first and having it checked rather than handed to him — the one thing Entries #121 through #123 all named as outstanding and none of them decided. New permanent section, “Declared — 15 September 2026,” under “Two goals, and the order they finish in”: round 6's targeted (not whole-codebase) re-audit satisfies the Constitution's literal bar; the stronger basis offered is the suite's positive, checked result across all three post-round-5 fix commits (466 passed, golden snapshot and code_hash as predicted) on a suite now hardened against the vacuous-pass failure mode, not merely an absence of suspicion; criterion 2 fully closed as of Entry #123's correction; criteria 3 through 6 already on record via Entries #119 through #121. This closes the two-goals-ordering question standing since 3 September: Goal B (backtesting) can now start without risk to what Viktor submits. Tagged portfolio-v1 at this commit. Docs-only; code_hash confirmed unmoved, no engine module or test touched; golden snapshot not applicable — no code path touched.

#125	docs/build/session_handover_check.py Added; Confidence-Scored Lesson Auto-Pruning Declined September 15, 2026	DECISION — ADOPTED

“Viktor read an unrelated agent-harness project and asked for critical feedback on two of its components, then delegated the final call: “It is your call, do what is best for the engine and the project.””

Commit 63d8620 (twentieth patch). Adopted, in modified form: the seven-item handover checklist in Working practice had been an instruction (“run these checks”) rather than a structural fact of the repo — Viktor's own standing preference is a structural fix over a reminder to be careful. The new script answers four of the seven mechanically: untracked/modified files, ignored files (filtered for routine build/cache noise so a real signal isn't buried, the same blind spot this repo's own .gitignore comments already document), loose .patch/*_commit_message.txt files anywhere in the tree, and a staged-but-uncommitted index. The other three — is this session's state written into PHASE7_NEXT.md, are today's rulings recorded, are the Engineering Notes current or the gap stated — need someone to read and judge content, so the script prints them as an explicit reminder rather than pretending to answer them. Not automatic, and said so rather than overclaimed: the source project wires this class of check to persistent-session hooks; this project's actual workflow is single commands in cmd.exe with no session to hook into, so the script runs by hand at session close, computed instead of recollected. Declined: confidence-scored auto-pruning of “Lessons carried forward” — conflicts directly with this project's own standing rule that wrong turns are recorded rather than quietly corrected, and recurrence count is not a proxy for importance (the git add -A trap, Entry #113, stopped recurring specifically because it was fixed structurally; a recurrence-based prune would have deleted that lesson first, precisely because it worked). All four detectors exercised against deliberately introduced test conditions before delivery and confirmed to actually fire, then the test scaffolding removed and the tree confirmed clean again. code_hash confirmed unmoved on the pristine tree, the pre-apply working copy, and a fresh clone with the patch applied — all three identical. Full three-configuration suite run on a fresh clone with the patch applied, matching baseline exactly, including the 32 pre-existing run_tests.py errors diffed by name, not just counted.

#126 README.md Brought Current; a Stray Delivery Leftover Found; Backtesting-Scoping Findings Carried Forward

PROCESS — RECORDED

September 15, 2026

“README.md had not been touched since 30 August and still said the release gate was shut, the day after Entry #124 declared it open — the first thing anyone visiting the repository sees.”

Commits ebb0a46, d5bb329. Found while starting to scope Goal B (backtesting): the status table, the four-Critical framing, and the test counts were all stale relative to the 15 September portfolio-ready declaration. Fixed by reading current source directly rather than carrying the old table forward — decision_log.py and lineage.py's hash-and-archive mechanism for the old Item 6, bias_score/confidence de-duplication for the old Item 11, three further Criticals a later round found and closed, current test counts, the release-gate declaration and portfolio-v1 tag. The repository-layout note on backtesting deliberately softened from a first draft that would have flatly repeated “blocked by the release gate”: the Constitution's own backtest-start condition (Items 2, 3, 6, 18) is a separate clause, evidence points to it already being satisfied but nobody has formally re-checked it since 15 September, stated as not-yet-confirmed rather than assumed. Viktor's own correction, recorded plainly: this was not a case of him failing to ask for document checks — the standing handover check had never covered README.md at all, only PHASE7_NEXT.md, the Engineering Notes, and git status; that checklist gap is named, not fixed, here (closed two commits later, Entry #127). Separately, a live instance of a known blind spot, found by hand-checking what the new handover script checks: phase7_mutant_escape_correction_commit_message.txt sitting at the repo root, gitignored and therefore invisible to git status --short — very likely Entry #123's own delivery file, never deleted per its numbered sequence's own step 11. Not deleted by Claude; flagged for Viktor to remove himself. Also recorded: today's backtesting-scoping research — the backtest-gate's separate Constitution condition, Item 17's actual blast-radius requirement, a real precedent already fixed for what it guards against, and two governance questions (what “known-good checkpoint” and “fixed evaluation dataset” concretely mean) put to Viktor and not yet answered — written down so a new session can resume it rather than only a closed chat window remembering it. code_hash confirmed unmoved on both commits, README.md and PHASE7_NEXT.md both outside core/code_fingerprint.py's file walk. Full suite re-run on the README fix specifically (466 passed, matching baseline exactly — “obviously safe” has been wrong on this project before); the pure head-block addition's own re-run skipped as disproportionate, per this document's own precedent for that class of change.

#127	README.md-Currency Check Added to the Handover Script (Item 8); Ignored-File-Sweep Question Closed With No Code Change September 15, 2026	PROCESS — RECORDED

“Viktor named the split himself — the checkpoint/dataset questions are his to write, this one is Claude's engineering call — then said “Yes. Fix this.””

Commit a441023 (twenty-second patch). New check_readme_currency() function, printed as “## 5.” in the handover script's output: README.md's last-touched commit (hash, date, subject), HEAD's own, and the commit count between them. Informational only — it does not feed the script's exit code, the same way the four “still manual” reminders don't, because whether README.md's content still matches what PHASE7_NEXT.md's head block declares is a judgment call, not a git fact; what the script removes is the reason that judgment call went unmade for sixteen days (Entry #126) — nobody ran git log on README.md to see how old it was. The other half of the open question needed no code change: whether the ignored-file sweep should run “every time” was already answered by the script's own existence — check_ignored_files() runs on every invocation, unconditionally, since Entry #125 created it; the stray commit-message file Entry #126 found slipped through because the script itself didn't run that session (no shell access to Viktor's machine), not because its sweep logic was incomplete. Working practice's checklist gets item 8 to match. Three scenarios run against a throwaway git repo — a clean tree with README.md one commit behind HEAD; a stray commit-message file plus an ignored file plus a staged index together, confirming the new check doesn't disturb the existing flagged/clean logic; and a repo where README.md has no commit history at all, confirming graceful degradation — all three produced the expected output and exit code. code_hash confirmed unmoved: docs/ excluded from core/code_fingerprint.py's file walk by directory name, checked directly in EXCLUDED_DIR_NAMES rather than assumed from precedent. Test suite not re-run, stated as a scoping call rather than a skip: this script has no pytest coverage and nothing under test imports it. First real Windows run confirmed working, closing the platform question every prior entry in this row could only leave unconfirmed.

#128	Goal B Specified and Ratified Before Any Backtesting Code Exists; Two of Claude's Own Errors Recorded Inside It September 15, 2026	DECISION — ADOPTED

“After four rounds of writing his own position first, Viktor delegated the remaining open points: “make the changes you want and are content with... it is your call.””

Commit dff7d00 (twenty-fourth patch), with f9e5127 (twenty-third patch) — the latter only the head-block note recording this document's own v1.30 regeneration. New permanent section, “Goal B — the backtesting phase, specified before it starts,” written into PHASE7_NEXT.md and since moved, verbatim, to docs/PHASE7_DECISIONS.md by the 18 September split (Entry #130). Same shape as the portfolio-ready declaration (Entry #124): Viktor wrote his position on each governance question first, Claude critiqued it, corrections were folded in, Viktor ratified — four rounds covering checkpoint, rollback and isolation rules; the fixed evaluation dataset; the verdict thresholds; and the completion boundary. Closes the two questions open since Entry #126 (what “known-good checkpoint” and “fixed evaluation dataset” concretely mean) and settles three that were never on the list. A verdict procedure fixed before any result exists: INCONCLUSIVE below 30 out-of-sample trades, then POSITIVE or NEGATIVE on a symmetric band of 0.30 absolute Sharpe units with drawdown guards, NEUTRAL otherwise, every verdict printed beside absolute return and time-in-market. A completion boundary that is outcome-shaped and accepts a negative verdict as completion. And a re-run clause: the first out-of-sample evaluation is the verdict of record, and any later re-run is a new phase with its own pre-registration, because pre-registration only works once. Three findings came from reading the code rather than the documents: PHASE7_PINNED_DATA silently falling back to live data (named a precondition here, fixed at Entry #131); the evaluation dataset being three series per run rather than one file, which makes truncating each series by candle open time the highest-risk line of the phase; and the pinned path already serving df.iloc[-limit:], so the harness is a time cursor on existing code rather than a new pipeline. Two of Claude's own errors, recorded in the section rather than dropped, both the shape Entry #123 exists to warn about: Claude said no lookahead sweep had ever run, and git showed 793e863 (30 August) had done it, with tests/test_no_lookahead.py pinning the result; and Claude called fc35a2f a precedent covering most of the Item 17 guard without reading it — it is a test-time path redirect, not a runtime guard. Both were caught only because Viktor asked for a critical review. Standing mitigation adopted: every claim Claude makes about the codebase's state says whether it came from the code or from a document. No Goal B code written. Docs-only; code_hash unmoved at 38458f20… on both commits.

#129	The Course Correction: Pace Slowed, and an Engine Review Before Any Goal B Work September 18, 2026 (recording a September 15 decision)	DECISION — ADOPTED

“Six audit rounds have checked whether this engine is honestly built. Not one has checked whether it is analytically sound.”

Commit 4a97c32. Decided in conversation at the close of 15 September and kept until then only in a handover document outside git; written into docs/PHASE7_DECISIONS.md as a subsection of Goal B so the reasoning survives in the repository, per the standing rule that a conversation is not a record. Two rulings by Viktor. First, the course correction: before any Goal B implementation, the engine itself is reviewed for logical soundness — is the decision logic coherent, are the constants justified, are the signals actually independent, and is the engine's thesis written down anywhere at all. Everything Entry #128 ratified stands; nothing is cancelled, only deferred. The reasoning: a backtest grades a decision process against outcomes, but it cannot show that the process's own inputs are non-redundant or that its constants are more than convenient — and backtesting an arbitrary threshold tests the threshold, not the idea. Of Goal B's preconditions only the PHASE7_PINNED_DATA fix was judged urgent on its own merits (Entry #131). Two scope conditions were agreed the same day and have not yet been acted on: the review needs its own completion boundary, written before it starts, since “review the whole engine” is unbounded by construction; and a boundary on what it can settle — internal coherence is reviewable, but whether the market thesis is correct is not something a review, or an AI, can establish; only data can. Second, the pace ruling, made by Viktor on his own judgment: 226 commits across 22 days, active on 20 of them, had become hasty, and the project slows down considerably. A session no longer opens by proposing work from the open items as a queue to clear; it asks what Viktor wants to do first. Split into an entry of its own, rather than folded into Entry #130, at Viktor's choice. The review's first two findings are Entries #132 and #134.

#130	PHASE7_NEXT.md Split Into NEXT, DECISIONS and HISTORY, Proven a Move and Not a Prune; the Bias Weights Confirmed Fingerprinted September 18, 2026	PROCESS — RECORDED

“Reassembling the eight extracted blocks in their original order reproduced the 5,831-line file byte for byte.”

Commit 4a97c32. docs/PHASE7_NEXT.md had grown to 57,833 words and was what every session read first. Split three ways: PHASE7_NEXT.md becomes a current-state entry point only, rewritten each session rather than appended to (1,059 words, a 98% reduction in what a session reads first); docs/PHASE7_DECISIONS.md holds standing material — the two goals, Goal B, the rulings, Working practice, the earned rules; docs/PHASE7_HISTORY.md holds the head-block archive and every dated session narrative, append-only. Proven programmatically rather than eyeballed: every one of the original's 5,831 lines was extracted by exact line range into one of eight blocks, and reassembling the blocks in original order matched the pre-split file by string equality. Four sections titled “Rulings” needed a judgment call — standing rulings went to DECISIONS, rulings inseparable from build narrative went to HISTORY — each named in the commit message so Viktor can override it. Resolved in the same commit by reading the code, rather than re-running the malformed grep that had left it unverified: the six bias weights are not in FINGERPRINTED_CONFIG, which holds config.py-level names only, but they are fingerprinted through the separate FINGERPRINTED_MODULES / module_snapshot() mechanism, and both feed run_hash — already pinned by tests/test_lineage.py::test_the_run_hash_moves_when_a_bias_weight_moves. No code change was needed. Named and deliberately not fixed: seven test files and three docs/build/ scripts, this document's own build script among them, still cite PHASE7_NEXT.md for facts that now live in HISTORY; none of them reads the file, so the citations are stale rather than broken. Earlier entries of this document that cite PHASE7_NEXT.md are left as written, under this document's no-silent-edits rule — they were true when written. One consequence surfaced a day later and was repaired at 5c73e24 (Entry #135): HISTORY had not been appended to after the split, so the 18 September session's account existed only in NEXT until that commit moved it across verbatim. code_hash unmoved at 38458f20…; full three-configuration suite unchanged (466 passed; 338 passed, 117 skipped; 395 passed, 0 failed, the same 32 errors by name).

#131	Three Tier-1 Fixes: PHASE7_PINNED_DATA's Silent Live Fallback, _save_state's Non-Atomic Write, and the Handover Filter's Routine Noise September 18, 2026	FINDINGS — FIXED

“Viktor: “Let us do Tier 1 and then start a new chat.””

Commit 5be5d82, with its landing update 2ebbd6b. Three defects, each confirmed real against source before the patch touched anything. (1) data/data_fetcher.py: with PHASE7_PINNED_DATA set to a path that did not resolve to a directory — a typo, another machine's path, a dropped mount — pinned_source() returned None, which get_tf() reads as “go live.” A broken pin request served live data with nothing in the output saying so, reintroducing the exact nondeterminism the pinned mechanism exists to remove. It now raises ValueError naming the resolved path, the choice set_pinned_source() already made for the sibling case; engine_core.py's existing top-level handler turns that into a clean decision-object error, traced before anything was written. (2) core/engine_core.py: _save_state opened the state file with open(path, "w"), which truncates before a byte of the new state is written, so an interrupted write left an empty or partial file — and _load_state read that silently as a first run, losing Exit Watch's comparison baseline. It now writes a temp file in the same directory and os.replace()s it over the real path; the degrade-not-halt behaviour of 29 August is unchanged. (3) docs/build/session_handover_check.py flagged logs/ and the generated docs/audit_package/round*/ content as unexplained ignored files every session an engine run or package rebuild had happened, burying a real surprise among expected clutter; both joined the routine-noise filter, with a negative control proving an unrelated ignored file still surfaces. Seven new tests; each fix's test fails against the pre-fix code, reproducing its defect. One wrong turn recorded: the atomic-write test was first written with pytest's monkeypatch fixture, which passed under pytest but moved run_tests.py's error count from 32 to 33 — caught only by running all three configurations, and rewritten fixture-free. The first engine-code change since the portfolio-v1 tag: code_hash moved from 38458f20… to 5080ccf0…, in exactly the two touched engine files' per-file digests, and had not moved on any commit between the tag and this one. Golden snapshot unmoved, as predicted, since neither changed path is reachable on the golden fixture. Suite 466 to 473 on the Linux sandbox. The landing update, 2ebbd6b, also removed a duplicate open-item bullet its own first editing pass had missed, found by rereading the whole file rather than trusting the diff of the block it had touched.

#132	Bias-Weight Independence: ADX's Second, Near-Duplicate Read Removed From continuation_strength September 19, 2026	FINDINGS — FIXED

“Asked to write a position on whether the six bias weights are defensible, Claude found 40% of bias_score's weight reading the same raw ADX value on two different curves.”

Commit f24a6e9, with its landing update 2fee78f. The engine review's first finding (Entry #129). trend_health (weight 0.30) and reversal_continuation (0.10) both read adx_val from indicators/trend_health.py and scored it on two different monotonic curves. Not Item 11's defect — that was a reused computed value; this was a reused raw input reaching bias_score through two formulas, which is why it passed Item 11's audit and six later rounds that checked for shared values, not shared raw indicators. models/risk_model.py's Item 14 comment had already named and accepted this exact fact on 11 September, under the narrower standard that mattered there. Ruling: offered (a) independence only, (b) a reasoned rationale for the weight magnitudes, or (c) both, Viktor chose (a). Removing ADX rather than RSI from continuation_strength was Claude's engineering call within that scope, not separately ruled on. The fix deletes continuation's adx_component; its ceiling honestly drops from 60 to 35 rather than being rescaled back up, the same standard as the 31 August health_component removal. RSI was deliberately left in both factors, on the grounds that one reads it as a symmetric question and the other as a directional one — an exemption Entry #134 later measured. Golden snapshot predicted to move in bias.score and what derives from it, and nothing else; the prediction held. Re-baselined with PHASE7_UPDATE_SNAPSHOT=1 and diffed at leaf-field level: exactly 16 fields, all downstream of bias_score or confidence (bias.score 78.70 to 77.10 on the TESTUSDT fixture); entry, exit, structure, degradation and provenance untouched. code_hash moved from 5080ccf0… to bb47ab53…, in indicators/trend_health.py's digest alone — the comment-only edits to bias_engine.py and risk_model.py cannot move it, because the fingerprint hashes the parse tree and # comments are not part of it. run_hash unchanged on the golden run, as designed: it captures input data and fingerprinted configuration, not formula code, and code_hash is what tells the two code versions apart. One new regression test. Not fixed, by Viktor's choice: the weight magnitudes remain unreviewed hand-picked judgment calls, recorded as open; the other four factors were not checked for hidden shared inputs in this pass.

#133 git add -A Nearly Sweeps an Untracked PDF Into a Commit; Files Are Now Added by Name

PROCESS — RECORDED

September 19, 2026

“Caught only by reading git status --short before committing — a carefulness step, not a structural one.”

Recorded in 2fee78f, the landing update for Entry #132. Delivering f24a6e9, the command sequence's git add -A swept an untracked file, Claude outputs/Phase7_Session_Handover_2026-09-15.pdf, toward staging. It did not reach the commit, because git status --short was read first. The project had paid for this exact trap before — at least four earlier occurrences are recorded in docs/PHASE7_HISTORY.md, among them commits bd79800 and 77d822f and a commit-message file swept into a commit — and each time the answer had been to be more careful. Viktor's standing preference for anything that cannot be undone is a structural fix over an instruction to be careful, and a commit to a public repository is that kind of thing. So the delivery procedure changed: on Viktor's repository every file is added by name, one git add per file, so an unrelated untracked file cannot be staged whatever else is lying in the tree. git add -A survives only inside the throwaway repository used to build a patch, where nothing else exists. The change lives in the patch-delivery skill, outside the repository; 2fee78f records it as proposed for Viktor's review, not yet confirmed saved. The PDF itself was left where it was, neither committed nor deleted, with its fate recorded as an open item for Viktor. Split into an entry of its own, rather than folded into Entry #132, at Viktor's choice.

#134	The Bias Blend Asks One Question Four Times — Recorded, Not Fixed; Two Errors in the Project's Own Record Corrected September 19, 2026	OBSERVATION — FILED FOR THE RECORD

“In a sustained trend the four agree because they must, and bias_score reports that agreement as four independent confirmations.”

Commit e2c6637, with its landing update 742b514. The engine review's second pass, tracing the four factors Entry #132 left unchecked. Trend health (0.30), structure regime (0.20), SuperTrend direction (0.15) and macro bias (0.10) — three quarters of the blend — are four transforms of one measurement, the recent direction of close: an EMA slope, the gap between two close means, an ATR-banded flip, and price against its 50-EMA one timeframe up. Different windows and arithmetic, so they genuinely disagree at turning points and this is not the ADX defect repeated; but all four are monotonic in the same quantity and enter the blend with the same sign convention. Only volume sentiment (0.15) and continuation's acceleration term read something other than price direction. Entry #132's RSI exemption was measured across RSI 0 to 100: it holds in downtrends (r = 0.37) and largely fails in uptrends (r = 0.83), worth at most 1.5 points of bias_score against the 2.5 the ADX fix removed — a note, not an alarm. Why nothing is fixed: acting means reweighting or dropping factors, a trading judgment about what corroboration between correlated trend reads is worth, which this project cannot yet evaluate — the golden baseline proves a change is attributable, never that it is correct. The decision is Viktor's and has not been made. Two errors in the project's own record corrected with it. models/bias_engine.py's dependency graph described structure_regime as a swing-based label; it never was — it is a close-mean gap with hysteresis, and swing structure reaches only the panel — so the first pass had cleared the factor as independent against a mechanism it does not use. And PHASE7_NEXT.md had claimed the session's test counts matched on Linux and Windows; that was false, and Claude's error: on any LF checkout test_manifest_hashes_match_the_files fails, because the pinned CSVs are normalised by .gitattributes while MANIFEST.json pins raw-byte hashes, so the test is green only where the manifest was generated. Recorded as a standing platform difference, not fixed here. Two new tests pin facts, not judgments — structure_regime responds to close alone, and both RSI correlations stay inside bands — each with a negative control. code_hash unmoved at bb47ab53…, the bias_engine.py edit being comment-only; golden snapshot untouched.

#135	Pre-Push Hook: the Handover Check Runs on Every Push and Stops It on a Finding September 19, 2026	DECISION — ADOPTED

“Viktor's first answer was that a finding should “only warn, then I consult with you and we push or not.””

Commit 2f5fdaa, with its landing update 5c73e24. githooks/pre-push runs docs/build/session_handover_check.py on every git push: exit 0 lets the push through, anything else stops it with the findings printed. Installed once per clone with git config core.hooksPath githooks, since git does not version .git/hooks. Proposed the session before and never built; Viktor raised it himself. Ruling: three options were put to him — (A) stop the push, consult, override with git push --no-verify; (B) a y/N prompt at the terminal; (C) warn and let the push through. Claude pointed out that warn-only could not deliver the second half of his first answer: the push would already have left before the consult. He chose (A), and --no-verify is his override after the findings have been read together, never a routine step. (B) was set aside not on merit but because reading the keyboard from a hook under cmd.exe through Git for Windows' sh could not be verified from the sandbox. The hook fails closed — no python on PATH, a missing script or a crash all stop the push too. The handover check gained a section 6 that flags a clone without the hook installed. Two findings from verification. Applying the patch in one step wrote the new hook with CRLF despite the same patch adding githooks/* eol=lf, because git apply works under the attributes in force before the patch — so a patch that changes .gitattributes and adds files now applies in two steps, .gitattributes first. And on Linux git 2.43 a hook that cannot be executed makes git push hang rather than fail; the first negative control hung the sandbox for five minutes, and the push tests now carry a timeout. Nine new tests, no fixtures; against the pre-change code 8 fail with FileNotFoundError, proving only that the file is new, and 1 fails behaviourally. The hook's first run on Viktor's Windows machine was the push of 2f5fdaa itself, clean. code_hash unmoved at bb47ab53…; golden snapshot untouched. Suite on the Linux sandbox, clone made with core.autocrlf=true: 485 passed, 0 failed with pandas_ta. The same count on Viktor's machine is recorded as his confirmation by proceeding, not a count Claude read. The landing update, 5c73e24, also found PHASE7_HISTORY.md had not been appended to since the 18 September split, and moved the previous account there verbatim before rewriting NEXT — proven by un-demoting the moved block and comparing it byte for byte.

#136 PHASE7_NEXT.md Brought Current at c7ced36; a Wrong Prediction About git status Recorded

PROCESS — RECORDED

September 19, 2026

“The prediction was made from memory, not checked; the patch-delivery skill itself says to confirm the delivery files are absent, which was right.”

Commits e431714 and e115272, docs-only. PHASE7_NEXT.md had been written at 2f5fdaa/5c73e24 and did not know about c7ced36 (this document's v1.31 regeneration), so two sessions had to open from a resume prompt that corrected it. e431714 made the entry-point file true on its own again: the previous version moved verbatim into PHASE7_HISTORY.md under a dated entry, headings demoted one level, and proven a move rather than an edit by un-demoting the block and comparing it byte for byte, with a one-character negative control. README.md, fifteen commits behind, had its test counts and layout tree brought current; two older omissions it found (Claude outputs/ and decision_log_backups/ missing from the tree, no pointer to the three PHASE7_* files) were recorded as an open item rather than fixed out of scope, and closed at Entry #140. e115272 recorded a wrong prediction from that delivery: Claude said git status --short would list the two delivery files as untracked, and they did not appear at all, because .gitignore ignores *.patch and *_commit_message.txt. No harm followed, but the error existed only in chat until Viktor asked for it to be recorded. Test counts unchanged at 485 / 0 with pandas_ta; on Viktor's machine recorded as his confirmation by proceeding. code_hash unmoved at bb47ab53…, recomputed on both commits' trees for this regeneration.

#137 The Pinned-Data Manifest Check Made Portable: One Hash Per Content, on Every Checkout

FINDINGS — FIXED

September 20, 2026

“One committed file, two byte forms — and a test that hashed the bytes.”

Commit b68de08, tests and tooling only. tests/test_pinned_source.py compared the raw bytes of the three pinned CSVs with MANIFEST.json. With * text=auto in .gitattributes, git checks those CSVs out CRLF on Windows and LF elsewhere; the manifest had been generated on a CRLF tree, so the test passed on Windows and failed on every default LF checkout (484 / 1). Recorded since Entry #134 as a standing platform difference and never quietly subtracted from a count. The fix hashes content with CRLF normalised to LF, in both places: make_pinned.py's sha256(), which writes the manifest, and a copy of the same function in the test, which cannot import make_pinned.py because that module writes files when imported. The manifest was regenerated by running the script itself, not by editing hashes in, and its CSVs came out byte-identical. A new test, test_manifest_hash_ignores_line_endings, builds both byte forms of each CSV itself and requires both to match, so a revert to raw-byte hashing fails on every platform — the existing test alone could not catch that on an LF checkout. Against the pre-fix manifest, 2 behavioural failures and no import errors; four negative controls, each restored and confirmed with cmp. Counts +1 everywhere: 486 / 0 with pandas_ta, 355 passed / 120 skipped without it, run_tests.py 415 / 0 / 32; a default LF clone now 486 / 0. make_pinned.py sits under docs/, excluded by directory, so code_hash stayed at bb47ab53… — computed on three trees, not assumed. Golden snapshot untouched.

#138	The Open-Items List Scrapped in Favour of an Independent Audit September 20, 2026	DECISION — ADOPTED

“We are gonna scrap that entire list. 1-7.”

Commit 16d3c1f, docs-only. Viktor's ruling, recorded as its own section in docs/PHASE7_DECISIONS.md. The seven judgment and review items open at e115272 — his written position on the bias-weight magnitudes and thesis, five points from the 15 September PDF, retiring the stale to-do PDF, the four-factor finding of Entry #134, RSI reaching bias_score through two factors, the engine review's missing scope and completion boundary, and the backtest-start condition never formally declared met — were scrapped as open items in favour of “a new audit by an independent model,” with three mechanical items first: the manifest test (Entry #137), the stale citations (Entry #139) and README's omissions (Entry #140). How it came about: asked whether the weights had been changed without a record, Claude checked git log rather than answering from memory — the six weights are unchanged since 83a9425 (26 August), which replaced the first commits' four-term blend wholesale; no commit ever changed a value, and the original choice predates the repository, which is why no reason for it exists. Claude's recommendation on items 4 and 5 — that fixing either means choosing new weights, item 1 by another name — was superseded by the full scrap rather than answered. What the ruling does not change: the findings stay on record in DECISIONS and HISTORY; the Constitution's backtest-start condition is neither declared met nor waived; the 15 September course correction (Entry #129) is not amended. Still undecided, and Viktor's: which model, which package, and whether the auditor sees the scrapped findings. The same commit moved the previous NEXT to HISTORY verbatim and found it headed a day ahead of the commits that wrote it — recorded in the new entry's heading note, not corrected in the moved text, since HISTORY is append-only. code_hash unmoved (bb47ab53…).

#139	Stale Citations of PHASE7_NEXT.md Fixed in Live Files; Dated Records Are Not Edited to Follow a Move September 20, 2026	DECISION — ADOPTED

“Viktor's first position was to update the dated records too: “the aim was to keep a real record.””

Commit 9a35f1c, comments, docstrings and docs only. After Entry #130's split, citations of PHASE7_NEXT.md written before 18 September sent readers to a file that no longer held what they cited. The item was recorded as seven test files and three build scripts; a repository-wide search found twenty-four files, including comments in six engine modules. Two rulings. (A) The item covers every stale pointer the search found — yes. (B) Whether dated records — this document's entries, session handovers, audit reports, constitution reviews — are edited too: Viktor's first answer was yes. Claude argued against: those citations were true when written, since DECISIONS and HISTORY did not exist before 18 September; this document is declared append-only in the project's own AI-Attribution Statement; and some of the records are other models' output. Viktor accepted a forward note instead, and Claude read his reply as that ruling and said so before building. Rule: a dated record is not edited because a file it cites was later split or moved; one sentence in PHASE7_NEXT.md's head block, plus the ruling in DECISIONS, tells a reader where such content now lives. The citations inside this document's own earlier entries are therefore left as written, as v1.31 had already done. Each live pointer now names the file and section that hold the content, each target checked to exist. Two citations were wrong from the day they were written, not merely stale: “rule 18” in models/entry_model.py and tests/test_decision_bar_integrity.py — the same commit that wrote them, 108cc9f, inserted a rule and moved that one from 18 to 20; it is now rule 22. Found and not fixed: DECISIONS' rules list has no rule 18, and is not renumbered because rule numbers are cited throughout. The rendered text of the Portfolio Document and the AI-Attribution Statement changed with it; both PDFs were regenerated with this document's v1.32. code_hash predicted and confirmed unmoved at bb47ab53… — comments and docstrings are stripped by the fingerprint's AST pass — and all 33 per-file fingerprints equal; a negative control (one real expression changed) moved the hash, and restoring the file restored it. Suite unchanged, 486 / 0; Windows confirmed by proceeding.

#140 README.md Brought Current: Layout Tree, a Pointer to the Project's Record, and Four Stale Statements

PROCESS — RECORDED

September 20, 2026

“Sixteen modules at ratification, twenty-two now — and the README still said sixteen.”

Commit 077d120, docs-only. Closes the omissions Entry #136 recorded: Claude outputs/ and decision_log_backups/ are in the layout tree, and a new paragraph describes PHASE7_NEXT.md (entry point), PHASE7_DECISIONS.md (standing rules and rulings) and PHASE7_HISTORY.md (dated record), noting that older documents cite the first for content now in the other two. Because the item's recorded list had come from a narrower search at Entry #139, README.md was read end to end and each factual claim checked against the tree and the code. Four more stale statements, fixed in the same commit: the test counts, which predated the test b68de08 added; Logs/ in the tree where the directory is logs/ — one directory on Windows, two on Linux, the defect class of 6 September; the module lists and “sixteen files”, which was the count at ratification (83a9425, checked with git ls-tree) and is now twenty-two; and the claim that run_tests.py needs “nothing but a Python interpreter,” which the runner's own docstring had withdrawn on 6 September — most test files import pytest at module scope. Left alone, because it is not Claude's call: README.md still cites docs/Phase7_Audit_Findings_Complete.pdf, which does not exist; whether the reference goes or the source material is restored has been Viktor's open decision since 6 September, and is now an open item in PHASE7_NEXT.md rather than only in HISTORY. code_hash unmoved at bb47ab53…; suite unchanged; Windows confirmed by proceeding.

#141	The Round-1 Audit Outputs Recorded as Lost; README.md and the 30 August Reinstall Entry Corrected September 20, 2026	DECISION — ADOPTED

“The pre-reinstall check looked at what the repository tracked and ignored, and could not see material that had never entered it.”

Commit 4629002, docs-only. README.md said all raw auditor output was published verbatim in docs/Phase7_Audit_Findings_Complete.pdf. Neither that PDF nor its source — the four round-1 auditor outputs of 27 August (Run 1, DeepSeek V4 Pro, blind; Runs A, B and C, Kimi K3), run through the OpenRouter chat room — was ever in the repository, and the same passage's claim that the later rounds' reports sat in docs/audit_reports/ “the same as the original four” was wrong for the same reason. The item had been Viktor's call since 6 September. Claude recommended restoring the source if a copy existed and correcting README.md if not; Viktor chose to search first. The repository on disk, git history on every branch, D:\Phase_7_Engine_Random_Files, the OpenRouter chat room (it keeps history only in the browser), OpenRouter's request logs (I/O logging was off: token counts, no text) and a Claude data export for 26–31 August were searched, and none held a copy. Claude's first guess, that an icon on some request-log rows meant the text was stored, was wrong; opening a row showed it. Viktor then ended the search and approved Claude's three recommendations as proposed. README.md states the loss and what survives — Phase7_Remediation_Plan.pdf, GLM 5.3's prioritisation of the findings, not the findings themselves — and says no copy “has been found” rather than that none survives, because Downloads, the Cowork task list and one unnamed project were not opened. The error in PHASE7_HISTORY.md's 30 August entry, “nothing lived only on disk except gitignored logs/”, is corrected by a new dated entry, not an edit, per the ruling Entry #139 records. build_findings_bundle.py is kept, as the record of how the PDF was built. Found and deliberately not fixed: README.md's “Every raw report from every round is in docs/audit_reports/” — round 2's folder holds a reasoning trace that ends before any report — recorded as an open item, since it lay outside this one. code_hash unmoved at bb47ab53…, computed on both trees. Suite unchanged, 486 / 0 with pandas_ta on the Linux sandbox; on Viktor's Windows machine recorded as his confirmation by proceeding, and the pre-push hook reported SUMMARY: clean on the push.

#142	Round 3 Ruled Not Independent — for Authorship, Not Exposure; README's Audit-Round Claims Corrected September 20, 2026	DECISION — ADOPTED

“The reason that holds is authorship: that session produced the Remediation Plan, and round 3 reviewed work done under it.”

Commit 53394ff, docs-only. Viktor chose four loose ends from PHASE7_NEXT.md's Open items. The first was README.md's “Every raw report from every round is in docs/audit_reports/”. That was narrowly overstated: round 2 produced no report, since every attempt on 2 September stopped before writing one, and its folder holds a reasoning trace. The same paragraph said every round ran “on a model reporting no prior exposure”. Checking that turned up a contradiction in the record about round 3 (GLM 5.3 Flash, 5 September). HISTORY's 5 September entry called it not independent under the lab-level rule. The 2 September ruling clears chat-session exposure, and calls applying the lab-level rule to it a category error. GLM 5.3's 28 August session was chat-session exposure, read in the provider export. The reason that does hold is authorship: that session produced the Remediation Plan, and round 3 reviewed work done under it. Viktor first proposed counting round 3. He remembered the 28 August run as minor, and as a deliberate compromise. Claude checked the record, which showed a Step 5 run with the full source, the Constitution and all four round-1 outputs, and no compromise ruling anywhere. Viktor then ruled that round 3 is not counted as independent. It is recorded as its own section of PHASE7_DECISIONS.md, including what it weakens. The independent rounds after the original audit go from five as published to three that are both independent and complete (rounds 4, 5 and 6). README's status row, release-gate paragraph and audit paragraph were corrected to match. The wording was delegated to Claude. code_hash unmoved at bb47ab53…, computed on both trees.

#143 The Suite's One Live Network Call Left in Place, and the Empty Rule Number Marked

September 20, 2026

DECISION — ADOPTED

“One rule ever changed number, and it changed three times.”

Commit 53394ff, the same four-item session as Entry #142. The api.mexc.com connection seen during test runs was traced by hooking every outbound connection (a sitecustomize hook on getaddrinfo and on HTTP tunnel setup). The traced call is tests/test_clean_checkout.py::test_main_runs_without_a_logs_directory, which runs python main.py in a temporary copy of the repository. It is the suite's only outbound connection, with pandas_ta, under both pytest and run_tests.py, and it writes only inside its temporary directory. Claude recommended pointing that subprocess at tests/fixtures/pinned/, which ran main.py to exit 0 with no network in the sandbox. Viktor ruled to leave it, and Claude was overruled; no test changed. Rule numbers: PHASE7_DECISIONS.md's rules list was rebuilt at all 98 commits that touched it. One rule ever changed number: “Fixing the instance you found does not close the item”, 18 → 20 → 21 → 22, all on 1–2 September. Every citation of rules 18–21 in tracked text files and in every commit message was checked against the list as it stood when written. All were correct except “rule 18” meaning today's rule 22, which 9a35f1c had already recorded; PDFs were not searched. At Viktor's approval, item 18 now says the number is empty and where the rule went. Found while adding it: Markdown numbers an ordered list by position, so from rule 19 on, every rule rendered one lower than its cited number. This was checked with a CommonMark renderer, not on GitHub itself; after the note, all 38 match. A wrong turn, sandbox only: while testing the pinned-data option, a failed cd let main.py run inside the scratch clone and create logs/ there. Nothing on Viktor's machine was involved. Found and not fixed here: the Portfolio Document's “no prior involvement” claim (fixed at Entry #146), and the Qwen-or-Kimi attribution of round 2's observations (settled at Entry #158).

#144 Landing Facts Filed Between Sessions, and Why the Notes Gap Cannot Fall Below One

PROCESS — RECORDED

September 20, 2026

“A commit that regenerates the Notes cannot cover itself, so one commit behind is the lowest the gap can go.”

Commits 92775ea, 982e70f and 65a0aef, docs-only. They follow 3a899b5, which regenerated this document at v1.33 and so has no entry of its own, as c7ced36 and c745a67 had none. 92775ea filed 3a899b5's landing facts: Windows counts confirmed by proceeding, and SUMMARY: clean from the pre-push hook. Claude laid out three options for filing them and pointed out that the third, a Notes regeneration, would not close the gap. Viktor chose to leave them for the next docs commit, then changed to a commit of their own. Claude's wording for that option said the Notes would “stay one commit behind”. That was wrong: the commit made it two. 982e70f filed 53394ff's landing. Its three recorded points were Viktor's fresh clone confirming the tip; Claude's expected git status output listing ?? lines for the two delivery files a second time, although .gitignore ignores them (lines 106–107, checked); and the to-do PDF deleted (Viktor's report, not checked by Claude). Claude noted that landing these ahead of the audit-planning commit departed from Viktor's own earlier instruction; the choice was his. 65a0aef did the once-per-session rewrite of PHASE7_NEXT.md. The previous version moved to HISTORY verbatim, with the move proven by un-demotion and a negative control. A first run of that proof reported a mismatch. The cause was the proof's own offset (+2 instead of +4 for the two separating blank lines), not the file. Ruled in 65a0aef: the Notes were not regenerated. Claude recommended waiting, because a regeneration then would have to be repeated after the audit-planning commit. Viktor agreed to regenerate once, after the audit rulings, and only if the audit package included the Notes. This v1.34 regeneration is Viktor's later call (Document History, v1.34). code_hash unmoved at bb47ab53… through all three commits.

#145	The Run Summary Named the Expected-Value Line, Not the Action's Own Reason September 20, 2026	FINDINGS — FIXED

“A refusal summarised as worth taking.”

Commit a9d4b1f. models/decision_model.py built the decision's one-line summary from reasons[-1], and _compute_ev appends to reasons last on every path through evaluate(). So every run's summary was the illustrative expected-value sentence, whatever the run had decided. The golden fixture recorded a NO-TRADE (RISK TOO HIGH) summarised as “positive -- worth taking on average”. The summary reaches the decision object and the decision log, but panel_render.py does not print it, which is why six audit rounds driven by the panel did not reach it. It was found, not by an audit, while answering the five questions from Viktor's 15 September PDF against the code. None of the five turned out to be a calculation error. Claude's prediction in chat that they were “mostly checkable arithmetic” was wrong, and is recorded as wrong. The fix makes the summary name the reason of whichever stage last set final_action, tracked in evaluate(). Claude proposed that over the simpler reasons[0], because two later stages can replace the action after the first reason is written. Viktor: “the summary fix can go into the patch also, whatever.” The EV sentence now says it is confidence restated (a linear rescaling that cannot disagree with it). It quotes the breakeven and positive levels, 33/100 and 43/100, derived from AVG_REWARD_R and the new named EV_BREAKEVEN_BAND_R. The VALIDATION line gained its “/100”. A bias_engine.py comment records where volume and macro each reach an output, and says whether three penalties is the right total weight is Viktor's judgment, unmade. Not fixed: whether the EV line belongs in the reasoning at all, the unreviewed bias weights, and the BTC adjustment scaling up with correlation. code_hash bb47ab53… → b0213766…. Golden re-baselined in exactly the two predicted fields. A mechanical consequence was predicted only after it happened: the Linux re-baseline wrote the fixture all-LF, so it was converted back to CRLF by hand. Tests 486 → 500 with pandas_ta, 14 new; against the pre-fix code, 9 behavioural failures and 1 AttributeError for the new constant (which proves only that it is new). Confirmed on Windows by Viktor's live run: VALIDATION 35.00/100 with its arithmetic checked, and the logged summary naming the risk reason.

#146	README and the Portfolio Document Stopped Calling the Four Original Runs Independent of the Build September 20, 2026	FINDINGS — FIXED

“We fix this”

Commit 6e1baba, docs-only. The pre-push hook's section 5 flagged README.md three commits behind. Read end to end against the tree, it held four stale or wrong claims. First, its two test-count lines predated a9d4b1f. Second, “four independent runs … on models with no prior involvement in the build” was contradicted by README itself seventy lines later: three models worked on the build through Aider, two of them DeepSeek's (V3 and R1), and Run 1 was DeepSeek V4 Pro. Independence is tracked at the lab level here, so the lineage counts; Runs A, B and C (Kimi K3) are unaffected. Third, “Every raw report from every round is in docs/audit_reports/” was wrong: round 2 produced no report. Fourth, docs/build/README.md listed build_findings_bundle.py's PDF with no note that neither the PDF nor its source ever reached the repository. The Portfolio Document carried the second claim in its own words. That had been a standing open item, and Viktor ruled that it rides along now (“We fix this”). The script is corrected and Phase7_Portfolio_Document.pdf regenerated. Each correction says in the text that the document used to say otherwise, and since when. A wrong turn, recorded: the first check that the PDF's text had changed extracted zero characters (a hand-rolled zlib walk that matched nothing) and reported “changed text lines: 0”. That reads as “nothing changed” when in fact nothing was read, the same vacuous-check defect this project has recorded before. It was redone with pdftotext: exactly one paragraph differs. A no-change rebuild differs in 74 bytes, all timestamps. Found and not fixed here: the TREND line's missing denominator (Entry #148). code_hash unmoved at b0213766….

#147	code_hash Is Comparable Only Within One Python Minor Version September 20, 2026	OBSERVATION — FILED FOR THE RECORD

“A hash that appears to have moved is a question about the interpreter before it is a question about the code.”

Commit b869a30, docs-only: the once-per-session rewrite of PHASE7_NEXT.md at 6e1baba. The previous version moved to HISTORY verbatim, proven four ways: un-demotion reproduces the old file byte for byte; a one-character negative control fails; exactly seven lines differ, each by one leading #; and old HISTORY is a byte prefix of the new. The new fact: the sandbox's default python3 was 3.11.15, and it reported code_hash 938720cd… on an unmodified 6e1baba tree. Under Python 3.12.3 the same tree gives b0213766…, matching Viktor's Windows decision-log record, and under 3.13.13 a third value again. core/code_fingerprint.py hashes ast.dump output, a CPython implementation detail, and already said so under “WHAT IT DOES NOT SURVIVE”. It is now a standing rule in PHASE7_NEXT.md: every code_hash claim about this project is computed under Python 3.12 (Viktor runs 3.12.10). The same thing happened again the next morning (a45366b3… under 3.11 at 119c8a3; Entry #149). The v1.34 row's code_hash values were all recomputed under Python 3.12.3. Also recorded in the rewrite: Viktor's position that the independent audit might be paused, “since we fixed all the issues now”, recorded as a position and not a ruling. Claude's objection sits beside it: an audit exists to find what nobody has found, and the TREND-line defect, found by Viktor on a live panel, is a small instance of exactly that. The hook's results on the pushes of 65a0aef, a9d4b1f and 6e1baba were not reported, and the file says so rather than carrying earlier clean runs forward. The ruling came the next day (Entry #156).

#148 The TREND Line Printed a Score With No Denominator — Rule 22 Collecting on a9d4b1f

September 21, 2026 — **FINDINGS — FIXED**

“Fixing the instance you found does not close the item.”

Commit 119c8a3. core/panel_render.py printed “TREND : BULLISH / HEALTHY (Score: 100.00)” directly above “VALIDATION : WEAK (Score: 35.00/100)”. A reader could not tell the top of a 0–100 scale from an unbounded total that happened to reach three figures. Viktor found it reading a live panel about two hours after a9d4b1f fixed the identical defect one line below. It was not found by an audit, nor by the work that produced a9d4b1f, which had gone looking for this class of defect and stopped at the instance it was asked about. That is rule 22 of PHASE7_DECISIONS.md, collecting. Every line render_panel() prints was then read, not sampled: after this fix, no score on the panel lacks its scale. The denominator was checked in indicators/trend_health.py, not assumed from the 100.00 on screen: the score is clamped with max(0.0, min(100.0, …)). The new test holds the scale from two independent directions, because either alone can pass vacuously. It reads the clamp bounds from the function's parse tree, and it runs the function on four generated frames plus the degraded path. Three negative controls: revert the panel line (3 fail); change the clamp to 10 (1 fails, the parse-tree half); inflate the value after the clamp so the parse tree still reads 0–100 (1 fails, the behavioural half). The third proves the two halves separately load-bearing. code_hash b0213766… → 35718f6b…; of 33 per-file digests exactly one moved. Golden unmoved: the rendered panel is discarded by every caller and never reaches the decision object. Tests +4 in all three configurations, run_tests.py's 32 errors unmoved. Confirmed on Windows by Viktor's live run at 00:02 on 21 September: “Score: 100.00/100” above “Score: 35.00/100”.

#149 The Entry-Point Document Carried a Wrong code_hash — Eleven Statements Corrected in Place

September 21, 2026 — **PROCESS — RECORDED**

“A wrong code_hash in the entry-point document is the one fact a later session would trust without rechecking.”

Commit 635a94e, docs-only. PHASE7_NEXT.md was written by b869a30, and 119c8a3 landed after it in the same session. A commit cannot name its own hash, and the once-per-session rewrite rule meant nothing corrected the file afterwards. So the document every fresh session reads first carried five false statements and six incomplete ones. They were found by the handover check at the session's close and by reading the file against 119c8a3 at the next opening. The worst was code_hash: the file gave b0213766… as unmoved, and 119c8a3 had moved it to 35718f6b…. The rest were the tip, the test counts, the Notes gap, and the TREND line still listed as open; plus 119c8a3's Windows live run (only in chat until then), the golden-snapshot line, and the pre-push hook. The hook reported SUMMARY: clean after a first attempt failed on DNS, but whether that push carried both b869a30 and 119c8a3 or only the latter was not established. It is written as unknown. Viktor was offered three scopes and ruled all eleven, as Claude recommended. The Notes count disagreed, nine against ten. Both were right: the file counted itself, and Viktor did not count the amendment. The file now states the convention. Tests against pre-change code were shown to be no evidence here: with both files replaced by the single word GARBAGE the suite still passed. So equal counts before and after are vacuous for this change, and the commit says so instead of offering them. The evidence was the code_hash on both trees, the byte match of the applied files, a clean reverse-apply, and a patch that refuses a drifted base. The 3.11 interpreter gave a45366b3… on an unchanged tree, a third interpreter and a third value (Entry #147). code_hash unmoved at 35718f6b….

#150	The Pre-Backtest Review: Thirteen Findings From Reading the Code, and the Order of Work Delegated September 21, 2026	DECISION — ADOPTED

“Organize a to do list and we start working, It is up to you.”

Commit ebb4e5c, docs-only. Viktor asked what to put to the engine before backtesting: whether its logic is correct, whether the panel shows the correct information, and whether it carries dead code. Claude read the decision path and the input side, and built an import graph of every engine module. The decision path was risk_model, entry_model, exit_model, decision_model, panel_render, and what they use from engine_core, signal_router and bias_engine. The input side was indicators, trend_health, structure and volume_profile. Thirteen findings resulted. Every one comes from reading the code, none was reproduced by running the engine, and each says how far its reachability was checked. Four are questions about what the engine should do, and are Viktor's to rule on before any backtest is designed. They are: the entry zone against a plan measured from the last close (4); a NEUTRAL bias still printing a plan (5); the stop pulled to the 75-day volume point of control (6); and indicator values beyond 5 sigma silently replaced by the previous bar's (7). The rest became Claude's work orders, B to F. Viktor delegated the order of work to Claude. He deferred the read of data_fetcher, validation, decision_log and lineage until after the work list (“we check it after”); that read is Entry #158. Two items were reported as Claude's claims rather than findings: every engine module is reachable from main.py except two by design (14), and _refuse_incoherent_plan cannot fire today, correctly (15). Also filed: SUMMARY: clean on 635a94e's push (Viktor's paste). The previous NEXT moved to HISTORY verbatim, proven by un-demotion. code_hash unmoved at 35718f6b….

#151	Work Order B: the Panel Printed AERO for Every Symbol, “nan” Scores, and 0.00 for Absent Ones September 21, 2026	FINDINGS — FIXED

“One function rather than six guards, so the next score line added cannot forget.”

Commit a530006, review findings 1–3. (1) Two sentences in the panel's BTC section named AERO whatever the symbol. On BTCUSDT the engine always skips BTC context, so every BTCUSDT run printed that the “AERO analysis” was unaffected. Sequence item 12 had fixed the same string in decision_model.py and missed these two. Both now use asset_name(), moved to module level so the panel and the model cannot drift apart. (2) TREND and VALIDATION could print “Score: nan/100”. The router sends NaN for a value it could not measure, and both lines formatted it directly. The price lines were guarded at Round 6 F3; these two were not, including in a9d4b1f and 119c8a3, which edited exactly these lines. Decision Reasoning's “trend strength nan/100” was the same defect. (3) Entry quality, confidence, trade quality and BTC-adjusted confidence printed 0.00/100 when absent, a score nobody computed; latent, since the router always sets the first three. The fix is one helper, _score_text(), used by every score line: the score with its scale, or “not computed”. Found while scoping and fixed in the same function: asset_name tried “USD” before “BUSD”, so ETHBUSD read as “ETHB”; longest suffix first now, with no recorded run affected. The new test renders a panel with every numeric field missing and asserts “nan” appears nowhere, so a later score line fails too. Against the pre-fix code 8 fail: 6 behavioural, 1 parse-tree, and 1 ImportError that proves only that asset_name is new. Nine negative controls each failed exactly their tests. A wrong turn: two controls were first contaminated by stale bytecode, because the broken suffix order was the same length and was restored within the same second. They were rerun with PYTHONDONTWRITEBYTECODE=1. code_hash 35718f6b… → ec88cf24…; golden unmoved, and the pinned run's whole panel is byte-identical. Confirmed on Windows by Viktor's live run: every score with its scale, no “nan”, AERO in the BTC section, and the record carrying ec88cf24….

#152	Work Order C: the Risk Plan's Direction Parameter Decided Nothing; a Wrong-Side Stop Now Refuses September 21, 2026	FINDINGS — FIXED

“A check at the only producer is the structural form of what D asked for.”

Commit e3f3d51, review findings 8, 9 and 10. (8) RiskModel.calculate_stop_targets tested its detailed_bias parameter for “LONG” or “SHORT”. Its only caller passes the bias state machine's “BULLISH CONFIRMED”, “BEARISH CONFIRMED” or “NEUTRAL”, so the test never matched, and the plan's direction has always been the sign of bias_score. It is the same wrong-vocabulary shape as entry_model's A1/A2 gates, in a second file. It gave no wrong answer, because raw bias comes from the same score. But the parameter read as the direction source, so it is removed, and passing it raises TypeError. (9) A fallback for a stop on the wrong side of price replaced the stop distance and left the stop where it was. It is unreachable on the engine's path, because bias_score is clipped to ±100. It is wrong if reached, shown on the pre-fix code with a bias_score of 400: a LONG with its stop at 101.04 above a price of 100, which validate_risk_parameters passed, since it measures distance with abs(). A NaN bias_score gave a NaN stop with a SHORT's targets. Both are refused now. (10) Nothing at runtime checked the stop's side. It was planned as work order D and folded in as Claude's call: calculate_stop_targets is the only producer of a stop, and it now refuses a wrong-side one. detailed_bias is also removed from lineage.risk_inputs, which it never fed. The new tests include a grid over every bias score, trend health, volatility state and structural level the engine can produce. It asserts the stop's side, so the unreachability claim is checked rather than argued. Against the pre-fix code, 5 fail through the signature change alone, which proves only that the signature changed, and 2 fail behaviourally. A wrong turn: the NaN control first failed to fail, because the new wrong-side refusal caught it with a different message. The test now asserts the message names bias_score. Golden re-baselined in exactly the one predicted line; no decision field moved. code_hash ec88cf24… → ac02a155…. Finding 6 was evidenced twice from the record: the live run's stop was the HVN, 18.2% below price, where the ATR stop sits 6.2% below; the golden fixture's was 19.9% below. Confirmed on Windows: the next live record carries ac02a155… with no detailed_bias in risk_inputs.

#153	Work Order E: the Last Fabricated Defaults, and the Forming Candle Found From a Wrong Prediction September 21, 2026	FINDINGS — FIXED

“The simulated order log is an obvious thing to build a backtester on, which is why it is fixed before that work starts.”

Commit afd8460, review findings 12 and 13. (12) live_trading._build_simulated_order recorded an absent entry zone, stop, target set or price as 0.0, and an absent risk reason as “OK”. It was the third copy of the shape Round 6 F3 removed from the router and the panel. Absent is now None, which is JSON null. Found while scoping: the order's timestamp called datetime.utcnow(), deprecated since Python 3.12. A 5 September note in the same file records fixing the same call twice elsewhere. This one was missed, and it was the source of the suite's two DeprecationWarnings. A test now sweeps every engine module's parse tree for utcnow. (13) structure.calculate_structure copied its analysis with .get() defaults of “NEUTRAL STRUCTURE” and 0.0 on keys that are always present. Had they fired, the HVN column would have held a price of zero. They are indexed directly now. Six negative controls each failed exactly their tests. The first draft of the counts said 527 and 392: the file has six tests and the arithmetic added five. The suite run caught it before delivery. code_hash ac02a155… → 3e76c1c5…; a docstring counts, since the hash is over the parse tree. The suite now prints no warnings line. Finding 16 was found from a wrong prediction. Claude said Viktor's next live run would repeat the 05:28 numbers, because it was the same candle. It did not: same stop, a new price, each target exactly 1R, 2R and 3R from it. The engine decides on the candle still forming. That was confirmed in data/data_fetcher.py, which keeps every kline and reads close_time only to discard it. Which candle counts is a rule, not a defect, so it is Viktor's to rule on with findings 4–7. Confirmed on Windows (filed at 49de810): the 06:18 run's record carries 3e76c1c5…, and pytest printed 528 passed with no warnings line.

#154	The Constitution Checked Against the Session's Work; Two Claims Reheaded as the Builder's Own September 21, 2026	PROCESS — RECORDED

“Pausing the audit is consistent with the Constitution; pausing it and then starting backtesting is not.”

Commit 49de810, docs-only. After afd8460 the end-of-session check found five things that existed only in chat, and this commit files them in PHASE7_NEXT.md. First, work order E's Windows confirmation (Entry #153). Second, SUMMARY: clean on afd8460's push. Third, a third live instance of finding 6: the 03:28, 03:51 and 04:18 UTC records all carry a stop equal to the HVN, 0.549596, 18.2%, 18.1% and 17.6% below price, and all three are NO-TRADE (RISK TOO HIGH). They were read from the records, not the panels. Fourth, Viktor's plan, stated that morning: about four more weeks of our own review, fixes and improvements before the next audit, with no time pressure. It is recorded as a plan in his words, not a ruling. Fifth, the Constitution check he asked for. Claude read the ratified PDF's rules and its Next Steps sequence, not its version history or glossary: nothing the session did breaks it. One condition bears on the audit decision: step 8 requires an independent re-audit of what changed before step 9, backtesting, and none of the session's fixes has had one. Where the session is thinnest is recorded too: two riders found mid-work (the BUSD suffix in B, utcnow in E), and finding 16 as a possible Item 3 or Item 8 gap, Unknown until ruled on. Following from the check, findings 14 and 15 were reheaded. Under “Checked and found sound” they read as the builder certifying his own compliance, which the Constitution does not allow. They are now “Claude's claims, open to the independent auditor”. code_hash unmoved at 3e76c1c5….

#155	Work Order F: the Entry Signals Now Confirm the Side the Decision Ladder Chooses September 21, 2026	DECISION — ADOPTED

“Make the adjustments you want and do what is best for the engine and the project”

Commit 3f263c2, review finding 11. long_signal and short_signal were computed on every run and carried into the decision object, the log and the simulated order, but read by nothing that decides. The golden fixture recorded long_signal: true beside NO-TRADE (RISK TOO HIGH). Claude set out three options (remove, confirm, relabel) and named the choice as Viktor's, because it changes which trades are taken. He chose confirm: a trade the ladder chooses is taken only if that side's signal confirms it, else NO-TRADE (SIGNAL UNCONFIRMED). He wrote the conditions after two rounds of Claude's critique. Macro is removed, as counted twice. The reversal check is ruled by its components: HVN proximity alone no longer blocks, divergence and exhaustion do, and only a reversal pointing against the trade blocks. The CONFIRMED requirement is dropped, so the signal follows raw_bias as decision_model does. When risk and confirmation both fail, RISK TOO HIGH stays the label and the confirmation failure is appended. He then delegated the rest, in the words above. Claude's calls under that delegation: one direction-aware divergence rule on every tier, replacing the upper tiers' direction-blind veto (CONSERVATIVE had none); trend health of at least 50 dropped from the signal, since every trading branch already requires it; and macro inside CONSERVATIVE left for its own commit, finding 17, work order G. Degradation overrides only actions naming a side. Claude had told Viktor during scoping that DEGRADED INPUT would override SIGNAL UNCONFIRMED; reading the code corrected that. The live log was checked first: of 29 records, no action would change. The one SHORT passes where the old signal refused it. Golden: three fields added, none changed, as predicted. code_hash 3e76c1c5… → 44f7296b…. 21 new tests; against the pre-fix code, only an ImportError, which proves only that the file is new. The behavioural evidence is seven negative controls. Found, not fixed: finding 17; finding 18 (the bias state machine now gates no trade); and the dependency that, if finding 6 stops pulling the stop to the HVN, nothing checks HVN proximity. Not backtested: nothing yet says the gate improves decisions.

#156	The Independent Audit Paused for Our Own Work — No Backtesting Before an Independent Re-Audit September 21, 2026	DECISION — ADOPTED

“We pause the audit. We work on the engine another four weeks.”

Ruled by Viktor and recorded in PHASE7_DECISIONS.md with work order F (commit 3f263c2). Asked to write his position first, he set out about four weeks of reviewing everything, analysing, fixing and evaluating our own work, with integrity and truth kept. After Claude's critique he added that we should not audit our own work, but we have fixes to make and should prepare a new audit as well as we can. What the ruling does not change: the Constitution's step 8 still requires an independent re-audit of what changed before step 9, backtesting. So no backtesting before an independent re-audit, and the open questions of the 20 September ruling (which model, which package, whether the auditor sees the scrapped findings) stay open until the audit is planned. Agreed with Viktor: nothing done in those weeks is written up as verification. Our own review, tests and negative controls are recorded as unaudited work awaiting the auditor, never as “found sound”. Preparing the audit means three things: a running list of every change since the last audit, which becomes the auditor's scope (Entry #158); Viktor's open rulings answered before the package is sent; and the deferred read done first. Claude made two points, and neither was adopted. First, the pause had no end condition: “about four weeks” is a plan, not a trigger, and Claude suggested a date around 19 October to decide again. Second, the no-backtest rule exists only as text, where a structural form would be a backtest entry point that refuses to run without a recorded re-audit. An end condition was ruled on 22 September. It is recorded in PHASE7_DECISIONS.md in the same commit that brings this document to v1.34, and enters this document at its next regeneration. The second point remains open.

#157	F's Live Run Came After Its Commit — From Then On the Command List Stops at the Live Run September 21, 2026	PROCESS — RECORDED

“Never predict live numbers; check the record.”

Commits aafded0 and cb659f1, docs-only. F's command sequence asked for a live run before git commit. It did not happen: at the push the decision log still held 29 records, the newest on the old code_hash. Claude caught it by reading the log, not from a paste. The record of Viktor's 19:24 run carries 44f7296b…, so F was confirmed on Windows after its commit, not before, and is recorded that way. The record shows F's fields as designed. The hook's output on 3f263c2's push was not pasted, so a clean result was inferred from the push landing, not seen. A wrong prediction: the expected git status put the new test file last, but git sorts by path. The once-per-session rewrite was deferred to the next opening, agreed at 25% of Viktor's weekly usage. cb659f1 did it. It condensed the fixed findings to one line each while keeping their numbers, because test docstrings cite them by number. It wrote the carried lesson: on any change that moves code_hash, the live run and the panel read happen before git commit, and Claude checks the newest decision-log record for the new hash before giving the commit steps. It also corrected “working tree clean” to “last seen clean”, since the hook's summary is not a git status. README's audit and backtesting rows were brought current in Viktor's first person and are his to reword. cb659f1's HISTORY move was proven by un-demotion but without the one-character negative control. That control was run afterwards (Entry #158): the proof holds, and the mutation is detected. The rule was followed at 2c7a7d1, 4e2b1c8, 486f1a5 and 3bfa6b7: each command list stopped at the live run, Claude read the record, and only then gave the commit steps. code_hash unmoved at 44f7296b… across both commits.

#158 The Deferred Read (Findings 19–27) and a Running Change List for the Next Auditor

September 21, 2026 — **PROCESS — RECORDED**

“we check it after”

Commit 1cc2142, docs-only. The read Viktor had deferred (Entry #150) covered data/data_fetcher.py, data/validation.py, core/decision_log.py and core/lineage.py, read in full at cb659f1 with their call sites. It produced nine findings, 19 to 27, all Claude's, none changing a decision, none fixed in this commit. They are from reading the code, except finding 19, checked against Viktor's live decision log: 1 of 30 records (6 September) carries a bare NaN token. Finding 16, the forming candle, was confirmed again. New file docs/audit_change_list.md: the running list of every change since the last independent audit, which becomes the next auditor's scope. Its baseline is e65a0f7. The nine files round 6's fix-verification was sent match every commit from ac0a211 to e65a0f7 by sha256, and code_hash is 13908bba… at both (Python 3.12, computed). Each line gives what changed, what it closes, which tests guard it, and whether output moved. Every test file named was checked to exist. Two engine changes had no test at all, the SETUP DIRECTION box and its NEUTRAL branch (Entry #159). From here on, each commit changing engine code, tests or tooling adds its line in the same commit. A wrong claim was caught before delivery: the first draft said F was the only listed change that alters which trades are taken. But f24a6e9 moves bias_score on every run (78.70 → 77.10 on the fixture), so it can change actions near a threshold too; corrected. Two record questions were settled. The 119c8a3 counts came from an autocrlf clone, and the note the question leaned on was the patch-delivery skill's, stale since b68de08. Round 2's eleven observations are Kimi's, since HISTORY's later same-day section supersedes the earlier Qwen attribution. SUMMARY: clean on cb659f1's push (Viktor's paste). code_hash unmoved at 44f7296b….

#159 The SETUP DIRECTION Box Gets Its First Tests

September 21, 2026 — **PROCESS — RECORDED**

“There is no pre-fix code — the box already behaves this way; these tests hold it there.”

Commit 9c4917c, tests only. It is the first of the six commits Viktor chose from the third session's list. The audit change list (Entry #158) found 66f1479 (the SETUP DIRECTION box, 14 September) and 39e0e79 (its NEUTRAL branch, Viktor's ruling the same day) to be the only engine changes since the last independent audit with no test. tests/test_setup_direction_box.py, six fixture-free tests, pins them. A BULLISH bias prints the LONG box and nothing else, and a BEARISH bias the SHORT box. A NEUTRAL bias prints no box and no SETUP DIRECTION line. NO-TRADE (PLAN CONTRADICTS ACTION) prints CONTRADICTORY and neither side, for both biases: that case keeps the 2 September “long label on a short plan” defect from coming back through this line. The box follows the bias on a run that takes no trade, which is why it was added, and it sits directly under the DECISION line. Against the current code all six pass, so the evidence is four negative controls on core/panel_render.py, each restored and confirmed with cmp: the CONTRADICTORY guard disabled (1 fails), a NEUTRAL line printed (1 fails), the box shown only when a trade is authorised (2 fail), and the box moved below Decision Reasoning (1 fails). code_hash unmoved at 44f7296b…, since tests/ is excluded by directory. Tests 549 → 555 with pandas_ta, run_tests.py's 32 errors unmoved. Confirmed on Windows by Viktor proceeding past the command sequence's stop conditions.

#160 **The Decision Log Wrote Bare NaN, Which Is Not JSON — and Its Reader Dropped Damaged Lines Silently**

September 21, 2026

FINDINGS — FIXED

“Python reads it; jq and JSON.parse reject the line.”

Commit 05a12c7, deferred-read findings 19, 20 and 21, core/decision_log.py only among engine files. No decision changes. (19) write() used json.dumps with its default allow_nan=True, so the router's deliberate “not measured” NaN was written as the token NaN. It was checked against the live log: 1 of 30 records. Now _json_safe() copies the record with every non-finite float as null, and json.dumps runs with allow_nan=False, without modifying the caller's object. Records already in the log keep their NaN, and read() still accepts them. A numpy.float32 keeps its old path, the string “nan”, which is valid JSON; that is out of scope and said so in the docstring. The archive's half of the same defect was left for the lineage commit (Entry #162). (20) read() skipped any unparseable line anywhere, silently, though its comment named only a torn final line. read_with_report() now returns the records and the skipped line numbers. read() keeps skipping, so one damaged line cannot make the history unreadable, but now logs a warning naming the file and the lines. (21) The module docstring named “the pinned directory” as the recorded source; the engine records the literal “pinned”. Against the pre-fix code, 6 of 8 new tests fail: 4 behavioural, and 2 AttributeErrors that prove only that read_with_report() is new. Five negative controls on the module each failed exactly their tests. code_hash 44f7296b… → f691c4d8…. It was not confirmed on Windows by itself: the one live run was taken on the next commit's tree, which carries this code (Entry #161). There the record was strict JSON, but held no non-finite value to convert, so the null conversion is evidenced on Linux only.

#161	The Readable Fingerprint Missed Three Constants — and a Scan Now Finds the Next One September 21, 2026	FINDINGS — FIXED

“This list has fallen behind the code three times before, each fixed by adding names.”

Commits 2c7a7d1 and 5d3a4b4, deferred-read finding 22, on its own because it moves run_hash and so the golden snapshot. FINGERPRINTED_MODULES listed DecisionModel.BTC_ADJUSTMENT_CAP, but not its sibling BTC_STRESS_PENALTY, nor AVG_REWARD_R and EV_BREAKEVEN_BAND_R. All three were inside code_hash, so a change was detected, but the record could not say which value a run used. They are added. The structural part is tests/test_fingerprint_names_every_constant.py. It does not name the three: it reads every module the list covers, and fails on any UPPER_CASE, numeric, finite constant at module or class level that is not listed. Its exclusions are stated in the test. Without pandas_ta the scan is skipped whole rather than narrowed, because a pass that silently left a module out would look like full coverage. Negative controls: one of the three names removed (2 fail), and a new constant added to risk_model.py (1 fails, naming it). The first attempt at the second control appended to a file with no final newline and broke its syntax, failing for the wrong reason. It was redone, and is recorded because a control that fails for the wrong reason proves nothing. The golden re-baseline changed exactly seven leaves: run_hash in two places, the archive name in two, and three added module_constants leaves. No decision field moved. An incomplete prediction is recorded: Claude had named only run_hash and the archive name, missing the three leaves that are the change itself, and that run_hash sits in two places. code_hash f691c4d8… → 82c14ef9…. Confirmed on Windows before the commit, by Viktor's 22:09 run, read by Claude from his disk: 31 records, the first 30 byte-identical to the earlier copy, the new one strict JSON on 82c14ef9… with the three constants at 15.0, 2.0 and 0.3. 5d3a4b4 filed that check and the hook's SUMMARY: clean, section 1 “none”, on the push of 2c7a7d1, and brought README's counts current.

#162	An Archive Can Now Be Checked Against Its Decision-Log Record, Not Only Against Itself September 21, 2026	FINDINGS — FIXED

“An edit that rewrote both the frame and its digest passed.”

Commit 4e2b1c8, deferred-read findings 23, 24 and the archive's half of 19, core/lineage.py only among engine files. No decision changes. (24) verify_archive() compared each stored frame with the digest stored beside it in the same file. So an edit that rewrote both passed, and its docstring overstated what it checked. The check that means something, the archive against the decision log's input hashes, existed only in the tests. New verify_against_record() re-hashes each archived frame, trusting no stored digest, and compares the result with the record's input hashes and run_hash. Run on Viktor's live log on Linux, 30 of 31 records matched in every frame and in run_hash. The remaining one, from 6 September, carries no input hashes. Not built: a command-line wrapper, and a comparison against re-fetched data. (23) The archive is named by run_hash, which excludes code_hash by design, so a rerun on the same candles under changed code overwrites the file. It is documented and pinned by a test, not renamed, because the name is pinned by the golden snapshot and the record keeps its own code_hash. (19, archive half) write_archive now uses the decision log's sanitiser with allow_nan=False. This is latent, since only meta can hold a float and none is non-finite today. All 22 of Viktor's archives re-serialise byte-identically under the old and the new call (Linux). Not changed: run_hash() still uses the default allow_nan, because changing it could move the pinned run_hash, and no non-finite value reaches it. Against the pre-fix code, 10 of 12 new tests fail: 8 AttributeErrors that prove only that the function is new, and 2 behavioural. Six negative controls each failed exactly their tests. code_hash 82c14ef9… → 4fe5084e…; golden unmoved. Confirmed on Windows before the commit, by Viktor's 22:54 run, the 32nd record: the earlier 31 byte-identical, and the new record strict JSON on 4fe5084e…. Its archive had the same hash, and both verify functions returned all True. SUMMARY: clean, section 1 “none”, on the push.

#163	Validation Refuses a Future-Dated Last Candle, Reads Timeframes Exactly, and Converts an Aware Clock September 21, 2026	FINDINGS — FIXED

“One bar, not zero: the exchange's clock and the local clock are two clocks.”

Commit 486f1a5, deferred-read findings 25 and 26, and 28, found while fixing 25; data/validation.py only among engine files. No decision changes. The commit also did the once-per-session rewrite of PHASE7_NEXT.md, with the move to HISTORY proven by un-demotion and a one-character negative control. (25) The staleness check bounded a series' age from above only, so a last candle dated after the reference time passed as current. A last candle more than FUTURE_TOLERANCE_BARS = 1 bar after now is now rejected as future-dated. Reachability on MEXC was not measured: the sandbox cannot reach api.mexc.com. (26) The timeframe table lower-cased its input, so MEXC's month “1M” became one minute, and MEXC's “60m” was unlisted, which switched the spacing and staleness checks off. The lookup is exact now, with “60m” and “1W” listed and “1M” deliberately unlisted, since a month is not a fixed number of minutes. The stated consequence: an upper-case spelling the table does not list, such as “4H”, now skips the checks where before it was lower-cased and checked. Nothing in the engine uses one. The evidence for MEXC's spellings is weak, and is said to be. (28) An aware now was stripped with tz_localize(None), which keeps the wall-clock reading, so a Stockholm now was read two hours late. It is converted to UTC first now. The engine's own caller passes naive UTC, but after 25 a relabelled now west of UTC would make a current series look future-dated, so the two land together. Recorded, not changed: an unlisted timeframe still skips the checks silently. Against the pre-fix code, 7 of 11 new tests fail, all behavioural. Five negative controls each failed exactly their tests. code_hash 4fe5084e… → 2c8ebe32…; golden unmoved. Confirmed on Windows before the commit, by Viktor's 23:17 run, the 33rd record: no fetch error, so the forming candle passed 25's check as predicted. The earlier records were byte-identical, and the new record and its archive strict JSON on 2c8ebe32…. SUMMARY: clean on the push.

#164 **A Failed Fetch Quotes the Exchange's Reply Instead of Discarding It — the Sixth of Six**

September 21, 2026

FINDINGS — FIXED

“Fixing only the non-list branch would have left the usual case as it was.”

Commits 3bfa6b7 and 5e55eb8, deferred-read finding 27, data/data_fetcher.py only among engine files. No decision changes. fetch_ohlc answered every reply that was not a non-empty list with “Empty or invalid API response.”, so MEXC's own code and message were read and thrown away. Now a reply that is not a list is quoted: MEXC's code and message as that pair, anything else as its repr. An empty list is reported apart. The fix was widened past the finding as written, as Claude's call: a 4XX or 5XX reply's body is quoted too. MEXC documents 4XX for a malformed request, so that is where its message usually arrives, and requests' HTTPError text holds only the status line and the URL. Quoted text is cut at 300 characters, and the unused import time is removed. The text reaches the panel's “Data fetch failed” line only, because a failed fetch writes no decision-log record. What MEXC actually returns for a bad symbol was not observed, since the sandbox cannot reach it. Against the pre-fix code, 6 of 8 new tests fail, all behavioural; one fails on any unused import in the module. Five negative controls each failed exactly their tests. code_hash 2c8ebe32… → b3c2308f…; golden unmoved. Confirmed on Windows before the commit, by Viktor's 23:29 run, the 34th record: strict JSON on b3c2308f…, the earlier records byte-identical, and both verify functions all True. The fetch succeeded, so the new error text is evidenced on Linux only. 5e55eb8 filed that check, the hook's SUMMARY: clean on the push, and the session's close: the six commits for findings 19–28 and the direction-box tests are done. Filed with this regeneration, from 22 September: the push of 5e55eb8 printed SUMMARY: clean, section 1 “none”, section 6 “installed”. Section 5 reported README.md one commit behind, which was expected. The previous session's sandbox fetched the tip afterwards, and the verified close patch reverse-applied to it cleanly (from Viktor's message opening 22 September). This regeneration's own clone of GitHub's tip is 5e55eb8.

#165 When the Audit Resumes: a Closed List, and What “Done” Means

DECISION — ADOPTED

September 22, 2026

“We do the audit after we have done G and my findings.”

Ruled by Viktor on 22 September and recorded in PHASE7_DECISIONS.md, “Ruling, 22 September 2026 — when the audit resumes”, in commit e804b64. That commit is v1.34's own regeneration, which by precedent has no entry; this entry is for the ruling it carried. It answers the first of Claude's two points in the 21 September pause ruling (Entry #156): the pause had no end condition. Viktor wrote his position first, in the words above. Claude's critique named three gaps. An open list could grow without end, since every new finding would push the audit back, which is the no-end-condition point again by another route. “Done” was undefined. And it was unclear whether the trigger replaced the four weeks. His answers closed all three. The list is closed: work order G, and findings 4, 5, 6, 7, 16 and 18. Anything found afterwards goes on a separate list for after the audit, kept in PHASE7_NEXT.md, and does not move it. His argument was that a fixed list is itself the end condition, and Claude accepted it and withdrew the objection. A finding is done when he has ruled it and any code his ruling calls for has landed. A ruling to leave something as it is counts as done. Claude proposed that definition, and he accepted it. The four weeks was his estimate with a buffer, not a rule, and no date is attached. Recorded but not ruled: a tension Claude named between “until … it actually generates profit” and Goal B's verdict-of-record and re-run clauses. Working toward a profitable engine fits Goal B one pre-registered phase at a time; re-running until a result passes does not. What the ruling weakens is stated in DECISIONS: a defect found during the work waits for after the audit, whatever its severity, because the ruling makes no exception. Claude raised that consequence before the patch was applied, and Viktor did not add an exception. Claude's second point, that the no-backtest rule exists only as text, is untouched and open. The Constitution's order is unchanged: an independent re-audit comes before backtesting.

#166 e804b64's Push Filed, and README's “About Four Weeks” Replaced by the Ruling

PROCESS — RECORDED

September 22, 2026

“A commit cannot file its own push.”

Commit c5dc4cd, docs-only, at Viktor's request for the session's finishing touches, nothing major. It filed the pre-push hook's result on the push of e804b64: SUMMARY: clean, section 1 “none”, section 6 “installed”, and section 5 README.md two commits behind, as predicted (Viktor's paste). The tip was fetched into the sandbox afterwards, and its six changed files matched the verified build byte for byte. README.md's Independent-audit row still said the next round was paused “for about four weeks of fixes and audit preparation”, stale since the ruling in Entry #165. It now says the round is paused until work order G and Viktor's rulings on six open findings are done, with anything found meanwhile waiting for after the audit. The row is in his first person, like the rest of README, and his to reword. PHASE7_NEXT.md was amended in place, because the session's rewrite and HISTORY move had already happened at e804b64. Filed with this regeneration: the push of c5dc4cd printed SUMMARY: clean, section 1 “none”, section 6 “installed”, and section 5 README.md 0 commits behind, as predicted (Viktor's paste). The tip fetched afterwards reverse-applies the verified patch cleanly. Each such filing leaves the next push's result owed to the commit after it. That regress cannot be closed, only kept one step long. code_hash unmoved at b3c2308f…, computed on the applied tree under Python 3.12.3. The suite was unchanged at 597, 460 with 126 skipped, and 526 / 0 / 32.

#167 The Round-1 Audit Outputs Recovered and Filed — Entry #141's Loss Corrected Forward

DECISION — ADOPTED

September 26, 2026

“Hashed where they were found, before anything was copied or opened.”

Commit 75682ee, documentation only. Entry #141 recorded the four round-1 auditor outputs of 27 August (Run 1, DeepSeek V4 Pro; Runs A, B and C, Kimi K3) as lost. Viktor found them on 23 September in D:\USB Backup\Phase7_Engine documents, a backup that ends 30 August, and hashed all six files there with certutil before copying or opening any. They are filed byte-identical, with a README of their own, in docs/audit_reports/round1_deepseek-v4-pro_kimi-k3_2026-08-27/. His hash table was matched three times: on copies staged from his disk, on the files and index blobs after the patch was applied to a sandbox clone, and on his disk after git apply. The folder is -text in .gitattributes, so git neither normalises nor rewrites the files. Checked first, because the repository is public: no key, token, e-mail address or local path in the six. The room export records OpenRouter's routing region as “OSL”; editing it would break the hash match, so the choice was six or five, and Viktor ruled all six. A file in the same backup folder whose name marks it as an API key was neither copied nor opened. Run A's findings match the export's final assistant message except for 21 newline characters; Runs 1, B and C have no primary source in the backup. A new HISTORY entry carries the account, and README.md and docs/build/README.md are corrected forward, each saying it used to say otherwise. Entry #141 and the v1.33 row stay as written. Also filed: the push of bc48f59 (SUMMARY: clean, README.md one commit behind); a correction that “To do list Claude Phase 7 Engine.pdf” was not deleted, as recorded at 982e70f; and the 22 September ruling in Entry #168. Found and not fixed: build_findings_bundle.py still reads the four files from docs/audit_raw/ and says they are not in the repository; repointing it is Viktor's call. The session ran under Claude Sonnet 5 until Viktor switched it to Claude Opus 5.5; everything in the commit was done under Opus 5.5. code_hash unmoved at b3c2308f….

#168 bias_score's Weighting Waits for After the Audit, and the Auditor Is Shown It

DECISION — ADOPTED

September 22, 2026

“The audit can show the weighting is implemented as written; it cannot say the weights are right.”

Ruled by Viktor in chat on 22 September, option 3, and filed in PHASE7_DECISIONS.md by commit 75682ee as owed. The weighting of bias_score stays off the closed list of Entry #165 and waits until after the independent audit, and the independent auditor is shown these findings. “The weighting” means four things: the six weights, 0.30 / 0.20 / 0.15 / 0.15 / 0.10 / 0.10, which have never been reviewed or given a written reason; the four-factor overlap (Entry #134); RSI reaching bias_score twice; and one volume/macro disagreement adding up to three penalties. The last is recorded in Viktor's words. A keyword search did not find it as a finding in the repository; the nearest is the 15 September PDF's “volume may be counted three times”. Options 1 and 2 were not seen by the session that filed the ruling, and are not recorded. It answers part of what the 20 September ruling (Entry #138) left open, whether the auditor is shown the scrapped findings: yes for two items of that list and for the weights half of a third, whose thesis half is not covered. The rest is still Viktor's call. Claude's note of 20 September applies to the part now decided: an auditor that has already seen a finding is less independent on that point. What it weakens, as DECISIONS states it: an unreviewed weighting sits at the centre of every decision the engine makes, and the audit runs without it being ruled.

#169 Any Model May Build or Review, at the Cost of Audit Eligibility; the Pre-Send Check Is Audit Preparation

DECISION — ADOPTED

September 26, 2026

“a model can be used to review and build but not for audit”

Commit 2582994, the eighth session's first, documentation only. Two rulings, both stated in chat on 22 September and ruled on 26 September; Claude drafted the wording and showed it before the commit. First, in Viktor's words above: no model is barred from build or review work on the engine, and a model that has done it may not audit it. Authorship disqualifies permanently; having read engine source disqualifies for independence. The 22 September wording, “spent models for build and review”, was Claude's framing in the roadmap, and wrong: Viktor asked why any model would be spent for building, and Claude agreed that “spent” is a fact about the audit role only. What goes with it: a ledger entry at the time of use, the normal patch gates with the model named in the commit message, and no build or review work for a family still clean for auditing. His sentence ended “... but not for audit, if a model.”; Claude read the trailing clause as the same condition and did not ask. Second: the pre-send token check, which measures the package with the chosen auditor's own tokenizer and refuses the send if the package and an output reserve do not fit, is audit preparation, not an engine item. It does not reopen the closed list, is built under docs/ so it cannot move code_hash, takes a line in docs/audit_change_list.md, and cannot be finished until the auditor is pinned. Not started. Also recorded, from Claude's reading of staged code, not a run: the stop-distance refusals return before the ladder work order G changes, and the effective stop ceiling is 8%, not 15%; finding 6's text was corrected in place, with a note. A wrong prediction: 75682ee's git status prediction missed ten untracked files in “Claude outputs/”, so its first push stopped at the hook's section 1. Built from a sandbox clone, it could not see files that exist only on his disk. From then on his repository is listed over the device bridge first. Also recorded: Grok's own account of what it read, and that Viktor stopped using Grok after its run reported 2 failures, on a suite that passes on his machine, which it did not resolve. code_hash unmoved at b3c2308f….

#170	The Stop-Distance Finding Is Evidence for Findings 4 and 6, Not a New Item September 26, 2026	DECISION — ADOPTED

“We can put it on the after audit list”

Commits 1861208 and eba6a2a, documentation only, the eighth session's second and third. Viktor's finding of 22 September was that the engine almost never trades; Claude's count from his decision log found 33 of 37 records refused as RISK TOO HIGH, 29 of them on the distance limit. Ruled on 26 September, by agreeing to Claude's recommendation rather than writing a position first: it is evidence for findings 4 and 6, not a new item and not part of work order G, and the closed list is unchanged. Claude said at the time that the only critique of the recommendation was its own, and that this was acceptable because the question was where a finding is filed, not what the engine does. Two parts of Claude's reading were ruled separately. The 8% ceiling, under which a stop 8–15% away is refused as EXTREME RISK, is part of finding 6 (“Ok”). The risk gate running before the weak-validation and lean checks goes on the list for after the audit, in the words above: it changes what the panel says (2 of the 33 refused runs would otherwise have been WAIT), not which trades are taken. Why not G: the risk verdict returns before the ladder G changes. What it weakens: findings 4 and 6 now carry the measurement of how often the stop vetoes a setup, which makes them heavier to rule. The count and a ladder replay are Claude's re-implementation, not engine output. eba6a2a filed 1861208's push (SUMMARY: clean, as predicted) and saved Grok's own account of what it read, unchanged, outside the repository in the Random Files folder, with its SHA-256 recorded: the evidence for Grok's ledger entry, then still owed. It also recorded that filing a push's result always leaves the filing commit's own push owed, a floor of one like this document's (Entry #144). Entry #178 records the ruling that ended it. code_hash unmoved at b3c2308f….

#171 The Repository Moved to E:\phase7_engine, Proven Before the Originals Were Renamed

PROCESS — RECORDED

September 26, 2026

“Every path the engine records is relative to the repository root — read from the code, not assumed.”

Commit 7d103d7, documentation only, the ninth session's first. Before phase 3 started, Viktor moved the repository from D:\phase7_engine to a new drive, E:\phase7_engine, and the files kept outside it from D:\Phase_7_Engine_Random_Files to G:\Phase_7_Engine_Random_Files. He chose Claude's recommended option on all three questions: keep the folder names; leave the USB backup folder where it is; rename the originals after verification and delete them himself later. Verified three ways. robocopy: 666 files in 83 directories and 25 files in 7, none failed. On Windows, in the new folder: git status --short empty, git fsck --full clean, HEAD eba6a2a, and pytest 597 passed. Over the device bridge: the 91 git-ignored data files (logs/ and the audit-package rounds beyond each MANIFEST) and the 25 Random Files, staged from both drives and compared by SHA-256, all identical, with a negative control (one byte appended changed a hash). The first rename failed because Viktor's cmd window was inside the folder; it succeeded after cd /d. Read from the code: code_hash records paths relative to the repository root, the hook path is relative (core.hooksPath = githooks), and logs/ is relative (core/config.py). Not checked: whether E: and G: share a physical disk with D:, and anything on his machine, such as shortcuts, editor workspaces or scheduled tasks, that names the old path. Not changed, deliberately: a usage example and a test docstring that name D:\phase7_engine, and every dated record. The previous NEXT moved to HISTORY verbatim, proven by un-demotion with a negative control. The push of eba6a2a filed (SUMMARY: clean). Outside the repository, the patch-delivery skill named the old path; Claude proposed its update the same day. code_hash unmoved at b3c2308f….

#172	Finding 16 Ruled: Decisions on Closed Candles; Phase 3 Starts in the Order 16, 7, 4, 6, 5, 18, Then G September 26, 2026	DECISION — ADOPTED

“Two runs inside one candle could decide differently.”

Commit bf2e802, documentation only, the tenth session's first; phase 3 of the roadmap to the audit starts here. Finding 16, found from a wrong prediction (Entry #153): data/data_fetcher.py read MEXC's close_time, threw it away and kept every row, so on a live run the decision bar's close, volume and indicators came from the candle still forming, and the staleness check, measured from the last candle's open, never called a forming candle stale. Claude re-read the fetcher off Viktor's disk and asked four questions. On two, Viktor's own position: decide on the last closed candle, on every series the engine fetches (the 4h symbol, the 1d macro, BTCUSDT 4h); and record in the log which candle each series decided on, and whether a forming one was dropped. On the other two he asked for Claude's suggestion and adopted it: the panel shows the live price on one line, as information only, with its distance from the decision close; and staleness is measured from the decision candle's close, with a short grace for exchange delay. Accepted with it: a decision can be up to one candle old, and a run just after a close may fail instead of deciding. current_price becomes the decision close, so the stop and targets are measured from it until finding 4 is ruled. The order of phase 3, with G last, was Claude's proposal under the 21 September delegation, agreed by Viktor. Also filed: a Windows test run before 7d103d7 was committed gave 590 passed and 7 errors in a new non-administrator window, all PermissionError on pytest's temp folder, and 597 once Viktor removed that folder; Claude's inference that an elevated window had created it is recorded as not confirmed. And the push of 7d103d7 (SUMMARY: clean). Its code had not landed, so the finding was not yet done by Entry #165's definition. code_hash unmoved at b3c2308f….

#173	Finding 16's Code: the Forming Candle Is Dropped and Recorded; CURRENT PRICE Becomes DECISION CLOSE September 26, 2026	FINDINGS — FIXED

“MEXC's close_time equalled open plus one bar exactly — the first evidence of MEXC's convention in this project.”

Commit add8540, the eleventh session's; data/data_fetcher.py, data/validation.py, core/engine_core.py, core/decision_contract.py and core/panel_render.py among engine files. fetch_ohlc keeps close_time, drops the trailing candle still forming before validation, and records what it dropped. A row opening more than one bar after now is not dropped, so finding 25's future-dated check (Entry #163) still rejects it. MEXC's close_time must equal open plus one bar within 1 s, or the reply is malformed. Validation rejects a forming last candle, calls a series stale once a newer candle has closed, and calls a decision candle that closed less than 60 seconds ago “not yet final”. Pinned data counts every row as closed and records forming_candles_dropped as null, not 0, since a pinned file cannot answer the question. The decision object gains provenance.decision_candles for each series. The panel's CURRENT PRICE is now DECISION CLOSE, with a LIVE PRICE line beside it, information only. Claude's reading of the ruling's point 3 (a run inside the grace fails) and the label were agreed by Viktor; the 60 s grace is Claude's under the ruling, set against his clock, measured at +0.2 s from time.windows.com. MEXC's own finalisation delay was not measured. The scope stated before the diff missed two things, both caught by the suite on the first full run: decision_contract.py must declare every new field, and a test's startswith('DECISION') matched the new label before the DECISION line. Against the pre-fix code, 22 of the new and amended tests fail: 13 prove only that the interface is new, 9 are behavioural. Seven negative controls, each failing as intended. Golden: one field added, provenance.decision_candles; no decision field moved, run_hash unmoved. The live run before the commit, read off his disk: every series decided one candle earlier than before, and the action was NO-TRADE (RISK TOO HIGH), as on 33 of the 37 earlier records, which is findings 4 and 6, not this change. code_hash b3c2308f… to 36d99565…. Tests 597 to 613 with pandas_ta.

#174	Finding 7: Indicator Values Beyond 5 Sigma Are No Longer Replaced; inf Still Becomes NaN September 27, 2026	DECISION — ADOPTED

“The mean and standard deviation spanned the whole series, so whether bar k survived depended on bars after k.”

Commit 57f0521, indicators/indicators.py only among engine files. clean_series set every value more than five standard deviations from its series' mean to NaN once the series had more than 10 values. The forward fill then put the previous bar's value there, and at the decision bar the value stayed NaN and was reported as an indicator failure. Nothing recorded either. Viktor ruled on 27 September, in a session that made no commit: remove the replacement, keep the step that turns inf into NaN, and put checking the candles for bad data on the list for after the audit. The repository does not record whether he wrote his position first. The evidence, sandbox and synthetic: a fresh SuperTrend flip of 16 bars or fewer in a 450-bar frame was erased, and so were an ATR spike of 2.19 times its median at the decision bar and a 10% move after a flat market. The rule reached EMA_20 and EMA_50, RSI, ADX, the SuperTrend and ATR, on the primary paths and the fallbacks. Measured before building: instrumented at add8540, the rule fired on no fixture value, which is why no test had noticed it and why the golden snapshot does not move. Run over all 30 lineage archives on Viktor's disk, 90 frames, it would have fired on no value in any, with a positive control that registers. So on the live runs recorded so far the change moves no value and no decision. Accepted with it, per DECISIONS: a bad price print that is internally consistent now reaches every indicator unaltered. Of nine new tests, 8 fail against the old code, all behavioural; the ninth pins the inf step the ruling keeps. A wrong turn caught before delivery: the first planted-value test asserted only that a value was present, which the old code passed, because an erased EMA fell through to a fallback that recomputed an ordinary number. Found and dealt with outside the commit: an untracked PDF on Viktor's disk in a tracked folder, which he moved out. code_hash 36d99565… to 84f5c4ac…; golden and run_hash unmoved. Tests 613 to 622.

#175	Finding 4: the Plan Is Measured From the Decision Close; the EMA Band Is Not an Entry September 27, 2026	DECISION — ADOPTED

“The panel printed an entry zone, and no entry price.”

Commit b0efe23, a panel change: core/panel_render.py and utils/plotting.py move code_hash, and two other engine files change in comments only. The panel printed an ENTRY ZONE, the band between EMA_20 and EMA_50, while the stop, T1–T3 and every R:R are measured from the decision candle's close, and no rule requires price to be in the band to trade. Ruled by Viktor on 27 September, not in the usual order: he asked for Claude's suggestions first. Of three options (the decision close made explicit; the band as a limit order with an expiry; the live price), Claude recommended the first and advised against the third, which would undo finding 16 and could not be reproduced under Item 2. Viktor chose the first: the decision close is the plan's single entry, printed and recorded; the band is labelled for what it is; a backtest fills at the next candle's open and records the gap; the band as a limit order goes on the list for after the audit. Accepted unmeasured: entering at the band might reduce the distance refusals behind most NO-TRADEs. Claude offered to measure it from the log, and the ruling was made without it. What the fill rule weakens, stated in DECISIONS and pointed to from Goal B's methodology: the printed R:R holds only at the decision close, and a next-open fill is optimistic about timing. The panel now reads EMA BAND and BAND DISTANCE, the band's lengths come from config, and the DECISION CLOSE line names itself the plan's entry. Deliberately unchanged: entry_status's values, which decision_model reads, so the panel still prints “STATUS : ACTIVE ENTRY ZONE”, and the data keys recorded in the log. Against the old code, 5 of the 7 new or changed tests fail, all behavioural; the other 2 pin the chosen behaviour, not a fix. Golden unmoved. The live run before the commit showed the new labels and carried the new code_hash. code_hash 84f5c4ac… to b5e6af9e…. Tests 622 to 627.

#176	Finding 6: the Stop Comes From ATR Alone; the HVN No Longer Pulls It September 27, 2026	DECISION — ADOPTED

“The HVN could only widen the stop, with no distance limit.”

Commit 76c8cde: models/risk_model.py and core/engine_core.py move code_hash, and five other engine files change in comments only. engine_core passed the HVN, the single highest-volume bin of the 450-candle frame (about 75 days on 4h), to calculate_stop_targets, which set a long's stop at the lower of it and the ATR stop and a short's at the higher. A trend that had moved away from its point of control got its stop there, and past 8% the risk check refused it. Ruled on 27 September by agreeing to Claude's suggestion, not by writing a position first, and recorded that way at Viktor's instruction: the stop comes from ATR alone and the HVN stays as information only; the 8% and 15% limits are unchanged; a check on HVN proximity goes on the list for after the audit; a swing-structure anchor was neither chosen nor measured. The evidence, corrected while building: the ruling's figures, 29 of 39 runs passing and 7 refused, summed to 36. Recomputed from the same log, 32 pass and 7 are refused. That is Claude's recomputation, not engine output; it reproduced the logged stop exactly on the runs the HVN had not taken over. Passing structural_level now raises TypeError. Against the old code 5 of the 6 new tests fail: 4 behavioural (on the pinned engine path the stop was the HVN, 0.6422, not the ATR stop, 0.7878) and 1 signature error. Claude first wrote “3 of 6” from reasoning; running the tests corrected it. Golden re-baselined, 11 fields exactly as predicted, among them the action, NO-TRADE (RISK TOO HIGH) to CONSERVATIVE LONG. The live run before the commit: the logged stop equalled the ATR formula recomputed from the record exactly, where the HVN would have put it at 0.54608. Volatility was EXTREME, so the run was still refused, the ruling working as specified. Found, not fixed: lineage.risk_inputs omits trend_health, which sets the stop's distance. code_hash b5e6af9e… to ed2710b6…. Tests 627 to 631.

#177 Finding 5: a NEUTRAL Bias Prints No Plan — Panel Only

September 27, 2026

DECISION — ADOPTED

“A score of exactly 0 printed a long.”

Commit 8b9ac1e, core/panel_render.py and models/exit_model.py. The plan's direction is the sign of bias_score (since work order C, Entry #152), so a NEUTRAL bias, inside ±20, which prints no direction box, still printed a long- or short-shaped stop, three targets and their R:R. Ruled on 27 September by agreeing to Claude's suggestion, not by writing a position first: panel only. Under a NEUTRAL bias the panel prints one PLAN line saying there is no plan because the bias is NEUTRAL, with the score and the threshold, and no stop, targets or R:R. The engine still computes and logs the plan, because the risk check needs a stop. The ruling asked for two checks before the diff. Where the plan prints: the DECISION CLOSE line's entry claim, STOP LOSS, the three TARGET lines, and the Exit Watch note naming Target 1. That NEUTRAL never reaches LONG or SHORT: from the code, two independent guards hold; from Viktor's log, 2 NEUTRAL runs of 41, one WAIT and one refused. Claude's readings, recorded for Viktor to correct: the DECISION CLOSE line stays but drops its entry claim; the Target 1 note is withheld; only NEUTRAL withholds the plan. A grid of 48,384 NEUTRAL inputs through the full evaluate returns only WAIT and NO-TRADE, with a negative control. Two wrong turns, recorded: the first ladder control failed to fail, because the minimum-bias guard still sent every such lean to WAIT; and a sed matched two lines, so one control run was redone with a counted replacement. Against the old panel 9 of the 14 new tests fail, all behavioural. Golden unmoved. What it weakens: a NEUTRAL record in the log carries a plan the panel did not show. The live run was BULLISH, so the NEUTRAL path did not run live. It was LONG, the first trade the log holds since a SHORT on 16 September. Found, not fixed: under NEUTRAL, entry and trade quality are still scored for the side the score's sign picks. code_hash ed2710b6… to c6361577…. Tests 631 to 645.

#178	A Push That Goes as Predicted Leaves Nothing Owed — the One-Step Regress Ended September 27, 2026	DECISION — ADOPTED

“Only a deviation becomes an owed item.”

Ruled by Viktor on 27 September, by agreeing to Claude's suggestion, not by writing his position first, and recorded in PHASE7_DECISIONS.md by commit 8b9ac1e with finding 5's code. Since the pre-push hook (Entry #135), each push's result was filed by the next commit, which left that commit's own push owed. Entry #166 called it a regress that cannot be closed, only kept one step long. Claude proposed on 26 September, as proposal (b), that a push whose result Viktor does not paste back be recorded once as going as predicted, and that only a deviation be filed as owed. Adopted. It rests on his standing practice of reporting only what differs from a prediction, and on Claude reading master and origin/master off his disk, and GitHub's tip, at the start of the next session, which is kept. What it weakens, per DECISIONS: a push that went wrong and was not reported now leaves no trace of being unconfirmed; before, the owed item kept asking. The same commit recorded a second confirmation made the same way: finding 6's “information only” means only the stop stopped reading the HVN. Its uses in the entry score, the reversal reading and the Exit Watch stay, and revisiting them belongs with the bias_score weighting after the audit (Entry #168). The same commit filed the push of 76c8cde (the refs on his disk and GitHub's tip all 76c8cde, SUMMARY: clean); the first push recorded under the new rule was 8b9ac1e's, at d76ddfd. One deviation inside that command sequence was caught by the git status prediction before the commit: models/signal_router.py was first left unstaged. This document's own floor (Entry #144) is not touched: a regeneration still cannot cover itself.

#179	Finding 18: the Bias “State Machine” Had No Memory; the Label Is Now a Function of the Score September 27, 2026	DECISION — ADOPTED

“CONFIRMED only ever meant a score beyond 30.”

Commit d76ddfd: models/bias_engine.py and core/engine_core.py move code_hash. Finding 18 asked whether the bias state machine's persistence requirement should gate anything, since work order F removed CONFIRMED from the confirmation gate. Read at 8b9ac1e, before Viktor wrote anything: there was no persistence requirement. transition() never read its previous state, and engine_core built a fresh instance for every run. So CONFIRMED only ever meant a score beyond 30, and the ladder already needs 30 to act. Nothing that decides reads the label. From the record: all 42 records in Viktor's log carry the label the current score gives. Ruled on 27 September by agreeing to Claude's suggestion, not by writing a position first. Of four options (label only; real persistence; the same labels from a plain function; a CONFIRMED gate), the third: behaviour exactly as it was, no new trading rule. “A side must hold for N closed candles” goes on the list for after the audit, computed from the run's own closed candles rather than a state file, and renaming CONFIRMED waits for after the audit too. The class became bias_label(raw_bias, bias_score). Its NEUTRAL band now reads RAW_BIAS_THRESHOLD instead of a literal 20. The 30 stays a literal, deliberately: naming it would put it among the fingerprinted constants and move run_hash and the golden snapshot for a number that decides nothing. Equivalence kept before the diff: the old class and the new function agree on 126,120 inputs, with fresh and reused instances. Six negative controls. The live run before the commit fell inside the same 4h candle as the previous run on the old code, so it was an equivalence run on live data: input hashes and run_hash identical, every decision field identical except the live prices. Two wrong turns, recorded: Claude opened by saying finding 5 was not yet built, from its own notes, and reading the refs showed it had landed minutes earlier; and Claude said a test used the class, when it only named it in a docstring. Recorded, not changed: the exactly-30.0 edge between the label (above 30) and the ladder (30 or more). With this commit all six of Viktor's findings are done. Golden unmoved. code_hash c6361577… to f4b23f94…. Tests 645 to 652.

#180	Work Order G: Macro No Longer Decides the CONSERVATIVE Tier — and the Closed List Is Done September 27, 2026	FINDINGS — FIXED

“The two upper tiers never had the clause; only the tier with the weakest case did.”

Commit 7d0024e, models/decision_model.py and models/signal_router.py. Finding 17, Claude's: both CONSERVATIVE branches of the ladder required macro_bias to agree with the side, but macro is already a 10% factor inside bias_score, so the clause counted it twice: the double count Viktor removed from direction on 2 September and from the entry signal at F (Entry #155). It was left out of F so that each change to which trades are taken lands alone. Claude's calls under the delegation, each reversible by Viktor, whose design he confirmed before building: the clause removed from both branches; macro_bias removed from DecisionModel.evaluate() and the ladder, since an unread parameter is the item-14 shape (the router still records it); and the later arguments made keyword-only, so an old positional call fails instead of putting macro into btc_context. Viktor's one ruling inside G, by agreeing to Claude's suggestion: the CONSERVATIVE sentence blamed entry quality even when trend strength had fallen short, a false claim in the line read first, and it is fixed inside G rather than listed for later. It now names what fell short. Before any diff Claude replayed Viktor's 43 logged runs through the router on the old and new code: the old reproduced 42, and G changes exactly 2 (15 September, bearish bias with bullish macro, WAIT becoming NO-TRADE (SIGNAL UNCONFIRMED)). A wrong turn: the first replay passed the degradation block in the wrong shape and reproduced 40; checking against the logged actions caught it before any figure was reported. Macro still reaches the decision in three other places, left for after the audit with the bias_score weighting. What it weakens, per DECISIONS: the CONSERVATIVE tier now admits setups whose higher timeframe disagrees or is flat. Against the old code 9 of the 11 new tests fail, all assertions. Golden: 2 fields, the sentence, as predicted. The live run was an equivalence run inside the same candle, LONG, and did not exercise the CONSERVATIVE tier. With G landed, the closed list of Entry #165 is done, and the independent audit is next. code_hash f4b23f94… to c6a44d4d…. Tests 652 to 663.

#181 What “Finished” Means: T1 Before the Stop, 100 Closed Trades, Spot Only

DECISION — ADOPTED

September 29, 2026

“A reading changes nothing.”

Commit 22cb39b, documentation only; ruled on 28 and 29 September. Until then the project had no claim a test could check, and no finish line except Goal B's verdict on a backtest. Before choosing any work, Viktor asked which questions the project should ask itself. Claude listed sixteen, he answered the four that were his, and the discussion became this ruling; the other twelve are open items. By Viktor's position first: a LONG should reach T1 before its stop, with no time limit on the trade; T1 must net at least 3% after fees; spot only, no leverage, and on spot SHORT means sell to USDT and wait; no weekly target or cap, and the count is never raised by loosening a threshold; results are read at 3, 6, 9 and 12 months, with no tuning between readings; the portfolio is secondary. By agreeing to Claude's suggestion: finished means one engine version and at least 100 closed trades on data it was not tuned on, with at least 60% reaching T1 first, positive average profit after fees, and a better result than random longs with the same stop and T1. Under 100 trades is not enough evidence yet. Viktor first proposed 51%; Claude showed it leaves almost nothing per trade before costs and takes about 10,000 trades to tell from luck. Also: the win is decided at T1, with half sold there and the stop moved to the entry; trends only, with a range mode later; paper trading in ten equal slices. Reconciled with Goal B on 29 September: its verdict stays as pre-registered and the criteria are reported beside it; spot only and the exit rules are added to its methodology, each stating what it weakens. Read from the code: T1 equals the stop distance and a stop past 8% is refused, so T1 cannot pay more than 8%, and the 3% floor needs a new filter, after the audit. Claude's errors, recorded: it said a trade needed a time limit, and it read 10% as per year on the account, where Viktor meant per trade on the position. It changes no code, does not amend the Constitution, and does not lift “no backtesting before an independent re-audit”. The push of 7d0024e was filed once. code_hash unmoved at c6a44d4d….

#182	Fable 5.1 Works on the List for After the Audit, Not Before It September 29, 2026	DECISION — ADOPTED

“The early warning is given up.”

Commit 68c6191, documentation only, the twentieth session's. Viktor has Claude credit for Fable 5.1, an Anthropic model, and asked whether it should review the engine and suggest improvements before the independent audit. Claude set out three options and recommended A, and Viktor chose it, by agreeing to Claude's suggestion, not by writing his position first. A: the audit first, then Fable works on the list for after the audit. Not taken: B, Fable reviews now, with everything it finds going onto that list; C, the closed list reopened. Claude's reasons, given before the ruling: the closed list already sends anything found before the audit to the later list, so Fable's suggestions could not be built before the audit without reopening it, and the auditor would then check code nobody had reviewed; suggestions and improvements are what the later list holds; and Fable, like Claude, is an Anthropic model, which Claude expects, untested, to share many of Claude's blind spots, making it a weak second opinion on Claude's work. Using it costs no audit candidate, since Anthropic is spent already. Cost was not the constraint: Fable's pricing was read from Anthropic's pricing page on 29 September; its context window was not confirmed. What goes with it: a ledger entry at the time of use, the Constitution and the Assistant Instruction read first, and its work reaching the repository as a patch, named in the commit message. What it weakens, in the words above: a serious defect Fable might have found before the audit is left for the auditor, or for Fable afterwards. Also in that session, and not ruled: Viktor asked which auditor, and whether to audit now with few clean labs left; Claude recommended Nemotron 3 Super, with Laguna S 2.1 second, and auditing now. The push of 22cb39b was filed once, as predicted. code_hash unmoved at c6a44d4d….

#183	Mistral Put Back Among the Clean Labs by Mistake, and Taken Off Again — the Gap Entry #48 Had Found September 29, 2026	FINDINGS — FIXED

“Five models ran through Aider; README named three.”

Commit 68c6191, documentation only. Viktor supplied OpenRouter's activity export (727 requests, 22 August to 14 September), and Claude reconciled it against the 22 September independence ledger, kept outside the repository, which said it was built from README.md and the AI-Attribution Statement. Five models ran through Aider during the build: Claude Sonnet 4 (194 requests), DeepSeek V3 (315), DeepSeek R1 (96), Mistral Nemo (71) and Claude 3 Haiku (4). README.md, since 29 August, and the Attribution Statement, written 14 September (Entry #120), named only the first three. Built from them, the 22 September ledger and the model roster put Mistral back among the clean labs. This document had it right on 2 September: Entry #48 found the same gap from an earlier export and named all five. Entry #146, on 20 September, repeated README's count of three; it is a dated record and stays as written. The repository had also rebuilt the ledger from a full export on 5 September, which the 22 September ledger did not use either. Claude repeated the roster's claim to Viktor in the session before the export was read. Viktor first held that Mistral was not spent (“I accept the risk”). Claude's critique: that reason does not hold; two stronger ones exist, since Mistral Nemo wrote nothing in the repository and its requests went to third-party hosts serving fixed weights; and accepting it would be a named exception to the lab-level rule. Viktor withdrew the position. Mistral stays spent, and nothing was ruled. Whether NVIDIA's co-development of Mistral NeMo counts against NVIDIA is open, and Viktor's call. Corrected forward: README.md's “How this was built” and the Attribution Statement's section 6 now name all five and say they named three until 29 September, and the statement's PDF was regenerated; a new HISTORY entry carries the account. The roster and the ledger were reissued outside the repository, dated 29 September, and the reissued ledger carries Grok's entry, owed since 26 September. Dated records that carry the old count or call Mistral clean, this document's included, are not edited.

#184 The Independent Audit Goes Ahead Now: Laguna S 2.1, the Full Package, Parts 1–6 Before Part 7

September 29, 2026

DECISION — ADOPTED

“Nothing is planted: the list for after the audit is the test.”

Commit e7a94d1, documentation only, the twenty-third session's. Two rulings made in chat in the twenty-second session, which landed no commit; this session could not read that chat and filed them from the record that session kept, the Phase 4 list of Viktor's roadmap document (outside the repository, revision 71), which paraphrases him, so the entries quote none of his words. Both by agreeing to Claude's suggestion. First: the audit now, preparation first and the send when the list is done, at his pace, by Laguna S 2.1 pinned to Poolside, the full package in one session; NVIDIA kept for a scoped round later. Claude had first recommended Nemotron 3 Super on the roster's 1M context; OpenRouter serves it at 262K, and the recommendation was withdrawn. The package was estimated at about 515K tokens with the commit messages cut to those since e65a0f7, by character count, not Laguna's tokenizer; no fallback was ruled for a package that does not fit. Second: two messages, Parts 1–6 first and Part 7 only after they are saved — the bias_score findings, the list for after the audit, any 15 September PDF point still live, and the commit messages since e65a0f7. Nothing is planted; the list for after the audit is the test, the method of 3 September. Point 1 of the 15 September PDF went back onto that list: models/decision_model.py still set win_rate = confidence / 100, read from the code at ec4e5fd.

#185 What Opens Backtesting: Four Items Rated Compliant, Then a Harness and a Known-Good Checkpoint

September 29, 2026

DECISION — ADOPTED

“The list for after the audit is not a condition.”

Commit e7a94d1, documentation only, filed with Entry #184's rulings from the same record. Viktor's position first, refined by Claude's critique, the merged version agreed. Backtesting may be built once Items 2, 3, 6 and 18 are rated Compliant, or fixed and confirmed by the same auditor; no Critical or Major finding is open and every other finding is ruled; every fix is confirmed by the same auditor; and Viktor declares it in DECISIONS. It may be run once Goal B's harness — the write guard, the look-ahead check, cut_checkpoint.py and the entry-point guard — is built and tested and a known-good checkpoint passes. If the auditor finds fewer than half the usable test bugs, it runs a second time first. Written by Claude at filing and marked so: what the ruling weakens — a backtest may run before the list's trade-changing rules are decided, and under Goal B's re-run clause the first out-of-sample verdict is then of the engine without them — and two points the record left open, which model runs the second time and what counts as usable and found. Entry #186 settles both.

#186 Six Questions Before the Audit's Send

DECISION — ADOPTED

September 29, 2026

“The instruction carries the ruled behaviour as plain requirements, not its reasons.”

Commit d5ced0c, documentation only. After e7a94d1 landed, Viktor asked whether every question needed before the audit had been answered. Claude: the large ones were ruled that day; six smaller ones were open, and none of the twelve preparation items had started. Viktor asked for Claude's suggestion on 1, 2, 4 and 5 and agreed to all of them; 3 and 6 are his own answers. Point 2 had been marked for his position first, and he chose to ask for the suggestion instead; it is recorded that way. 1: the commit the package is built from is tagged, and no commit touches engine code or tests until the report is triaged. 2: the instruction carries the ruled behaviour as plain requirements and the files changed since e65a0f7, and asks where a requirement itself looks wrong; DECISIONS and the reasoning stay out. What it weakens: the auditor may check against the intent instead of questioning it. 3: the second pass is Laguna S 2.1 again, as often as needed, each run a fresh session scored on its own. 4: a test bug is a defect, usable only if nothing shipped gives it away, found only with the same place and the same fault in Parts 1–6; under half found means a rerun; locked by a commit before the send; Claude scores, Viktor checks. Two were out already, read from the code: the 30.0 boundary and point 1 of the 15 September PDF. 5: the instruction does not say the ratings open backtesting. 6: no fallback decided in advance for a package that does not fit. Checked the same day from OpenRouter's model page, not its live API: Laguna S 2.1 at 1,048,576 tokens of context, Poolside its only provider.

#187 The Pre-Send Token Check: the Package Measured With the Auditor's Own Tokenizer

PROCESS — RECORDED

September 29, 2026

“A changed template is a refused check, not a wrong count.”

Commit 59b747a, docs/build/package_token_check.py and its 15 tests; tooling, no engine code. Item 2 of the preparation list, under Viktor's ruling of 26 September that the pre-send token check is audit preparation. For each turn it renders the request the way Laguna S 2.1's chat template writes it, counts the whole string with Laguna's own tokenizer, counts every earlier reply at the full output reserve — the worst case, not a forecast — and checks prompt plus reserve against the context. It sends nothing and writes nothing. The tokenizer files are Poolside's, downloaded by Viktor at revision e80da38 and kept outside the repository; the tool pins them by SHA-256 and refuses a folder where any is missing or different, since the tokenizer was changed after release and a sibling model ships files under the same names. Claude's design choices under the delegation of 21 September: not yet wired into the send (item 3, Entry #189); the template not run but rendered by the tool, checked equal to jinja2's rendering on four cases. Measured in the Linux sandbox on a trial package: 833,821 of 1,048,576 tokens at the worst case, 20.5% to spare; the content alone about 11% more than the estimate from four characters a token. Seven negative controls, each failing at least one test. Not covered by the suite, deliberately: loading the real tokenizer, whose files are outside the repository; covered instead by Viktor's run on Windows of the Constitution audit copy, its count stated in advance (23,045 tokens).

#188 The Test Bugs Classified: Four Usable, and Four Is Enough

September 29, 2026 — **DECISION — ADOPTED**

“Three of the four change no trade.”

Commit f256937, documentation only; items 7 and 6 of the preparation list, in that order, as the only items whose result could reopen a ruling. Item 7 traced points 3–5 of the 15 September PDF against the code at 59b747a: point 3, the BTC confidence boost, live but deciding nothing, and point 4, macro's ×0.90, live and never ruled, both moved onto the list for after the audit; point 5, what the validation score is out of, not live. Item 6 checked each entry of that list against what ships, taken from round 6's builder: five new trading rules and one weight's size not counted; five given away by a shipped line, each cited; four usable — the risk gate runs before the bias checks; lineage.risk_inputs leaves out trend_health; under a NEUTRAL bias the entry score is measured for a side nothing chose; the weak-validation WAIT decides nothing. Viktor checked the classification. Claude then put three options with no recommendation — four is enough; reopen “plant nothing”; leave it open until the send — and Viktor chose the first, with no written reasoning, recorded that way. None or one found means a rerun; two or more, a pass. What it weakens: three of the four change no trade, so a careful auditor may leave them out as cosmetic and a miss is weak evidence; the fourth is a record gap on Item 6, one of the items that open backtesting.

#189 The Send in Two Requests, the Token Check Inside It — and Rounds 5 and 6 Confirmed to Have Shown Part 7 Early

September 30, 2026 — **PROCESS — RECORDED**

“A run is not paid for whose Part 7 could not follow.”

Commit 12b483e, items 3, 4 and 5 of the preparation list: tooling, tests and docs, no engine code. send_audit_round.py was repointed to Laguna S 2.1, pinned to Poolside, fallbacks off, and split in two: --send sends the first message after measuring the whole conversation with the first reply at the full reserve; --send-part7 sends Part 7 only if the first reply finished, was served by Poolside, has content, and its files and the first message are committed and unchanged. Both refuse before any network call when the request does not fit, so the check needs no second command to remember. OpenRouter's prompt compression is pinned off. The package builder writes one folder per request, and the commit messages are cut to those after e65a0f7. Ruled by agreeing to Claude's suggestion: the second request carries the first reply's reasoning, not only its answer; item 5 lands here. The commit between the two requests is Claude's design, put to Viktor before the build. The correction owed since 29 September confirmed: rounds 5 and 6 each sent the Part 7 file in the same message as Parts 1–6, shown by rebuilding each request as it was sent and matching its recorded payload hash; whether that changes those rounds' weight is Viktor's. Found and fixed on the way: a stale largest-reply figure (36,085, now 41,861, round 4's) and a docstring counting three failures and listing four. Twenty-two new tests; seventeen negative controls, each failing on a behavioural assertion; against the old code the new file is one collection error and the two added tests fail on names that do not exist yet, which proves only that the code is new.

#190	Rev 8 of the Audit Instruction, Approved — It Says Nothing About This Round's Test September 30, 2026	DECISION — ADOPTED

“Rev 7 told the auditor twice that nothing was held back to test it.”

Commit 01f2892, documentation and a file name in two docs/build scripts; item 9 of the preparation list, approved by Viktor (“I read it, sounds great. I approve.”). Rev 7 edited, not rewritten. Its counted disclosure was re-run on a trial round-7 package: fifteen rules carry a stated verdict, severity or attributed finding where rev 7 counted eight, and seven AI parties are named where it counted five. A statement that the audit narrative was moving out of code comments described a plan, not a fact, and was corrected. Ruled by agreeing to Claude's suggestion, Viktor having leaned the same way: for round 7 the sentence that nothing was held back is untrue, since the list for after the audit is the test; option B — rev 8 says nothing about this round's test, the old sentences are narrowed to the rounds they describe, and the Part 7 document says it all when it arrives. Also approved: the requirement lines cover the rulings since e65a0f7 plus three older ones the code carries. Claude's, under the delegation: the hand-written Part 7 file renamed part7_material_PART7_ONLY.md, since its name, once committed, appears in the first message's history and “findings” would have told the auditor what rev 8 describes neutrally. Checked: no requirement line names any of the four test bugs, and rev 8 does not say the ratings open backtesting. Found and not changed: “ROUND 6” in seven test comments means round 5's sixth requested run — put on the list for after the audit, not usable as a test bug, and not edited in the code the auditor reads.

#191	The Part 7 Document, Approved — With a “Found Means” Sentence for Each Test Bug September 30, 2026	DECISION — ADOPTED

“Every entry was known, is unfixed, and was held back until Part 7 on purpose.”

Commit 8b6fd00, documentation only; item 13 of the preparation list, drafted by Claude and approved by Viktor (“I approve if you are satisfied also.”). docs/audit_package/part7_material_PART7_ONLY.md, the document the second message carries beside the commit messages: why the list for after the audit exists and why nothing was planted; the scoring rule in full; every finding on the list, classified — S1–S4 scored, N1–N5 named by a shipped line, O1–O3 not scored for other reasons, five new trading rules; the four bias_score findings; all six questions of the 15 September PDF. Claude's under point 4 of Entry #186, put to Viktor as the part to check most closely: a “found means” sentence for each of the four, narrowing the 29 September wording, and two entries added on 30 September classified as not scored. Four usable stays four. Corrected after the approval, before building, and shown to Viktor before he applied: the opening said the first message pointed to none of the entries, and shipped comments name some; Section 1 said every entry was found 26–30 September, and three come from the 15 September PDF. Found writing it: models/entry_model.py compares macro_bias with trade_direction, BULLISH/BEARISH/NEUTRAL against LONG/SHORT, always true; the result is right for every value a caller passes. Not fixed: on the list for after the audit. Every file:line the document cites was checked by script against the code, 40 checks, with a shifted citation as the negative control.

#192 The Roster and the Ledgers Fused Into One Document; Two Draft Copies Caught Before a Commit

PROCESS — RECORDED

September 30, 2026

“Neither reached a commit.”

Commit e186423, documentation only, the commit round 7's package was built and tagged at (round7-sent-2026-09-30). It files what was owed before the send, so the tag leaves nothing owed across the freeze. Item 8 of the preparation list: the twenty-ninth session, which made no commit, reissued the model roster outside the repository with Nemotron 3 Super at the 262K OpenRouter serves, NVIDIA now for scoped rounds only, and then, at Viktor's request, fused the roster and the ledgers into one document in place of five, which are kept as they were; both SHA-256 hashes are in PHASE7_NEXT.md. Two deviations, both a draft copy left in “Claude outputs\” — the folder whose files become public on commit: at 01f2892 Viktor's git status showed the rev 8 draft as one untracked file more than predicted, and he deleted it before committing; at 8b6fd00 a copy of the approved Part 7 draft was there, and he moved it to the superseded folder on G: before pushing. The first was owed to 8b6fd00 and missed, because that commit said the push of 01f2892 had no deviation, true of the push; the deviation came earlier, at git status. The superseded Part 7 draft differs from the committed document in exactly the three places corrected after the approval (Entry #191). Nothing was ruled; this commit before the send was Claude's suggestion, agreed by Viktor.

#193 Round 7, Run 1: Laguna S 2.1 Rates All 44 Rules Compliant in 52 Seconds — 0 of 4

AUDIT ROUND 7, RUN 1 — RECORDED

September 30, 2026

“It names itself Poolside's Muse Spark 1.3.”

Commit 86f30f1, documentation only, saved the day of the send as item 11 requires. The first message went at 09:58 UTC on 30 September from the tag at e186423. Five files written on Viktor's machine, none by Claude: the build's manifest, 111 files hashed; the reply, its reasoning file — empty, 0 bytes — and its metadata, as the send wrote them; and OpenRouter's record of the request, which Viktor fetched with a read-only request after the send's own lookup ran out of tries. The reply: finished after 52 seconds, 3,626 completion tokens, 0 reasoning tokens. All 44 rules rated Compliant, no finding, the release gate called met; Item 14 is Compliant in Part 1 and not verifiable in Part 3. It names itself Muse Spark 1.3, round 6's reviewer, a Meta model. It names none of S1–S4: 0 of 4, Claude's scoring, Viktor's check owed. Under half, so by Entry #186 its ratings do not count. The request did not ask for reasoning; why none happened is not known. Served by Poolside, cost $0.0416. Part 7 was not sent. The token check, tested: it counted 455,359 prompt tokens for the first request, and Poolside counted 455,360. Checked before the send, in the sandbox and on Viktor's build: the 111 hashes equal the files at e186423; his build equals the sandbox's except build times and four path lines printing Windows backslashes; nothing that ships had changed since 12b483e. Nothing was ruled; what follows is Entry #194.

#194	What Follows Round 7's First Reply: Part 7 Withheld, a Second Run With Reasoning, Two Runs in All September 30, 2026	DECISION — ADOPTED

“The freeze gets its first exception.”

Commit 3e8191a: docs/build/send_audit_round.py, one test file and docs, no engine code. Viktor asked for Claude's suggestion on the three questions run 1 left open and answered “Agreed.” to all three and to a fourth Claude raised; points 1–4 are ruled by agreeing to Claude's suggestion, and point 5 is his choice between two options, the one Claude recommended. 1: Part 7 does not go to run 1. 2: run 2 asks for reasoning and changes nothing else, and the send is refused unless a probe shows reasoning; tests/test_send_audit_round.py is the one exception to the freeze, since it tests a script that is not the engine and is not in the package. 3: run 1 counts as a run. 4: two runs in all; if run 2 also finds fewer than two, the next step is another auditor, not a third run. 5: a run is any reply on record, including one cut off at the output ceiling; a send that fails before any reply is not. Viktor checked Claude's scoring of run 1, 0 of 4, re-scored from the committed report: two near misses both fail the rule. The script's design is Claude's under the delegation, put to Viktor before it was built: a probe with nothing from the package before every real send, a limit of two runs, the first message checked against the hash every earlier run recorded, run 1 closed, a longer lookup. Found on the way: --out-dir with --send had skipped the check that refuses to overwrite a paid run. What it weakens: tests/ on master now differs from the tag in one file. Nine tests; nineteen negative controls, each failing the test it targets.

#195	Round 7, Run 2: Reasoning Requested, Item 11 Found Partially Compliant — Still 0 of 4; Round 7 Ends September 30, 2026	AUDIT ROUND 7, RUN 2 — RECORDED

“Its reasoning says it would focus on the most critical parts of a package this size.”

Commit 1523e0c: the reply's four files, written by the send on Viktor's machine and none by Claude; DECISIONS' record that round 7 ended; run 2 closed in the send script; and the one test file the freeze allows. Nothing new was ruled: this applies point 4 of Entry #194. Sent by Viktor at 11:43 UTC from the same package, its printout matching the prediction before sending; the probe replied with 1,102 reasoning tokens. The reply: Poolside, finished after 45 seconds, 4,415 completion tokens, 2,166 of them reasoning; 455,328 of its 455,360 prompt tokens were served from the provider's cache of the identical first message, which run 1 had sent less than two hours earlier — the prompt, not run 1's reply. Cost $0.0049. 42 rules Compliant; Item 11 Partially compliant (Major), four of the six bias factors reading the direction of close; the release gate called met. It named itself “Poolside's Meta-Muse Spark 1.3”. Scored 0 of 4 by Claude and checked by Viktor (“Yes, checked”): none of S1–S4 is named, and its one near miss on S3 is in the reasoning file, which is not Parts 1–6. Its Item 11 finding is named by a shipped comment and is one of the bias_score findings; it is not a test bug. With two runs under half, round 7 ended: no Compliant rating from it counts, Part 7 went to neither run, and the next step is another auditor. Claude's, under the delegation: run 2 added to the closed runs, since the script cannot know a score and would otherwise send Part 7 to it.

#196 The Next Audit Delegated to Claude Outright; Round 8 to Xiaomi's MiMo-V2.6-Pro

DECISION — ADOPTED

October 4, 2026

“If you have to send 3 or 4 packages instead of one full audit so be it.”

Commit ff12e17, documentation only. Viktor's delegation of 4 October: the next independent audit is Claude's to run, every decision included (“you do all the work, take all decisions and see through that the Audit take place”). Accepting or rejecting the result, running the commands, and anything the delegation does not name stay his. Filed from Claude's notes of the session that received it, not from its chat, which this session could not read; DECISIONS says so. Claude's four decisions under it, each with what it weakens: round 8 goes to MiMo-V2.6-Pro, the full package in one session, if four checks pass, pinned to Xiaomi with three named providers after it and data collection never allowed to reach an endpoint, two runs; the package is round 7's tag with only rev 9 of the instruction and a corrected Part 7 document, since a build from master would ship round 7's scores — the commits after the tag record “0 of 4”; the freeze holds until round 8's report is triaged, with the token check's test file joining the exception; the fallback is scoped packages to Upstage, NVIDIA, Cohere and AI21 on Items 2, 3, 6 and 18 and the test bugs' code. The checks as far as they could be made: served context and output read through a summarising web tool, not raw JSON; no lineage disqualifier found, though Xiaomi does not publish its distillation teachers; no Xiaomi or MiMo row in Viktor's OpenRouter export. The own-tokenizer count was not made: the sandbox's proxy refuses huggingface.co, so Viktor downloads the pinned files.

#197 Round 8's Tooling: Round 7's Built Files Reused by Hash, the Provider Order Run by Hand

PROCESS — RECORDED

October 4, 2026

“The package rests on files git ignores, which exist only on Viktor's disk.”

Commit 4a4c6ce, tooling and tests only, no engine code; Claude's under the delegation (Entry #196), the design filed in DECISIONS with what each choice weakens. The token check gains a MiMo-V2.6-Pro entry, its tokenizer files pinned by SHA-256 and a renderer for its chat template; --model no longer defaults to round 7's model. send_audit_round.py is repointed to round 8: two runs, the highest of the four endpoints' prices, and --provider naming one of the four, pinned with fallbacks off, a 404 naming the next in the order; --send-part7 goes only to the provider that answered the first message. New, build_round8_package.py makes both messages' folders from round 7's built files, each checked against the hash both round 7 runs recorded, plus rev 9, and refuses the Part 7 document while it is still the tag's text; nothing is written until everything is checked. build_audit_package.py now refuses to rebuild round7/, which would overwrite the files round 8 is built from. A refinement recorded: “built from the tag” is done by reusing round 7's files, not by a new build at the tag, which would carry a new build time and the tag in its history file; so the auditor sees a manifest reading “Round: round7”, and rev 9 has to say why. The two test files the freeze allows grow from 17 to 21 and 30 to 43 tests.

#198	Rev 9 of the Audit Instruction: Rev 8 Was Short by Six Rules October 4, 2026	DECISION — ADOPTED

“Twenty-one rules carry a prior judgement, not fifteen.”

Commit 217ede6, documentation only; Claude's under the delegation, neither document shown to Viktor before landing. Rev 8 edited, not rewritten. Its counted disclosure was re-run on the package as built, from round 7's files staged off E: and each hashing to what both runs recorded: rev 8 had searched for a rule's number beside a verdict word, and six more rules carry a defect attributed with no verdict — Items 4 and 12, three Tier 2 principles and T3-5, all in rev 8's package. Rev 8's other counts were re-run and held, and Section 4a's list of changed files regenerated the same. New: why the manifest reads “Round: round7”; round 7 produced no report the project could use; a disclosure that no rating, finding or wording was carried over from it; the addressed lab described, not named, its knowledge cutoff stated with its one source — a sample system prompt on Xiaomi's model page, called an example, not a specification. The Part 7 document corrected in place where it named round 7, and one change beyond that letter: its threshold said “if run after run stays under half”, made untrue by the ruling of 30 September, and now says two runs in all; entries, classifications and the four “found means” sentences unchanged. Left out on purpose, and recorded: MiMo's tokenizer family and its team lead's past employer, since rev 9 does not name the lab and either would point the auditor at an identity.

#199	Round 8, Run 1: MiMo-V2.6-Pro Finds 0 of 4 — and Names Itself “Claude (Anthropic)” October 4, 2026	AUDIT ROUND 8, RUN 1 — RECORDED

“The prompt count matches MiMo's tokenizer exactly.”

Commit d27743f, the run's files and tooling; Claude's under the delegation, the decision filed in DECISIONS with what it weakens. Before the send, on Viktor's machine: master and GitHub's tip both 217ede6; the round 8 package's nine files equal in size and SHA-256 to the sandbox's build; no shipped .py file changed since 01f2892 except the two test files the freeze allows; OpenRouter's endpoints for the model read as raw JSON for the first time, by Viktor's own request, the sandbox's proxy refusing openrouter.ai, and kept outside the repository. Tag round8-sent-2026-10-04 at 217ede6, then the send: Xiaomi's endpoint; 430,820 prompt tokens, the token check's count to the token; 47,832 completion tokens, 40,434 of them reasoning; 623 seconds; $0.229. Scored by Claude, Viktor to check: 0 of 4 — _determine_final_action is named only for other things, risk_inputs is not named, calculate_entry_quality only for the macro multiplier, the weak-validation branch not at all. Under half, so its ratings (37 Compliant, 5 Partially compliant, 2 Not verifiable) do not count; its five findings are received, not triaged. The reply names itself “Claude (Anthropic)”; OpenRouter's record names Xiaomi's endpoint and MiMo's model, and the prompt count matches MiMo's tokenizer; recorded for Viktor to weigh. Run 1 is closed to Part 7: before this, once its files were committed, --send-part7 would have found it as the one run waiting, shown in the sandbox.

#200 Round 8, Run 2: 2 of 4, a Pass — Its Ratings Count Toward Nothing Until Viktor Rules

October 4, 2026

AUDIT ROUND 8, RUN 2 — RECORDED

“Claude's first draft said 20 / 21 / 3, written before counting.”

Commit a0276dd, documentation and the run's files, no .py change. The same first message, by the send's own hash check; Xiaomi's endpoint; 430,820 prompt tokens, 430,720 from the provider's cache, which holds the prompt, not run 1's reply; 80,023 completion tokens, 68,464 reasoning; $0.071. Claude's score, Viktor to check: S2 found — F1 names risk_inputs and says trend_factor scales the stop while risk_inputs leaves it out; S4 found — F10 names the weak-validation branch of _determine_final_action and says it can change only the reason string, not the action; S1 and S3 not found. Two of four is a pass. Its ratings, counted from its table: 29 Compliant, 13 Partially compliant, 2 Not verifiable, none Non-compliant. Claude's first draft of the DECISIONS entry gave 20 / 21 / 3, written before counting, and was corrected before delivery. Five Major and five Minor findings, no Critical; the release gate met on its own findings, with a caveat on Items 3 and 6. Run 2 again names itself “Claude (Anthropic)”. Also corrected: two stale self-references carried in PHASE7_NEXT.md from earlier versions. Not changed: the send script's largest-reply figure, 41,861, which run 2's 80,023 exceeds while still under the 131,072 cap; left so the script stays as it was for both runs.

#201 Round 8's Part 7 Reply: Its Own Reading Agrees, 2 of 4 — With Three Caveats on the Measure

October 4, 2026

AUDIT ROUND 8, PART 7 — RECORDED

“The scorer is the fixer.”

Commit 040a529, documentation and the reply's files, no .py change. --send-part7 picked run 2, run 1 being closed, and went to Xiaomi's endpoint, which had answered the first message. 658,197 prompt tokens by the provider's count — exactly the token check's count with the first reply's reasoning, so the reasoning reached the model, although MiMo's template reads a differently named field; the question the send script left open is answered. 63,761 completion tokens, 58,054 reasoning; $0.156. What the reply says, summarised; the reply is the record: its own reading of its Parts 1–6 is 2 of 4, S2 and S4, the same as Claude's score. Nothing retracted; two findings refined; its fixture-error count corrected to 32. Six gaps between stated intent and shipped code, among them that work order G did not take macro off the LONG/CONSERVATIVE boundary, which the entry score's macro multiplier still moves. Three caveats on the measure: all four scored entries are named in the commit messages it saw in Part 7; the two-run threshold was ruled after a score existed; and the scorer is the fixer. It would classify O3 as a defect. What became Viktor's: check the two scores, accept or reject the result, and triage if accepted.

#202	Round 8 Accepted Narrowly: Findings to Triage, Ratings Not Counted, Backtesting Closed October 4, 2026	DECISION — ADOPTED

“I approve the narrow version, backtesting remain closed.”

Commit 9e42d30, documentation only. Viktor's ruling, by agreeing to Claude's suggestion, not by writing his position first. Asked to check the two scores and write his position first, he said he would not read the report and would approve it regardless, and asked whether Claude approved. Claude recommended accepting the findings and not letting the ratings count; said its approval could not stand in for his, since it chose the test bugs, scored both runs and wrote many of the graded fixes — the auditor's third caveat; reminded him of his own rule to pause rather than decide without the work; and offered the narrow version or a pause. He chose the narrow version. He did not check the scores; they rest on Claude's scoring and the auditor's own reading, which agrees. Run 2's findings F1–F10 and the points of its Part 7 reply go to triage; its ratings go with them as the auditor's view, and none counts toward opening backtesting. Before recommending, Claude checked four points by reading the code: F5, the band distance measured from the band's midpoint; F6, clean_series keeping an interpolate branch that reads later rows, called from nothing outside tests/, so latent; a golden-snapshot test that returns instead of skipping without its file; and a watched error count given as 29 against 32. The rest were not checked. What it weakens: the adjudication was made without reading the report or checking the scores, which rest on the fixer alone. The gate to backtesting is stricter, not weaker. Not ruled: what would make a rating count.

#203	Round 8 Triaged: the Stop From ATR Alone, No Trend No Trade, Bitcoin Reference Only — the Freeze Lifts October 5, 2026	DECISION — ADOPTED

“Conviction decided the 8% risk verdict.”

Commit eb86846, documentation only. Viktor triaged round 8 on the night of 4 to 5 October, beside a review Claude wrote that night at his request, kept outside the repository; that session landed no commit, and this one filed the rulings from its resume note, not its chat, which it could not read. Every point by agreeing to Claude's suggestion. 1: the stop is ATR × 1.2 × the volatility factor; trend_factor and bias_factor go — a break of Item 14, since conviction decided the 8% risk verdict; confirmed by Viktor in this session's chat, Claude having read his “OK” on point 2 as covering point 1. 2: the release gate closes until point 1's fix lands and round 8's auditor verifies it; it reached live decisions, so it is Critical. 3: the risk regime from volatility alone. 4: ADX under 20 blocks both sides, replacing the exhaustion flag in the confirmation gate. 5: the trend score's scale waits for after the audit. 6: test runs write to a log, archive and Exit Watch state of their own. 7: the BTC-adjusted confidence removed; Bitcoin is reference only. 8: the label and record fixes are made now. 9: F4 stays on the list; rev 9's two wrong sentences get a correction note. 10: the check of Claude's score, declined. 11: triage complete, the freeze lifts. Claude's, under the delegation of the order of work: point 4 before point 3, so no commit leaves a window in which a choppy market could be called AGGRESSIVE. Re-derived before filing from Viktor's live log, by arithmetic, not engine output: on the 37 live runs the median stop goes from 6.65% to 5.25%, and ADX under 20 would have changed no live action. Claude's wrong turns, recorded: a first draft said live stops narrow “by about a third” — computed, by 4% to 28%; and one said round 8 questioned none of rev 9's other requirement sentences, when it questioned two. Found before the commit: three untracked PDFs in “Claude outputs\”, the third time a file meant for G: had landed there.

#204

Fix 1: the Stop Is ATR × 1.2 × Volatility — Conviction No Longer Sets It

October 5, 2026

FINDINGS — FIXED

“The same market passed or failed the 8% line by how convinced the engine was.”

Commit c2b07fb, engine code, tests, the golden fixture and documents; round 8's triage, point 1, and round 8's F1 with it. The stop multiplier had been ATR_STOP_MULT × (1 + trend_health / 200) × (1 – |bias_score| / 300) × the volatility factor, so trend health widened the stop by up to ×1.5 and a strong bias score narrowed it to ×0.667; the risk check refuses a stop wider than 8% of price, so the same price, ATR and volatility passed or failed it by conviction. The three LONGs of 27 September passed at 7.15%, 9.53% without the bias factor. Viktor's ruling of 5 October, by agreeing to Claude's suggestion: a break of Item 14, Critical, the release gate closed until round 8's auditor verifies the fix (Entry #203). The order of the fixes is Claude's. Now stop_mult = ATR_STOP_MULT × the volatility factor; trend_factor, bias_factor and their two divisors are removed; calculate_stop_targets refuses trend_health and takes keyword-only parameters, Claude's call, so a call in the old order cannot bind trend health to the price. Six new tests, five changed. Against eb86846's code 18 fail: 2 behavioural, 3 on the signature or call, 13 only because the old signature needs trend_health. Seven negative controls. A wrong turn caught by reasoning: the verdict test first used one market, on which the bias factor alone would have left it green; it became a 25-market sweep. Found and fixed on the way: the non-finite-ATR test asserted any Exception and was vacuous; it now asserts the guard's ValueError. The replay of the live log: 3 of 37 live runs change, NO-TRADE (RISK TOO HIGH) to WAIT; median stop 6.65% to 5.25%; the triage had put all three WAITs down to MIN_ACTION_BIAS, and record 15's is its NEUTRAL bias. Golden snapshot: 11 leaves predicted, 11 moved. The live run on Windows: LONG, its stop and targets rebuilt exactly from lineage.risk_inputs. Found, not fixed: a long whose ATR × 1.2 × the volatility factor exceeds the price gets a stop below zero, refused downstream. Owed to the next commit: the tests ran after the live run, not before it (Entry #205).

#205	Fix 2: No Trend, No Trade — ADX Under 20 Blocks Both Sides in the Gate October 6, 2026	FINDINGS — FIXED

“In practice the exhaustion flag asked whether the last candle was wider than the one before.”

Commit 1b25f36, engine code, tests, the golden fixture and documents; round 8's triage, point 4. The confirmation gate blocked both sides on the exhaustion flag, which, trend health almost never being under 35, asked whether the last 4h candle was wider than the one before in a low-ADX market — on the review's count, not re-derived, it fired on 94 of 185 bars with ADX under 20 — while the engine is to trade trends only. Viktor's ruling of 5 October, by agreeing to Claude's suggestion: the flag leaves the gate; ADX under 20 blocks both sides, with the reason printed. Claude's proposal, Viktor's to reverse: an ADX with no valid reading — None, NaN, a bool, anything outside 0–100 — blocks both sides too, so a degraded run whose ladder chose a side is now NO-TRADE (SIGNAL UNCONFIRMED). MIN_TREND_ADX = 20.0 is new and fingerprinted; adx is keyword-only with no default; the value prints rounded down, so 19.96 never reads as “20.0, under the 20”. The flag is still computed and read elsewhere. Thirteen new tests; against c2b07fb's code the new file fails to import, and with the import shimmed 33 fail: 4 behavioural, 3 on the signature, 26 only because adx is required. Twelve negative controls. Golden snapshot: 6 leaves predicted, 6 moved. The replay: no recorded action changes. Wrong turns: decision_model first imported MIN_TREND_ADX by name, a second copy of a fingerprinted constant, caught by the suite; the gate's comment stated the replay's result before it had run. Correction of c2b07fb's message: pytest's cache on Viktor's disk showed pytest had not run after c2b07fb's apply, so its counts were not confirmed by proceeding; fix 1's tests first ran on Windows in this session. From now on Claude reads that cache before giving the commit steps. Found, not fixed: calculate_dynamic_bias's trend_exhaustion parameter is dead, on the list for after the audit.

#206	Fix 3: the Risk Regime From Volatility Alone — ADX Leaves the Risk Check October 6, 2026	FINDINGS — FIXED

“A high ADX raised conviction and lowered the assessed risk together.”

Commit fcb7823, engine code, tests, the golden fixture and documents; round 8's triage, point 3 (round 8's F2, the review's B8). classify_risk_regime read raw ADX twice: under 20 the regime was HIGH VOLATILITY RISK whatever the volatility, and LOW VOLATILITY reached LOW RISK only at ADX 25 or more. Trend health spends up to 40 of its 100 points on the same raw ADX and is 0.30 of bias_score, so a high ADX raised conviction and lowered the risk together — the coupling Item 14 forbids, one level below the one removed on 11 September; the constants block's update of 19 September had called the read independent. The chop test also named the wrong cause: LOW VOLATILITY printed beside HIGH VOLATILITY RISK. Viktor's ruling of 5 October, by agreeing to Claude's suggestion: the regime comes from volatility alone, both ADX tests go. Claude's proposals, Viktor's to reverse: LOW VOLATILITY is LOW RISK; a volatility state the regime does not know is UNKNOWN RISK, not NORMAL RISK, which keeps AGGRESSIVE out — unreachable on the engine's path, read from the code. Claude's engineering call: validate_risk_parameters takes three required parameters and no **kwargs, so a stale or misspelt keyword is refused before the body runs. lineage.risk_inputs no longer records adx. The labels change, and no action can: since Entry #205 a run with ADX under 20 takes no trade. Eleven new tests, five removed. Against 1b25f36's code 15 fail: 5 behavioural, 5 on the signature or call, 5 only by TypeError. Twelve negative controls. Golden snapshot: 7 leaves predicted, 7 moved. The replay: all 48 verdicts, reasons and regimes the same. On Windows pytest's cache shows step 3 ran before the live run. Wrong turns: the first search for the risk check's callers looked for the keyword and missed two positional calls; a probe wrote a decision log into the pristine copy, cleaned before anything was built. Found, not fixed, on the list for after the audit: two docstrings in risk_model.py miscount return paths; calculate_stop_targets still defaults volatility_state to “NORMAL”.

#207	B6: Runs on Pinned Data Keep Their Own Log, Archive and Exit Watch State October 6, 2026	FINDINGS — FIXED

“The first live run that day compared live AERO with a synthetic series.”

Commit 9e507b1, engine code, tests, the golden fixture, the audit-package builder and documents; round 8's triage, point 6, the first patch of the label and record batch. Every run wrote its decision record, its archive and its Exit Watch state into config.LOG_DIR whatever its candles were, so the live log on Viktor's machine holds nine runs on the synthetic fixture and one hand-built record, and the first live run of 6 September read its prior state from a pinned run 27 seconds earlier. Found by Claude's self-review of 5 October, on the real log; Exit Watch is advisory, so no action was affected. Viktor's ruling of 5 October, by agreeing to Claude's suggestion: pinned runs write to their own records, chosen by the code from the run's source; the pinned records already in the live log stay. Claude's proposals, Viktor's to reverse: logs/pinned/ holds them; a record goes in the live log only when its provenance says live in both places, so a record with none, or one that contradicts itself, fails safe to the pinned log, and no caller can override it; a broken PHASE7_PINNED_DATA counts as not live; each source prunes its own archives. records_dir is built from string operations so its answer is the same on Windows and Linux. The audit-package builder now redacts both spellings of the temp directory, which would otherwise have let the Windows username into the next package; found by reading every consumer of the printed path. Ten new tests, seven changed. Against fcb7823's code, with a shim giving it the new names, 19 fail, 5 of them behavioural. Golden snapshot: 3 leaves predicted, 3 moved, the paths; run_hash unmoved. The live run: record 51 in the live log, no logs/pinned/ on his disk. Wrong turns: a negative control passed, because every not-live shape had two defects (the 27 September lesson again), and the shapes were rewritten; the golden updater was run again after CRLF was restored; records_dir was first built with os.path.join, which would have failed only on Windows. Not separated, his only if he wants it: the chart, the process log and LiveSim.

#208	N5: Bitcoin Is Reference Only — the BTC-Adjusted Confidence Removed October 9, 2026	FINDINGS — FIXED

“A number computed and recorded that nobody reads is the shape this project keeps removing.”

Commit 6469a67, engine code, tests, the golden fixture and documents; round 8's triage, point 7, the second patch of the label and record batch. DecisionModel._compute_btc_adjusted made a second confidence figure, the run's confidence moved by up to 20 points by BTC's bias strength and the pair's correlation and by 15 down under market stress, printed as “BTC-ADJUSTED CONFIDENCE: x/100 (vs y/100 unadjusted)”. It ran after the action, confidence, trade quality and EV were set; nothing that decides read it, and nothing had tested whether it predicted anything (the 15 September PDF's point 3, round 8's N5). Viktor's ruling of 5 October, by agreeing to Claude's suggestion: Bitcoin is reference only, the line removed, the BTC section kept. Claude's proposal, the reading the ruling named, Viktor's to reverse: the computation and its record field go with the line, and with them the two constants, evaluate()'s btc_context and symbol parameters, the panel's validation-status line and the BTC bias score put in the block for it alone. The BTC block is now available whenever BTC was measured; before, any exception in the adjustment hid the whole section. Nine new tests, nine removed with their subject, each named where it stood. Against 9e507b1's code all 9 new tests fail, 7 on behavioural assertions. Golden snapshot: 8 leaves predicted, 8 moved. No decision moves. A wrong turn: the method's last line lay outside the range Claude first read, and the removal's own check stopped the first try. Found, not fixed: the Constitution's Item 7 note now describes a feature the engine no longer has; amending it needs a reviewer who is not Claude. README's paragraph on what is validated is rewritten in Viktor's voice, his to change. The same commit files Viktor's position of 6 October (Entry #209).

#209	All the Work and the List for After the Audit First, Then One Full Audit by Amazon — a Position, Not Ruled October 9, 2026	OBSERVATION — FILED FOR THE RECORD

“If the whole list is fixed first, nothing is left to score Amazon with.”

Commit 6469a67, filed in PHASE7_DECISIONS.md as a position, not a ruling, from the resume note of the session of 6 to 9 October, which landed no commit, checked against that session's chat. Viktor's position, 6 October: do all the work and the list for after the audit, then audit with Amazon through Bedrock, whose billing he will sort; “we have the best one left”. Claude's critique: it agrees with the order — Amazon is the last lab clean for a full round, best spent once on the engine as it will be backtested — but the four defects that score an auditor come from that list, so fixing all of it first leaves the under-half rerun blind; the options are to hold some back, reopen “plant nothing”, or audit without the check. Amazon stays clean until the send; whether Nova 2 Pro is reachable on his account and fits the package was not checked since 22 September. Who decides: the plan and the scoring question are Viktor's; running the audit is Claude's under the delegation of 4 October, as Claude reads it, and Claude's reading that “all the work” includes round 8's fix verification is not confirmed. Viktor's thought of 7 October, not ruled: tune the engine after the paper-trading year to match what happened. Claude's critique: the trades used for tuning cannot judge the tuned engine, which is a new version needing data it has not seen; tune few settings and name them before the year starts. Open, Viktor's, his position first on four of them.

#210	F3: a Frame With No Readable Time Axis Is Rejected by the Data Validator October 10, 2026	FINDINGS — FIXED

“The validator's gap was real; the refusal was an accident one frame further on.”

Commit ae37795, one engine file, a new test file and documents; round 8's triage, point 8 (round 8's F3), the third patch of the label and record batch. validate_ohlcv accepted a frame with neither a DatetimeIndex nor a timestamp column, and one whose timestamp column it could not read; every check on time sat behind that early return, so a pinned file with the column removed passed as clean. No caller in the engine hands over an untimed frame, read from the code. Viktor's ruling of 5 October, by agreeing to Claude's suggestion: the label and record fixes now; he chose to start F3. Round 8's reproduction was run, and its outcome was wrong: an untimed copy of the AEROUSDT file does not give a complete panel with row number 419 as the decision candle, but crashes later in _decision_candle with a generic error, no record, no panel. Now rejected: no time axis (F3 itself), and — Claude's proposals, Viktor's to reverse — an unreadable timestamp column, and a missing timestamp, until now reported as out of order at index 0, the wrong defect at the wrong place. Not changed, deliberately: the pinned path checks no staleness, by the ruling of 30 August, so F3's “five years stale analysed as current” describes that ruling, not this gap. Six tests; against 6469a67's code 5 fail, all behavioural; three negative controls. Golden snapshot not re-baselined, run_hash unmoved. On Windows pytest's cache was rewritten after the apply and lists the new tests; the live run, record 53, carries the new code_hash.

#211	F5: the Band Distance Is Labelled From the Band's Midpoint October 11, 2026	FINDINGS — FIXED

“A positive distance “away from the band” printed under ACTIVE ENTRY ZONE.”

Commit 9d00d99, the panel line and comments in two files, tests and documents; round 8's F5, the fourth patch of the batch. The panel's BAND DISTANCE line read “x% away from the band” while the number is the distance from the band's midpoint as a percentage of the close, so with the close inside the band it printed a positive distance away from it under STATUS : ACTIVE ENTRY ZONE; Claude's review of 5 October reproduced it on a real run (1.40%). Viktor's ruling of 5 October, the batch now; the repair is Claude's, Viktor's to reverse: the line reads “x% from the band's midpoint”, the number, record key and score unchanged. Not chosen: an edge distance, which would change what the record's key means from one record to the next. Three new tests; an existing guard that “0.00% away from the band” is not printed could, after the relabel, never fail, and now looks for “0.00%”. Against ae37795's code 4 fail, all on the label; three negative controls. Not changed, deliberately: ACTIVE ENTRY ZONE reaches half a band-width outside the band and scores the zone's full 30 points, which can change a tier — a decision-path question, put on the list for after the audit. code_hash moves through the panel only; run_hash and the golden snapshot unmoved.

#212	F6: clean_series Accepts Only Methods That Read No Later Row October 11, 2026	FINDINGS — FIXED

“One keyword argument from a look-ahead leak.”

Commit 27d3ed5, one function, a new test file and documents; round 8's F6, the fifth patch of the batch. clean_series accepted four methods; all eight engine calls pass forward_fill. interpolate filled each interior gap from the next valid row, a bar in its future (Item 2), and carried the last value across the decision bar that forward_fill refuses to fill; drop was unused and shortened the series; an unknown name, or fill_value without a value, returned the series silently. Latent, not live, and live the day a backtest walks the decision bar backwards. Viktor's ruling of 5 October; Claude's proposals, Viktor's to reverse: interpolate and drop removed, and an unknown method or a fill_value without a number raises ValueError. Checked, and right: round 8's statement that test_no_lookahead.py cannot see the branch. Four tests, the accepted methods read from the function's own refusal so the list cannot drift; the check run on the removed interpolation catches it (the control that it can fail). Against 9d00d99's code 3 fail, all behavioural. Not changed, deliberately: whether a gap at the decision bar is refused depends on whether a later bar exists — no value comes from a later row, but a backtest computing indicators once over its whole history would decide on a carried value where the live engine reported a failure; on the list for after the audit, Goal B's harness being where it matters. code_hash moves through indicators.py only; run_hash and the golden snapshot unmoved.

#213	F8: One Name, One Meaning — VWMA Names Only the Engine's VWMA October 11, 2026	FINDINGS — FIXED

“The two can point opposite ways in the same run.”

Commit 9051ffd, two engine files, a new test file and documents; round 8's F8 (Item 10), the sixth patch of the batch. “VWMA” named two calculations: the engine's 20-candle VWMA, scored by entry_model, and structure.py's vwma_recent, vwma_prev and vwma_slope, two 5-candle volume-weighted means of the close feeding the volume sentiment and through it the bias score. And the panel printed momentum_mode twice, after TREND's slash and as the whole MOMENTUM line. Viktor's ruling of 5 October; the new names, and removing the MOMENTUM line rather than the word after the slash, are Claude's proposals, Viktor's to reverse. The three locals become vw_mean_recent, vw_mean_prev and vw_mean_slope; the MOMENTUM line goes. No name containing “vwma” in structure.py was printed or recorded. Five tests on the parse tree of every file code_hash reads: a function assigning a “vwma” name touches the engine's VWMA, with controls that the walk reaches the scoring and catches the old function. Against 27d3ed5's code 3 fail, all behavioural. Equivalence: the new function's parse tree with the names mapped back equals the old, and the two give the same label on 1,434 windows of the four fixture files. code_hash moves through the two files; run_hash unmoved, the live run reading the same closed candle with the same run_hash. Not changed, deliberately: the word after TREND's slash carries no name of its own; nobody has ruled on one.

#214	F9: Four Claims the Code Did Not Support, and the Labels Ruled in With Them October 12, 2026	FINDINGS — FIXED

“It persisted nothing: the only caller builds a fresh engine per call.”

Commit e5ac20a, six engine files (one comments only), a new test file and documents; round 8's F9 (Item 8), the seventh patch of the batch. Round 8 named four statements the code beside them did not support, all four still present on 11 October: StructureEngine's “regime persistence to reduce whipsaws” persisted nothing, since calculate_structure builds a fresh engine per call (the shape of finding 18); trend health's failure labels were NEUTRAL, a reading no classifier produces; UNKNOWN VOLUME was reported as “Volume sentiment is neutral.”; and the ADX failure note said “25 of its 100 points”, where the component is worth up to 40 — round 5 had reported the 25 on 12 September and it was not fixed then. Viktor's ruling of 12 October, by agreeing to Claude's suggestion: F9(1) is option C, a plain function with the same output on the engine's path; REGIME_HYSTERESIS_THRESHOLD keeps its name, a rename going on the list for after the audit since it would move run_hash for a word; trend_direction's failure label and a frame under 15 rows in scope; the TREND line's fallbacks in scope if they cannot fire. Wording Claude's, Viktor's to reverse. Failure labels become UNAVAILABLE; a short frame reads UNKNOWN STRUCTURE; the TREND line falls back to UNKNOWN, checked unreachable and kept. A wrong turn, recorded: the first build made a short frame raise and mark the run degraded, against a recorded decision and the ruling's words; changed before any test ran, and the suite could not have caught it. Eleven tests, all behavioural; against 9051ffd's code 7 fail. Equivalence on 20,436 frames, no difference. A prediction wrong as made: Claude predicted the live record's run_hash would stay, but run_hash moves with the candle; recomputed on the record's own inputs, the code side of it had not moved. Found, not fixed: eleven more panel fallbacks of the same kind, not checked for reachability, Viktor's call (Entry #215).

#215	The Panel's Fallback Readings and Two Repeated Readings Join the Label and Record Batch October 12, 2026	DECISION — ADOPTED

“The ruling does not repeat those two conditions, and the record does not read them in.”

Commit ef65f46, documentation only, filed at Viktor's instruction in the session's first commit. Viktor's ruling of 12 October, 08:24, in his own words, not by agreeing to a suggestion: the eleven panel fallbacks found building F9 join round 8's batch, and so do two repeated readings found on F9's live run — REGIME prints the same _detect_regime value as STRUCTURE, and bias.volatility is printed on both the STRUCTURE line's Vol field and the VOLATILITY line. Claude's suggestion had been the batch after the rest of it, if a check showed the fallbacks cannot fire; the ruling did not repeat those conditions, and the record does not read them in. Not ruled, Viktor's before any code: each item's shape and their place in the order. Claude checked, from the code at e5ac20a, that REGIME and STRUCTURE are one value on every run where calculate_structure returned its frame, and that both volatility fields print the same key. “Eleven” was not numbered in NEXT; Claude's count, pairs as one and engine_core's trend_direction as the eleventh, is recorded as Claude's reconciliation, Viktor's to correct. Found filing it, not covered: STATUS falls back to ACTIVE ENTRY ZONE, the one panel fallback asserting a positive reading; ACTION to WAIT; the BTC call's three defaults, those the 2 September note removed from the asset's own path. code_hash unchanged at 3f8f218f….

#216 The Fallback Sweep: an Absent Value Reads UNKNOWN or Fails, Never a Reading

October 13, 2026

FINDINGS — FIXED

“None could fire on the engine's path, which is how they survived.”

Commit 797bdc9, three engine files, tests and documents; patch 1 of the 12 October additions to round 8's batch (Entry #215). The panel and the router read most of the decision with .get(key, default), and many defaults were readings: BIAS NEUTRAL, VOLATILITY LOW, RISK REGIME NORMAL RISK, STATUS ACTIVE ENTRY ZONE, ACTION WAIT, an entry score of 0.0, a reward multiple of 2.0 that was a second copy of AVG_REWARD_R, nothing degraded. Found in pieces and then by an inventory of every default in the two files. Viktor's ruling of 13 October, before any code, by agreeing to Claude's suggestions and on two questions choosing the option Claude recommended, filed in DECISIONS by this commit: a value only printed or recorded falls back to the router's own value for an absent one (UNKNOWN and its kin); a value the engine computes with, or one whose type has no UNKNOWN, is read directly and a missing one fails as a defect; patch 1 is a sweep of every default; patch 2 takes out the repeated readings, then B4's labels, B7, B9, B10 and Part 4's test shapes. Read into the rule by Claude, Viktor's to reverse: a recorded number falls back to NaN; the correlation label to NOT MEASURED; a Bitcoin block missing a key is reported unavailable rather than failing the asset's analysis; an absent exit watch prints that it was not computed; BROAD MARKET STRESS prints UNKNOWN for an absent flag. Sixteen new tests, thirty hand-built router inputs given the keys with the values the router used to invent. Against ef65f46's code 12 of 16 fail, 10 behavioural, 1 an ImportError, 1 the list's own control. Eight negative controls. The golden run's whole decision object and its 86-line panel identical but for code_hash. Claude's errors, recorded: its first lean, to remove the fallbacks, was wrong and corrected before the ruling; a line cited as :181 is :192; its question said a missing key becomes the router's error record, untrue for a section missing from the engine's output, which leaves no record — on Open items, Viktor's. Found, not changed: the router's own parameter coercions, three degraded_inputs defaults not checked for reachability, and a supertrend direction of 0.0 on an indicator failure, a degraded value feeding the bias.

#217 Repeated Readings: Each Reading Is Printed Once on the Panel

October 13, 2026

FINDINGS — FIXED

“One value on two lines; they could differ only when a fallback fired.”

Commit 69c92d6, the panel, tests and documents; patch 2 of the 12 October additions, under the rulings of 12 and 13 October (Entries #215 and #216). REGIME printed bias['regime'], read back from the STRUCTURE column that calculate_structure fills with the same value STRUCTURE printed; the STRUCTURE line's “| Vol:” field and the VOLATILITY line both printed bias['volatility']. The REGIME line and the Vol field are removed; STRUCTURE and VOLATILITY print each reading at its source. BTC REGIME keeps its Vol field, Bitcoin's volatility being printed nowhere else. Read into the ruling by Claude, Viktor's to reverse: the decision object and its record keep both keys, the ruling having named panel lines, so bias.regime is now recorded and not printed. Five new tests with sentinels; against 797bdc9's code 4 fail in the new file and 4 in the changed one, all behavioural; two negative controls. Golden snapshot unchanged; the golden run's panel goes from 86 lines to 85, no other line different. code_hash moves through panel_render.py only. The live run on Windows printed no REGIME line and STRUCTURE with no Vol field.

#218	B4's Labels: the Divergence Sentences Say What Is Compared; FADING VOLUME October 14, 2026	FINDINGS — FIXED

“Renaming the label in structure.py alone would have dropped its 25 validation points with no test failing.”

Commit 22a9a35, seven engine files, three tests, a new test file, the golden fixture and documents; round 8's triage, point 8 (the 5 October self-review's B4, the labels only), next in the order ruled on 13 October. Momentum divergence compares this candle's close and RSI with those four candles back, but its comments said “new highs” and three sentences called it momentum divergence; a close above the close four candles back can sit far below the highs in between, and it vetoes the side it points against. VOLUME DIVERGENCE fired when volume was below the volume four candles back and not expanding, whichever way the close moved; nothing compares a trend in volume with one in price. Found reading it: volume_agreement matched the label by part of a word, so a rename in structure.py alone would have dropped its 25 validation points with no test failing. Viktor's ruling of 14 October, before any code, by agreeing to Claude's suggestion, option B of three: the sentences say what is compared; VOLUME DIVERGENCE becomes FADING VOLUME; the part-word match goes; every score unchanged; the keys keep their names, the rename going on the list for after the audit. Wording Claude's, Viktor's to reverse; whether the divergence should veto is his, not ruled. Eight new tests; against 69c92d6's code 6 fail, all behavioural. Six negative controls. Golden snapshot: 4 leaves predicted, 4 moved; run_hash unmoved. The log's vocabulary changes here: of 59 earlier records, 44 carry VOLUME DIVERGENCE; records are not edited, and from this commit FADING VOLUME is the same condition and score. Found, not fixed: bias_engine scores an unknown volume label as neutral, a weighting question, Viktor's, on the list for after the audit.

#219	B7: the Stop-Distance Limits Record the Risk Regime October 15, 2026	FINDINGS — FIXED

“The 15% limit refused nothing on its own: it chose the message.”

Commit d77a8c4, one engine file, two tests, a new test file and documents; the self-review's B7, next in the batch, which Viktor chose. validate_risk_parameters tested the 15% distance limit and the 0.2% floor before classifying the regime, and both returned the regime as UNKNOWN; any stop over 15% is over 8% and so EXTREME RISK and refused anyway, so the 15% limit chose the message and made the record say UNKNOWN for a regime that was EXTREME RISK. Found reading it: the 0.2% floor has the same shape and, unlike the 15% limit, is reachable on the engine's own stop. On a refused run the regime decides nothing. The live log: 32 records carry the 15% message, all with UNKNOWN. Viktor's ruling of 15 October, 02:12, by agreeing to Claude's suggestion: option A of two — classify the regime first and have both limits return it, messages and precedence unchanged, the 15% limit kept as a backstop should the 8% ever be raised past it; option B, removing the limit, would have amended the ruling of 27 September and moved run_hash. Five new tests, one with an inline copy of the pre-B7 order to show no verdict or reason changes on any input. A sweep test that found the floor's refusal by regime == “UNKNOWN” would have passed for the wrong reason after the fix, and now finds it by its reason. Against 22a9a35's code 6 fail, all behavioural. code_hash moves through risk_model.py only; run_hash and the golden snapshot unmoved. Wrong turns: the docstring's first draft dated the finite-number check wrongly, corrected by git log -S; Viktor re-ran the apply after it had landed, and “patch does not apply” meant only that nothing was left to apply. Found, not fixed: two words, UNKNOWN and UNKNOWN RISK, for a regime that was not read; Viktor's whether it joins the batch.

#220	B9: Bitcoin's SuperTrend Is Read Directly October 15, 2026	FINDINGS — FIXED

“The panel's BTC BIAS printed a reading one factor short with nothing to say so.”

Commit 64ad625, one engine file, a new test file and documents; the self-review's B9, next in the batch, which Viktor chose. Read again before asking: the fallback sweep (Entry #216) had already made all of B9 but the SuperTrend, without naming B9, and four versions of NEXT kept B9 as still to come. Checked now: ST_Direction is missing exactly when SuperTrend's calculation failed; for Bitcoin that reached only a log line, and bias_engine scored the 0.0 as a neutral SuperTrend, so BTC BIAS printed one factor short, unmarked. Bitcoin is reference only; no trade follows. The asset's own read differs: a failed SuperTrend degrades the run, and a degraded run authorizes no trade. Viktor's ruling of 15 October, 07:55, by agreeing to Claude's suggestions: Bitcoin's SuperTrend now, read directly, so a failure makes the BTC section unavailable for the run, as any BTC failure does; the alternative, recording it in the BTC section's degraded_inputs, was not chosen, since the panel does not read that field. The asset's 0.0 after the audit, with what was checked recorded. Five tests that fail pandas_ta's supertrend for one series, told apart by price, so the failure takes the real path; against d77a8c4's code 2 fail, both behavioural, the 3 controls passing. Three negative controls. code_hash moves through engine_core.py only; run_hash and the golden snapshot unmoved. Claude's errors, recorded in DECISIONS: the sweep made most of B9 unnamed; it filed Bitcoin's SuperTrend for after the audit though the triage had put B9 in the batch; the review said the main path had removed these defaults on 2 September, and the asset's read kept its 0.0. On the way, two more: a draft cited two phrases to the wrong record, and the script writing the record stopped at a percent sign after writing HISTORY, which was restored and re-proven. Found, not fixed: a comment says BTC indicator failures are recorded inside btc_context, and only what BTC's trend health reads is; the asset's 0.0 is recorded and saved as the next run's prior state, changing no trade.

#221	B10: Two Outputs Agree With Finding 4 and the Action October 15, 2026	FINDINGS — FIXED

“Exit Watch advised a stop move at the target of a plan the engine did not suggest.”

Commit 9ea1cf3, two engine files, two tests, a new test file and documents; the self-review's B10, next in the batch. live_trading.py's simulated order recorded the EMA band as entry_zone, where finding 4 ruled the entry is the decision close and the band information; and on a WAIT or NO-TRADE under a leaning bias, Exit Watch still advised moving the stop to breakeven at Target 1, the target of a plan not suggested — record 62, 64ad625's live run, was one. No simulated order file exists on Viktor's disk. Viktor's ruling of 15 October, by agreeing to Claude's suggestions: B10 first, then the three test shapes as one patch; the band renamed ema_band, the panel's label; the note withheld on the panel unless the action names a side (LONG or SHORT, the test decision_model makes of its own action), kept in the record. At 11:13, after Claude found the order already records the decision close as current_price, the decision log's name for the plan's entry: no field is added. Ten new tests; against 64ad625's code 5 fail, all behavioural; eight negative controls. The live log's 62 records rendered through both panels: 44 lose the note's line and nothing else. code_hash moves through the two files; run_hash and the golden snapshot unmoved. Claude's errors, recorded: the suggestion added entry_price without reading line 168, which would have recorded one value under two names, found building and put again; it did not name that finding 5's control expects the note on a WAIT fixture; and the chat prediction left out that a section whose only flag was the note then prints “No exit-watch flags are active right now.” Found, not fixed: the simulator writes an order file on every run, WAIT and NO-TRADE included.

#222	Part 4's Three Test Shapes: Three Tests Check What They Name October 15, 2026	FINDINGS — FIXED

“The test it replaces passed under all eight.”

Commit f00bef6, tests, run_tests.py's docstring and documents, no engine code; the last item of round 8's label and record batch. Round 8's Part 4 named tests whose shape lets them pass whatever the code does; read and run again: the snapshot-coverage test printed SKIP and passed when the golden fixture was missing; the non-finite-panel test asserted only that “math.isfinite” appeared in panel_render.py, and stayed green with safe_float's check deleted; its sibling, named in the same sentence, asserted only that a phrase appeared in a function's source, and stayed green with the guard's raise deleted. And the watched error count was written as 29 in three places where run_tests.py reports 32. Viktor's ruling of 15 October, 14:28 and 15:51, by agreeing to Claude's suggestions: the coverage test fails on a missing baseline; the non-finite test renders the golden decision with each of its 120 float fields set to +inf, –inf and NaN in turn, 360 renders; the SuperTrend test asserts what only the level's guard does, SuperTrend failing whole; the count is no longer written as a number outside PHASE7_NEXT.md, and counting fixture-taking tests apart goes on the list for after the audit. The 14:28 turn was stopped by the usage limit before anything reached his disk; its sandbox work was discarded and rebuilt. Against the old code the changed tests pass, the engine being unchanged; the evidence is the negative controls — eight guard removals each fail the non-finite test, which the old test passed under all eight. code_hash unchanged, run_tests.py's change being a docstring the fingerprint strips. Claude's errors, recorded: where the counts are kept; a suggested assertion that the final sweep would have kept green with the guard removed, found building; and calling Viktor's a question the 22 September ruling already settles. Found, not fixed: the non-finite test cannot reach LIVE PRICE, null in the golden decision.

#223 Round 8's Fix Verification Decided: What Is Checked, What Is Sent, to Whom

October 16, 2026

DECISION — ADOPTED

“After this push nothing about the decision exists only in a chat.”

Commit 36de35d, documentation only. Round 8's fix verification is Claude's under the delegation of 4 October, as Claude reads it; this commit files the decision in DECISIONS before the send it governs. Checked: the Critical's fix, the stop (Entry #204), which is the release gate's question; round 8's own fixed findings against its own words — F1, F2, F3, F5, F6, F8, F9 and Part 4's three test shapes; and every other change to code it read, for new defects. Declared not fixed: F4, F7, F10, Part 4's other named tests, and two sections of the Part 7 reply, one marked as not checked in the code. To MiMo-V2.6-Pro through OpenRouter, Xiaomi's endpoint pinned, fallbacks off, data collection denied, as in round 8. One single-turn request: round 8's run-2 reports and the rulings since 5 October verbatim, the whole engine source, the new test files and the diffs, the Constitution's audit copy and the instruction — close to full size, because the fixes touched 17 of the 33 engine files; estimated at 330,000–350,000 tokens from round 8's bytes per token. One run, under the cap of two; the commit tagged first; no freeze. The gate reopens on a report that says the fix holds and that no change adds a Critical, once Viktor accepts it. Viktor ruled nothing; his “Do it.” agreed to a plan that files it as Claude's reading, and whether “all the work” includes this verification stays open. Why it lands alone: on 16 October, with 17% of Viktor's weekly usage left, Claude split the stage so the budget could not leave it half done. Claude's error, recorded: the scoping message called F1 the Critical; F1 is round 8's traceability finding, Major, and the Critical is the review's B2, under Item 14.

#224 Round 8's Fix Verification: the Send Script, Its Instruction and Its Tests

October 17, 2026

PROCESS — RECORDED

“The send is refused unless the probe reasoned, came from Xiaomi and was served by that checkpoint.”

Commit 78ea37d, tooling, an instruction and documents; nothing sent. docs/build/send_round8_fix_verification.py builds one request to MiMo-V2.6-Pro from git objects at the tagged commit, never from the working tree: Xiaomi's endpoint only, fallbacks and context compression off, data collection denied; measured with MiMo's pinned tokenizer before any call; a probe asks first for the dated checkpoint xiaomi/mimo-v2.6-pro-20260921, and the send is refused unless the probe reasoned, came from Xiaomi and was served by that checkpoint or one named on the command line; two runs at most, the second only after a commit to DECISIONS. The instruction, sections A to F: the gate question on the stop, a verdict per fixed finding, other changes and every ruling's claims checked against the code, the declared items, anything else, identity and recognition. Twenty-five tests, no network; against the old tree the file fails at import, which proves only that the script is new. Thirteen negative controls. Claude's decision of 17 October, filed in DECISIONS, records seven places the tooling goes past the decision of 16 October, each Claude's and Viktor's to reverse — among them diffs from e186423, the code round 8 read, with the script refusing if the engine differs from 217ede6's, and the checkpoint checked before the send. Measured in the Linux sandbox: 56 files, 1,410,498 bytes, payload SHA-256 aa9b9241…, 341,082 prompt tokens, 472,154 with the reserve, of 1,048,576. code_hash unchanged. Claude's errors, recorded: the instruction's first draft based both diffs on 217ede6; the script first pinned diff.orderFile to an empty value, which git refuses; two first-draft tests were wrong, caught by running them.

#225	Round 8's Fix Verification Answered: the Stop's Fix Holds, No New Critical October 17, 2026	AUDIT ROUND 8, FIX VERIFICATION — RECORDED

“Nothing found wrong anywhere, on 7,309 reasoning tokens.”

Commit d3188f2, documents and the run's four files. Sent once from 78ea37d, tagged round8-verify-sent-2026-10-17. Viktor's dry run on Windows printed the sandbox's payload hash and token count exactly; probe and send were served by xiaomi/mimo-v2.6-pro-20260921, from Xiaomi, finish_reason stop; the provider counted 341,082 prompt tokens, the token check's count; 11,634 completion tokens, 7,309 of them reasoning, 240 seconds, $0.158. The report: A, the stop's fix holds and no change adds a Critical; B, F1, F2, F3, F5, F6, F8, F9 and Part 4's three test shapes fixed; C, no new defect in the other changes, every ruling claim it checked supported; D, no declared item worse; E, three notes, none a defect; F, it does not know its checkpoint and recognised no code. Claude's check, reading the code at 78ea37d, not running it: twenty-three claims checked and the code each describes is there; one quotation merges two passages; line numbers approximate; whether it read in full what it says is not checkable; its weight set out for Viktor — 7,309 reasoning tokens on 341,082 read, against round 8's run 2 at 68,464 on 430,820. No second run, Claude's under the decision of 16 October. Accepting the report is Viktor's, his position first (Entry #226). Claude's error, recorded and corrected: 78ea37d's rewrite of NEXT left three passages saying “the commit that writes this line” about 36de35d; the rewrite replaced named lines and did not search the file for the phrase. code_hash unchanged.

#226	Round 8's Fix Verification Accepted on the Stop's Answer — the Release Gate Reopens October 17, 2026	DECISION — ADOPTED

“Nobody independent of Claude weighed the report.”

Commit 67f6ec1, documentation only. Viktor's ruling of 17 October, 09:37, by agreeing to Claude's suggestion. He first wrote “I accept the report.”; Claude asked for a sentence each on the basis, the scope and the weight; he asked for Claude's suggestions instead (“You are the engineer.”), asked whether to send a more focused run or choose another auditor, and agreed. Accepted on A.1 alone, the stop's fix holds, which is what point 2 of the triage asked. A.2, no new Critical, recorded, not relied on: it says it checked every change, but names eight of the seventeen engine commits it was sent. Sections B to F triaged: noted, no action; B's “fixed” counts toward nothing until he rules what makes a round-8 rating count. No second run and no other auditor now, Claude having advised against both for this gate. The next full audit names the stop's fix as its own question and checks the seven fixes in B for new defects. The release gate reopens; backtesting stays closed. What it weakens, in the entry: the fixer suggested accepting the check of its own fixes; the gate rests on one run, by an auditor whose round-8 runs found 0 and 2 of 4; the fixes in B go unchecked independently until a full audit not yet planned. Claude's errors, recorded: the check filed at d3188f2 missed A.2's list, and the no-second-run decision rested on A.2; the decision of 16 October checked B's fixes only as fixes and did not list that as a weakness; DECISIONS never recorded that the instruction sent the other answers to triage; four stale passages in NEXT and the change list, corrected.

#227 The Forward Measurement: main.py Takes the Pairs and Works From Its Own Folder

October 17, 2026

DECISION — ADOPTED

“Claude said the measurement might need no code at all, before reading main.py.”

Commit 0491b84, main.py, a new test file and documents. Viktor's ruling of 17 October, by agreeing to Claude's suggestion: “finished” needs 100 closed trades from one version, and nothing on record said how often the engine trades — the live log held 63 hand-started records across 42 code versions. So the engine runs every four hours on ten pairs, logging only, started by Windows Task Scheduler: AEROUSDT, ETHUSDT, SOLUSDT, XRPUSDT, BNBUSDT, DOGEUSDT, ADAUSDT, LINKUSDT, AVAXUSDT, SUIUSDT, the pairs Claude's suggestion. Claude's error, recorded in the ruling's entry: it first said no code might be needed, before checking; main.py ran only config.SYMBOL and wrote logs/ into whatever folder it was started from. Now “python main.py ETHUSDT SOLUSDT” runs each pair with a router of its own and config.SYMBOL set for its run; a malformed or repeated pair refuses the whole command (exit 2); a failed pair does not stop the rest (exit 1); run as a script, main.py changes to its own folder first. Claude's design calls, his to reverse: config.SYMBOL per run, the exit codes, refusing a malformed list whole, one task for all ten pairs. Nine tests, all needing pandas_ta; against the old main.py all 9 fail, 6 behavioural and 3 only because parse_pairs is new. Nine negative controls. code_hash bc1b9c2c… → 816d4951…, main.py's digest alone; golden unmoved. Eleven live runs on Windows before the commit, each record its pair's own; a pythonw.exe run started from C:\ wrote into the repository's logs/. The 63 earlier records intact. Viktor answered the ledger's question: no Amazon model had been sent anything. Four findings went on the list for after the audit, none deciding a trade. The model roster and ledger were reissued the same day, outside the repository, with their SHA-256 recorded.

#228 README's Status Rows Brought Current; the Forward Measurement's First Checks

October 17, 2026

PROCESS — RECORDED

““Nothing currently on record contradicts it” had been false since 4 October.”

Commit a235783, documentation only, Claude's work first at Viktor's request. README's rows Independent audit, Release gate and Backtesting had been stale since 29 September; they now name rounds 7 and 8, the Critical from round 8's triage — found by Claude's review of 5 October, not by round 8 — with its verified fix, the gate's closing and reopening, and the ruled conditions for a backtest, not met; Viktor approved the wording as drafted. Other stale README passages went on Open items at his choice. The forward measurement's first checks, read from his disk (Windows evidence): a run at 14:00:00.9 refused all ten pairs “not yet final”, by design, and what started it is not in the logs; the first scheduled run, 14:05, wrote ten records, each its pair's own, on code_hash 816d4951…, no LONG and no failure; the task, from his own schtasks output, runs every four hours from 14:05. Readings due 31 October and 14 November. How the count is made was set before any reading, a DECISIONS entry Claude's under the ruling and his to change: from the task's first run, one decision per pair per candle per version, refusals, failures and missed candles counted apart. B1, the trend score, seen on ten pairs: 83.71 to 100.00 on 21 records, the slope term 44.95 to 45.00 of its 45 points on every one; each score recomputed from its record's own inputs, all 21 matching. Corrected and recorded: three carried sentences that said “the commit that writes this line” or “this commit” of an earlier one, and the forward measurement's ruling missing from NEXT's “Ruled — in force”. code_hash unmoved.

#229	Amazon's Audit Is Scored on Four Defects Held Back From the List, and Its Ratings Are the Ones That Count October 17, 2026	DECISION — ADOPTED

“Nothing is planted: four known defects are left in through the audit and fixed after it.”

Commit 8d5d64a, documentation only. Viktor's rulings of 16:58 on questions 2 and 3 of his position of 6 October (Entry #209). Question 2 by agreeing to Claude's suggestion, which he asked for instead of writing first; question 3 by his position first — “Yes” — then Claude's critique, whose suggestions he agreed to. Ruled: Amazon's audit is scored on four known defects held back from the list for after the audit, the method of 29 September; each minor, touching none of Items 2, 3, 6 and 18, not pointed at by the code's own comments, with a “found means” sentence committed before the send; chosen before work on the list starts, left in through the audit, fixed after it, and the fixes verified by Amazon. Amazon's ratings are the ones that count toward opening backtesting, if the check passes as committed, the provider's record (not the model's own account) shows the pinned model answered, the run finishes normally, and Viktor checks the score himself. Round 8's verified fixes do not count toward condition 1 by its second path. Claude's reading, his to correct: the plan of 6 October, as changed, is ruled. Left open, his: what happens if Amazon fails the check twice (Entry #230). a235783's push checked from his disk: its pytest step ran after the apply. code_hash unmoved.

#230	If Amazon Fails the Check Twice: a Scoped Round on Items 2, 3, 6 and 18, Then a Wait for a Clean Lab October 17, 2026	DECISION — ADOPTED

“No lab would then be clean for a full round.”

Commit a48d9e5, documentation only. Viktor's ruling of 17:44, by agreeing to Claude's suggestion, which he asked for; it settles what Entry #229 left open, and that entry is not edited. If Amazon finds fewer than half of the four held-back defects in both runs: first, a scoped round on Items 2, 3, 6 and 18 — all that condition 1 of what opens backtesting asks to be rated — by a lab clean for scoped rounds (Upstage, NVIDIA, Cohere or AI21), scored on the same four defects, its ratings counting on the same four conditions, Amazon's findings still going to triage; second, if that fails too, backtesting stays closed until a lab that has never seen the project is available. Not chosen: Poolside or Xiaomi again. It adds a fifth criterion for the four defects: they sit in the files that scoped package would carry, still breaking none of those Items. Not checked: whether any of the four scoped labs can be reached, or fits even a scoped package. 8d5d64a's push checked from his disk. code_hash unmoved.

#231 The Scheduled Task's Second Run Checked: Nothing Ran at 18:00

October 18, 2026 — **PROCESS — RECORDED**

“What the gap rules out is a trigger at every close.”

Commit fbd7877, documentation only, at Viktor's request. Read from his disk (Windows evidence): the process log goes from the 14:05 run's last line, at 14:05:30.215 on 17 October, to the 18:05 run's first, at 18:05:01.300, with no line between, so whatever started the 14:00 run that refused all ten pairs did not fire at the next close. Viktor reads that run as a one-off; what the gap rules out is a trigger at every close, and one firing only at 14:00 would show at the next 14:00 close. The 14:00 run took the ten pairs in the task's order, and main.py runs them only when given them (parse_pairs, read at a48d9e5), so whatever started it was given the task's pairs; what it was is still not known. The second scheduled run, 18:05:01 to 18:05:39: ten records, each its pair's own, live from MEXC on the candle that closed at 16:00 UTC, code_hash 816d4951…; no LONG, seven on the SHORT side, one WAIT, two NO-TRADE, DOGEUSDT's DEGRADED INPUT on the same 11.8x volume candle; no failed run. a48d9e5's landing read from his disk: its pytest step ran after the apply. Reworded in the rewrite: two accounts that said “this session” and “this commit” of earlier ones. The four held-back defects were proposed in the message that delivered this patch and filed only with his answer (Entry #232). code_hash unmoved.

#232 The Four Held-Back Defects for Amazon's Audit: H1, H2, H3 and H5

October 19, 2026 — **DECISION — ADOPTED**

“Four runs of rounds 7 and 8 missed H4, and a shipped test asserts its behaviour.”

Commit 8311167, documentation only. Viktor's ruling of 18 October, under Entry #229. On 17 October Claude proposed H1 to H4 and said only H1 and H2 were strong; at 07:56 he asked for suggestions and at 07:58 agreed to them — accept H1 and H2, replace H4 through a search of the files a scoped package on Items 2, 3, 6 and 18 would carry, keep H3 unless something clearly better turned up. The search found H5, which he accepted at 09:28. The four, each with its “found means” sentence in DECISIONS: H1, the weak-validation WAIT decides nothing (decision_model.py; round 8's F10, rated Minor); H2, calculate_dynamic_bias takes trend_exhaustion and never reads it (bias_engine.py); H3, the simulator writes an order on every run, never reading the action (live_trading.py, run_once); H5, new, a flat candle can be credited to the wrong volume bin, because neighbouring bins share an edge and the strict “<” gives a tie to the bin below — in the sandbox a flat candle at 4.2 landed in (3, 4]; the repository's real candle files hold no flat candle. Not chosen: H4, the risk gate before the bias checks (S1): four runs missed it, a shipped test asserts it, and a docstring states it as Viktor's ruling; it stays on the list. The search, over every engine module's parse tree, also found two unread parameters in _compute_trade_quality and an unused import, which go on the list, and one loop variable, which is style and does not. fbd7877's landing read from his disk. code_hash unmoved.

#233 Nothing Ran at the 14:00 Close of 18 October — the Question Needs a Later Close

October 19, 2026

OBSERVATION — FILED FOR THE RECORD

“Viktor's PC was not running then.”

Commit d2dddf7, documentation only, the forward measurement's line on the 14:00 close of 18 October corrected at Viktor's request. The line said a trigger firing only at 14:00 “would show at the 14:00 close of 18 October”. Nothing ran at that close: his PC was not running (his word, 19 October). From his disk (Windows evidence), the process log has no line between 10:06:01 on 18 October and 02:05:06 on 19 October, so the question needs a later 14:00 close with the PC on. Not filed in NEXT: the log has no line between 18:05:39 on 17 October and 10:05:15 on 18 October either, so three closes had no run; the readings count missed candles from the logs. Found in the rewrite and corrected: the roster item said its SHA-256 are in “the message of the commit that writes this line” — right at 0491b84, wrong from a235783's rewrite on, carried through four commits; it now names 0491b84. HISTORY's entry for the file at a48d9e5 keeps the wrong wording, HISTORY being append-only, and this commit's entry names the exception. 8311167's landing read from his disk. code_hash unmoved.

#234 The List for After the Audit Sorted: Kind, Who Decides, Size and Order

October 19, 2026

DECISION — ADOPTED

“Not ruled, since the notes suggested nothing there.”

Commit aa9e7ac, documentation only. Viktor asked for the sort first; Claude read the list's forty-five entries at d2dddf7, with the rulings they cite, and put a table and six notes to him. At 11:03: “I agree with your suggestion.” Ruled by agreeing to Claude's suggestion, with Claude's reading of his words written into the entry, his to correct: the table as drafted, and the two things the notes suggested in so many words — no commit that moves code_hash before the forward measurement's reading of 14 November, and entry 1 (S1) worked after the audit with the four held-back defects. Not ruled, since the notes suggested nothing there: whether bias_score's weighting joins the work before Amazon's audit, whether Fable 5.1 or Claude builds the list, and the severity of F4's third path, which the table only puts first. Read from the code at d2dddf7 for the notes, by parse tree. Three corrections to the list from the same reading: entry 13 had missed an eighth “ROUND 6”, in a test docstring; entry 18's keys are named in thirteen test files and the golden fixture, not “more than 15” — and Claude's “14 test files” in chat had counted the fixture among them; the unused loop variable found on 18 October joins the list as entry 46. Corrected in the rewrite: an Open item said every point the triage left open was on the list, which was not true of the ten questions of 1 October or of F2's severity. d2dddf7's landing read from his disk, and the four points the previous session left in chat filed. code_hash unmoved.

#235	F4's Third Path Is Critical: Macro Moving the Tier Breaks Item 11, and the Release Gate Closes October 20, 2026	DECISION — ADOPTED

“The suggestion is the code's author's.”

Commit 391c338, documents and README; the first of its three rulings of 20 October, each by agreeing to Claude's suggestion, not by Viktor's position first. The facts were put to him first without a view. Round 8 rated F4 Major; Claude suggested Critical under the Constitution's rubric — macro's multiplier moving the action's tier breaks Item 11 — and said that the suggestion is the code's author's. Round 8's example arithmetic (76.9 and 65.6) is slightly off; Claude's is 76.4 and 65.5. The release gate closes again, three days after it reopened (Entry #226); goal A stands and portfolio-v1 is not moved; the fix lands after the forward measurement's reading of 14 November and before Amazon's send. Open, his: the fix's shape, and who re-audits it. README's Engine code and Release gate rows and the paragraph under the table, drafted by Claude in his voice and approved by him as drafted; the Engine code row now also counts round 8's eighth Critical, which it never had. The 14:00 reading of 19 October, read from his disk: nothing ran at that close, all five scheduled runs of the day complete; a trigger every day at 14:00 is ruled out, a weekly one is not, since 17 October was a Saturday. aa9e7ac's landing read from his disk. code_hash unmoved.

#236	Claude Builds the List for After the Audit, Fable 5.1 Waits; 4h Only for Now, With a Guard to Come October 20, 2026	DECISION — ADOPTED

“Claude suggested itself and said so.”

Commit 391c338, the second and third of its rulings of 20 October, filed in PHASE7_DECISIONS.md, each by agreeing to Claude's suggestion. Entry 38: Claude builds the list for after the audit, and Fable 5.1 waits until after the audit; Claude suggested itself and said so, and how Viktor's Fable credit can be used was not checked. Entry 22, B5: the engine is 4h only for now, and a guard will refuse any other timeframe — code, in wave 2. Read at aa9e7ac: config.TIMEFRAME is “4h”, main.py refuses only a missing or empty one, and a search of the .py files outside tests/ found no refusal of another, so the guard does not yet exist. No code changes with either ruling.

#237 Wave 1, Entries 13, 21 and 28: Three Comments That Said What the Code Does Not

October 20, 2026

FINDINGS — FIXED

“Counted from the parse tree, it has seven.”

Commit 8360923, comments and docstrings only, the first of wave 1 of the sort of 19 October (Entry #234); Viktor asked Claude to decide what came next and said go. Entry 13: eight comments and docstrings said “ROUND 6” for round 5's requested-run 6, a day before round 6 was sent; each now names it, and git grep finds none left. Dated records keep the old wording. Entry 21: core/engine_core.py's comment said a BTC indicator failure is recorded in btc_context; only btc_trend's own degraded_inputs is, and the comment now says so. Read from the code: a failed BTC structure sub-routine is neither recorded nor logged — entry 47, new, not changed, BTC being reference only. Entry 28: read_risk_verdict's docstring said validate_risk_parameters has six return paths; the parse tree counts seven, at the commit that wrote the sentence and since. A wrong turn, recorded: the first build added a four-line note to that docstring, and one behavioural assertion failed in all three configurations — a test reading models/risk_model.py from the working tree needs line 400 to hold the stop's formula. Never delivered; the change went in place, no line moving, and the dependency became entry 48. Also corrected: a sentence 391c338 left stale. code_hash unmoved at 816d4951…, each changed file's digest identical; a negative control, 100.0 → 100.1 in risk_model, moves it. Counts unmoved.

#238 Wave 1, Entry 48: the Location Check's Control Reads the Tree Its Check Reads

October 20, 2026

FINDINGS — FIXED

“The control broke on any line added above line 400 while the check held.”

Commit ed005a0, one test and documents. Claude decided the order of wave 1's test gaps — 48, then 36, then 37 — as Claude's under the sort, putting 48 first as the one that fails the suite whenever a change moves line 400 of models/risk_model.py, and the smallest. test_the_location_check_can_fail is the negative control of the round 8 fix verification instruction's location check: a citation one line off must not find the stop's formula and the line above must. It read the file from the working tree, while the check itself reads through _instruction_tree() — the commit that last changed the instruction once that is committed. Found when 8360923's first build added four docstring lines (Entry #237). Now the control reads through _instruction_tree() too, with a comment saying why. Negative controls (Linux sandbox): one line added at the top of risk_model.py fails the control at 8360923 and passes with the fix; with the instruction also changed in the working tree, both tests read the working tree and both fail, before and after — so the control still reads real content and moves with its check. Correcting the chat's “byte for byte” in 8360923's landing: each file on his disk is byte for byte an autocrlf checkout of the tip, and differs from the raw blob by its CRs alone. code_hash unmoved; counts unmoved.

#239	Wave 1, Entry 36: the Non-Finite Panel Test Reaches LIVE PRICE October 20, 2026	FINDINGS — FIXED

“The golden run is on pinned data, so the sweep never reached the LIVE PRICE line.”

Commit e88e0dd, one test and documents, Viktor's go at 13:19. test_the_panel_never_prints_a_non_finite_number renders the golden decision with each float field set to +inf, –inf and NaN in turn; on pinned data the struct candle's live_price is null, not a float, so the sweep skipped it. Found landing Part 4's three test shapes (Entry #222). Now the test asserts the fixture's live price is null, puts in a finite one before anything renders, and its precondition asserts the panel prints one LIVE PRICE line with a price; the sweep then sets it like every other float. The fixture is read, not written. Of the golden decision's nine null leaves, struct.live_price is the only one the panel reads. Negative controls: the line's isfinite check weakened to “is not None” passes before and fails after, printing “LIVE PRICE : $nan”; the line removed fails the precondition; the fixture's price set to 2.0 fails the seed's assertion. A wrong turn, recorded: the first try at that last control put 2.0 into another candle and the test passed; done again by key, it fails as it should. code_hash unmoved; counts unmoved. Also corrected: a sentence ed005a0 left stale.

#240	Wave 1, Entry 37: Part 4's Other Named Tests Checked Against Its Description October 21, 2026	FINDINGS — FIXED

“The scan has had nothing to read since 21 September.”

Commit d508b29, six test files, README's counts and documents; the last of wave 1's three test gaps. Round 8's Part 4 read first, then every other test it names checked against it, the test read and then run against a changed engine. Shape 1, absence guards that pass if the scan finds nothing, was right on every one. test_execution_surface.py's four scans had no control, and writing them found three blind shapes: “import json, ccxt”, importlib and __import__ unseen; python-binance, binance-connector and alpaca-trade-api, the names pip installs two of the four libraries under (read from each wheel on PyPI); and requirement operators other than == and >=. The risk_valid guard was a text match on three sites, missing a default without its space, a default of 1 and any other module — and, parametrized, run_tests.py had counted it as a fixture error and never run it. The swing-structure guard missed the price returned through a local; structure.py has had no .get call since afd8460 (21 September), so that scan had nothing to read. Shape 4 is right too: the router never builds a decision with no score keys, but the engine's own default responses return a finite 0.0, and forced through the router entry_model's printed 0.00/100 on a run recorded as not degraded, trading authorized, CONSERVATIVE LONG — entry 24's fix, a code change, not made here. A milder one entry 37 left out: a test asserting “95” absent had been unable to fail since c4dfcc7 (1 September). Now each scan's matcher is a function its control feeds; the risk guard reads every engine module's parse tree. Twenty negative controls: fourteen leave the old tests green and fail the new; six fail both. Ten matchers or file lists made to find nothing: ten of ten controls fail. Found, not changed: entry 49, conftest.ENGINE_MODULES leaves out three modules; entry 50, README writes the watched error count, which the ruling of 15 October keeps in NEXT. Qwen's reasoning in docs/qwen_reasoning_3.txt had seen the first gap; nothing was filed then. Tests 875 → 883; run_tests.py 804 / 0 / 32 → 815 / 0 / 31, the risk guard no longer a fixture error. code_hash and golden unmoved.

#241 d508b29's Landing, Three Missed Closes, and the Document on the Engine's Indicators

OBSERVATION — FILED FOR THE RECORD

October 24, 2026

“The document is not independent: Claude built most of the engine.”

Commit 9f2f3d7, documentation only, filing what was only in chat. d508b29's landing read from Viktor's disk: its pytest step ran 2 minutes 49 seconds after the apply, holding the new ids, and he went on to the commit, which confirms 883 by proceeding. The forward measurement's runs of 20 and 21 October, from the process log (Windows evidence): complete runs at 22:05 on 20 October and at 02:05, 06:05 and 10:05 on 21 October, no error line; the closes of 14:00, 18:00 and 22:00 on 21 October missed, his PC off, counted as missed candles. The document on the engine's indicators, written for Viktor on 23 and 24 October and kept outside the repository, in two versions. Measured with the engine's own functions on the ten pairs' archived candles, 3 September to 23 October, never reading a price after the decision bar, so inside the ruling that keeps backtesting closed. ChatGPT and Gemini reviewed version 1 at his request; their points were checked against the code and the data, and version 2 corrects what they got right — among them agreement read without base rates, macro's independence (the daily trend bullish on 86% of bars) and “circular reasoning” for readings that merely agree. It adds a decision-impact test: the full decision path per bar with each disputed factor removed; its baseline matched the engine's logged action on all 196 comparable live records, no removal flipped the side, and removing the four disputed bias factors cut trade actions by 37%. Nothing was ruled. The scripts were in Claude's sandbox only (Entry #242). code_hash unmoved.

#242 The Overlap and Decision-Impact Tool: the Document's Scripts Made Repeatable, With Its Tests

PROCESS — RECORDED

October 24, 2026

“The scripts behind version 2 left out the daily candle that closes with a 20:00 bar.”

Commit d410bfa, documentation tooling and a test file; no engine code. Viktor said start at 18:50. docs/build/indicator_overlap.py, new, runs the engine's own functions bar by bar on each pair's newest archived run, on the candles up to each bar only, and reports base rates, agreement corrected for chance (Cohen's kappa), correlations, threshold pass rates, and how many actions change, and how, with each disputed factor removed (ten variants). With --logs it checks its baseline against the engine's own logged actions, live, non-degraded, at the current code_hash. It reads no price after the bar it judges, so it cannot say whether a change is better — that is backtesting, closed. Claude's error, found building it: the scripts behind version 2 of the document (Entry #241) read macro from the daily candles closed by the bar's open, leaving out the one that closes with a 20:00 bar; the engine reads it, and so does the tool. On the same archives every figure moves a little and no conclusion changes: 2,980 bars where version 2 has 2,970; trade actions 1,333 to 860 with the four disputed factors removed, –35.5%, where it has –37%; no side flipped; 197 of 197 logged actions matched. Version 2 is not changed; the tool's docstring and NEXT record the difference. Thirteen tests, none taking a fixture, two needing pandas_ta: kappa against hand-worked cases, the weights restored after a failure, the change counts, the 4h guard, the macro rule, the golden decision reproduced, and a bar's row unchanged when every later candle is changed beyond recognition. Eight negative controls, each failing its test; the first macro control stayed green on the golden set, so the macro test was added. Tests 883 → 896. code_hash and golden unmoved.

#243 d410bfa's Landing: Its Test Steps Ran After the Commit

October 24, 2026

PROCESS — RECORDED

“The commit came 28 seconds after the apply.”

Commit 34d8e3b, documentation only, NEXT edited in place. d410bfa's landing read off Viktor's disk and GitHub after its push: the apply at 19:39:43, the commit from 19:40:11, the push from 19:40:22; GitHub's tree byte for byte the patch applied, and its message the one delivered. Its two test steps had not run before its commit: the commit came 28 seconds after the apply, and the pytest cache had last been written at 18:50:42, by 9f2f3d7's run. Claude asked for both steps after the push, under their original numbers; the cache was rewritten at 19:43:56 and holds the 13 new tests' ids, and Viktor reported “Done. All good.”, which confirms 896 and 828 / 0 / 31 on Windows. GitHub's tip for 9f2f3d7: both 9f2f3d7 and d410bfa said it was not fetched, because the repository had been made private and the sandbox could not reach it; Viktor made it public again that evening, and fetched then, 9f2f3d7's tree is byte for byte the patch applied. The sentence is corrected. Two wrong turns in the sandbox, caught before delivery: a patch first diffed against the LF blob rewrote the whole file, and a stray git stash stashed the edit in both sandbox trees. code_hash unmoved.

#244 Wave 1, Entry 49: conftest.ENGINE_MODULES Held to the Tree

October 24, 2026

FINDINGS — FIXED

“Every module in the engine” held 21; the engine has 24.”

Commit ad34c02, tests and documents; Viktor agreed to Claude's suggestion that it go next. core.lineage, core.code_fingerprint and utils.decision_log_backup were never added to tests/conftest.py's ENGINE_MODULES, so test_every_module_imports never imported them: an error at module level in any of the three was caught only where something else imports it. Shown at 34d8e3b: a RuntimeError at the end of either file, and the import check passes. Now ENGINE_MODULES names the three, and NOT_ENGINE_MODULES, new, names the two .py files code_hash reads that are not the engine, each with its reason — run_tests, and test_live, whose module body runs the engine once and writes a simulated order, so importing it is running it. A new test reads code_fingerprint.iter_source_files(): every file code_hash reads must be a listed module, a package marker with only a docstring, or a listed script, and every name a file, once, in one list. Rule 11: a list of names decays, so it is checked rather than extended and trusted; the test does not build the list from the walk, because test_live.py shows what importing everything can do. A new module fails the suite until someone decides its list (rule 28). Against the old code: with the old conftest whole, an ImportError, proving only that the name is new; with NOT_ENGINE_MODULES added, one behavioural failure naming exactly the three. Controls on every rule the test holds, each restored. Not read before in this project: pytest's lastfailed, {} since 26 September. Found, not fixed: entry 51, the suite's file walk excludes fewer directories than code_hash does. A wrong turn: Claude told Viktor before building that 49 was wave 1's last work of Claude's; building it found 51. Tests 896 → 897. code_hash and golden unmoved.

#245	The Engine's Indicators Ruled: One Direction Source, Gates for Weights, and One Redesign Before Amazon's Audit October 25, 2026	DECISION — ADOPTED

“Every suggestion is its author's, and the entry says so, with the case against.”

Commit 89d2367, documents and README. Viktor's eight rulings of 25 October, at 16:34, answering the eight questions of section 6 of the document on the engine's indicators (Entry #241). He first kept his rule of his position first, was given each question's options without a pick, then asked for Claude's suggestions and accepted all eight. In short: the trend-pullback hypothesis; the EMA slope as the one direction source, with ADX, the structure regime and divergence as gates; macro a gate on LONG only; trend health removed; entry quality as distance from the EMA band in ATRs and HVN nearness; confidence and the EV line dropped for now; one redesign, after the forward measurement's reading of 14 November and before Amazon's audit, with bias_score's weighting in it — reversing the ruling of 22 September, which is marked reversed, not edited — and F4's third path (Entry #235) fixed by removing the multiplier, re-audited by Amazon within its audit. Claude built most of the engine and wrote the document; the entry says every suggestion is its author's, and gives the case against. Found and not fixed: the entry score's structure multiplier (BOS 1.05, CHOCH 0.90) is not covered by the eight and stays unless he rules otherwise. Claude's errors, recorded: at 16:30 it suggested ruling 7's “leave the four held-back defects in place” beside rulings that reach three of them; and 391c338 had left README's sentence that every Critical has a verified fix, false from 20 October until this commit. README's rows, in his voice, drafted by Claude under his delegation. ad34c02's landing read from his disk; GitHub's tip read with no credentials once he made the repository public again. code_hash unmoved.

#246	The Trend Multiplier Is Critical by F4's Precedent; Held-Back Defects the Redesign Reaches Are Replaced October 25, 2026	DECISION — ADOPTED

“The trade's side comes from that slope on 99.6% of bars.”

Commit 89d2367, two decisions Viktor delegated to Claude in the same session, each his to reverse, both found filing the rulings of Entry #245. First, read at ad34c02: models/entry_model.py multiplies the entry score by 1.05 or 0.90 for the trade agreeing with trend health's EMA-slope label, while the trade's side comes from that slope on 99.6% of bars (version 2, measured); it moved the tier on 61 of 2,970 bars. The same shape as F4's third path, ruled Critical on 20 October, so Critical; entry 52, fixed by ruling 5 in the same change. Two Criticals are now open, and the release gate, already closed, reopens when both fixes have landed and been re-audited. Claude's error, recorded: version 2 of the document had suggested removing the multiplier without asking whether it breaks Item 11. Second: the held-back defects the redesign reaches are replaced or judged again, not protected. Ruling 4 removes trend health, which deletes H1's branch; rulings 2 and 8 rewrite H2's function; ruling 5 raises HVN's share, so H5's one-bin shift may stop being minor. The redesign is designed on its merits; H1 replaced under the 18 October ruling's own clause; H2 kept only if the design keeps its parameter for its own reasons; H5 judged again once the weights are set; H3 kept. Replacements are chosen when the design is fixed, before its code is built; Claude proposes, Viktor accepts or rejects. Both multipliers stay in the code until the redesign.

#247	Wave 1, Entry 51: the Suite's Walks Read the Files code_hash Reads October 25, 2026	FINDINGS — FIXED

“Two lists of the same directories, drifted apart.”

Commit b387c03, tests and documents; Viktor agreed to Claude's suggestion that it go next. tests/conftest.py's all_python_files, the walk the compile check and four source scans read the engine through, skipped a list of directory names of its own; core/code_fingerprint.py's EXCLUDED_DIR_NAMES, what code_hash counts as the engine, skips four more — .pytest_cache, env, Backup and Claude outputs (rule 11). Building the fix found a second walk with a list of its own, the utcnow scan, skipping five fewer; fixed here too, Claude's call, said to Viktor before building. Shown at 89d2367: a utcnow() call under Claude outputs/, a syntax error under Backup/, a place_order(call under Backup/ each failed a scan, though none of those files is the engine in any record. Nothing is read differently today: no .py file sits in any of those directories, in the repository or on his disk. Now all_python_files skips EXCLUDED_DIR_NAMES, read when it is called, and gives docs/ back to the compile check alone — one list, so the two cannot drift (rule 28). A new test plants a .py file in a directory of every excluded name, at three levels, beside two named like them (environment/, Backups/), and holds each walk to what code_hash's reads. Against the old code: a TypeError proving only that root is new; then one behavioural failure naming exactly the 20 planted files. Controls on every rule, and on the three scans. Found, not fixed: entry 54, the audit package builder's own list. A wrong turn: the sandbox's first clone was shallow, and five history-reading tests failed before the history was fetched in full. Tests 897 → 898. code_hash and golden unmoved.

#248 Wave 1, Entry 54: the Audit Package Reads the Engine code_hash Reads

October 26, 2026

FINDINGS — FIXED

“A .py file under Claude outputs/ went into the engine bundle.”

Commit 9862568, audit tooling, a test and documents; Claude's under the delegation of 4 October, Viktor agreeing that it go next. docs/build/build_audit_package.py builds the bundles an auditor reads as the engine and its tests, and its walk skipped a list of its own, not code_hash's: shown at b387c03, a .py file under Claude outputs/, env/ or Backup/ went into the engine bundle, one under tests/Backup/ into the test bundle, and one under node_modules/ into neither. Nothing is read differently today. Now the engine bundle is code_hash's own walk, iter_source_files, less run_tests.py and test_live.py, which the package puts with the tests on purpose; the test bundle is those two and every .py under tests/ in no directory code_hash skips; both loaded from the packaged tree's core/code_fingerprint.py by path. One list (rule 28). A new test plants files in every skipped directory at three levels, beside four named like them, and holds both bundles to the rule, then on the repository. Against the old script, one behavioural failure naming exactly the 17 misplaced files, and three bundles out of order. A consequence predicted: the engine bundle is now sorted by path, 17 of its 31 places changing. Two parts of b387c03's landing record had said more than its checks showed and are recorded as checked: the message was compared with its CRs removed, not byte for byte; and a failed comparison was put down to GitHub's files being CRLF, where they are LF. Claude's error, corrected: NEXT said twice that 89d2367's message was read byte for byte; it was compared with its line endings made the same. GitHub asked for credentials that day, so the sandbox's repository was built from Viktor's own .git objects, 350 commits, fsck reporting nothing missing. Found, not fixed: entry 55. Tests 898 → 899. code_hash and golden unmoved.

#249 Wave 1, Entry 55: Each Bundle's Header Says the Order Its Files Follow

October 26, 2026

FINDINGS — FIXED

“A small untrue sentence in what an auditor is told about its own material.”

Commit 3d7e6cb, audit tooling, a test and documents; Claude's under the delegation of 4 October, Viktor agreeing that it go next. The package builder's _bundle wrote “N files follow, in path order” above each bundle, and the engine bundle is the modules sorted, then the five project files — pytest.ini, the two requirements files, .gitattributes and .gitignore — so as a whole it is not in path order: the kind of small dishonesty _bundle's own comment names when it keeps the delimiter out of the header. Found reading the script for entry 54 (Entry #248). Now _bundle takes the sorted files and the files that follow them as two arguments, refuses unsorted first files rather than describe them, and makes the header's sentence from the two, so it cannot say more than the bundle does. Kept: the order itself, a grouping an auditor can use, and the bundles of rounds already sent, which are records. A new test holds that the call that wrote the untrue header is refused, then reads both bundles back as main() builds them. Against the old script, one behavioural failure at its first assertion. Six controls. Built with the old and the new script, the engine bundle's files, order, sizes and SHA-256s are the same and its text differs in one line, the header's. Wrong turns caught before any count: the first build of the test broke a string literal and failed to collect; the first control changed a name the header never used, so it changed nothing, and was redone. 9862568's landing read off his disk and GitHub, read with every credential source switched off. Tests 899 → 900. code_hash and golden unmoved. No entry found.

#250	Entry 26: Goal B Keeps a Candle Only Once It Has Closed at t October 26, 2026	DECISION — ADOPTED

““I don't understand. Just fix it as you see fit.””

Commit 78204ff, documentation only. Entry 26 of the list for after the audit, decided by Claude under Viktor's delegation, his to reverse. Claude suggested it as next, his and position first, or stopping; he said go, was shown what the entry asks, asked what leak, was told, and at 19:02 answered: “I don't understand. Just fix it as you see fit.” Entry 31, Goal B's other ruling, was not delegated and stayed his (Entry #251). The defect, in the ratified text: precondition 2 truncated each series at “open time strictly before t”. A daily candle opens at 00:00 UTC, so at a 12:00 decision that rule keeps a candle that has not closed, which a pinned file holds finished: the backtest would read the rest of the day. The ratified negative control, noise on every bar after t read by open time, never reaches that candle, so it would have passed. Decided: at t a candle of any series is kept only when its open time plus one bar of its own timeframe is at or before t — the live engine's own rule, read from the code at 3d7e6cb (_split_forming drops such a trailing candle; _load_pinned drops none). The control's noise now covers every candle the rule drops, the forming daily candle among them. Precondition 2 is amended in place and marked; the entry quotes the sentences it replaces and gives the case against. What it weakens: macro reads the previous day's candle all day, as the live engine does; and a ratified specification changed by Claude under delegation is a precedent. Reasoned from documents and code, not tested: no harness exists. 3d7e6cb's landing read from his disk, the loose commit object hashing to its name. code_hash unmoved.

#251	Entry 31: a Backtest Decides Through the Live Path October 27, 2026	DECISION — ADOPTED

“Writing the text, Claude read RSI's caller and found a missing last value goes to a fallback first.”

Commit 0d86fda, documentation only. Viktor's ruling of 27 October, by agreeing to Claude's suggestion, which he asked for instead of writing a position. The finding: clean_series fills a gap from the row before it at every row but the last; read again at 78204ff and run in the sandbox, [1, 2, NaN, 4] becomes [1, 2, 2, 4] and no value comes from a later row. At the last row the gap is reported, and for RSI that raises and sends the engine to its fallback. So a backtest computing indicators once over its whole history would decide on a carried value the live engine, deciding at that bar, did not have. Ruled: precondition 2 gains a paragraph — the harness reaches the engine only through the single-decision path a live run takes, on the files truncated at t, never computes an indicator once over the whole history, and records what that path returns, a decision, a fallback or a refusal. Its test, once the harness exists: a missing indicator value planted at sampled decision bars, at least three seeds, each bar's record byte for byte the live path's. clean_series is not changed; nothing ratified is removed. Two corrections to the suggestion as put in chat, recorded in the ruling: it said the backtest “records a refusal”, but RSI's fallback comes first; and it said precondition 1 puts recomputing into the budget, where it is precondition 0, the timing benchmark. What it weakens: the faster design is ruled out, by how much was not estimated. The other indicators' paths were not read for this entry. Reasoned from documents and code, not tested. 78204ff's landing read from his disk. code_hash unmoved.

#252	A Commit Names the Tree Verified Before Delivery October 27, 2026	DECISION — ADOPTED

“The first commit to carry a Verified-Tree line.”

Commit b8f88d8, documentation only. Viktor's ruling of 27 October, by agreeing to Claude's suggestions. At 11:29 he asked how to solve, or move past, what a landing check could not read again; Claude answered each point, and at 11:43: “We go with your suggestion.” He took Claude's suggestion on the two points left open too. Ruled: every delivered commit message carries a Verified-Tree line, the git hash of the tree verified before delivery, and a landing check compares it with the commit's own tree, from GitHub alone; a commit-msg hook refuses a commit whose staged tree is any other — its own patch, next, as entry 56; a later session no longer compares the message with the one delivered, but reads that it holds no CR, ends with its attribution lines and is .git/COMMIT_EDITMSG byte for byte; run_tests.py writes a last-run record after the reading of 14 November, being inside code_hash's walk (entry 57). Claude's, stated with it and not overruled: each read of GitHub dated, and a read with no credentials carrying a probe of a repository that does not exist. Tested before it was suggested: 0d86fda's CRLF checkout, staged into an empty index, gives 0d86fda's tree hash exactly, and Viktor's .git/index lists the same 293 paths, modes and blobs. The probe, run that day with every credential source removed, got a request for credentials, so none reached GitHub; that it answers “Repository not found” when credentials do arrive is known behaviour, not tested. What it weakens: the comparison of the message with the one delivered. 0d86fda's landing read from his disk. This commit carries the first such line, 9f69d94d…. code_hash unmoved.

#253	Wave 1, Entry 56: the Commit-msg Hook Checks the Verified Tree October 27, 2026	FINDINGS — FIXED

“Nothing compared it with the staged tree before the commit existed.”

Commit a1c4218, a hook, a test file and documents. Entry 56 of the list for after the audit, ruled by Viktor on the morning of 27 October by agreeing to Claude's suggestions, the hook Claude's to build. Since b8f88d8 a landing check compared the Verified-Tree line with the commit's tree only after the push, when a wrong tree cannot be undone without rewriting history. githooks/commit-msg, new (sh, LF, 100755), collects every line that looks like the tree's line — its key in any case, indentation or spacing before the colon — and stops the commit on none, on two or more, or on one not written exactly as the key, a space and forty lowercase hex digits; otherwise it compares the hash with git write-tree on the index git names to the hook, so a commit -a or a partial commit checks the tree git commits from. Equal prints one line and exits 0; anything else stops the commit, and it fails closed. Viktor ruled its one open point at 13:08, by agreeing to Claude's suggestion after being asked for his position first: a message with no such line is stopped, and git commit --no-verify is the way past after a consult. Two lines, or one not in the exact form, are stopped too: Claude's, his to overrule. tests/test_commit_msg_hook.py, 11 tests, no fixture: against the tree before, 10 fail with FileNotFoundError, proving the hook new and nothing more, and 1 passes, the LF rule already existing; the behavioural evidence is eight negative controls, each restored and checked with cmp. Claude's wrong turns, caught before any count: its first comparison of his index with the tree split a “Claude outputs/” path at its space; two tries at the fail-closed branch for a failing write-tree did not reach it, so that branch is untested, and the test file's docstring says so. b8f88d8's landing read from his disk. Linux sandbox evidence. code_hash unmoved.

#254	Entry 50: README Does Not Write the Watched Error Count October 27, 2026	DECISION — ADOPTED

“A second copy of a number that has gone stale before would go stale again.”

Commit 713e084, documentation only. Entry 50 of the list for after the audit, found on 20 October: README wrote run_tests.py's watched error count, 31, twice — its Tests row and its paragraph “All tests currently pass” — though point 5 of the ruling of 15 October keeps that number in PHASE7_NEXT.md only; the paragraph also said the count moves only when a fixture-taking test is added on purpose, and entry 37 had moved it by making one fixture-free. The shape was Viktor's, position first; asked for a suggestion at 17:49 and offered one, he answered at 17:52, “Do it.” Claude took that as handing it the shape, as with entry 26 (Entry #250), and did not ask again: Claude's decision under that delegation, filed in DECISIONS, his to reverse. README's row and paragraph now name every way the count moves and point to NEXT's “Test suite” for the number; the other counts README writes stay. New wording in Viktor's voice, put to him before his commit. Point 5 is not amended: naming README in it would change his ruling, which was not delegated. Left: README's other stale passages, an Open item of their own, and counting fixture-taking tests apart from real errors, still on the list. a1c4218's landing read from his disk, its pytest cache holding the 11 new tests' ids. code_hash unmoved.

#255	README's Audit History and Layout Brought Current October 27, 2026	PROCESS — RECORDED

“That sentence would have published a false correction.”

Commit 134f85b, documentation only. README's other stale passages, an Open item since 17 October, Claude's to draft and Viktor's to approve; he said “Go.” at 18:42. “Five more rounds have run since” had stopped being true when round 7 ran on 30 September; both lists of models in “How this was built” stopped at round 6, missing Laguna S 2.1 and MiMo-V2.6-Pro, and, checked against the roster and ledger of 17 October, had left out GLM 5.3 Flash, round 3's auditor, since 15 September; the layout's logs/ line was silent on logs/pinned/, and its build/ line named only the PDF scripts. Now the paragraph counts seven rounds since the original audit and names rounds 7 and 8 and the fix verification; fourteen models have seen the Constitution and ten have graded the engine source, with a sentence saying the lists stopped at round 6 until 27 October. A wrong turn, caught before building: Claude's first draft, approved at 18:52, took the second list's “GLM 5.3” for the Flash and said README had named it in error; README has named GLM 5.3 there since 7158a7f, on 29 August, a week before round 3 ran, for its own session of 28 August. Told to Viktor with the paragraph corrected — the Flash added beside it, ten graders, not nine — he took that suggestion at 19:03 and said “Go.” at 19:09. Whether that session graded the source is README's claim, not found in the record, and left. Qwen3.8-Max stays out of both lists: on the record, no round-2 attempt received the Constitution's text. 713e084's landing read from his disk. code_hash unmoved.

#256 A Landing Check Does Not Report Its Standing Gaps

October 28, 2026

DECISION — ADOPTED

“They recur by the design of the report, not as findings.”

Commit fad0588, documentation only. Viktor's ruling of 27 October, 19:35, by agreeing to Claude's suggestion. Every landing check since that morning's ruling (Entry #252) had reported the same four parts as not read again — whether run_tests.py ran, the commit-msg hook's line at his commit, the commit message as delivered, and the session before's read of GitHub's tip — and so had every landing record. At 19:34 he asked whether that was an issue, whether it affected the project and why it happened. Claude answered that none of the four leaves a gap in what is verified, and suggested that a part not read again be reported only when it is new or bears on whether the landing holds, the standing four named once in DECISIONS; “I agree.” Ruled so, run_tests.py's gap leaving the list when entry 57 lands; point 4 of the ruling of Entry #252 amended in place with a dated note. A slip, caught before anything was reported: the comparison of his index with the tree split a “Claude outputs/” path at its space, as at b8f88d8's landing check. The patch-delivery skill, outside the repository, still told a landing check to report those parts; an updated skill was put to him to save. 134f85b's landing read from his disk. code_hash unmoved.

#257 A Ruling That Changes the Delivery Procedure Changes the Skill in the Same Session

October 29, 2026

DECISION — ADOPTED

“The ruling of 27 October reached DECISIONS only.”

Commit 3dbd8dc, documentation only. Viktor's ruling of 29 October, by agreeing to Claude's suggestions. The landing check of the session from 18:23 on 28 October reported two of the four standing gaps against the ruling fad0588 had filed (Entry #256): a landing check follows the patch-delivery skill, which is in his account, outside the repository, and does not read DECISIONS, and the skill's section 10 still said to report what cannot be read again. At 13:08 he asked what the issue was, beyond the obvious, and said he wanted it prevented. Claude suggested comparing the whole skill with DECISIONS at once, and giving the skill a marker, the commit it was last compared at, so that each session finds what was filed since by one command; “Go.” at 13:44. The comparison found nothing else contradicting the skill and four rulings absent from it. He said “Go.” to a skill carrying all of it, marked fad0588, at 13:46, saved it at 13:51, and at 13:52 agreed to filing it. Ruled: a ruling that changes how a patch is delivered or checked changes the skill in the same session; every delivery and landing check begins with the diff of DECISIONS since the marker, which moves once the commit filing a ruling has landed. Two slips in the draft, caught before delivery: the ruling's numbered points indented unlike every list in the file, and a reason saying none of the four absent rules had yet caused a wrong step, which had not been checked — it now says each was a step a session could have got wrong. fad0588's landing read from his disk. code_hash unmoved.

#258 A Candle Touching Both the Stop and T1 Counts as the Stop

October 29, 2026

DECISION — ADOPTED

“Claude first suggested a rule for a case it had not yet counted.”

Commit a2ca918, documentation only. Viktor's ruling of 29 October, by agreeing to Claude's suggestions: the fifth question of his position of 6 October and the first item of the Freqtrade comparison, the one open item that must be ruled before any backtest result exists. Set out for his position first, he asked for Claude's suggestion instead, and Claude said that this reversed the order he asks for on his own decisions, then gave it; “Go.” at 14:58 to a count of the candles and, given the count, “Agreed.” at 15:01. Ruled: before T1, a candle reaching both the stop and T1 counts as the stop, the trades it decided reported beside the verdict with the result had they reached T1, not graded; at T1 the stop moves to the entry at the close of T1's candle, not within it, the T1 candles that also reached the entry reported with the profit had those second halves closed there; paper trading and the random-longs benchmark follow the same rule; Goal B's data stays 4h. The count (Linux sandbox, from logs/archive/ on his disk; candles, not trades, and not Goal B's pinned dataset): of AERO's 1,218 4h candles, 4.8% could reach both the stop and T1 at a 3% stop, 0.9% at 5% and 0.2% at 8%; both T1 and the entry, 29.8%, 9.3% and 2.5%. Claude's error, recorded in the ruling: it first suggested that such a T1 candle close the second half at the entry, before that case was counted; up to 29.8% was enough to move the profit total, and the suggestion changed to moving the stop at the close. Goal B's methodology gains one line. 3dbd8dc's landing read from his disk. code_hash unmoved.

#259 What May Be Tuned After the Paper-Trading Year: at Most Three Thresholds, Judged on Fresh Trades

October 31, 2026

DECISION — ADOPTED

“”You are the engineer and the brain of this project. I am just the manager.””

Commit d7d3ae0, documentation only. Question 6 of Viktor's position of 6 October. On 29 October Claude set it out for him to write his position first; at 16:06 he answered that these were highly technical questions he could not answer, and that session closed with nothing filed. In the session from 23:21 on 30 October he asked for it to be filed as Claude's decision, with one line first on whether any part needed his: point 3 excluded anything already ruled, and its examples, the stop's ATR multiple 1.2 and the ADX floor 20, are values in his rulings of 5 October (read from DECISIONS, not from code). Asked for a suggestion, he said “B” at 23:58. Decided, Claude's under his delegation of 29 October, his to reverse: at most three settings, named when the redesign's design is fixed, before Amazon's audit; only thresholds whose effect the year's records show directly, not indicator lengths and not what a ruling decides; a threshold moves for the result, never for the trade count; the tuned version is judged on its own fresh 100 closed trades paper traded forward, a held-out backtest window allowed as an earlier check once Goal B's harness exists, as a new phase under the re-run clause; one tuning round a year. His, point 3: a number inside one of his rulings may be named, what the ruling decides stands, and any ruled value on the list comes to him for a yes or no. Claude's errors, recorded: an engineering question set out for his position on 29 October; point 3, as first given, contradicting itself, found before filing. a2ca918's landing read from his disk. code_hash unmoved.

#260 Volume Enters the Redesign as Two Gates: a Liquidity Floor and Volume Confirming the Trigger

October 31, 2026

DECISION — ADOPTED

““Sounds good, we add in the next patch.””

Commit b233c4a, documentation only. Viktor's ruling of 00:55 on 31 October, by agreeing to Claude's suggestion. He asked whether volume should feed T1, T2 and T3; Claude said it should not, and he agreed. He then said volume should be implemented and asked which kinds. Claude had said that choice was its own; he calls the ruling his; either way the gates are new trading rules, and their design comes to him. Ruled: a liquidity floor, a minimum of USDT traded per candle, from MEXC's quote_volume, which the fetcher drops today; and the trend-resuming candle's volume above its recent median. Not a buy/sell split (MEXC's candles as fetched carry none; whether MEXC publishes one elsewhere was not checked) and not OBV-type indicators, which would count direction twice. The gates' cutoffs are candidates for the at-most-three tunable settings (Entry #259); the median's window, a length, is not — Claude's reading, added in the filing. No code before 14 November. Read from the code at d7d3ae0: data/data_fetcher.py drops quote_volume before the frame is returned, the decision contract declares volume as a label, and the candle archive holds no quote_volume as far as read, not traced end to end. Claude's error, recorded in the ruling: it said the reading of 14 November could show whether trades on higher volume did better because every decision record logs volume; the measurement records decisions, not trade results, and the record logs a label, not the numbers — corrected in that session. d7d3ae0's landing read from his disk. code_hash unmoved.

#261 The Forward Measurement's Reading of 31 October

November 1, 2026

OBSERVATION — FILED FOR THE RECORD

“The missed candles are mostly the night closes.”

Commit 6d073f9, documentation only. The first of the measurement's two readings, due on 31 October and Claude's to count under the ruling of 17 October, by Claude's decision of that day on how it is counted, set before any reading; Viktor said “Go” at 12:17, and again at 12:21 to filing it after the 14:05 run. Counted by a script from the bytes of logs/phase7_engine.log and the ten decision logs, staged off his disk (Linux sandbox). The window: the 84 candles that opened from 08:00 UTC on 17 October to 04:00 UTC on 31 October; per pair 39 read and 45 missed; 10 refusals, all at 14:00 on 17 October, their candle read at 14:05; no failed run. All 390 records are the task's, live, on code_hash 816d4951…, one per pair per candle. Actions: 56 LONG, 143 SHORT, 60 WAIT, 131 NO-TRADE; 3 degraded runs, all DOGEUSDT, from one volume candle. 45 of the 56 LONGs fall on eight candles in a row, 20 to 21 October, all ten pairs LONG on the last of them. The counts are for the hours the machine ran, and the reading of 14 November reads with that stated. Nothing ran at 14:00 on that Saturday, nor at 13:00, the close by the winter clock; with the read of 19 October, a trigger at 14:00 daily or weekly is ruled out, which fits Viktor's reading of the run of 17 October as a one-off. Read, and not a defect: the one degraded run labelled NO-TRADE (SIGNAL UNCONFIRMED) follows the precedence work order F set. Nothing ruled. b233c4a's landing read from his disk. code_hash unmoved.

#262 A Session Does Not Ask Viktor What to Do Next

November 1, 2026

DECISION — ADOPTED

“”Stop asking me what to do next and take your own initiative and do that which should be done first””

Commit 8c810b1, documentation only. Viktor's ruling of 00:28 on 1 November, written when the landing check of 6d073f9 ended, as the patch-delivery skill's section 10 then said, by asking him what he wanted to do. The words are his; points 2 to 5 of the ruling, and what “first” means, are Claude's reading, his to reverse. After its opening a session chooses the work that should be done first and starts it, saying in one line what it is and why; this replaces only the last sentence of the pace ruling of 15 September, and the Open items are still not a queue, one piece at a time. What is marked his stays his to decide, and the rulings in force bind the choice, among them no commit that moves code_hash before the reading of 14 November. The case against, recorded with it: the question at a session's opening was the brake the pace ruling set, and keeping the pace is now Claude's. The skill's section 10 changes to match, put to him with the marker after the push (Entry #257). 6d073f9's landing read from his disk. code_hash unmoved.

#263 8c810b1's Landing, and the Skill Saved With Its Marker

November 1, 2026

PROCESS — RECORDED

“One command meant to run run_tests.py also carried a git stash nothing called for.”

Commit 248037d, documentation only. Records 8c810b1's landing, read off Viktor's disk and GitHub — applied at 15:41:31, the pytest cache rewritten 3 minutes 1 second later, committed at 15:48:29, pushed at 15:48:44, its tree the one its trailer names; nothing differed — and his save of the patch-delivery skill with its marker at 8c810b1, its section 10 carrying the ruling of 1 November (Entry #262). No ruling is filed, so the marker stays. A wrong turn, recorded: in the sandbox one command meant to run run_tests.py also carried a git stash of the applied tree and the working copy, which nothing called for; it was stopped before the stash, both trees were confirmed unchanged, and run_tests.py was run again on its own. One of the five commits of 1 to 3 November that recorded only a landing (Entry #276). code_hash unmoved.

#264 Entry 58: a Reasoning Sentence Rounds Across Its Line

November 1, 2026

OBSERVATION — FILED FOR THE RECORD

“”trend strength 75/100 (down), below the 75 the SHORT tier needs””

Commit e748b0b, documentation only. Found on Viktor's hand run of 16:09 on 1 November (AEROUSDT) and filed as entry 58 of the list for after the audit, not fixed. Decision Reasoning printed a trend strength of 75/100 as below the 75 the tier needs, on a trend health of 74.52: the ladder compares the unrounded value with AGGRESSIVE_TREND_HEALTH_MIN (75.0), and models/decision_model.py:634 prints it with :.0f. The same rounding is at :512 and :519, against AGGRESSIVE_ENTRY_SCORE_MIN, and, found by Claude reading the file, at :687–688, against MIN_ACTION_BIAS; :680 and :774 print a rounded value beside a comparison they do not name. Checked in Python 3.12: 74.515…, 69.5 and 29.5 print as 75, 70 and 30, and 74.5 as 74, round-half-even — hence “can”. It changes no action and no recorded number. From his disk: models/decision_model.py byte for byte the tip's, and the decision log's last three records his three hand runs, on code_hash 816d4951…, the same closed candle and a trend health of 74.51544804066864, each with that sentence first, so all three printed it. Not fixed: a fix moves code_hash before 14 November, and entry 53 removes trend health. Claude's suggestion, his to change: a label fix within entry 53, whose design gains the rule that no reasoning sentence rounds a compared value across its threshold. Wrong turns, recorded: the first virtualenvs were started under Python 3.13 and rebuilt under 3.12.3 before any test ran; the first check that the suite left logs/ alone was vacuous, the clone's logs/ being empty, and was repeated with three sentinel files; the index comparison split on spaces once more. 248037d's landing read from his disk. code_hash unmoved.

#265 e748b0b's Landing

November 1, 2026

PROCESS — RECORDED

“Predicted, every line of the diff.”

Commit f6aa126, documentation only. Records e748b0b's landing, checked at 17:52 on 1 November: GitHub's tip, read with no credentials, e748b0b, its tree the one its trailer names; Viktor's disk the tip, the pytest cache rewritten 3 minutes 29 seconds after the apply; nothing differed. No ruling is filed, so the skill's marker stays at 8c810b1. By the ruling of 1 November the session then starts the redesign's design draft, which the session of 16:26 had named as its next work and Viktor's opening of 17:49 named as still to start. One of the five commits of 1 to 3 November that recorded only a landing (Entry #276). code_hash unmoved.

#266 Entry 53's Design Draft, Version 0

November 1, 2026 — PROCESS — RECORDED

“The golden action would move from CONSERVATIVE LONG to WAIT.”

Commit e69d5bb, documentation only, the session's second commit, so NEXT is edited, not rewritten. Files the redesign's design draft, version 0, as docs/redesign_entry53_draft.md, so that the draft lives in the repository and not only in chat. Claude's, started by the ruling of 1 November and read at e748b0b; nothing in it is ruled, and no code is written from it until Viktor rules the design, after the reading of 14 November. It covers the decision ladder, the gates, entry quality's weights, the strength floor, the two volume gates (Entry #260) and entry 58's rule (Entry #264); for each part it names the tests and golden fields the part moves, and which parts are his to decide; its gaps are marked “to measure” and wait for the data of 14 November. It proposes a trigger candle, which no ruling defines and the code does not have (searched at e748b0b). It finds that the pinned golden run's entry quality comes out at 0 under the proposed formula — computed in the sandbox from the golden record's own close, band, ATR and HVN, not by the engine. f6aa126's landing recorded. code_hash unmoved.

#267 Entry 53's Design Draft, Version 1: Every Golden Field and Every Test It Moves

November 1, 2026 — PROCESS — RECORDED

“Version 0 was wrong in three places, recorded in its section 15 rather than quietly changed.”

Commit e8b42b8, documentation only. Version 1 of the draft, Claude's and not ruled: the parts Viktor asked for that need no data from 14 November. Section 13 takes every golden field one by one — 227 leaves: 125 unchanged, 81 removed, 15 that move, 2 renamed, 4 his, which go only with the structure multiplier — and the pinned run becomes a WAIT at gate 8: its decision candle closed under its open and under the previous close, and its volume (84,693.73) is under the twenty-candle median (179,247.89). Section 14 takes every test it moves: 467 of 874 test functions — 95 go, 124 rewritten, 2 with their expectation reversed, 118 expected to hold, 128 moving only through a helper — classified by a script from names, docstrings and the names each body and helper uses; the bodies were not read line by line, and the draft says so. Entries 14 and 19, read against section 3.2, are settled by the design. Found on the way: indicators/trend_health.py also supplies the ADX and divergence gates; the exit watch would lose its exhaustion and reversal flags (flagged as Viktor's); the bias-flip flag would fire once on every pair unless its saved state changes key; BTC's reference line on the pinned run moves from BEARISH to BULLISH; quote_volume would move every live input hash. Version 0's three errors: it named gate 9 as where the pinned run stops, though gate 8 fails first, written without reading the candle; it gave entry 14 the wrong reason and put entry 19 in the wrong section; and it gave an unmeasured market the label of a thin one. e69d5bb's landing recorded. code_hash unmoved.

#268 e8b42b8's Landing; What Comes First Is an Item of Viktor's

November 1, 2026 — **PROCESS — RECORDED**

“No piece of Claude's can land before the reading of 14 November.”

Commit d6fba83, documentation only. Records e8b42b8's landing, nothing differing, and Claude's choice of what comes first under the ruling of 1 November: waves 2 to 5 move code_hash and the redesign's draft waits for the reading, so no piece of Claude's can land before 14 November, and what comes first is an item of Viktor's, the Bedrock billing, the step of the plan after which the check of reach and fit is Claude's. Recorded in NEXT, his to redirect. One of the five commits of 1 to 3 November that recorded only a landing (Entry #276). code_hash unmoved.

#269 d6fba83's Landing; the Draft's Measurements to Be Fixed Before Their Data Is Read

November 2, 2026 — **PROCESS — RECORDED**

“Named, but not each with its computation and the result that would count against the draft.”

Commit 7c23886, documentation only. Records d6fba83's landing, nothing differing. At the session's opening Viktor said the Bedrock billing is his and that the draft, version 1, waits for the reading of 14 November. Claude's choice of what comes first, then: the reading is the work due by its date, and the draft's five measurements marked “to measure” were named, but not each with its computation and the result that would count against the draft; Claude fixes those before any of that data is read for them, as the count's own method was fixed before any reading. Claude's, as the draft is, his to change, documentation only. One of the five commits of 1 to 3 November that recorded only a landing (Entry #276). code_hash unmoved.

#270 How the Reading of 14 November Measures the Draft, and What Result Counts Against It

DECISION — ADOPTED

November 2, 2026

“The thresholds are Claude's judgment, set now, not derived.”

Commit e939734, documentation only. Claude's decision, Viktor's to change, filed in DECISIONS rather than in the draft, because the draft is rewritten and a method fixed in advance must not be; set before any of the reading's data was read for it. One population — the records the reading counts under the decision of 17 October, degraded ones apart — the record's own values, and the run's own archive only for what a record lacks, checked against the record. Six computations, each with the result that counts against the draft: direction (opposite sides on more than 2% of the records where both have one, or NEUTRAL on more than half of today's trade actions); the tiers (a class empty or above 90%); F (the 5% per pair; against it, the stand-in's bounds disagreeing by more than 2 points, or F under 1% of the median on nine pairs); the trigger and its volume (blocking 90% or more of what passes every other gate); H5's reach; and the whole ladder (entries projected under 100 a year on the ten pairs). Each says what it weakens. Read for it at 7c23886: the record's fields, generate_entry_signals's blocker lists, validate_risk_parameters, write_archive and compute_volume_profile; in the sandbox (pandas 3.0.5, not his venv) the pinned AERO frame gives the golden HVN, 0.6421571. Not blind on one point, and the entry says so: the document of 24 October measured the bias side against the slope on archives that overlap the window's first week. Found at the landing check: no run at 02:05, 06:05 or 10:05 on 2 November; Viktor, at 11:30: “When something doesn't run, the computer is off.” 7c23886's landing read from his disk. code_hash unmoved.

#271 The Reading of 14 November: the Script for Its Six Computations

PROCESS — RECORDED

November 2, 2026

“Not run on the live logs; it runs at the reading.”

Commit 1b82ed0, a script under docs/build/, its tests and documents; no engine code. docs/build/measure_draft_gaps.py, new, runs the six computations of Entry #270, written from that entry before any of the reading's data was read, and prints each with what counts against the draft. Its population: live 4h records at one code_hash, of candles that opened from 08:00 UTC on 17 October, the scheduled task's only — logged within five minutes of its first start at or after the close, Swedish time from the EU rule rather than a time-zone database, so that Windows needs no tzdata and gives Linux's answer — one per pair per candle, degraded records apart, the archive checked against the record before it is used. Its docstring lists the eleven choices the entry left open. One bears on the draft, found writing it: the draft calls AGGRESSIVE's “close inside the band” today's rule, but today's AGGRESSIVE reads ACTIVE ENTRY ZONE, which reaches a whole band width either side of the band's midpoint; the script counts the band, as the entry words it, and prints today's reading beside it, for the draft's next version to choose. tests/test_measure_draft_gaps.py, new: 20 tests, no fixtures, on synthetic records only, the live logs never read, the expected values worked by hand; the golden test checks the script against numbers the draft and the entry computed apart from it. Against the code before, all 20 fail on FileNotFoundError, proving the file new and nothing else; twelve negative controls, each restored byte for byte. Suite 911 to 931, 755 / 165 without pandas_ta, 863 / 0 / 31 under run_tests.py. e939734's landing read from his disk. code_hash unmoved.

#272 1b82ed0's Landing, Its run_tests.py Deviation, and Entry 59

November 2, 2026

OBSERVATION — FILED FOR THE RECORD

“A guard that fails on a normal event teaches its reader to dismiss it.”

Commit 7ff8128, documentation only. Viktor reported that run_tests.py's first run at 1b82ed0's landing gave 862 / 1 / 31, the failure test_the_real_log_directory_was_not_written_to_or_deleted, and his re-run the predicted 863 / 0 / 31. Read again off his disk: each of the ten decision logs holds exactly one record logged after 22:05 on 1 November, the scheduled run of 18:05 (live, code_hash 816d4951…), their mtimes 18:05:22 to 18:06:02, so the task appended to five logs between conftest's snapshot and the guard's read, and no test reached the record; the 14:05 candle was missed, the computer off. The guard cannot tell the scheduled task's write from a test's: filed as entry 59 of the list, not in the sort. Claude's suggestion, his to change: a test fix, Claude's, XS, tests only, moving no code_hash. Found while rewriting, and corrected: a sentence in the account of the session from 17:49 on 1 November about e69d5bb's patch had come to name each latest commit in turn, through e8b42b8, d6fba83, 7c23886, e939734 and 1b82ed0; it now names e69d5bb, with a note of the five. What comes first next, by the ruling of 1 November: entry 59's fix. code_hash unmoved.

#273 Entry 59: the Real-Log Guard Watches the Suite's Processes, Not the Folder

November 3, 2026

FINDINGS — FIXED

“A folder cannot say which process wrote to it.”

Commit 5fe8522, two patches, tests and documents; the fix's shape was Claude's to settle. Since 17 October the scheduled task writes into logs/ every four hours, and the size snapshot tests/conftest.py took at the suite's start failed the guard on the task's own records (Entry #272). tests/suite_guard/sitecustomize.py, new, installs a Python audit hook (PEP 578) that records every open for writing, delete, rename, mkdir, rmtree, truncate, utime, link and symlink a process makes inside a logs-named directory of the root, in any casing; conftest installs it in the suite's process and puts its folder on PYTHONPATH, so every Python child the suite starts imports it; the scheduled task is not started by the suite, so nothing it writes is recorded. The size snapshot goes, and with it its two blind spots, a file rewritten to the same size and a directory deleted and rebuilt. What the hook cannot see is stated in its docstring: a child that is not Python, a Python child started with -I, -E or -S or without PYTHONPATH, an operation relative to a dir_fd, anything before conftest is imported. Evidence (Linux sandbox): 7ff8128, with a stand-in process appending to logs/ every two seconds, reproduced 862 / 1 / 31; patched, the stand-in running, 933, 757 / 165 and 865 / 0 / 31. Against pre-fix code the new test file fails only at import. Six negative controls; a first try at one left two variables stripped, broke nothing and passed — caught and redone. At his pytest step on Windows, 932 / 1: the capitalised-folder control expected Logs/ and the hook reported logs/, os.path.normcase lowercasing on Windows, which the Linux sandbox could not show — the shape of 6 September, the defect in the control's expectation, not the hook. A second patch, built on his applied files, makes the control expect the path as normcase reports it. 7ff8128's landing read from his disk. code_hash unmoved.

#274 5fe8522's Landing

November 3, 2026

PROCESS — RECORDED

“By the ruling of 1 November a commit owed comes first, and this is the one owed.”

Commit 7b93769, documentation only. Records 5fe8522's landing, two patches: GitHub's tip, read with no credentials at 16:51:28, 5fe8522, its tree the one its Verified-Tree line names; off Viktor's disk, the seven files each byte for byte an autocrlf checkout of the tip, four with the first apply's times (16:21:06) and three with the second's (16:39:44), the pytest cache rewritten after the second apply and holding the four new controls' ids, lastfailed {}; nothing differed. The one deviation, his first pytest run's 932 / 1 before the second patch, is recorded by 5fe8522 itself. Made as the commit owed under Claude's reading of the ruling of 1 November at the time, which the next commit corrected (Entry #276): one of the five commits of 1 to 3 November that recorded only a landing. code_hash unmoved.

#275 HISTORY Keeps What Each Rewrite of PHASE7_NEXT.md Removes, Not the Whole File

November 4, 2026

DECISION — ADOPTED

“Each commit added 400 to 440 KB.”

Commit b990c93, a script, a test file and documents. Entry 60 of the list for after the audit, found by Viktor and verified at 7b93769 in the session of 00:21 on 4 November, which made no commit, on purpose: every commit since 19 September appended the whole of PHASE7_NEXT.md to PHASE7_HISTORY.md, though the patch-delivery skill's section 9 said the session account. Measured on GitHub's blobs: HISTORY 16,679,872 bytes at 7b93769 against 12,223,465 at 8c810b1, eleven commits before. GitHub warns on a file over 50 MiB and refuses a push holding one over 100 MiB; at the pace of 1 to 3 November, an estimate and not a measurement, the warning in about 55 to 60 commits. Claude read what the dated record keeps as Viktor's to rule and the mechanism as Claude's; at 01:02 he asked for Claude's suggestion and answered “Go”. Ruled: each commit that changes NEXT appends one entry holding every line the rewrite deletes or replaces, word for word and in order, in fenced blocks labelled with their line numbers in the previous version, under a header pointing to the whole version in git; the 16.7 MB already in HISTORY stays. docs/build/check_history_move.py, new, checks it between a base commit and the working tree — HISTORY append-only, the blocks exactly the runs deleted or replaced, the labels and the git pointer right — and --emit prints the blocks to paste; this commit's entry, the first, holds 46 lines in 23 runs. tests/test_check_history_move.py, new: 18 tests, all failing against the tree before only because the script is absent; thirteen negative controls, one of which first passed because its test accepted any message containing “fence” — the test now expects the refusal's exact sentence. Not done, on purpose: the check is not run from the pre-push hook, which would need a mode comparing two commits. Suite 951, 775 / 165, 883 / 0 / 31. 7b93769's landing read from his disk. code_hash unmoved.

#276 A Landing That Went as Predicted Is Never Itself a Commit Owed

November 4, 2026

DECISION — ADOPTED

“This commit is made for its own reason, the fix.”

Commit b990c93, filed with the ruling of Entry #275 and carried from the session of 00:21 on 4 November: Claude's correction of its own reading of “first” under the ruling of 1 November (Entry #262), Viktor's to reverse. That reading had put first a commit owed, a landing to record among them, so a session with nothing else to do made a commit to record a landing. Corrected: a landing that went as predicted is recorded in the next commit made for its own reason and is never itself a commit owed, since a landing-only commit creates a new landing to record — the floor the ruling of 27 September removed for pushes. A deviation is still owed, and so is a ruling of his to file. Found by it, from the DECISIONS entry it files: five commits since 1 November recorded nothing but a landing — 248037d, which also recorded the skill's save, f6aa126, d6fba83, 7c23886 and 7b93769 (Entries #263, #265, #268, #269 and #274); they stand as made. What it weakens: a landing that went as predicted waits, unrecorded in the repository, until the next commit made for its own reason. b990c93 itself is made for the fix and records 7b93769's landing. The patch-delivery skill's sections 9 and 10 change with it, put to Viktor after its push, with the marker moved to b990c93. code_hash unmoved.

#277 Round 8's F2 Is Critical: the Risk Regime Read ADX

November 7, 2026

DECISION — ADOPTED

“The rubric's Critical has no exemption for a break that errs only toward caution.”

Commit 706e858, documentation only. Viktor's ruling of 7 November, at 00:33, by agreeing to Claude's suggestion: round 8's F2 breaks Tier 1 Item 14 (Risk Is Not Conviction) and is Critical. Round 8's run 2 had rated it Major; the triage of 5 October fixed it (fcb7823, verified fixed by round 8's auditor on 17 October) and left its severity unruled, the one item of the triage the audit's record needs before the send. Claude put the facts to him at 00:30 without a view, his position asked for first; he asked for a suggestion at 00:32. Read at fcb7823~1, the code before the fix: classify_risk_regime returned HIGH VOLATILITY RISK for an ADX under 20 whatever the volatility, and LOW RISK for LOW volatility with ADX 25 or more; AGGRESSIVE is allowed only under NORMAL or LOW RISK, so the chop test could change a live action and the LOW RISK test could not. From the evidence filed on 5 October, read from DECISIONS and not re-derived: volatility was HIGH or EXTREME on all 37 live runs, so no live action changed. Claude's reasons: the rubric's Critical is a Tier 1 break that could reach a live decision, which the chop test could, and ruling Major because it errs toward caution would add an exemption without an amendment; F4's third path was ruled Critical on 20 October on the same reading. The case against, given with them: round 8's own rating, since no live action moved. 0c7dec5, which put ADX into the regime, is co-authored by Claude, so the suggestion is its author's. README: F2 counted among the Criticals fixed and verified before it was ruled, and nine engine-code findings resolved instead of eight. The release gate does not change. bc250fc's landing read from his disk. code_hash unmoved.

#278 The Forward Measurement's Count: a Filed, Tested Script

November 7, 2026 — PROCESS — RECORDED

“That count cannot be re-run, and the reading of 14 November would have had nothing filed to be checked against.”

Commit 4ab59d7, a script under docs/build/, its tests and documents; no engine code. The reading of 31 October (Entry #261) was counted by a script in the sandbox that was never committed. docs/build/count_forward_measurement.py, new, written from the decision of 17 October on how the forward measurement is counted, its rules 1 to 4; its records are measure_draft_gaps.py's population, imported from it, so the two readings count the same records. Refusals and failures are read from main.py's “Engine returned error for” lines in the process log, never from ERROR lines; a “not yet final” one is a refusal; a candle with no counted record is missed. Ten choices the decision left open are in its docstring, among them that the window's last candle is given, never guessed, and that a second code_hash among the task's records stops it. Checked against the reading of 31 October on Viktor's ten decision logs and process log, staged off his disk: every figure that reading filed came out the same — 84 candles, 390 records, 450 missed (54%), 10 refused, 0 failed, LONG 56, SHORT 143, WAIT 60, NO-TRADE 131, the three degraded DOGEUSDT records. tests/test_count_forward_measurement.py, new: 24 tests, no fixtures, on synthetic records only. Against the code before, all 24 fail on FileNotFoundError, proving the file new and nothing else; twelve negative controls. Two defects of the first draft were found by its own tests before delivery. Found while rewriting NEXT, and corrected: 706e858's rewrite had left four of bc250fc's sentences saying “the commit that writes this line”. Interrupted at 01:34, the laptop off with the work in the sandbox only; finished from 04:03, the base files staged off his disk again. Suite 951 to 975, 799 / 165 without pandas_ta, 907 / 0 / 31 under run_tests.py. 706e858's landing read from his disk. code_hash unmoved.

#279 Entry 53's Design Draft, Version 2: Its Direction and Its Band Cannot Meet

November 7, 2026 — PROCESS — RECORDED

“A candle the draft calls BULLISH always closes above the band.”

Commit 94b7e96, documentation only. Version 2 of the draft, Claude's and not ruled, written before any of the reading of 14 November's data is read, so that the case against cannot be fitted to it: a new section 16, what the design as a whole weakens and the case against it (section 12, point 5). Found writing it, read in indicators/indicators.py at 4ab59d7: EMA_20 and EMA_50 are recursive and each slope is the one-candle change of its EMA, so a slope is above zero exactly when the close is above the previous candle's EMA, and the close is then above this candle's EMA too. A candle the draft calls BULLISH always closes above the band, BEARISH below it: section 4's AGGRESSIVE, the close inside the band, is never reached on either side, nor section 3.1's full distance points on a candle with a side, and the tier partly reads the EMAs' spread. Version 1 had also called AGGRESSIVE's condition today's rule, which reads ACTIVE ENTRY ZONE, not the band. Checked in the Linux sandbox on the pinned candles, through the engine's indicators with pandas_ta and through ewm alone: the identity held on all 400 rows of AEROUSDT and of BTCUSDT. Not fixed there: reconciling sections 2, 3.1 and 4 became section 12's point 7. Found while rewriting NEXT, and corrected: its “Before this commit” bullet still gave bc250fc's sentence. 4ab59d7's landing read from his disk. code_hash unmoved.

#280 The Redesign's Direction Is EMA50's Slope and the EMAs' Order

DECISION — ADOPTED

November 7, 2026

“We go with B”

Commit a11ffc4, filed in DECISIONS with the draft's version 3, a documentation script and its test; no engine code. Viktor's ruling of 7 November, at 05:34, by agreeing to Claude's suggestion: in the redesign's draft, direction is BULLISH when EMA50_Slope is above zero and EMA_20 is above EMA_50, BEARISH when both point down, else NEUTRAL; EMA20's slope is not read. It settles the draft's strength floor and reconciles the sections Entry #279 found cannot meet. Claude put two routes to him without a view, his position asked for first — A, keep the floor and measure entry quality's distance from EMA_20; B, take the floor from EMA50's slope and the EMAs' order — and suggested B when he asked, at 05:33. Due by a date: a rule the decision of 2 November measures had to be filed before the reading's data is read, so that decision's computation 1 (Entry #270) was amended in place, marked so and naming the new entry, which quotes the sentences it replaces; measure_draft_gaps.py and its direction test follow. Checked through the engine's indicators on the pinned candles: the sign of EMA50_Slope equals the sign of the close minus the previous EMA_50 on every row, so no close with a side is on the wrong side of EMA_50; sides AERO 342, 64 inside the band, BTC 307, 52 inside, against version 2's 278 and 255, none inside. Synthetic candles: they say nothing about the live pairs. Claude's errors, recorded: the counts first put to him in chat came from pandas' ewm alone, not the engine's ta.ema, stated with that method; and a first comparison of the index with the tip's tree split paths at their spaces, caught before it was reported. Against the pre-change script the rewritten test fails on one behavioural assertion; two negative controls. Suite unchanged, 975, 799 / 165, 907 / 0 / 31. 94b7e96's landing read from his disk. code_hash unmoved.

#281 Entry 53's Design Draft, Version 4: Which Distance STATUS Reads

PROCESS — RECORDED

November 7, 2026

“One distance is read in three units today, and the draft's d adds a fourth.”

Commit 542085f, documentation only. Version 4 of the draft answers the first half of its section 12, point 6, a question §3.1 had marked Claude's. Version 1's error, found scoping it and corrected: §3.1 said that entry.distance_from_zone sets STATUS. Read at a11ffc4 in models/entry_model.py, it does not: STATUS counts the close's distance from the band's midpoint in band widths — ACTIVE ENTRY ZONE within 1, NEAR ZONE within 2, APPROACHING ZONE within 3.5, else AWAY FROM ZONE — three numbers no ruling sets; distance_from_zone is the same distance as a percent of the close and decides nothing; and AGGRESSIVE reads the STATUS word itself. Proposed, Claude's and Viktor's to reverse: the tier reads a true-or-false field in the entry block, never the STATUS word, and STATUS is printed from that field, so a panel word cannot change an action. The field follows his reading of AGGRESSIVE, then still his: branch (a), close_inside_band, STATUS INSIDE, ABOVE or BELOW THE BAND, with no threshold; branch (b), in_active_entry_zone, today's words kept as display only. STATUS does not read d. Not first by a date: computation 2 of the decision of 2 November counts both readings. a11ffc4's landing read from his disk. code_hash unmoved.

#282 Entry 53's Design Draft, Version 5: the Second Pinned Decision, and Entry 61

PROCESS — RECORDED

November 7, 2026

“Each passes the first gate only because the short average goes unrecorded.”

Commit b874ec1, documentation only. Version 5 of the draft answers the rest of section 12, point 6, the second pinned decision's candle, as far as it can be answered before the reading of 14 November. Claude scanned first and put the finding, the scope and an estimate to Viktor before writing, by his rule of 6 September; he said go at 07:48. The scan, in the Linux sandbox with pandas_ta: the golden fixture cut to its first k candles, k from 60 to 450, each cut run as the golden test runs it and the draft's gates applied, gate 7 taken as passing, its column being synthetic; at k 450 it gives the draft's own path, gates 8 and 9 failing. Of the 336 undegraded cuts three pass every gate — k 171 and 173, SHORT, CONSERVATIVE, and k 234, LONG — none AGGRESSIVE under either reading, and each has 29 or 39 daily candles behind it, fewer than the macro average's 50, where ta.ema returns None and the ewm fallback averages the candles there are, recording nothing. None of the 156 cuts with 50 daily candles or more passes every gate. So none is proposed; instead a second series from a fixed seed by make_pinned.py's synth, with at least 300 4h candles before its decision candle, its seed chosen once F, the tier lines and his reading of AGGRESSIVE are set, the first seed stated before the search. The scan's script was not committed; §10.1 states its method. Filed as entry 61 of the list for after the audit, found, not fixed, not ruled: an EMA on a frame shorter than its length goes unrecorded; a live run asks for 100 daily candles, so AERO does not reach it, and a pair listed less than 50 days ago would — reasoned, not run. 542085f's landing read from his disk. code_hash unmoved.

#283 The Redesign's AGGRESSIVE Reads the Close Inside the Band

November 7, 2026

DECISION — ADOPTED

“We go with A then.”

Commit 82ea40e, filed in DECISIONS with the draft's version 6; documentation only. Viktor's ruling of 7 November, at 19:47, by agreeing to Claude's suggestion, not by writing his position first: the draft's §4 AGGRESSIVE needs the close between EMA_20 and EMA_50, ends included, not today's ACTIVE ENTRY ZONE, within one band width of the band's midpoint. With it §3.1's branch (a) of version 4 (Entry #281): the tier reads entry.close_inside_band, never the STATUS word, and STATUS prints a position, INSIDE, ABOVE or BELOW THE BAND; NEAR, APPROACHING and AWAY go with their 2 and 3.5. The tier lines, 70 and 50, and the rest of the design stay unruled. Claude's reasons: it is the pullback hypothesis ruled on 25 October, stated with no distance in it; it adds no number, where (b)'s band width is set by no ruling; (b) labels as AGGRESSIVE closes outside the band, on a LONG up to half a band width above EMA_20; and ruled before any of the reading's data is read, it cannot be fitted to its counts, computation 2 counting both readings. What it weakens: a close in today's ACTIVE ENTRY ZONE beyond the fast average would read LONG or SHORT, not AGGRESSIVE; AGGRESSIVE may be rare when the band is narrow; the tier changes no position size, every trade being one of ten equal slices. Claude's error, recorded in the entry: the third reason as first sent also said (b) would label AGGRESSIVE a LONG closing below EMA_50; under the direction of Entry #280 a BULLISH close never does, which was stated without being checked against the draft's §16.4; corrected in chat at 19:49, before the filing, after his ruling. The session had chosen these Notes first and named this question his, after Claude's work; a ruling of his to file is a commit owed, so it came first. The decision of 2 November is unchanged: its computation 2 already reads AGGRESSIVE so. b874ec1's landing read from his disk. code_hash unmoved.

#284 The Redesign's Entry Quality Has No Structure Multiplier

November 7, 2026

DECISION — ADOPTED

“Okay, go with A as your recommend it.”

Commit 7dac328, filed in DECISIONS with the draft's version 7; documentation only. Viktor's ruling of 7 November, at 22:19, by agreeing to Claude's suggestion, not by writing his position first: the redesign's entry quality is D + H with no structure multiplier. BOS ×1.05 and CHOCH ×0.90 go, and with them structure_multiplier, CONFLUENCE_BOOST_MULT, CONFLUENCE_PENALTY_MULT and SCORE_CEILING; D + H cannot leave 0 to 100, so no clip is needed. The rest of the design stays unruled. It settles the draft's §3.3, which the ruling of 25 October had left standing unless he ruled otherwise. Claude's reasons: the tiers come from entry quality only, and a ×0.90 that moves a setup across 50 or 70 is the shape of the two multipliers already removed; a multiplier changes a decision only when 10% happens to carry a score under a line, where a gate is the honest form if a CHOCH should stop a trade; a CHOCH gate now would be a rule with nothing measured behind it; and, ruled before any of the reading of 14 November's data is read, the choice cannot be fitted to its counts. Claude's error, recorded in the entry: the case against, as first sent, said structure is already a gate in the draft and that a CHOCH bar the gate let through would now score as well as a clean one, the gate taken from the draft and flagged as not re-read. Read before filing, at 9db4ae2: gate 4, signal_blockers in models/entry_model.py, compares only the structure regime and never the swing sequence, so the gate never sees a CHOCH, and with this ruling nothing in the design reads one. Claude said so and asked again. His first “We go with A.” (22:07) came before the correction; “Go with A.” (22:17), then “Stop” and a request for Claude's suggestion, and the ruling (22:19) after it. It came first as an item of his, no commit being owed and all of Claude's work waiting on the reading: the decision of 2 November's computation 2 computes entry quality with no multiplier, so the question had to be settled before the data is read. What follows: the draft's version 7, §13.2's four leaves that were his now removed (85 removed); a CHOCH count for the reading, Claude's (Entry #285); whether a CHOCH against the side becomes its own gate, Viktor's, with the whole design. 9db4ae2's landing read from his disk. code_hash unmoved.

#285	The Reading of 14 November Counts the CHOCHs Against the Side: Computation 7 November 7, 2026	DECISION — ADOPTED

“A threshold set here would make that ruling in advance.”

Commit 340ee8d, an amendment to the decision of 2 November (Entry #270), with docs/build/measure_draft_gaps.py and its tests; no engine code. Claude's, Viktor's to change, made in the session that opened at 22:50 on 7 November, before any of the reading's data was read; it carries out the second point of what follows from the ruling of Entry #284. With the multiplier gone nothing in the redesign reads a CHOCH, and the ruling asked that whether one against the side becomes its own gate come to Viktor with a number fixed before the data is read. Computation 7: of the records with a side, how many have a swing sequence (structure.sequence) that is a CHOCH against that side — CHOCH BEARISH (POSSIBLE REVERSAL) against LONG, CHOCH BULLISH (POSSIBLE REVERSAL) against SHORT, compared whole, as today's multiplier compares them — and how many of those pass every other gate, gates 0 to 9, which is a draft trade action. UNKNOWN, which the structure engine writes when no sequence was computed, and a missing or non-string sequence are not measured; NONE and every other value are a sequence that is no CHOCH. Unlike computations 1 to 6, nothing counts against the draft: the count informs a ruling on a gate the draft does not have. What it weakens: it reads today's swing detection only; telling a reversal the EMAs have not yet turned on from noise needs what price did afterwards, which is backtesting, closed; one population of about four weeks, the count likely small; and it is its author's, Claude having written the draft and suggested removing the multiplier. The script reports 7 CHOCH; computations 1 to 6 unchanged. tests/test_measure_draft_gaps.py: two tests, no fixtures, neither needing pandas_ta; against the script before, both fail, an AttributeError and a KeyError, proving the code new and nothing else; five negative controls. Found while rewriting NEXT, and corrected: a sentence on the session of 3 November from 16:50 said “the commit that writes this line” from 7b93769 to 7dac328; it now names 7b93769. Suite 975 to 977, 801 / 165 without pandas_ta, 909 / 0 / 31 under run_tests.py. 7dac328's landing read from his disk. code_hash unmoved.

#286	The Redesign's Exit Watch Drops Its Exhaustion and Reversal Flags November 9, 2026	DECISION — ADOPTED

“Go.”

Commit 6bb82cc, filed in DECISIONS with the draft's version 8; documentation only. Viktor's ruling of 9 November, at 08:38, by agreeing to Claude's suggestion, not by writing his position first: in the redesign the exit watch has no exhaustion flag and no reversal flag, and nothing replaces them. Its other flags stay as they are: the divergence, the CHOCH against the side, the HVN and LVN proximity, the SuperTrend flip, the bias flip and the Target 1 note. The rest of the design stays unruled. It settles an item the draft's §0 had marked his since version 1. The exit watch (build_exit_watch in models/exit_model.py, read at 82b824d) prints advisory flags for whoever holds a position and decides nothing; the two flags read outputs of compute_trend_health, whose score the ruling of 25 October removes. Exhaustion fires when the candle's range is wider than the previous candle's and ADX is under 20, or when the trend health score is under 35 and ADX under 15. Reversal fires against the raw bias when reversal_strength reaches 40: 25 for a divergence, 20 for an HVN within 1.5% of the close (10 within 3%) and 20 for exhaustion, the sum cut to 40% unless momentum is EXTENDED or EXTREME or the trend is exhausted. Claude put three options: (A) both go, as the draft proposed; (B) the reversal flag goes and exhaustion's first clause stays, relabelled; (C) both rebuilt with new weights. At 08:36: “Suggestion?” Claude gave (A) and said it would file it if he went with A. At 08:38: “Go.” Read as (A), Claude's reading, his to correct. Claude's reasons: every firing of the reversal flag needs at least two of its parts, and two of them, the divergence and an HVN within 1.5%, are flags of their own, so what it adds is a weighted sum whose weights and line at 40 no ruling or evidence backs; half of the exhaustion flag needs the trend health score, which goes, and the other half, ADX under 20 on a wider candle, is a weak trend, not stretched momentum, so its text does not say what it measures, and relabelled it would be a new flag, which belongs with the whole design; neither flag decides anything, so no action, tier or gate moves, the golden exit_watch holding only the Target 1 note and no test naming either flag's text; and it was ruled before any of the reading of 14 November's data is read. Claude's error, recorded in the entry: reading the exit watch in the code found two of its own statements wrong. The draft's §2, versions 1 to 7, named only the divergence and SuperTrend flags as what the exit watch keeps. And the draft's §2, §3.3 and §16.2, the ruling of 7 November and the amendment of 7 November say nothing in the design reads a CHOCH, where the exit watch's CHOCH flag reads the swing sequence and stays; it holds for the decision, where no gate, tier or score reads one, and not for the panel. The two dated entries are not edited; the draft's version 8 and NEXT's present-state lines are corrected. Entries #284 and #285 of these Notes carry the same statement from those entries, and stand as written with this correction beside them. What it weakens: whoever holds a position loses two warnings; the divergence and an HVN within 1.5% stay on flags of their own, and what goes is their combination, the HVN band from 1.5% to 3%, momentum mode's part and the one warning of a weak trend on a wider candle, which no other flag gives; whether the panel prints ADX elsewhere was not checked; and the suggestion is its author's, Claude having written the draft. It came first as an item of his, no commit being owed, Claude's work waiting on the reading and the Notes current: of his items in the draft, the only one needing none of the reading's data. 82b824d's landing read from his disk. code_hash unmoved.

#287 Amazon's Audit Goes Through OpenRouter to Amazon Bedrock

November 9, 2026

DECISION — ADOPTED

“It is not set up. Won't be for some while.”

Commit e4363d5, filed in DECISIONS; documentation only. Claude's decision of 9 November, under the delegation of 4 October, Viktor's to reverse; its second point he replaced the same day (Entry #288). With no commit owed, Claude's work waiting on the reading of 14 November and every item of Viktor's in the redesign draft needing the reading's data, Claude brought him the Bedrock billing, an item of his on the audit's path. At 09:42: “It is not set up. Won't be for some while.” At 10:23 he asked whether OpenRouter has an alternative to Amazon's Bedrock, and until 12:02 the session read what Amazon serves and where. Nothing from the project was sent to any Amazon model, and none was called: every read was of a public listing. What was found: Nova 2 Pro, the roster's model for Amazon since 30 September, is still a preview, on Bedrock since 2 December 2025 with no general availability announced; Nova 2 Lite is generally available, on Bedrock and through OpenRouter. Viktor's curl of OpenRouter's endpoints API for amazon/nova-2-lite-v1 at 12:01, pasted whole: two endpoints, both Amazon Bedrock, a context of 1,000,000, max_completion_tokens 65,535, reasoning supported, the tokenizer named “Nova”. Vercel's AI Gateway serves it from Bedrock alone. Ruled out: Nova Pro 1.0, too small; Nova Premier 1.0, without reasoning by one listing and, by one secondary site, retired on Bedrock; “Nova-3”, Deepgram's speech-to-text model. Except the endpoints read, the facts came from secondary sites through a search tool, not AWS's own pages, and the entry says which. What was decided: (1) the route is OpenRouter with the provider pinned to Amazon Bedrock and no fallbacks, as round 8's pin is written; the model is still served by Amazon from Bedrock and only the billing goes through OpenRouter, so the Bedrock billing is off the audit's path, and Vercel's gateway is the fallback route if OpenRouter's endpoint returns no reasoning tokens. (2) The model is chosen at the send: Nova 2 Pro if it can be reached then on a route the send script uses, with reasoning and an output cap no smaller than Nova 2 Lite's; otherwise Nova 2 Lite. (3) The checks before the send are round 8's four: the served context and output cap, lineage, reasoning by the send script's probe, and the count with the model's own tokenizer. Why: Amazon is the last lab clean for a full round, and spent on Nova 2 Lite while Nova 2 Pro might be served by the send it would go to the weaker model; the send cannot come before the redesign is built, so the rule costs no time. Claude's wrong turns, recorded in the entry: in the chat it first put Nova Premier second, then Nova 2 Lite first; the rule replaced both. Not checked, Claude's before the send: whether reasoning tokens count against the 65,535 (round 8's run 2 used 68,464, and a cut-off reply counts as a run); the count with Nova's own tokenizer; and the endpoint's data policy. Viktor's: how the fit is shown if no count is possible without the package reaching Amazon, which would change the ruling of 26 September, and accepting the audit's result. What it weakens: if Nova 2 Pro is not served by the send, Amazon's one full round goes to the smaller model; an output cap half of round 8's 131,072; facts from secondary sites; and the decision is its author's, Claude having written most of the code Amazon will audit and chosen its examiner, as the delegation allows. 6bb82cc's landing read from his disk. code_hash unmoved.

#288 Amazon's Audit Sends Nova 2 Lite, Through OpenRouter's API

November 9, 2026

DECISION — ADOPTED

“We are to use Nova 2 Lite with Openrouter through API”

Commit a60d66b, filed in DECISIONS; documentation only. Viktor's ruling of 9 November, at 12:38, his own position, given while Claude was checking e4363d5's landing: the model of Amazon's audit is Nova 2 Lite (amazon/nova-2-lite-v1), sent through OpenRouter's API. It replaces point 2 of Claude's decision of the same day (Entry #287), which chose the model at the send, Nova 2 Pro if it could be reached then and otherwise Nova 2 Lite, and was his to reverse; Nova 2 Pro is not sought at the send. His reasons, at 12:40, in his words: “We use Nova 2 Lite since the Pro version is not released, we will use openrouter because it hosts Nova 2 Lite, Vercel was just a backup.” Claude had recommended the rule and was overruled. Its feedback, given at 12:39, before his reasons. Against: Amazon is the last lab clean for a full round, and the ruling sends that round to the smaller model for certain, where the rule cost no time, the send not coming before the redesign is built. For: certainty now, the package and the send planned against Nova 2 Lite's limits today with no choice left for the send; and a Nova 2 Pro reachable only as a preview through his own AWS account would not have met the rule anyway. Unchanged: the output cap of 65,535, half of round 8's, and whether reasoning counts against it, unchecked. His reason agrees with the decision's reads, which found Nova 2 Pro still a preview; those came from secondary sites through a search tool and were not repeated. What stands from the decision: the route, OpenRouter with the provider pinned to Amazon Bedrock and no fallbacks; Vercel's gateway as a backup route and no more, “just a backup” read as keeping it as that, Claude's reading, his to correct; the checks before the send; and everything the decision left not checked, Claude's before the send. The Bedrock billing is off the audit's path entirely, Nova 2 Pro no longer being sought. The decision's entry is not edited. What it weakens: Amazon's one full round goes to Nova 2 Lite even if Nova 2 Pro becomes generally available before the send; a fail then takes the project to the fallback of 17 October, a scoped round and then a wait for a clean lab. The roster of 17 October, outside the repository, names Nova 2 Pro for Amazon; independence is counted by lab, so Amazon stays clean, and the roster's next reissue carries the change. Viktor's still: how the fit is shown if no count is possible without the package reaching Amazon, and accepting the audit's result. It came first as a ruling of his to file. e4363d5's landing read from his disk. code_hash unmoved.

Document History

Version	Date	Notes
v1.0	August 25, 2026	Converted from the standalone Engineering Note #1 into this standing log, at Viktor's request. Entries 1–3 carried over unchanged, split from one combined note into three individually tagged entries.
v1.1	August 26, 2026	Entries 4 through 9 added, tracking the Constitution's Rev 3 through Rev 5 and the new Tier0 Companion document. No earlier entry rewritten — Entry #9 corrects scope drift introduced in Entry #8's phrasing without altering Entry #8 itself, per this document's own no-silent-edits rule.
v1.2	August 26, 2026	Entry #10 added, closing Entries #1 and #6 — the credential gap is adopted into Constitution Rev 6 as Tier 1 Items 18 through 21, with operational detail in the new Phase7_Credential_Security_Protocol.pdf, and the audit now has a named independent auditor. New DECISION — ADOPTED status tag introduced for entries that record a decision actually made rather than a candidate awaiting one. Entries #1 and #6 left unedited; Entry #10 is where their closure is recorded.
v1.3	August 26, 2026	Entry #11 added: Viktor named Grok as the independent auditor for Steps 3, 4, and 8, once released, recorded in Constitution Rev 8. Answers the authorship-vs-judgment-independence question Rev 7 had just sharpened via Reviewer 4's identity.
v1.4	August 26, 2026	Entry #12 added, and format changed to hold it: a new gold-framed highlighted_entry_box treatment, used only when Viktor asks for an entry to be marked especially important. Documented in “How This Document Works” as a rare, deliberately non-standard marker rather than a new status tag. First use: a reflection on the audit loop's design, and the two places its soundness still depends on discipline rather than mechanism.
v1.5	August 26, 2026	Entry #13 added: criteria for evaluating any future auditor substitution — independence as a gate, power and cost as the actual trade-off, cost mattering because Step 8 makes the audit a recurring expense rather than a one-time one. Extends Entry #11 without editing it; Grok remains the named plan.
v1.6	August 26, 2026	Entry #14 added: Viktor ratified Constitution Rev. 8. Recorded in the Constitution's own Version History as a new RATIFIED row, not a Revision 9 — no rule text changed, only status. Scope freeze now in force for real; Step 2a of the audit sequence may begin.
v1.7	August 26, 2026	Entry #15 added: Gemini's assessment of the finished, ratified Constitution, logged and gold-highlighted at Viktor's request — second use of the highlighted_entry_box treatment, still rare. Not a structured finding and not the audit; flagged as such in the entry itself.
v1.8	August 26, 2026	Entry #16 added, recording how Entry #15's assessment was obtained — solicited by the party being evaluated — and the critical counter-review Viktor then commissioned. Entry #15 left unedited. Deliberately not gold-framed: the highlight marks Viktor's request for emphasis, not epistemic weight, and a counterweight entry that competed for attention would defeat its own point.
v1.9	August 26, 2026	Entry #17 added: Step 2a complete. Claude assembled the audit package from the actual working repository (19 source files, 4 evidence files not present in the public repo) at Viktor's instruction. Zero findings written, per Step 2a's own rule. Notes that Entry #3's file list is now superseded by the package's manifest, without editing Entry #3 itself.

Version	Date	Notes
v1.10	August 26, 2026	Entry #18 added: before running Step 3, Viktor noted Grok had already seen the Constitution in an earlier, separate conversation. Named as an independence risk on the same grounds as Entry #16, resolved by running Step 3 in a fresh conversation with no shared memory of the earlier one. The public GitHub repository was considered and deliberately left published.
v1.11	August 27, 2026	Entries #19, #20 and #21 added. #19 logs Grok's review of the Constitution as an external assessment, not Step 3, with its factual claims verified and its weaknesses named. #20 records Step 3 being paused and the auditor plan changing from Grok to a panel of models with no prior involvement, run pay-per-use; no revision of the Constitution was needed, since Rev. 8 named Grok as a plan rather than a requirement. #21 records an internal contradiction found in Claude's own drafting of the Minimum Viable Audit gate, deliberately recorded rather than repaired. Two new Constitution Version History rows (AUDITOR, DEFECT); no rule text touched and the register still frozen at 21 / 7 / 10 / 6.
v1.12	August 27, 2026	Entries #22, #23 and #24 added, covering the audit actually being run. #22 records Run 1, the blind source-only review by DeepSeek V4 Pro, whose ten findings were each checked against the real source and held. #23 records the first Step 3 attempt truncating inside its own reasoning at a default token ceiling, and the resulting split of the register across three scoped runs; the corresponding wording change is recorded as a new Section 6 in Phase7_Audit_Execution_Instructions.pdf rather than made silently. #24 records Run A, the Minimum Viable Audit gate: Items 2 and 18 Compliant, Items 3 and 6 Non-compliant, the Entry #21 DEFECT contradiction resolved by the auditor in favour of the four-item gate, and a false claim found in Claude's own Step 2a package. #24 also records, and withdraws, a false accusation Claude made against the auditor — a citation Claude called invented turned out to be accurate and located in the other evidence file, which Claude had not searched. Two genuine disagreements between Claude and the auditor are recorded unresolved and go to Viktor. No rule changed; register still frozen at 21 / 7 / 10 / 6, and the freeze stays in force until every Tier 1 item has a finding.
v1.13	August 27, 2026	Entries #25 and #26 added. #25 records Run B — the seventeen Tier 1 invariants outside the gate — with every substantive claim verified against the source, two precision errors in the auditor's evidence noted, two Compliant ratings disputed by Claude over the indicator cache, and the observation that two material Item 14 defects appear in Kimi's truncated first attempt but in neither run that completed. #26 records the scope freeze lifting: all 21 Tier 1 items now have findings, meeting the mechanical trigger Reviewer 4 forced into Next Steps at Revision 7. A new AUDITED row was added to the Constitution's Version History recording the same fact; no rule text was touched and the register stands at 21 / 7 / 10 / 6, no longer frozen. Four adjudications remain open for Viktor.

Version	Date	Notes
v1.14	August 27, 2026	Entry #27 added, recording Run C — Tiers 2, 3 and 4 — which completes the register: all 44 rules now carry a finding. Twenty-three items graded; four Tier 2 Non-compliances including a module that rewrites its caller's DataFrame in place, and a dependency manifest that omits two packages the code imports while declaring one nothing imports. Recorded as a method difference: Run C had web access and fetched the public repository, which Runs 1, A and B did not. Claude repeated that fetch independently and closed two Unknowns — no tags or releases (T3-7 becomes Non-compliant) and twenty-three commits on master (T3-9's depth question resolved). One contradiction between Run B and Run C is recorded with Run B's reading preferred, on source evidence.
v1.15	August 27, 2026	Entry #28 added, recording a hostile external review of the Constitution commissioned by Viktor from an uninvolved model. Two of its findings adopted into the Constitution the same day — a release gate, since the freeze-lifting definition established that the audit ran but attached no consequence to what it found; and amendment control, since the now-live amendment process would permit a Tier 1 invariant to be weakened by a short note reviewed only by Claude and Viktor. A front-matter statement of what the engine is was added alongside them. A third finding, Item 20 not covering crash-reporter capture of the process environment, is recorded and left unfixed: it requires amending a Tier 1 invariant, which the newly adopted rule says Claude may not do without external review. The review's remaining nineteen proposed invariants are declined as calibrated to a production trading system this engine is forbidden to become, and go to Future Amendment Candidates. Register unchanged at 21 / 7 / 10 / 6.
v1.16	August 29, 2026	Entries #29 through #33 added, covering the work between the audit closing and the remediation sequence arriving. #29 records Phase A — eighteen tests built and run, five passing, thirteen failing as written acceptance criteria, and the golden-path test that was shipped unverified and turned out to be wrong. #30 records the three defects that first run found, which four audit passes across three models had missed: a severity escalation on the clone failure, and two findings nobody had. #31 records the model-independence table in one place for the first time, and the rule that independence is tracked at the lab rather than the checkpoint. #32 records machine learning put on ICE with five written conditions for revisiting. #33 records Step 5 running on GLM 5.3, nine of nine source claims verifying, and Claude coming one sentence from a third false assertion of absence. Register unchanged at 21 / 7 / 10 / 6; the adjudications Step 5 asks for remain open.

Version	Date	Notes
v1.17	September 2, 2026	Entries #34 through #42 added, covering the six days between the audit closing and this row — a period this log had not recorded at all. #34 the machine rebuild and the golden baseline surviving it across a numpy major version. #35 Step 8, the independent re-audit: five Criticals, ten further findings, release gate not met. #36 remediation batches 1 and 2, including 46 tests that returned where they should have skipped. #37 items 3, 11 and 14 delegated by Viktor and ruled. #38 the macro degradation, dependency pinning, the four unguarded tests and two stale docstring labels. #39 Finding 3 — a Critical that had never been entered into any roadmap, found by reading the audit report rather than the summary derived from it, with a defect class wider than the two instances the report named. #40 the status sweep: three findings already closed by commits that never named them, and a count wrong by three. #41 the Step 8 auditor not being on the clean list, recorded rather than quietly noted. #42 the full document set read end to end, correcting four things no working document held. Register unchanged at 21 / 7 / 10 / 6, no longer frozen. Two amendments owed, neither adopted; the release gate stands closed on Item 6.
v1.18	September 2, 2026	Entries #43 through #49 added, covering the day the last Critical closed and a new one opened. #43 Findings 6 and 7 — input hashing, a raw-candle archive pruned at ninety days on Viktor's ruling, and the Item 6 lineage chain persisted. #44 a defect that shipped: a path separator that made the golden snapshot match only on the platform it was baselined on, caught on Viktor's machine on the first run. #45 the direction-source Critical — a CONSERVATIVE LONG printed over a stop above price and three descending targets, found by running the engine on live data, and missed by an independent 44-rule audit plus four passes across three models. #46 an unwritable log directory destroying a whole run, with Claude recommending refusal and being overruled. #47 a generated fixture that lets the suite reach the state that hid #45. #48 the session / training / lineage distinction ruled, and the independence ledger corrected from the billing log — Mistral had been wrongly listed clean for a week. #49 the round-2 audit package generated rather than assembled, and eight tools left enabled on the first attempt at it. Register unchanged at 21 / 7 / 10 / 6. Every Critical the Step 8 audit raised now has a fix that has landed; the release gate stays shut until a re-audit has seen them.
v1.19	September 3, 2026	Entries #50 through #56 added. #50 the re-audit's third failed attempt and the Constitution shipped as text, because two providers on two days could see its filename and not its contents. #51 eleven defects returned by an audit that produced no report, every one checked against source before it was believed. #52 the run hash that did not cover the risk multipliers — Item 6 broken inside the fix written to satisfy it. #53 a test whose failure message described a live defect it could not detect. #54 the volume-node fabrication reached through the producer after Finding 3 closed it at the consumer. #55 two defects found by reading the panel in one morning, one of them shipped by Claude an hour earlier and passed by all 226 tests. #56 the handover check, and the project goal split into a tagged portfolio milestone and a later finished engine. Suite 196 to 226. Register unchanged at 21 / 7 / 10 / 6. Earned rules 30 through 34 written the same day.

Version	Date	Notes
v1.20	September 5, 2026	Entries #57 through #70 added, covering 4–5 September — the three open rulings built and the whole of the “What is left to fix” table closed, across eight patches. #57 ruling 3, the bias engine inferring trend direction from the sign of continuation strength and reading a real trend as flat, measured across 9,800 live bars before the ruling rather than after. #58 rulings 1 and 2 — trend health, a magnitude, offered in reason strings as directional support, and a minimum bias strength kept as its own fingerprinted constant rather than folded into an existing threshold. #59 three patches diffed against a base that had already moved, and earned rule 30 violated in a test written to guard its own class. #60 the usage discipline Viktor delegated: reasoning goes where it survives, verification is not cut. #61 two fabrication constants that survived two passes because they were unreachable, one of which would have killed the run rather than degraded it. #62 a fallback whose docstring claimed an equivalence it did not have, and the test that had asserted on that claim for as long as it existed. #63 a degraded-dependency control that reported two errors and executed nothing, and a module-level object that defeated a deliberate lazy import. #64 correlation and beta paired by position after both timestamp indexes were discarded — correct by accident, and the printed label moving in 4 of 31 measured positions once the series differ by one bar. #65 an entry zone printed with its lower bound above its upper, found by reading the panel, and a close price that could not be read becoming $1.00. #66 five entry sub-scores summing to 39 under a printed total of 45.18, and the clip that was hiding a component list adding to 102 against a docstring saying 100. #67 the entry block's shape declared in four places, recorded rather than restructured inside a patch about something else. #68 the four-for-four generalisation: the fabrications that survive audits are the ones on paths a healthy run never takes. #69 two rejected CONNECTs to an exchange from a suite in which every route was pointed at a dead port, source unidentified, recorded rather than chased. #70 suite 226 to 319 passing and the fix table empty, with the release gate still shut because unresolved means fixed and re-audited. Register unchanged at 21 / 7 / 10 / 6. Earned rules 35, 36 and 37 written across the same two days.
v1.21	September 5, 2026	Entries #71 and #72 added, both corrections to this project's own record rather than to its code. #71 applies earned rule 29 a second time: the provider's billing log shows three Qwen calls followed by Kimi K3 on 2 September, so Entry #50 has the last two attempts of round 2 in the wrong order and the eleven observations of Entry #51 came from a Qwen run that preceded Kimi rather than followed it. The same log establishes that no attempt at round 2 ever received the UPLOAD_THESE folder — it was written after the last of the four calls — and corrects Entry #22, which names a DeepSeek checkpoint the bill says was never called. Claude's own inference from file modification times is withdrawn in the entry. Luna Pro's call count and Kimi K3's spent status are recorded as open questions, not resolved. #72 records a defect found by searching the round-3 package before sending it: Section 2 of the reviewer instruction told the auditor to stop if the code bundles contained prior-audit outcomes, and they contain a stated verdict for seven of the forty-four rules, three counts of Criticals, and two named prior reviewers. Viktor ruled the condition scoped to the Constitution file and the bundles' contents counted and disclosed, with the audit narrative to move out of code comments into this document later — not inside the audit it would affect. Register unchanged at 21 / 7 / 10 / 6. The release gate remains shut.

Version	Date	Notes
v1.22	September 5, 2026	Entries #73 through #82 added, covering the afternoon and evening of 5 September — the day's second half, in which nothing in the engine changed and most of what this project believed about its own audit did. #73 the package-search step created by Entry #72 finding, on its first run, that Entry #72's own disclosure had under-counted what the bundles leak, in three separate ways. #74 a commit count read off a generated file three hours after it was generated: the third instance of a figure treated as current because the artifact holding it looked trustworthy. #75 the largest correction this project has made — the four transcripts named for Qwen are Kimi's, identical for 57,631 characters, and the passage where the model says so never reached the repository, so Entry #71's own correction was valid reasoning to a false conclusion. #76 an Auto Router left on by accident sending the round-3 package to GLM 5.3 Flash, which returned a complete report; ruled to be round 3 rather than discarded, with one finding verified as under-scoped in a way that reaches the trade-authorization gate. #77 the ledger rebuilt from all 723 provider rows: Z.ai not clean and so the reviewer that ran was not independent, Kimi cleared on the check that mattered with its exposure thirteen times larger than recorded, and several paid-for outputs that were never saved. #78 the handover check's structural blind spot — ignored files do not appear in git status — and reviewer responses moved to a tracked directory where the safe location is the only one on offer. #79 rev 5: the instruction stops telling the reviewer who it is, and Part 8 is removed because its comparison document was the reviewer's own reasoning. #80 three PDFs checked against the repository before deletion, one holding a paragraph of ruling reasoning that existed nowhere else. #81 a commit message stating an impossibility where an undone task belonged, found because Viktor asked why. #82 Kimi ruled for round 4 on two grounds, with a third — that it is an important model — argued out and recorded as rejected. Register unchanged at 21 / 7 / 10 / 6. The release gate remains shut, and the round-4 package is built and waiting to be sent.

Version	Date	Notes
v1.23	September 6, 2026	Entries #83 through #90 added, covering the sixth session of 6 September — the round-4 confirmation run and the four fixes it and the two reviewers' reports led to, plus the two defects found by running the suite and reading the document set rather than by review. #83 the Section 11 confirmation run, made before any fix, confirming Finding 1 end to end and finding a second Item 6 gap neither reviewer filed. #84 Finding 1 fixed: a producer's not-measured value now passes through instead of being defaulted, and a failed merge's record keeps its lineage. #85 GLM F-7 fixed as extended: a risk block that was never assessed could authorize the engine's most aggressive action, closed by a validator that reads the verdict rather than assumes one. #86 Kimi Finding 5 item 6 fixed, and the live run it and F-7 were both owed: a swing-structure field invented its own current price as a location, the third time this class of defect was patched at a consumer while the producer kept inventing. #87 Kimi Finding 3 fixed — a source hash over the parse tree of every engine file, replacing a list of names nothing of that shape could hold — and a live run at 10:51 that found a second, independent route to the false-sentence class Kimi Finding 2 describes, filed as input to a decision rather than fixed. #88 the suite was destroying its own decision record: four routes, one of them a directory delete, found by running a clean clone rather than by reading, with the 10:04 live run's raw record among what could not be recovered. #89 the document set audited end to end: the audit-package record verifies byte-for-byte, no build script could write into the repository, and the semantic half — the eight PDFs read against current state — remains undone. #90 the build scripts fixed: seven of nine now write into the repository directly, the other two fail early and say why, and publishing the entries in this row became one command. Suite 319 to 412 passing across the batch; the golden snapshot moved once, in exactly the five fields predicted. Register unchanged at 21 / 7 / 10 / 6. The release gate remains shut — nothing that landed this batch has been re-audited, and eight decisions remain open, all Viktor's.
v1.24	September 9, 2026	Entries #91-#93: fourteen-item sweep, Kimi Finding 2 fix, Claude outputs fingerprint exclusion.

Version	Date	Notes
v1.25	September 11, 2026	Entries #94 through #97 added, covering the fourth commit batch of 11 September. #94 the four qwen_reasoning_*.txt transcripts moved from the repo root into docs/ (bb6d222), discharging the hold noted at Entry #93. #95 this document's own body reconciled with 76380ec a second time, in a new shape (4705d03): the 7-9 September Grok session had updated the head block and this builder but left “Open - decisions” #3 and “Open - work” #2 describing the sweep and Kimi Finding 2 as open after both had landed, docs-only and code_hash-neutral by directory exclusion. #96 docs/Phase7_Sweep_Session_Commands.pdf committed after sitting untracked on Viktor's machine (e75f7be). #97 Item 14 / decision 2 ruled and closed (0c7dec5): classify_risk_regime now reads ADX instead of trend_health, so the risk gate no longer shares its largest input (0.30 of bias_score) with the conviction number it is meant to check independently; REGIME_LOW_TREND_HEALTH/REGIME_HIGH_TREND_HEALTH renamed to REGIME_CHOP_ADX = 20.0 / REGIME_STRONG_ADX = 25.0, values already load-bearing in trend_health.py for the same two market states. code_hash moved and was confirmed on three trees, 47bc557119bb0350... to 2741062fae070f61..., matching Viktor's own Windows value. Suite 412 to 421 passed under pytest (0 failed both sides), 345 to 354 passed under run_tests.py with the watched 29-error count unmoved. Golden re-baselined at six sites: four predicted (module constants, risk_inputs, run_hash twice) and two not (the archive filename, built from run_hash, which should have followed the run_hash prediction mechanically and did not) - recorded as an incomplete prediction rather than a defect. No decision field moved on the pinned fixture, which exits before classify_risk_regime is reached; a live run the same day (AEROUSDT 4h, 22:50) did not exercise the new branches either, for two independent reasons, and is recorded as that rather than as validation. Found by reading the panel and left unfixed: RISK REGIME prints without the ADX number now deciding it, a transparency regression of the same class as GLM F-5. Register unchanged at 21 / 7 / 10 / 6. The release gate remains shut - nothing landed since round 4 has been re-audited.

Version	Date	Notes
v1.26	September 11-12, 2026	Entries #98 through #103 added, covering the round-5 send through all three Criticals landing as fixes. #98 (4d99cdb) GPT-6 Astra ruled for round 5 after Viktor's billing-export check cleared both of Luna Pro's exposed sessions as variant=standard, no training routing; send_audit_round.py and build_audit_package.py repointed, Part 7 folded into the single upload set per ruling 3. #99 (de7d135) the 2 September hedge closed for good: all nine Luna Pro OpenRouter calls checked individually, all standard tier, “No data training” -- evidence upgraded from “probably clean” to provably. #100 (a12640a) round 5 sent to GPT-6 Astra ($12.21636, 435,612 prompt tokens, confirmed twice over); three new Criticals returned, none previously on record -- F1 (volatility never reaching stop construction), F2 (entry quality scored for the wrong direction), F3 (a stale decision-bar VWMA presented as current) -- all three independently reproduced against live code this session and RULED to fix, no limitation accepted for any. #101 (83e334d) F1 fixed: the production call to calculate_stop_targets() omitted volatility_state, defaulting to NORMAL regardless of measured conditions; now passes the same volatility_mode its sibling call already receives. Golden snapshot unmoved -- the pinned fixture's own volatility state produces the same multiplier the bug's default did. #102 (3f92230) F2 fixed: eq_trade_direction was derived from short_signal, which zeroes for reasons unrelated to direction, letting a fully bearish setup score as a long entry; now derived from raw_bias directly, the same signal decision_model.py's own action selection reads -- the identical ruling already made for that module on 2 September, one caller downstream. #103 (7fcf646) F3 fixed, closing round 5: VWMA's own forward-fill had no trailing-edge exception, so a decision-bar miss was silently carried forward from the previous bar; fixed at the source and VWMA added to the critical_indicators sweep, with three new regression tests. Accepted consequence: some currently-clean runs will now show DEGRADED, naming VWMA. code_hash moved once per fix, confirmed on independently built trees at each step: e774b93d... to f8e4cc7d... to 8b0e97fd... to 00c8d4ea.... Suite 429 to 432 passed under pytest, 362 to 365 under run_tests.py, watched error count unmoved. Golden snapshot unmoved across all three fixes, confirmed by running it each time. Register unchanged at 21 / 7 / 10 / 6. The release gate remains shut: fixed is not re-audited, and unresolved means both.
v1.27	September 13, 2026	Entries #94 through #103 reconstructed from the v1.25/v1.26 Document History rows' own text into entry_box() form — a document-integrity gap found where those two rows claimed the entries existed as numbered log entries and they did not.
v1.28	September 13, 2026	Entries #104 through #111 added, covering 13 September: Round 6 (Meta Muse Spark 1.3) sent and graded clean, Round 6's F1/F2/F3 fixed, the docs and handover-file commits, F4 and F5 fixed. code_hash and the golden snapshot addressed per fix within the entries above; nothing in this row's own change touches either.

Version	Date	Notes
v1.29	September 13-14, 2026	Entries #112 through #118 added, closing the gap this document had carried since Entry #111 -- PHASE7_NEXT.md's own ninth-patch note had flagged this file as stale, last regenerated before Round 6. #112 (7fbf77c, 9c5f954, 4e0110b) the encoding defect that corrupted round 6's report.md found and fixed at the source -- OpenRouter's response carries no explicit charset and requests silently fell back to Latin-1; report.md corrected, three more evidence files caught untracked the same day. Appended out of strict date order, deliberately: it precedes Entries #104-#111 chronologically, matching this document's own Entry #71/#72 precedent for corrections filed when found rather than renumbering. #113 (214c5d8) the git add -A trap staging two of Claude's own delivery files a third time, closed structurally -- *.patch and *_commit_message.txt gitignored -- rather than a fifth restatement of rule 38. #114 (2fc578b) send_fix_verification.py built: a narrower tool for verifying round 6's own findings were fixed, deliberately not sharing send_audit_round.py's repointed-every-round MODEL constant. #115 (59de938, e65a0f7) round-6 fix-verification sent and graded: F1 through F5 all Fixed, independently checked against code, plus one bonus finding (a duplicate @staticmethod) found and not yet fixed at that point. #116 (a70fc1b) the bonus finding fixed. #117 (66f1479) Viktor's request: the setup's LONG/SHORT direction stated in words under the DECISION line, sourced from the same field the actual decision logic uses so it cannot disagree with an authorized trade, with a deliberate non-answer for the one case (PLAN CONTRADICTS ACTION) where stating a side would repeat the 2 September wrong-direction-label defect. #118 (39e0e79) Claude's own unrequested NEUTRAL-case line overruled by Viktor -- silence instead, since a neutral bias is not a failed measurement. Of the ten commits in this row, seven -- #112's three, #113, #114, and #115's two -- are docs/tooling-only and left code_hash and the golden snapshot untouched, confirmed rather than assumed in each case; the other three (#116, #117, #118) each moved code_hash by exactly the amount predicted in their own entries, and none moved the golden snapshot.

Version	Date	Notes
v1.30	September 14-15, 2026	Entries #119 through #127 added, closing the gap since Entry #118 -- ten commits landed without a fresh regeneration, batched per the rule Entry #122 itself records. #119 (69f1972) the Portfolio Document built, the fourteen-section deliverable drawn from current source rather than the (found-stale) README.md, with a Luna-Pro/GLM-5.3 misattribution caught and fixed before Viktor ever saw it. #120 (eee218e) the AI-Attribution Statement built, closing item 5 and -- with #119 -- all six portfolio-ready criteria; its own entry-count table caught and corrected (111 vs the true 118) by parsing this file's AST directly rather than trusting a regex. #121 (59b915a) this document's own entry-point brought current for #119/#120, a handover-check gap otherwise. #122 (48757fe) five token-saving workflow rules recorded in Working practice, already adopted in the patch-delivery skill and memory the same day. #123 (f7f2e51) a correction, not new work: the five round-6 mutant-escape findings were already closed at 044b055 three days before this session existed; Claude had reported them open from this file's own narrative rather than from git. #124 (99e022e) Viktor's ruling: the release gate declared open, the project declared portfolio-ready, tagged portfolio-v1 -- closing the question Entries #121 through #123 all left open. #125 (63d8620) docs/build/session_handover_check.py built, answering four of the standing checklist's seven items mechanically; confidence-scored auto-pruning of the Lessons list considered and declined, on the project's own record-don't-hide-errors rule. #126 (ebb0a46, d5bb329) README.md brought current after sixteen days stale -- it still said the gate was shut the day after Entry #124 declared it open -- plus a stray gitignored delivery file found and left for Viktor to remove, and Goal B's backtesting-scoping findings written down rather than left only in chat. #127 (a441023) item 8 added to the handover checklist closing the README-currency gap #126 exposed, and the ignored-file-sweep question closed with no code change since #125's own script already ran that sweep every invocation. All ten commits in this row are docs/tooling-only or pure PHASE7_NEXT.md prose; code_hash confirmed unmoved throughout (38458f20779d2709ed6402c69407339f797c69089ace3b18922ca5c8bb289fbb, computed on the tip this document was regenerated from, matching every commit's own individual confirmation) and the golden snapshot was never in play. The release gate itself moved this row -- shut at the start of it, open and tagged by #124 -- the first row in this document's own history to record that.

Version	Date	Notes
v1.31	September 15-19, 2026	Entries #128 through #135 added, closing the gap since v1.30 -- eleven commits landed without a regeneration, batched per the rule Entry #122 records and regenerated at Viktor's call. #128 (dff7d00, f9e5127) Goal B specified and ratified before any of its code exists, with two of Claude's own errors of asserting from the documents instead of the code recorded inside it. #129 (4a97c32) the course correction -- an engine review for analytical soundness before any Goal B work -- and Viktor's pace ruling. #130 (4a97c32) PHASE7_NEXT.md split into NEXT, DECISIONS and HISTORY, proven a move by byte-for-byte reassembly, and the bias weights confirmed fingerprinted by reading the code. #131 (5be5d82, 2ebbd6b) three Tier-1 fixes, PHASE7_PINNED_DATA's silent live fallback first among them. #132 (f24a6e9, 2fee78f) ADX's second read removed from continuation_strength on Viktor's independence-only ruling. #133 (2fee78f) git add -A nearly sweeping an untracked PDF, and files added by name from then on. #134 (e2c6637, 742b514) four of the six bias factors found to measure one thing, recorded and not fixed, with two errors in the project's own record corrected. #135 (2f5fdaa, 5c73e24) the pre-push hook, on Viktor's ruling to stop and consult. #129 and #133 were split into entries of their own at Viktor's choice. The first row in this document's history in which code_hash moved, and it moved twice: 38458f20779d2709ed6402c69407339f797c69089ace3b18922ca5c8bb289fbb, unchanged on every commit from portfolio-v1 through 4a97c32, to 5080ccf0490502ad72dadaf6c3c525d19792139ea52d7fb15ee28a923ab5d189 at 5be5d82, to bb47ab537314953e16b2db2fcf24003ead64eb5538a30221a6578d518bb7d34d at f24a6e9, unchanged since -- each recomputed on that commit's own tree for this regeneration rather than taken from the commit messages. The golden snapshot was re-baselined once, at f24a6e9, 16 leaf fields as predicted. Entries before the 18 September split that cite PHASE7_NEXT.md for material now in PHASE7_DECISIONS.md or PHASE7_HISTORY.md are left as written, per this document's own no-silent-edits rule.
v1.32	September 19-20, 2026	Entries #136 through #140 added, closing the gap since v1.31 -- six commits landed without a regeneration, batched per the rule Entry #122 records and regenerated once the three mechanical items before the independent audit had all landed. #136 (e431714, e115272) PHASE7_NEXT.md made current at c7ced36, and a wrong prediction about git status recorded. #137 (b68de08) the pinned-data manifest check made portable, closing the standing LF/CRLF platform difference. #138 (16d3c1f) Viktor's ruling scrapping the seven open items in favour of an independent audit. #139 (9a35f1c) stale citations of PHASE7_NEXT.md fixed in live files, and the ruling that dated records -- this document included -- are not edited to follow a move. #140 (077d120) README.md brought current. code_hash bb47ab537314953e16b2db2fcf24003ead64eb5538a30221a6578d518bb7d34d on every one of the six commits, each recomputed on its own tree for this regeneration. Golden snapshot unchanged since f24a6e9. Phase7_Portfolio_Document.pdf and Phase7_AI_Attribution.pdf regenerated in the same commit, their text differing from the previous build only by 9a35f1c's pointer changes (compared word by word).

Version	Date	Notes
v1.33	September 20, 2026	Entry #141 added for 4629002, the one commit since v1.32, regenerated at Viktor's call rather than left for the next batch. #141 (4629002) the round-1 audit outputs recorded as lost, and README.md and PHASE7_HISTORY.md's 30 August reinstall entry corrected. code_hash bb47ab537314953e16b2db2fcf24003ead64eb5538a30221a6578d518bb7d34d on 4629002's tree, recomputed for this regeneration. Golden snapshot unchanged since f24a6e9. Phase7_Portfolio_Document.pdf and Phase7_AI_Attribution.pdf were rebuilt from 4629002 and their extracted text compared with the committed PDFs: identical, because 4629002 changed neither script. They were therefore not regenerated, and PHASE7_NEXT.md's statement at 4629002 that they were one commit behind is corrected there.
v1.34	September 22, 2026	Entries #142 through #164 added, closing the gap since v1.33 -- twenty-seven commits landed without a regeneration, by Viktor's choice under the batching rule Entry #122 records, and regenerated at his call on 22 September, before work order G and his rulings on findings 4, 5, 6, 7, 16 and 18. 65a0aef had planned the regeneration for after the audit rulings and only if the audit package included the Notes (Entry #144); the audit is not yet planned, so this regeneration is Viktor's later call and not that condition met. 3a899b5, v1.33's own regeneration, has no entry, as c7ced36 and c745a67 had none. Twenty-three entries for twenty-six commits: #142 and #143 (53394ff) round 3 ruled not independent; the suite's live call left in place and the empty rule number marked. #144 (92775ea, 982e70f, 65a0aef) landing facts, and why the Notes gap cannot fall below one. #145 (a9d4b1f) the run summary named the EV line. #146 (6e1baba) the four-runs independence claim corrected. #147 (b869a30) code_hash comparable only within one Python minor version. #148 (119c8a3) the TREND denominator. #149 (635a94e) the entry-point document's eleven corrections. #150 (ebb4e5c) the pre-backtest review and the delegated order of work. #151 to #153 (a530006, e3f3d51, afd8460) work orders B, C with D folded in, and E, with finding 16 found from a wrong prediction. #154 (49de810) the Constitution check. #155 and #156 (3f263c2) work order F, and the audit paused. #157 (aafded0, cb659f1) the live run moved before the commit. #158 (1cc2142) the deferred read and the audit change list. #159 (9c4917c) the SETUP DIRECTION tests. #160 to #164 (05a12c7, 2c7a7d1 with 5d3a4b4, 4e2b1c8, 486f1a5, 3bfa6b7 with 5e55eb8) findings 19 to 28. Each written from its commit's own message, and #155 and #156 also from PHASE7_DECISIONS.md's two 21 September rulings; none re-checked against the code or the running engine. code_hash recomputed under Python 3.12.3 on every one of the twenty-seven trees for this regeneration, and every value matched its commit message: bb47ab53... through 65a0aef; b0213766... at a9d4b1f; 35718f6b... at 119c8a3; ec88cf24... at a530006; ac02a155... at e3f3d51; 3e76c1c5... at afd8460; 44f7296b... at 3f263c2; f691c4d8... at 05a12c7; 82c14ef9... at 2c7a7d1; 4fe5084e... at 4e2b1c8; 2c8ebe32... at 486f1a5; b3c2308f8f3e05981af25ee82468c071f7bf0b9d49519b6e75bd78c37b5fb365 at 3bfa6b7, unchanged at 5e55eb8 -- eleven moves. The golden snapshot was re-baselined four times (a9d4b1f, e3f3d51, 3f263c2, 2c7a7d1), read from git log of the fixture. Phase7_Portfolio_Document.pdf and Phase7_AI_Attribution.pdf not regenerated: their scripts last changed at 6e1baba, which regenerated the Portfolio Document, and neither script reads another file.

Version	Date	Notes
v1.35	September 22, 2026	Entries #165 and #166 added, at Viktor's call the same day as v1.34. #165 is the 22 September ruling on when the audit resumes, carried by e804b64; e804b64 itself is v1.34's regeneration and, by the precedent #144 states, has no entry of its own, so the entry is for the ruling it carried, not the commit. #166 (c5dc4cd) e804b64's push filed and README's audit row brought current. The gap goes from two commits to one, the floor. Both entries were written from their commits' own messages and PHASE7_DECISIONS.md's 22 September section, not re-checked against the code or the running engine. code_hash b3c2308f8f3e05981af25ee82468c071f7bf0b9d49519b6e75bd78c37b5fb365 at both commits, recomputed under Python 3.12.3 on c5dc4cd's tree; e804b64's was computed on its applied tree before delivery. Golden snapshot unchanged since 2c7a7d1. The portfolio PDFs were not rebuilt, as at v1.34.
v1.36	September 29, 2026	Entries #167 through #183 added, closing the gap since v1.35 -- sixteen commits landed without a regeneration, by Viktor's choice under the batching rule Entry #122 records (no time pressure, 26 September), and regenerated at his call on 29 September, with the closed list done and before the audit package is built. bc48f59, v1.35's own regeneration, has no entry, by the precedent #144 states. Seventeen entries for fifteen commits: #167 (75682ee) the round-1 audit outputs recovered and filed. #168 the 22 September ruling on bias_score's weighting, filed by 75682ee. #169 (2582994) build and review open to any model at the cost of audit eligibility, and the pre-send check as audit preparation. #170 (1861208, eba6a2a) the stop-distance finding ruled evidence for findings 4 and 6. #171 (7d103d7) the move to E:. #172 (bf2e802) finding 16 ruled, and the phase-3 order. #173 (add8540) finding 16's code. #174 to #177 (57f0521, b0efe23, 76c8cde, 8b9ac1e) findings 7, 4, 6 and 5. #178 the ruling that a push going as predicted leaves nothing owed, filed by 8b9ac1e. #179 (d76ddfd) finding 18. #180 (7d0024e) work order G, and the closed list done. #181 (22cb39b) what “finished” means. #182 and #183 (68c6191) Fable 5.1 after the audit, and Mistral taken off the clean labs again. Each written from its commit's own message and, for the rulings, PHASE7_DECISIONS.md's sections; #183 also from this document's Entries #48 and #146; none re-checked against the code or the running engine. code_hash recomputed under Python 3.12.3 on every one of the sixteen trees for this regeneration, and every value matched its commit message: b3c2308f... through bf2e802; 36d99565... at add8540; 84f5c4ac... at 57f0521; b5e6af9e... at b0efe23; ed2710b6... at 76c8cde; c6361577... at 8b9ac1e; f4b23f94... at d76ddfd; c6a44d4d7ab1c8d36f2a07f6f2c78a1967ecfc135259cbe828b3daa275e571c9 at 7d0024e, unchanged at 22cb39b and 68c6191 -- seven moves. The golden snapshot was re-baselined three times, at add8540 (one added field, provenance.decision_candles), 76c8cde (11 fields) and 7d0024e (2 fields), each checked for this regeneration by a field-by-field diff of the fixture between consecutive commits. Phase7_Portfolio_Document.pdf and Phase7_AI_Attribution.pdf were not rebuilt: each was built from its script at 68c6191 and its extracted text compared with the committed PDF, identical.

Version	Date	Notes
v1.37	November 4, 2026	Entries #184 through #203 added, the first of three regenerations closing the gap since v1.36 -- eighty-eight commits had landed without one, by the batching rule Entry #122 records, and Claude chose to close the gap in three parts so that no part is more than a session's work. This part covers the nineteen commits from e7a94d1 to eb86846; c2b07fb onward are still owed and stated in PHASE7_NEXT.md. ec4e5fd, v1.36's own regeneration, has no entry, by the precedent #144 states. Twenty entries for nineteen commits: #184 and #185 (e7a94d1) the audit to Laguna S 2.1 with Part 7 after Parts 1-6, and what opens backtesting. #186 (d5ced0c) six questions before the send. #187 (59b747a) the pre-send token check. #188 (f256937) the test bugs classified, four enough. #189 (12b483e) the send in two requests, and rounds 5 and 6 confirmed to have shown Part 7 early. #190 (01f2892) rev 8. #191 (8b6fd00) the Part 7 document. #192 (e186423) the roster and ledgers fused, and two draft copies caught. #193 to #195 (86f30f1, 3e8191a, 1523e0c) round 7's two runs, 0 of 4 each, and the ruling between them. #196 (ff12e17) the audit delegated to Claude, round 8 to MiMo-V2.6-Pro. #197 (4a4c6ce) round 8's tooling. #198 (217ede6) rev 9. #199 to #201 (d27743f, a0276dd, 040a529) round 8's two runs and Part 7. #202 (9e42d30) round 8 accepted narrowly. #203 (eb86846) round 8 triaged. Each written from its commit's own message, which for the rulings summarises the PHASE7_DECISIONS.md entry it files; none re-checked against the code or the running engine. code_hash recomputed under Python 3.12.3 on every one of the nineteen trees and on 68c6191 before them: c6a44d4d7ab1c8d36f2a07f6f2c78a1967ecfc135259cbe828b3daa275e571c9 on all twenty, no move, as every commit's message records; tests/fixtures/golden_decision.json the same blob on all twenty, never re-baselined. Phase7_Portfolio_Document.pdf and Phase7_AI_Attribution.pdf were not rebuilt; neither script changed in these commits.

Version	Date	Notes
v1.38	November 4, 2026	Entries #204 through #226 added, the second of three regenerations closing the gap since v1.36. This part covers the twenty-two commits from c2b07fb to 67f6ec1, round 8's fixes and their verification; 0491b84 onward are still owed and stated in PHASE7_NEXT.md. Twenty-three entries for twenty-two commits: #204 to #206 (c2b07fb, 1b25f36, fcb7823) the three fixes - the stop, the trend gate, the regime. #207 to #214 (9e507b1, 6469a67, ae37795, 9d00d99, 27d3ed5, 9051ffd, e5ac20a) the label and record batch - B6, N5, F3, F5, F6, F8, F9 - with #209 Viktor's position of 6 October, filed by 6469a67. #215 (ef65f46) the fallbacks and repeated readings ruled into the batch. #216 to #222 (797bdc9, 69c92d6, 22a9a35, d77a8c4, 64ad625, 9ea1cf3, f00bef6) the fallback sweep, the repeated readings, B4's labels, B7, B9, B10 and Part 4's three test shapes. #223 to #226 (36de35d, 78ea37d, d3188f2, 67f6ec1) round 8's fix verification decided, built, answered and accepted. Each written from its commit's own message, which for the rulings summarises the PHASE7_DECISIONS.md entry it files, and #209 also from that entry; none re-checked against the code or the running engine. code_hash recomputed under Python 3.12.3 on every one of the twenty-two trees and on eb86846 before them, and every value matched its commit message: c6a44d4d... at eb86846; 2c6365fb... at c2b07fb; b7e7bc05... at 1b25f36; 952dbb02... at fcb7823; a1381630... at 9e507b1; e539bebc... at 6469a67; d7145d86... at ae37795; 8a807866... at 9d00d99; 80a781b7... at 27d3ed5; cb87efd8... at 9051ffd; 3f8f218f... at e5ac20a, unchanged at ef65f46; 76b17282... at 797bdc9; 7324ee98... at 69c92d6; 515adca8... at 22a9a35; 6e287836... at d77a8c4; e2f6b2ad... at 64ad625; bc1b9c2c517215d0b16eec4209462fa14664a89fbf7b24fcd565c255c45800d7 at 9ea1cf3, unchanged through 67f6ec1 - sixteen moves. The golden snapshot was re-baselined six times, at c2b07fb (11 leaves), 1b25f36 (6), fcb7823 (7), 9e507b1 (3), 6469a67 (8) and 22a9a35 (4), each checked for this regeneration by a field-by-field diff of the fixture between consecutive commits, and each the leaves its message predicted. Phase7_Portfolio_Document.pdf and Phase7_AI_Attribution.pdf were not rebuilt; neither script changed in these commits.

Version	Date	Notes
v1.39	November 6, 2026	Entries #227 through #252 added, the third of four regenerations closing the gap since v1.36. v1.37 and v1.38 planned three; the third part, forty-eight commits, was more than a session's work, so Claude split it in two on 4 November and Viktor did not object: this part covers the twenty-four commits from 0491b84 to b8f88d8, the forward measurement, Amazon's audit plan and its four held-back defects, the sort of the list for after the audit and wave 1, the rulings on the engine's indicators and Goal B, and the ruling that a commit names the tree verified before delivery; a1c4218 to f8805a6 are the fourth part, still owed and stated in PHASE7_NEXT.md. Twenty-six entries for twenty-four commits: #227 (0491b84) the forward measurement and main.py taking the pairs. #228 (a235783) README's status rows and the measurement's first checks. #229 and #230 (8d5d64a, a48d9e5) Amazon's audit scored on four held-back defects, and its fallback. #231 (fbd7877) and #233 (d2dddf7) the scheduled task's runs read. #232 (8311167) the four defects, H1, H2, H3 and H5. #234 (aa9e7ac) the sort of the list. #235 and #236 (391c338) F4's third path Critical and the release gate closed, and entries 38 and 22 ruled. #237 to #240 (8360923, ed005a0, e88e0dd, d508b29) wave 1's entries 13, 21, 28, 48, 36 and 37. #241 to #243 (9f2f3d7, d410bfa, 34d8e3b) the document on the engine's indicators, its tool and the tool's landing. #244 (ad34c02) entry 49. #245 and #246 (89d2367) the eight rulings on the indicators, and the trend multiplier ruled Critical with the held-back defects the redesign reaches. #247 to #249 (b387c03, 9862568, 3d7e6cb) entries 51, 54 and 55. #250 and #251 (78204ff, 0d86fda) entries 26 and 31, Goal B. #252 (b8f88d8) the Verified-Tree ruling. Each written from its commit's own message, which for the rulings summarises the PHASE7_DECISIONS.md entry it files; none re-checked against the code or the running engine. code_hash recomputed under Python 3.12.3 on every one of the twenty-four trees and on 67f6ec1 before them, and every value matched its commit message: bc1b9c2c... at 67f6ec1; 816d4951bb0322763ec4abc4b9ed8161d5b902372823350e5272f2fb7176e380 at 0491b84, unchanged through b8f88d8 -- one move. The golden fixture was not re-baselined in this range: its blob is the same on all twenty-five trees, last changed at 22a9a35. Phase7_Portfolio_Document.pdf and Phase7_AI_Attribution.pdf not rebuilt: neither script changed in these commits.

Version	Date	Notes
v1.40	November 7, 2026	Entries #253 through #276 added, the fourth and last of four regenerations closing the gap since v1.36. This part covers the twenty-five commits from a1c4218 to 0a518fd: the commit-msg hook and the rulings on how a landing is checked and recorded, README brought current, a candle touching both the stop and T1, what may be tuned after the paper-trading year, volume's two gates, the forward measurement's reading of 31 October, the ruling that a session does not ask what to do next, entries 58, 59 and 60, the redesign's design draft, how the reading of 14 November measures it, and that reading's script. v1.39's row named the fourth part as a1c4218 to f8805a6, the range Claude proposed on 4 November, when f8805a6 was the tip; the part runs to 0a518fd, v1.38's regeneration, which landed after. f8805a6 and 0a518fd, v1.37's and v1.38's own regenerations, have no entry, by the precedent #144 states. Twenty-four entries for twenty-three commits: #253 (a1c4218) the commit-msg hook, entry 56. #254 (713e084) entry 50, README and the watched error count. #255 (134f85b) README's audit history and layout. #256 and #257 (fad0588, 3dbd8dc) a landing check's standing gaps, and a ruling on the delivery procedure changing the skill. #258 (a2ca918) a candle touching both the stop and T1. #259 (d7d3ae0) what may be tuned. #260 (b233c4a) volume's two gates. #261 (6d073f9) the reading of 31 October. #262 (8c810b1) a session does not ask what to do next. #263, #265, #268, #269 and #274 (248037d, f6aa126, d6fba83, 7c23886, 7b93769) the five commits that recorded only a landing. #264 (e748b0b) entry 58. #266 and #267 (e69d5bb, e8b42b8) the design draft's versions 0 and 1. #270 and #271 (e939734, 1b82ed0) how the reading measures the draft, and its script. #272 and #273 (7ff8128, 5fe8522) entry 59 found and fixed. #275 and #276 (b990c93) entry 60, HISTORY keeping what each rewrite removes, and the correction that a landing is never itself a commit owed. Each written from its commit's own message, which for the rulings summarises the PHASE7_DECISIONS.md entry it files, and #276 also from that entry; none re-checked against the code or the running engine. code_hash recomputed under Python 3.12.3 on every one of the twenty-five trees and on b8f88d8 before them: 816d4951bb0322763ec4abc4b9ed8161d5b902372823350e5272f2fb7176e380 on all twenty-six, no move, as every commit's message records. The golden fixture was not re-baselined in this range: its blob is the same on all twenty-six trees, last changed at 22a9a35. Phase7_Portfolio_Document.pdf and Phase7_AI_Attribution.pdf not rebuilt: neither script changed in these commits. a0ca299, v1.39's own regeneration, has no entry by the same precedent, so with this version the gap since v1.36 is closed.

Version	Date	Notes
v1.41	November 7, 2026	Entries #277 through #283 added: seven entries for the seven commits from 706e858 to 82ea40e, one regeneration, as v1.40 left no commit behind. #277 (706e858) round 8's F2 ruled Critical. #278 (4ab59d7) the forward measurement's count, a filed, tested script. #279 (94b7e96) the redesign draft's version 2, and the finding that its direction and its band cannot meet. #280 (a11ffc4) the redesign's direction ruled, EMA50's slope and the EMAs' order, with computation 1 of the decision of 2 November amended to it. #281 (542085f) version 4, which distance STATUS reads. #282 (b874ec1) version 5, the second pinned decision, and entry 61. #283 (82ea40e) the redesign's AGGRESSIVE ruled, the close inside the band. Each written from its commit's own message, which for the rulings summarises the PHASE7_DECISIONS.md entry it files, and #283 also from that entry; none re-checked against the code or the running engine. code_hash recomputed under Python 3.12.3 on bc250fc and every one of the seven trees: 816d4951bb0322763ec4abc4b9ed8161d5b902372823350e5272f2fb7176e380 on all eight, no move, as every commit's message records. The golden fixture was not re-baselined in this range: its blob is the same on all eight trees, last changed at 22a9a35. Phase7_Portfolio_Document.pdf and Phase7_AI_Attribution.pdf not rebuilt: neither script changed in these commits. bc250fc, v1.40's own regeneration, has no entry, by the precedent #144 states.
v1.42	November 9, 2026	Entries #284 and #285 added: two entries for the two commits from 7dac328 to 340ee8d, one regeneration, as v1.41 left no commit behind. #284 (7dac328) the redesign's entry quality ruled to have no structure multiplier. #285 (340ee8d) the reading of 14 November's computation 7, the CHOCH count, by Claude's amendment to the decision of 2 November. Each written from its commit's own message, which for the ruling summarises the PHASE7_DECISIONS.md entry it files, and both also from their entries, read this session; none re-checked against the code or the running engine. code_hash recomputed under Python 3.12.3 on 9db4ae2 and both trees: 816d4951bb0322763ec4abc4b9ed8161d5b902372823350e5272f2fb7176e380 on all three, no move, as every commit's message records. The golden fixture was not re-baselined in this range: its blob is the same on all three trees, last changed at 22a9a35. Phase7_Portfolio_Document.pdf and Phase7_AI_Attribution.pdf not rebuilt: neither script changed in these commits. 9db4ae2, v1.41's own regeneration, has no entry, by the precedent #144 states.

Version	Date	Notes
v1.43	November 10, 2026	Entries #286 through #288 added: three entries for the three commits from 6bb82cc to a60d66b, one regeneration, as v1.42 left no commit behind. #286 (6bb82cc) the redesign's exit watch ruled to drop its exhaustion and reversal flags, with Claude's correction that the exit watch's CHOCH flag stays. #287 (e4363d5) the route of Amazon's audit, OpenRouter pinned to Amazon Bedrock, by Claude's decision under the delegation of 4 October. #288 (a60d66b) its model ruled, Nova 2 Lite, Viktor's own position, replacing the decision's second point. Each written from its commit's own message and from the PHASE7_DECISIONS.md entry it files, read this session; none re-checked against the code or the running engine, and none of the entries' reads of Amazon's or OpenRouter's listings repeated. code_hash recomputed under Python 3.12.3 on 82b824d and all three trees: 816d4951bb0322763ec4abc4b9ed8161d5b902372823350e5272f2fb7176e380 on all four, no move, as every commit's message records. The golden fixture was not re-baselined in this range: its blob is the same on all four trees, last changed at 22a9a35. Phase7_Portfolio_Document.pdf and Phase7_AI_Attribution.pdf not rebuilt: neither script changed in these commits. 82b824d, v1.42's own regeneration, has no entry, by the precedent #144 states.

Future entries get appended below this line, in order, each with its own number, date, and status tag — this document grows with the project rather than being rebuilt each time.